# Tax Heterogeneity, Duration Risk, and the Term Structure: An Application to the Municipal Bond Puzzle*

Kevin Zheyuan Cui†

October 20, 2026

**Abstract**

I develop a term-structure model that jointly prices taxable and tax-exempt bonds. Heterogeneous taxes and balance sheets endogenously segment ownership by maturity under short-sale constraints. The model explains the municipal bond puzzle: short municipal yields closely match the after-tax taxable benchmark of top-bracket households, who face the highest taxes, while long yields exceed it. Lower-tax investors whose liabilities offset interest-rate risk bear duration risk more cheaply and hold long municipal bonds despite their smaller exemption benefits. Yields reflect both holders' taxes and risk-bearing costs. Without targeting municipal yields, the calibrated model explains 83 percent of the 118-basis-point thirty-year gap. Empirically, household tax effects weaken with maturity while institutional effects strengthen, as the mechanism predicts. Through portfolio substitution, investor tax policies affect bond yields in both taxable and tax-exempt markets.

**Keywords:** Term Structure, Investor Heterogeneity, Taxes, Endogenous Segmentation, Asset Ownership, Municipal Bonds
**JEL Codes:** G11, G12, G22, H24, H71, H74

*I am grateful to Moritz Lenel, Wei Xiong, and Motohiro Yogo for their guidance. I also thank Yacine Ait-Sahalia, Caio Almeida, Markus Brunnermeier, Fousseni Chabi-Yo, Aditya Chaudhry, Daniel Chen, Zhiyu (Julie) Fu, Nicolae Garleanu, Sergei Glebkin, Zhouzhou Gu, Nicolas Hommel, Paymon Khorrami, Jane Li, Shiqi Li, Ziang Li, Jonathan Payne, Yinan Qiu, Julia Selgrad, Amy Huber Wang, Lingxuan (Sean) Wu, Yinjie Yu, Yao Zeng, William (Yinhong) Zhao, Haonan Zhou, and Haoxing Zhu, and participants at the MPSS Conference and Princeton University Seminars for constructive feedback.

†Princeton University, Department of Economics.

# 1 Introduction

Whose taxes determine bond yields when investors pay different tax rates on the same bonds? U.S. municipal bonds provide a natural setting to study this question because the same tax exemption offers different benefits to investors facing different tax rates. These bonds finance state and local governments through a market of about $4 trillion. Municipal bond interest is exempt from federal income tax, and the exemption is worth most to top-bracket households, who face the highest tax rates. The standard benchmark in the literature therefore uses these households' tax rate to determine municipal yields: at every maturity, the municipal yield should equal their after-tax yield on a comparable taxable bond (Schwert, 2017; Fleckenstein and Longstaff, 2025). After callability and the issuance premium are removed from municipal yields, the household benchmark fits short maturities, but thirty-year municipal yields exceed it by 118 basis points. The tax rate implied by these municipal yields and comparable taxable yields falls from about 38% at two years, close to the top-bracket household rate, to just 21% at thirty years. This municipal bond puzzle has been documented for over thirty years (Feenberg and Poterba, 1991). Under the U.S. tax code, however, investors face very different tax rates, and the data suggest that top-bracket households hold only about half the municipal market. Lower-tax households, banks and insurers hold large amounts even though the exemption is worth less to them. Understanding relative yields therefore calls for a model that accounts for this heterogeneity.

Motivated by this heterogeneity in ownership, I explain the puzzle with an equilibrium term-structure model in which multiple investor groups differ in their taxes and the duration of their liabilities. They choose portfolios jointly across taxable and tax-exempt markets and maturities, subject to short-sale constraints. Their portfolio choices endogenously segment ownership and determine both yield curves. Tax differences draw high-tax investors toward tax-exempt municipal bonds and lower-tax investors toward taxable bonds. However, long-term liabilities allow some lower-tax investors to bear interest-rate risk more cheaply than high-tax households. When rates rise, the value of these liabilities falls, offsetting losses on their bonds. This advantage grows with maturity and draws them into long municipal bonds despite their smaller exemption benefits. Because the exemption saves these investors less tax, they accept a smaller reduction in yield relative to taxable bonds than high-tax households do. Long municipal yields therefore reflect their smaller tax benefits, explaining why yields exceed the household benchmark. Calibrated to U.S. taxes, balance sheets and holdings, without targeting municipal yields, the model explains 83% of the thirty-year puzzle. Untargeted empirical tests support its prediction that household tax effects weaken with maturity while those of banks and insurers strengthen. The framework applies beyond municipals because it jointly determines relative yields across markets. In the counterfactuals, investor substitution transmits tax reforms to the financing costs of both municipal and taxable issuers.

Household mortgage debt, banks' assets and deposit funding, and insurers' long-term obligations create different net exposures to rate movements. A two-investor illustration with common risk aversion shows how these balance-sheet differences interact with exemption benefits to determine

ownership across maturities. The high-tax household gains more from the exemption, while the lower-tax investor's stronger liability hedge lowers its cost of bearing rate risk. The resulting duration benefit grows with a bond's rate exposure and hence with maturity, whereas the difference in exemption benefits stays constant. Beyond a maturity threshold, the duration benefit outweighs the exemption disadvantage. In this illustration, the household therefore holds shorter municipal bonds, while the lower-tax investor holds longer municipal bonds and risky taxable bonds.

This sorting raises municipal yields above the household benchmark through three channels. First, lower-tax investors value the exemption less and therefore accept a smaller yield reduction, keeping the long municipal yields they price closer to taxable yields than the benchmark predicts. Second, the household alone holds shorter risky municipal bonds and bears rate risk at a higher cost than the lower-tax investor holding comparable taxable bonds. The household's higher required compensation raises municipal yields above the benchmark even at these maturities. Third, taxes scale risk as well as returns. Under the household benchmark, municipal yields and their sensitivity to interest rates equal the corresponding taxable quantities multiplied by the fraction of taxable interest the household keeps. Lower-tax investors keep a larger fraction, so the long municipal bonds they price are more sensitive to rates than this benchmark predicts. Holders require compensation for that additional rate risk, and household portfolio substitution carries the effect to shorter municipal bonds. All three channels vanish at the riskless shortest maturity, where the household's tax advantage determines relative yields.

Other motives can also lead lower-tax investors to hold long municipal bonds and generate the same qualitative pricing pattern. Lower trading costs or better default diversification can do so when the advantage grows with maturity and applies equally to taxable and tax-exempt bonds (Section 4.7). I focus on liability hedging because investors' widely differing liability durations provide an economically important and measurable source of differences in the cost of bearing interest-rate risk. Observed balance sheets and maturity holdings discipline these differences in the quantitative model.

The quantitative application asks how much of the puzzle observed taxes and portfolios can explain. It includes high- and lower-tax households in each state, banks, property-and-casualty insurers, life insurers and tax-neutral investors trading in state municipal markets and a common taxable market. I measure tax treatment, bond supplies, and investor balance-sheet sizes directly from data and estimate the rate process from historical short rates. Observed ownership and the maturities investors hold mainly discipline their unobserved exposure to interest rates, while taxable yields help discipline risk aversion. Municipal yields and empirical test coefficients are not targeted in calibration. The model accounts for 97 of the 118 basis points at thirty years, and measured default, liquidity and capital-gains differentials explain most of the remainder.

I evaluate the model against the household benchmark using three empirical tests, none targeted in the calibration. The model predicts that variation in top-bracket households' exemption benefits is fully reflected in implied taxes at the shortest maturity and less strongly at longer maturities. Changes in lower-tax investors' exemption benefits have no effect at the shortest maturity

and matter more at longer maturities. First, most states exempt residents' interest on their own municipal bonds from state income tax. Cross-state variation in state tax rates is reflected almost one-for-one in implied state taxes at short maturities and progressively less at longer maturities. Second, bank-qualified designation gives banks an additional tax benefit on a subset of municipal bonds issued by small borrowers. This bank-only benefit is barely reflected in short municipal yields and increasingly reflected at longer maturities. Third, I compare municipal yield changes with changes in the treasury reference yield at the same maturity. This relative response rises with maturity, consistent with lower-tax holders making long municipal yields more sensitive to interest rates.

Two counterfactual reforms show how removing different tax benefits changes financing costs and tax receipts. The first removes a state's exemption for residents' interest on its own municipal bonds, as Illinois, Iowa and Wisconsin generally do. Residents retain the federal benefit of holding municipal bonds but, as in states without an income tax, no longer have a tax incentive to favor home-state over other municipal bonds. High-tax households therefore mainly switch between municipal markets, with relatively modest changes in the maturities they hold. The state taxes the interest on these replacement investments, so selling local bonds does not substantially erode its tax base. Meanwhile, long municipal yields also reflect the valuations of institutions and nonresidents whose tax treatment does not change. Long yields therefore respond only partly to the lost household exemption, consistent with the partial pass-through of state taxes at long maturities documented in the first empirical test. Across seven state experiments averaged over 2000–2017, added state receipts exceed local issuers' added financing costs: own-state issuers incur \$0.68 to \$0.88 in added annual financing costs per dollar of added state receipts. Portfolio substitution also raises borrowing costs for other states' municipal issuers and taxable issuers.

The second reform removes the federal exemption as well, eliminating the tax advantage of municipal over taxable interest for every investor. Taxable municipal issuance has a precedent in the 2009–2010 Build America Bonds program, which subsidized issuers' interest payments. High-tax investors can no longer preserve that exemption benefit by switching between state municipal markets. They reduce their municipal holdings and shift toward shorter bonds or cash, which generate less interest than long bonds and limit the revenue collected from repeal. Tax-neutral investors, who pay no tax on the added interest, take over more long bonds, while lower-tax households and life insurers remain substantial holders. The shift of rate risk to tax-neutral investors raises required yields in both markets. The increase in financing costs consequently exceeds added tax receipts. Averaged over 2000–2017, municipal and taxable issuers together incur \$1.20 in added annual costs per dollar of added federal and state receipts, including \$0.13 borne by taxable issuers. Annual financing costs rise by \$46.3 billion, while annual tax receipts rise by \$38.7 billion. The excess becomes higher aggregate after-tax interest income, distributed unevenly across investors. Lower-tax households, life insurers and tax-neutral investors gain, while high-tax households, banks and property-and-casualty insurers, which previously benefited more from the exemption, lose on average.

The framework is not only about municipal bonds. It explains relative bond pricing across taxable and tax-exempt markets, with investor portfolio substitution transmitting tax changes in one market to yields in the other. I apply it to municipal and taxable bonds because municipal tax rules provide clean variation to test the mechanism and address a longstanding pricing puzzle. Its applications extend to other settings where tax treatment differs across investors and markets, including state-tax differences between treasury and corporate bonds, tax-favored accounts and cross-border withholding. More broadly, heterogeneous taxes, together with balance-sheet differences and limits to arbitrage, shape endogenous market segmentation and are important for understanding bond yields and the effects of tax policy.

Sections 2–3 review related work and present the tax, ownership and yield evidence. Section 4 develops the model, Sections 5–6 evaluate its quantitative and empirical implications, and Section 7 studies tax policy. Formal results, proofs and supporting evidence are in the appendix.

## 2 Related Literature

The paper contributes to research on how taxes affect portfolio choice and asset prices. Theoretical work studies how tax differences shape portfolios and equilibrium prices (Miller, 1977; Auerbach and King, 1983; Dybvig and Ross, 1986; Dammon and Green, 1987), and Sialm (2009) provides evidence that taxes are reflected in equity valuations. I study how these differences shape the bond yield curve. Under short-sale constraints, investors choose jointly across markets and maturities. Their taxes and costs of bearing rate risk determine who holds each bond and whose tax treatment its yield reflects.

Research on market segmentation links the yield curve to bond supply, investors' maturity preferences and their capacity to bear risk (Vayanos and Vila, 2021; Greenwood and Vayanos, 2014; Guibaud et al., 2013). Portfolio adjustment also transmits supply shocks across asset classes and currency bond markets (Greenwood et al., 2018, 2023).[1] I explain how taxes generate partial segmentation across markets and maturities. Participation responds to taxes and balance-sheet exposures, so tax changes alter both yield curves through joint portfolio choice.

The model's explanation of participation connects to research on institutional balance sheets and asset pricing. Intermediaries' capital and exposure to interest-rate movements affect the risk premia they require (He and Krishnamurthy, 2013; Haddad and Sraer, 2020). Kekre et al. (2024) show how losses on arbitrageurs' bond portfolios can amplify the effect of monetary tightening on long yields. Insurers' long-term liabilities also shape their demand for long bonds and their capacity to bear credit risk (Domanski et al., 2017; Li, 2026).[2] I show how a balance-sheet advantage in bearing rate risk can offset a smaller tax benefit, explaining why some lower-tax investors hold long tax-exempt bonds.

[1] Constraints on arbitrage can limit this adjustment (Gromb and Vayanos, 2018; Siriwardane et al., 2025). Related empirical work measures investors' asset demand and household rebalancing (Koijen and Yogo, 2019; Gabaix et al., 2026).

[2] Municipal offering spreads reflect banks' rate exposure and funds' liquidity needs (Pati, 2023). Liability structures also help explain investors' responses to credit-market downturns (Coppola, 2025).

The municipal bond puzzle has motivated a longstanding literature on relative bond yields. Early explanations emphasized tax benefits and investor sorting across maturities (Mussa and Kormendi, 1979; Hendershott and Koch, 1980; Kidwell and Koch, 1983; Skelton, 1983). The Tax Reform Act of 1986 reduced banks' tax advantages and municipal holdings (Hines and Horton, 2024), yet long municipal yields remained high relative to after-tax taxable yields (Green, 1993; Chalmers, 1998). Subsequent research examined the taxation of bond cash flows, tax risk, default risk, liquidity, convenience premia and callability (Ang et al., 2010; Longstaff, 2011; Schwert, 2017; Fleckenstein and Longstaff, 2025; Brown, 2011).[3] I assess bond-side adjustments on a common household after-tax scale and find that they leave a substantial long-maturity gap. My model explains most of this remaining gap through endogenous participation by investors with different taxes and costs of bearing rate risk.

Empirical research on municipal tax policy links state tax benefits to local ownership, financing conditions and risk sharing across states (Hendershott and Kidwell, 1978; Pirinsky and Wang, 2011; Babina et al., 2021; Moussawi, 2022, 2024). Changes in banks' tax benefits likewise affect municipal investment and access to financing (Hines and Horton, 2024; Dagostino, 2025; Benson and Marks, 2024).[4] I test how household and institutional tax effects vary with maturity. Household state-tax pass-through declines, while the yield reduction associated with bank qualification rises. This opposite pattern tests the participation mechanism against a single household tax rate applied across the curve.

Fiscal analyses of municipal tax exemption examine its revenue cost and benefits to borrowers. Poterba and Ramírez Verdugo (2011) show that estimates of forgone revenue depend on how households would reinvest without the exemption. Garrett et al. (2023) show how auction competition affects the borrowing-cost savings from tax benefits. I evaluate reforms by running counterfactuals that incorporate the equilibrium impact of investor reallocation, capturing changes in tax receipts and financing costs.

## 3 Taxes, Ownership and the Municipal Puzzle

This section establishes the empirical setting for the model. I describe the data and the tax benefits available to different investors, then show that municipal ownership extends well beyond high-tax households. I measure the municipal puzzle and examine whether bond characteristics or the choice of taxable benchmark explain its maturity pattern.

[3] Mankiw and Poterba (1996) link the two bond markets through equity holdings. Related pricing, tax-policy and maturity-clientele analyses include Fama (1977); Livingston (1982); Poterba (1986); Brown and Sirmans (2011). Trading-cost evidence informs my liquidity adjustment (Harris and Piwowar, 2006; Green et al., 2007; Bagley et al., 2023). For surveys, see Cestau et al. (2019); Spreen and Gerrish (2022). Appendix C.4 compares the model with Green's benchmark.

[4] Municipal fund flows and fragility also affect yields and financing (Li et al., 2024; Adelino et al., 2026).

### 3.1 Data

I combine data on bond issuance, investor holdings and taxes. Mergent provides issuance records, ratings, bond characteristics and call schedules. Insurers' filings in Schedule D of the National Association of Insurance Commissioners (NAIC) and a pension holdings panel collected through Freedom of Information Act requests allow me to examine institutional portfolios at the level of individual bond holdings. I use the Financial Accounts to measure aggregate holdings, the Survey of Consumer Finances to measure household balance sheets, and IRS Statistics of Income to document household tax-exempt interest income across income brackets.

The main yield comparison covers 2000–2020. Its taxable benchmark is FRED's A-or-above corporate zero-coupon curve. Federal Reserve treasury curves (Gürkaynak et al., 2007) provide the reference for the interest-rate response test. I measure trading and liquidity using municipal transactions from the Municipal Securities Rulemaking Board (MSRB) and corporate transactions from the Trade Reporting and Compliance Engine (TRACE) Enhanced database. Markit provides municipal and corporate credit default swap quotes, and the Tax Foundation provides state tax schedules. Appendix D.6 details the sources and their coverage.

### 3.2 Municipal issuance

Table 1: Municipal issuance: descriptive statistics (sample means).

| Variable | Full sample | Tax-exempt, non-AMT | |
|---|---|---|---|
| | | AA and above (score $\leq 3$) | A or above (score $\leq 7$) |
| Offering yield (%) | 3.11 | 2.82 | 2.96 |
| Corporate benchmark yield (%) | 4.47 | 4.33 | 4.41 |
| Maturity at issuance (years) | 10.15 | 10.76 | 10.86 |
| Offering price (per $100 par) | 102.88 | 104.68 | 103.85 |
| Issue rating (1 = AAA, 3 = AA, 7 = A−) | 3.80 | 2.18 | 3.55 |
| Implied tax rate $1 - y^M/y^C$ | 0.29 | 0.35 | 0.33 |
| Share callable | 0.46 | 0.51 | 0.51 |
| Share tax-exempt, non-AMT | 0.92 | 1.00 | 1.00 |
| Sample size (CUSIPs) | 3,014,420 | 984,210 | 1,990,791 |

*Notes.* The table reports bond-level means for issues of 2000–2020 before the call and issuance-premium adjustments, with one observation per municipal issue (CUSIP). The full sample contains all issues. The AA-and-above and A-or-above columns restrict to federally tax-exempt, non-AMT issues with average agency score $\leq 3$ and $\leq 7$, the latter matching the A-or-above corporate credit band. The implied-tax row averages the bond-level ratio $1 - y^M/y^C$, winsorized at the first and ninety-ninth percentiles, so it need not equal one minus the ratio of the displayed yield means. Each row uses non-missing observations, so the full-sample rating mean uses the 2,292,092 rated issues. The corporate benchmark is the corporate zero yield for the issue month and maturity.

Table 1 describes municipal bonds issued during 2000–2020, with one observation per bond identifier (CUSIP). The full sample contains more than three million issues, with an average maturity of about ten years. Most pay federally exempt interest: taxable municipal bonds account for 5.8% of issues, bonds subject to the alternative minimum tax (AMT) account for 1.9%, and the remaining issues are tax-exempt, non-AMT bonds. Call provisions allow issuers to redeem bonds

before maturity and are common in the data. Callable bonds account for 46% of all issues and 51% of the high-grade, tax-exempt samples, making callability important for comparing yields.

The table also reports two high-grade samples of tax-exempt, non-AMT bonds. The broader sample includes A-or-above issues (average agency score $\leq 7$ on the twenty-two-notch scale) and matches the credit band of the corporate benchmark used for the main yield curves. The main empirical specifications use the AA-and-above sample ($\leq 3$) to minimize the impact of default risk.

### 3.3 Bond taxation and the representative-tax benchmark

For a taxable household, interest on taxable bonds is ordinary income, while interest on tax-exempt municipal bonds is federally exempt. Interest includes both coupon payments and annual original-issue-discount accrual. Institutions generally do not receive the full exemption benefit from holding municipal bonds. For example, property-and-casualty (P&C) insurers pay an effective federal tax of 5.25% on municipal interest. Let $\tau_i^C$ and $\tau_i^M$ denote investor $i$'s effective tax rates on taxable and municipal interest. I define the relative after-tax share as the fraction of taxable interest retained divided by the fraction of municipal interest retained:

$$\xi_i \equiv \frac{1-\tau_i^C}{1-\tau_i^M}. \tag{1}$$

Investor $i$'s relative municipal tax benefit, the exemption benefit, is $1-\xi_i = (\tau_i^C - \tau_i^M)/(1-\tau_i^M)$. A smaller $\xi_i$ means a larger benefit. When municipal interest is fully exempt, $\xi_i = 1-\tau_i^C$, so the benefit equals the tax rate on taxable interest. Appendix D explains the underlying tax rules.

The representative-tax benchmark uses the top-bracket household's tax treatment at every maturity. Write $y_T^M$ and $y_T^C$ for municipal and taxable yields at maturity $T$, and $H$ for this household. Under the benchmark, the two bonds offer the household the same after-tax yield, $y_T^M = \xi_H y_T^C$. I call $\xi_H y_T^C$ the household benchmark and $1-\xi_H$ the household wedge. The benchmark's yield-implied tax is therefore $1 - y_T^M/y_T^C = 1-\xi_H$. With fully exempt municipal interest, it equals $\tau_H^C$. Investors with different $\xi_i$ cannot all be indifferent at this yield ratio. Their participation in the municipal bond market is the central question of the model.

### 3.4 Investor taxes and ownership

The tax benefit of municipal interest varies substantially across investors in the roughly $4 trillion municipal market. Table 2 summarizes these differences before the 2017 Act. For households and insurers, the net exemption benefit is $1-\xi_i$ in (1). Households in higher tax brackets receive a larger benefit because they pay more tax on taxable interest. Exemption from home-state income tax adds to the benefit on in-state municipal bonds. Appendix D gives the detailed tax treatment of each investor group and the changes after the 2017 Act.

Insurers receive smaller benefits than high-tax households because exempt interest reduces their deductions. The reduction is less severe for P&C insurers than for life insurers, giving P&C insurers

Table 2: Investor heterogeneity in the municipal market before the 2017 Tax Cuts and Jobs Act

| Investor type | Federal treatment | State treatment | Net exemption benefit |
|---|---|---|---|
| High-tax household | Interest exempt | Exempt on in-state bonds only | 0.35–0.43 |
| Middle-tax household | Interest exempt | Exempt on in-state bonds only | 0.25–0.33 |
| Low-tax household | Interest exempt | Exempt on in-state bonds only | 0.10–0.15 |
| P&C insurer | Exempt, less proration | Premium taxes, no distinction | 0.31 |
| Life insurer | Exempt, less proration | Premium taxes, no distinction | 0.21 |
| Bank | Exempt, funding cost disallowed | Varies by state | 0.15–0.33 |
| Other (pensions, etc.) | Varies | Varies | $< 0.05$ |

*Notes.* For households and insurers, the net exemption benefit is $1 - \xi_i$, from statutory rates and proration over 2000–2017. For banks it also deducts the financing-tax cost described below. Appendix Table 52 gives the values after the 2017 Act. Household values are federal marginal rates in three groups (35% and above, 25% to 33%, below 25%), and the top value includes the 3.8% net investment income tax from 2013. The quantitative model pools middle- and low-tax households into one lower-tax group, whose benefit with home-state taxes is 0.21 to 0.27 across states in 2015–2017, against 0.43 to 0.50 for the high-tax group (Table 3). Proration cuts a P&C insurer's losses-incurred deduction by 15% of its exempt interest and a life insurer's reserve deduction by the policyholders' share. Section 265(b) disallows the share of a bank's interest expense allocable to exempt holdings, of which 80% stays deductible for bank-qualified bonds. The bank value, net of that disallowance, is the break-even implied tax, equal to the bank's tax rate minus its annual financing-tax cost per dollar of municipal holdings divided by the taxable yield. That cost is 1.02% in 2000–2002 and 0.11% in 2015–2017, and the bank range covers the six three-year periods of 2000–2017 and maturities of two to thirty years. States generally tax insurers through premium taxes, which do not distinguish municipal from taxable interest, and banks through franchise or income taxes.

the larger benefit. Banks generally lose the interest deduction allocated to municipal holdings, with partial relief for bank-qualified bonds, so the table reports their benefit net of this financing-tax cost.[5] Tax-exempt institutions such as qualified pension funds and endowment funds already receive taxable interest free of current income tax and gain no additional benefit from the municipal exemption. Public pension funds in the sample hold almost no municipal bonds (Appendix C.8).

Figure 1 shows that municipal ownership spans these tax groups. Panel (a) reports aggregate holdings: households and municipal funds hold the largest amounts, while banks and P&C insurers also hold substantial portfolios despite smaller exemption benefits. P&C insurers allocate about 40% of their bond portfolios to municipals and maintain a larger municipal portfolio share than life insurers throughout the sample (Appendix C.8). Panel (b) documents household tax-exempt interest income across income brackets. Assigning municipal-fund holdings to households and allocating household holdings using tax-exempt interest shares puts roughly half the market with high-tax households.[6] The remaining holdings by investors with smaller exemption benefits motivate a model in which investors with different exemption benefits choose municipal and taxable bonds jointly. The analytical benchmark has two investors. The quantitative model includes household groups by state of residence, insurers, banks and a residual group (Section 5.1).

[5] The quantitative model charges the bank financing-tax cost separately. Table 3 therefore reports the bank's exemption benefit $1 - \xi_i$ before that cost.

[6] I use IRS Statistics of Income shares of tax-exempt interest on taxable returns, including fund distributions, to approximate household holdings by income group in 2010–2017. Adjusted gross income of at least $200,000 proxies for the high-tax group.

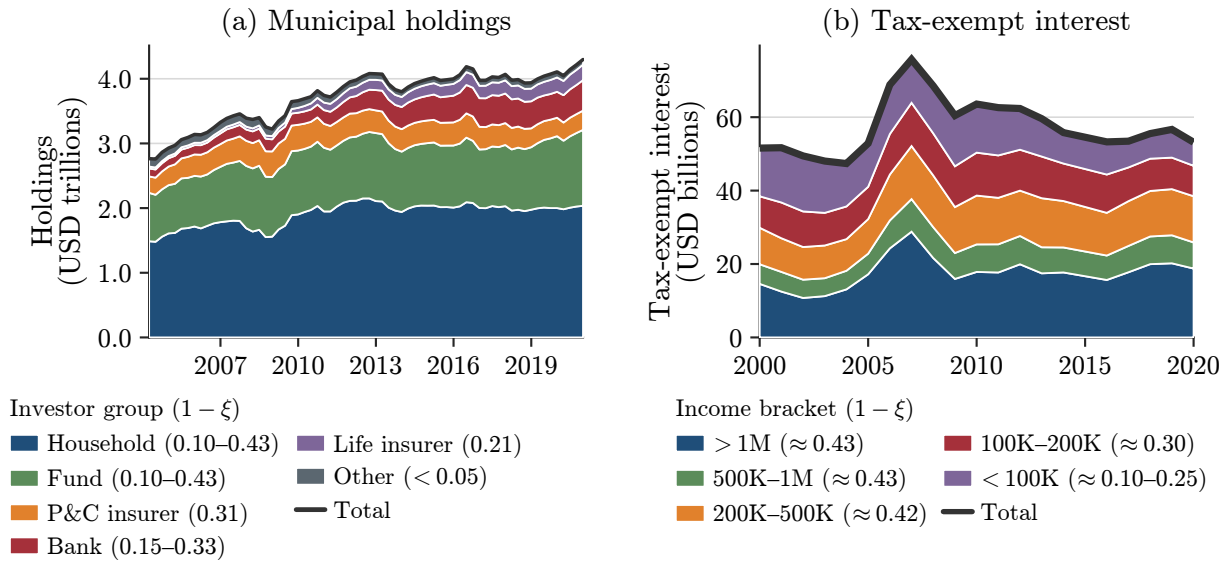



Figure 1: Municipal holdings and tax-exempt interest by investor group. Panel (a): direct holdings of taxable and tax-exempt municipal securities, 2005–2021 (USD trillions, Financial Accounts, with households including nonprofits). Panel (b): tax-exempt interest, an income flow, on taxable individual returns by adjusted-gross-income bracket, 2000–2020 (USD billions, IRS Statistics of Income). 2015 lacks comparable bracket detail and is omitted. Legend values in parentheses are net exemption benefits: in panel (a), the values before the 2017 Act from Table 2, with funds assigned their household shareholders' range, and in panel (b), approximate values by bracket.

### 3.5 Limits to arbitrage

Tax rules and trading frictions limit the trades that could exploit yield differences or equalize investors' costs of bearing rate risk. Interest and short-sale payments incurred to carry exempt bonds generally lose their deduction. Consider a household that shorts a long-maturity corporate bond to hedge its municipal holdings. With the deduction denied, an equal-principal position earns coupon carry equal to the pretax municipal-minus-corporate yield difference. This carry is negative along the sample-average par curves and about $-95$ basis points a year at thirty years. Municipal shorts are scarce and costly because substitute payments to the lender do not retain the interest exemption. Off-exchange swaps face eligibility restrictions, while futures losses can face deduction limits and deferral. The model summarizes these tax asymmetries and trading frictions through nonnegative bond positions and cash balances, ruling out short sales and borrowing. Appendix D.5 examines the tax treatment, carry and residual risk of these trades.

### 3.6 The municipal puzzle

The municipal puzzle is a maturity pattern: short municipal yields are close to the household benchmark, while long municipal yields substantially exceed it. To measure this pattern, I use the seasoned, option-free A-or-above corporate zero-coupon curve as the taxable comparison. High-grade corporate bonds are closer to municipals in credit and liquidity than treasuries. Treasury yields also contain a premium for safety and liquidity (Krishnamurthy and Vissing-Jorgensen, 2012),

which would enter the measured gap if treasuries supplied the benchmark. Studies using treasury yields (Chalmers, 1998) or swap rates (Longstaff, 2011) as the taxable benchmark also find the municipal puzzle.

I construct municipal zero-coupon yields in the same credit band from offering prices and contractual cash flows. Issuance data provide broad coverage and a common measurement date. Offering yields also avoid the direct tax effects of capital gains and market discounts arising from prior secondary-market trades. I remove callability and the issuance premium to make municipal yields comparable with the seasoned, option-free corporate curve. The issuance premium is the extra yield at issuance relative to comparable secondary-market trades.[7] Let $y_t^C(T)$ denote the taxable yield and $y_{s,t}^M(T)$ the municipal yield after these two adjustments, where $T$ is maturity, $s$ is the issuer's state and $t$ is the issue date. The municipal puzzle is the gap between $y_{s,t}^M(T)$ and the household benchmark $\xi_H y_t^C(T)$. The yield-implied tax expresses the same comparison in tax-rate units:

$$ImpliedTax_{s,t}(T) \equiv 1 - \frac{y_{s,t}^M(T)}{y_t^C(T)}. \tag{2}$$

For a given household relative after-tax share $\xi_H$, the two measures satisfy

$$\underbrace{y_{s,t}^M(T) - \xi_H y_t^C(T)}_{\text{municipal yield gap}} = \underbrace{\left[(1-\xi_H) - ImpliedTax_{s,t}(T)\right]}_{\text{shortfall from the household wedge}} y_t^C(T). \tag{3}$$

For positive taxable yields, a positive gap means an implied tax below the household wedge $1-\xi_H$. In Figure 2, the municipal yield is near the household benchmark at the short end but about 118 basis points above it at thirty years. The implied tax falls from about 38% to 21% while the household wedge remains near 41% (Figure 7(d)).

The declining implied-tax pattern also appears against treasuries (Chalmers, 1998), within issuer-matched tax-exempt and taxable municipals (Section 3.8), and in the AA-and-above and unlimited-tax general-obligation samples (Appendices C.11.1 and C.3).

### 3.7 Comparability and bond-side adjustments

Studies such as Fleckenstein and Longstaff (2025) use secondary-market prices of option-free bonds to examine the municipal puzzle. Consistent with that comparison, I remove callability and the issuance premium before defining the municipal puzzle as the gap over the household benchmark. I then assess default risk, liquidity and capital-gains taxation as explanations of this gap. Each adjustment measures the difference in compensation between municipal and corporate bonds, expressed in the household's after-tax yield units. Compensation common to both markets therefore cancels. The household wedge, one minus the ratio of averaged after-tax to pretax corporate yields, averages 41.3% and is nearly flat across maturities. Appendix C explains the adjustments, and

[7] The market-discount tax applies to later secondary purchasers. Its expected effect on resale values can also affect offering prices. The capital-gains adjustment includes this effect as well as taxes on realized gains.

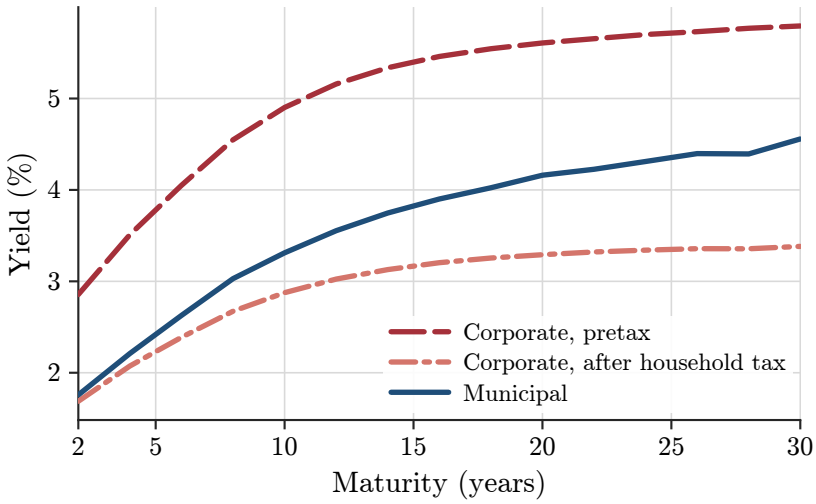



Figure 2: The municipal puzzle: A-or-above zero-coupon yields, 2000–2020. Municipal yields remove callability and the issuance premium. Averaging follows Appendix C.1.1. Appendix C.11.1 gives the AA-and-above counterpart.

Appendix Table 30 reports their contributions by maturity. Appendix C.4 discusses the tax timing of taxable bonds in Green (1993).

At thirty years, the call and issuance-premium adjustments reduce the 160-basis-point spread over the household benchmark to the 118-basis-point municipal puzzle. The default, liquidity and capital-gains differentials explain 40% of the puzzle at ten years but only 12% at thirty. Subtracting them leaves 103 basis points at thirty years for comparison with the model in Section 5.4. At two years, their sum exceeds the small puzzle, leaving a residual of −7 basis points, consistent with the small convenience premium discussed in Fleckenstein and Longstaff (2025). The adjustments leave the central pattern intact: the residual widens with maturity.

### 3.8 Issuer-matched evidence

Taxable and tax-exempt bonds issued by the same municipal borrower provide a further comparison that does not rely on corporate bonds. During 2009–2010, the Build America Bonds (BAB) program allowed municipal issuers to sell taxable bonds with a federal subsidy equal to 35% of their interest payments. The top federal household income-tax rate was also 35%. I leave tax-exempt offering yields unchanged and multiply BAB yields by 0.65 to put both on the household's after-tax scale. I then compare bonds from the same issuer and estimate how their yield difference varies with maturity:

$$y_b \;=\; \alpha_1\, EXMP_b \;+\; \beta\,(EXMP_b \times n_b) \;+\; \gamma_C\, y^{C}_{n_b,t_b} \;+\; \lambda_{j(b)\times n_b\times q_b} \;+\; \varepsilon_b, \tag{4}$$

where $y_b$ is the offering yield for a tax-exempt bond and 0.65 times the offering yield for a BAB. The indicator $EXMP_b$ equals one for a tax-exempt bond and zero for a BAB, and $n_b$ is maturity in years. The fixed effect $\lambda_{j(b)\times n_b\times q_b}$ compares bonds from issuer $j(b)$ at the same maturity and

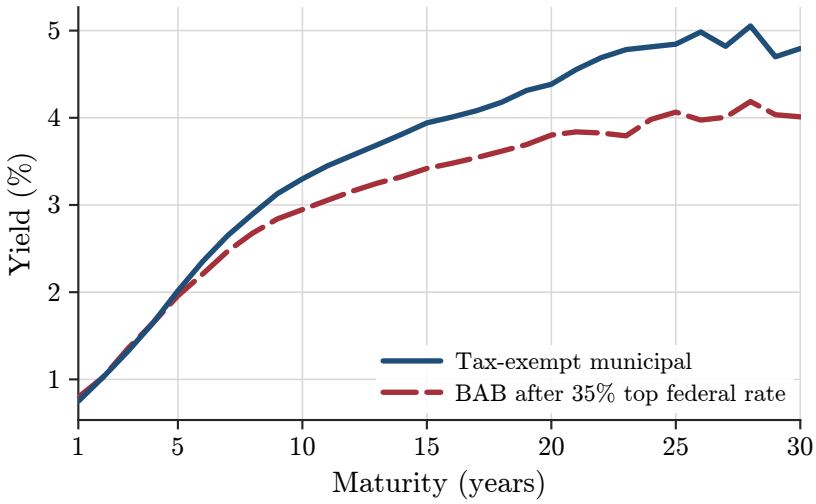



Figure 3: Tax-exempt and Build America Bond offering yields by maturity, AA-and-above issues, April 2009–December 2010. BAB yields are multiplied by 0.65, one minus the 35% top federal rate in 2009–2010. Appendix Table 32 reports regressions with and without call-status matching. Appendix C.2 describes the sample.

in the same issue quarter $q_b$. The call-matched specification also requires the same call status. The corporate yield $y^C_{n_b,t_b}$, measured at the bond's maturity and issue month $t_b$, controls for market yield movements within that quarter. The coefficient $\beta$ therefore measures how the tax-exempt yield minus the household after-tax BAB yield changes with maturity. Standard errors are clustered by state. Comparing bonds from the same issuer at the same maturity and with matched call status holds borrower identity and these bond characteristics fixed.

The estimated gradient $\beta$ is 3.4 basis points per year in the AA-and-above sample and 3.3 in the A-or-above sample. Matching call status gives gradients of 3.1 and 3.0 basis points. The binned estimates also rise monotonically: in the AA-and-above sample with call matching, the $[25, 30]$-year coefficient is 63.9 basis points above the one-to-five-year bin. Figure 3 shows mean yields by maturity before issuer matching. Appendix Table 32 reports the issuer-matched estimates. Both show a widening gap, and the regressions establish that it persists within the same issuer, using offering yields for both bond types and matching call status.

# 4 Model

I develop a term-structure model in which investors differ in tax treatment and balance-sheet exposure to interest rates. They choose bonds jointly across markets and maturities, subject to short-sale constraints. The model answers two questions: why investors with smaller exemption benefits hold tax-exempt bonds, and how their participation changes yields relative to the household benchmark.

I first present the portfolio-pricing framework. A two-investor, two-market affine benchmark then isolates the mechanism. The two propositions below summarize the return-gap decomposition

and rate-risk transmission, with formal conditions and proofs in Appendix A. Section 5 extends the framework for calibration.

### 4.1 Bonds, investors and portfolio choice

Time is discrete. Each investor solves an after-tax mean–variance portfolio choice problem across all assets and maturities, subject to short-sale constraints.

*Interest rates.* A single rate state $z_t$ drives the short rate $r$, the pretax return on cash, which is subject to a zero lower bound:

$$z_{t+1} = \rho z_t + \sigma_r \varepsilon_{t+1}, \qquad r(z_t) = \max\{\bar{r} + z_t,\, 0\}, \qquad 0 < \rho < 1. \tag{5}$$

Here $\varepsilon_{t+1}$ is standard normal, $\rho$ measures persistence, $\sigma_r$ the size of rate shocks and $\bar{r}$ the mean rate before the floor. Taxable cash earns this known log gross return $r(z_t)$ over the next period.

*Bonds and returns.* Bonds are zero-coupon and grouped into markets $a$. Taxes are investor-market specific. For a given investor, all maturities in a market receive the same tax treatment. A bond of market $a$ with $T$ periods to maturity has price $P_T^a(z)$, pays one unit at maturity and becomes a $(T-1)$-period bond next period. Cash is the one-period taxable asset. Dollar supplies $Q_T^a$ are fixed.

An investor who buys a $T$-period bond today and sells it next period earns

$$\text{gross holding return} = 1 + R_{T,t+1}^a = \frac{P_{T-1}^a(z_{t+1})}{P_T^a(z_t)}, \qquad P_0^a = 1.$$

The price ratio is the resale value per dollar invested. It carries rate risk because next period's price depends on next period's rate state.

With one period to maturity, the payment is known, so the bond is fully riskless. A zero-coupon bond's yield is $y_T^a = -\log P_T^a / T$. The holding return above covers one period, from purchase to resale next period. Investors optimize over this one-period horizon.

Consistent with the discrete-time return approximations used in the literature, I approximate the net holding return for portfolio choice as[8]

$$R_{T,t+1}^a \approx \mathcal{R}_{T,t+1}^a = \mu_T^a(z_t) + u_{T,t+1}^a. \tag{6}$$

Here $\mu_T^a(z_t) = \log \mathbb{E}_t[1 + R_{T,t+1}^a]$ is the mean of the approximated return used in the portfolio problem. The shock $u_{T,t+1}^a$ is the unexpected component of the log holding return, with conditional mean zero. Thus $\mu$ determines expected portfolio income, while $u$ determines risk. The approximation becomes exact in the continuous-time limit.

[8] I follow the discrete-time return approximation in Greenwood et al. (2024), Section 2.1. See also Greenwood et al. (2023), Section III.A, and Hanson and Stein (2015), footnote 18. Appendix A.1.1 gives the precise return convention and wealth equation.

*Investors.* Investor $i$ has funding $W_i$ that can be invested in any bonds or cash, risk aversion $\gamma_i$ and outside rate exposure $\ell_i$ from assets and liabilities beyond its bond portfolio. A positive $\ell_i$ raises net wealth when rates rise, as does a long-term liability whose value falls. A negative $\ell_i$ adds to bond losses. Investor $i$ keeps the after-tax share $\xi_i^a = 1 - \tau_i^a$ of returns in market $a$, where $\tau_i^a$ is its tax rate on bonds in market $a$. Cash is taxed like taxable bonds.

The investor chooses dollar positions $x_{i,T}^a \geq 0$ across all markets $a$ and all maturities $T$, keeps the rest of its funding in cash, $x_i^0 = W_i - \sum_{a,T} x_{i,T}^a \geq 0$, and maximizes expected next-period wealth less a penalty for its variance. With pretax returns measured by $\mathcal{R}$, next-period after-tax wealth is $\widehat{W}_i'$, and the problem is

$$\begin{aligned} \max_{x_i \geq 0,\, x_i^0 \geq 0} \quad & \mathbb{E}_t[\widehat{W}_i'] - \frac{\gamma_i}{2} \operatorname{Var}_t(\widehat{W}_i'), \\ \text{subject to} \quad & \widehat{W}_i' = W_i + \xi_i^C r(z_t)\, x_i^0 + \sum_{a,T} \xi_i^a x_{i,T}^a \mathcal{R}_{T,t+1}^a + \ell_i \sigma_r \varepsilon_{t+1}. \end{aligned} \tag{7}$$

Next-period wealth adds after-tax cash income, after-tax bond returns and the change in outside wealth to current funding. Investors choose portfolios each period, taking prices as given and conditioning on today's rate state. Nonnegative positions rule out short sales and borrowing (Section 3.5).

*Required returns.* The portfolio problem gives the Karush–Kuhn–Tucker conditions

$$\underbrace{\xi_i^a \mu_T^a - \xi_i^C r}_{\text{after-tax return over cash}} \quad \leq \quad \underbrace{\gamma_i \xi_i^a \operatorname{Cov}_z\big(u_{T,t+1}^a, \widehat{W}_i'\big)}_{\text{cost of the risk the bond adds}} + \underbrace{\kappa_i}_{\text{shorting constraint on cash}}, \tag{8}$$

with equality whenever $x_{i,T}^a > 0$. The covariance measures the risk the bond adds to total wealth, conditional on today's state. The multiplier $\kappa_i \geq 0$ on the shorting constraint on cash is zero when the investor holds cash and can invest another dollar without borrowing. A held bond must compensate for both its added risk and, when this constraint binds, the dollar it displaces from other bonds. If it offers less than an investor's required return in equilibrium, that investor holds none. Differences in taxes and balance sheets therefore allow the same bond to suit one investor but not another in equilibrium.

*Equilibrium.* Prices clear every market at every rate state:

$$\textstyle\sum_i m_i x_{i,T}^a(z) = Q_T^a, \tag{9}$$

where $m_i$ is the mass of type $i$, converting its positions and funding into aggregate holdings $m_i x_{i,T}^a$ and funding $m_i W_i$. In the calibration, one unit of mass represents one dollar of baseline portfolio assets (debt securities and cash). Thus $m_i$ measures a type's portfolio size, and $W_i$ is the share available for the bonds and cash included in the model. The rate process, supplies, taxes, funding, risk aversion and outside exposures are exogenous inputs. Equilibrium prices imply holding returns for which investors satisfy (8) and all markets clear. Who holds each bond is determined by these conditions. Appendix Table 6 collects the notation.

### 4.2 Taxes and the cost of bearing rate risk

To isolate the mechanism, I use a special case with two investors, two markets and rate risk alone. Consider a high-tax household $H$ and a lower-tax investor $I$ choosing municipal bonds $M$, taxable bonds $C$ and cash. Municipal maturities begin at $T = 1$, while risky taxable bonds begin at $T = 2$. The two determinants of required returns are the investor's tax treatment and the cost of adding risk to its portfolio.

The relative after-tax share of Section 3.3 summarizes the tax comparison:

$$\xi_i = \frac{\xi_i^C}{\xi_i^M} = \frac{1 - \tau_i^C}{1 - \tau_i^M}, \qquad \xi_H < \xi_I.$$

The household's smaller $\xi_H$ means a larger exemption benefit, $1 - \xi_H$. I measure each investor's payoffs in units of its after-tax municipal income. Municipal returns are then kept in full, and taxable returns and cash income are multiplied by $\xi_i$. This change of units preserves portfolio choices.[9]

In the full model, prices are nonlinear in the rate state for three reasons: the short rate cannot fall below zero, the shorting constraint on cash can bind, and which investors hold which bonds can change as rates move. A bond's price today depends on its resale price next period, so these future nonlinearities also affect today's prices. For tractability, the illustration fixes participation: prices are computed as if the same investors held each bond at every future rate state. I call such a fixed holding pattern a participation branch. The illustration also removes the rate floor, setting $r = \bar{r} + z$, and the shorting constraint on cash, so $\kappa_i = 0$. Each log price is then a constant plus a constant multiple of the rate state, $\log P_T^a(z) = -A_T^a z - C_T^a$, where $A_T^a$ is the bond's rate loading defined below. The quantitative model of Section 5 imposes none of these simplifications: participation changes with rates, and the rate floor and the shorting constraint on cash are kept. Fixing participation abstracts from the possibility that rate movements over a bond's remaining life change who holds it, and hence how strongly its price responds to rates (Appendix equation (44)).

Define a bond's rate loading as $A_T^a \equiv -\partial \log P_T^a / \partial z$, which measures its proportional price decline when the rate state rises. It is determined by equilibrium prices, as Section 4.5 shows. Rate shocks affect the holding return through next period's resale price, when the bond has $T-1$ periods remaining. Define the investor's net rate exposure as

$$\widehat{D}_i \equiv -\frac{\partial \widehat{W}_i'}{\partial z_{t+1}} = \sum_T x_{i,T}^M A_{T-1}^M + \xi_i \sum_T x_{i,T}^C A_{T-1}^C - \ell_i. \tag{10}$$

The derivative measures how much next-period wealth falls when the rate state rises, holding today's portfolio fixed. Liabilities whose value also falls offset losses on bonds, so a larger liability hedge $\ell_i$ lowers this net rate exposure.

[9] With payoffs divided by $\xi_i^M$, risk aversion $\gamma_i \xi_i^M$ and outside exposure $\ell_i / \xi_i^M$, the objective is $1/\xi_i^M$ times the original, up to a portfolio-independent constant. Below, $\gamma_i$ and $\ell_i$ denote these normalized values.

The investor's cost per unit of additional rate exposure is $\gamma_i \sigma_r^2 \widehat{D}_i$, which falls with the liability hedge at given bond holdings. The portfolio first-order conditions are

$$\begin{aligned} \mu_T^M - \xi_i r &\leq \gamma_i \sigma_r^2 \widehat{D}_i \, A_{T-1}^M, \\ \mu_T^C - r &\leq \gamma_i \sigma_r^2 \widehat{D}_i \, A_{T-1}^C, \end{aligned} \tag{11}$$

with equality for a positive position. The municipal condition compares $\mu_T^M$ with after-tax cash, $\xi_i r$. The taxable condition, divided by $\xi_i$, is on a pretax scale and compares $\mu_T^C$ with $r$. In both lines the right side is the return the investor requires for the rate risk added by the bond. The quantitative model uses common risk aversion in these normalized payoff units, so its investors differ in tax treatment $\xi_i$ and net rate exposure $\widehat{D}_i$, which depends on portfolio choice and outside exposures. The analytical benchmark allows $\gamma_i$ to differ.

These conditions separate the two forces behind participation. The household's larger exemption benefit lowers its required municipal return. A lower unit cost of rate risk also lowers the required return, increasingly so for a bond with a larger loading.

### 4.3 Participation at longer maturities

At the riskless maturity, only the cash alternative matters. With $A_0^M = 0$, a municipal holder requires $\mu_1^M = \xi_i r$. The household accepts the lowest return because it keeps the smallest share of taxable interest. With a positive short rate and a slack shorting constraint on cash, it holds the municipal supply and

$$y_1^M = \mu_1^M = \xi_H r, \qquad y_1^C = r. \tag{12}$$

At the riskless short end the municipal yield equals the household benchmark exactly.

At a risky maturity, compare the returns the two investors require on the same municipal bond. The lower-tax investor requires less when its duration benefit exceeds the difference in exemption benefits:

$$\underbrace{\sigma_r^2 A_{T-1}^M (\gamma_H \widehat{D}_H - \gamma_I \widehat{D}_I)}_{\text{duration benefit}} > \underbrace{(\xi_I - \xi_H) r}_{\text{difference in exemption benefits}}. \tag{13}$$

All quantities are evaluated at equilibrium. A positive $\gamma_H \widehat{D}_H - \gamma_I \widehat{D}_I$ means that the lower-tax investor requires less return per unit of rate exposure. With the same $\gamma_i$, this happens when its outside liability exposure $\ell_I$ is large enough to give it a lower equilibrium $\widehat{D}_I$. Where the duration benefit is smaller than the exemption difference, the household requires less return and holds the bond. Where it is larger, the lower-tax investor requires less and holds the bond. Both can hold a maturity at which the two terms are equal. Thus participation is determined by the comparison at the equilibrium allocation, including the risk each investor bears.

In the illustration, the municipal rate loading $A_{T-1}^M$ is endogenous and increases with maturity (Section 4.5). At a given equilibrium, the difference in unit risk costs is common across maturities. The duration benefit therefore grows with maturity, while $(\xi_I - \xi_H) r$ stays fixed. Lower-tax par-

ticipation can begin at the long end. For taxable bonds, bond returns and cash income receive the same tax treatment, so there is no difference in exemption benefits to overcome. The investor with lower $\gamma_i \widehat{D}_i$ requires less compensation for rate risk at every risky maturity.

Figure 4 shows sharp sorting when rate risk is the only risk. Adding a small bond-specific shock creates a separate risk cost that grows with each investor's holding of that particular bond, allowing both investors to hold the same maturity. In the otherwise identical illustration of Appendix A.2.8, joint holdings extend across several long maturities while municipal yields change by less than 0.4 basis points. The quantitative model includes this component and multiple investor types, generating overlapping holdings and partial segmentation.

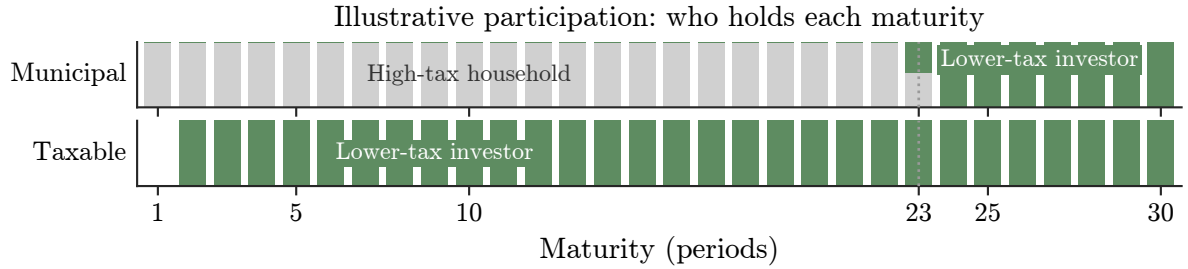



Figure 4: Participation in the two-investor illustration without bond-specific risk. Bars show supply shares at the mean rate state. Both investors hold municipal maturity $T^* = 23$ (dotted line). The high-tax household holds shorter municipal maturities, and the lower-tax investor holds longer municipal maturities and risky taxable bonds. The one-period taxable claim is cash. Appendix A.2.7 gives the parameters and checks.

### 4.4 Three channels in required returns

Given this allocation, I compare municipal and taxable bonds of the same maturity, the relevant comparison for the municipal puzzle. The municipal return above the household's after-tax taxable return is

$$g_T = \mu_T^M - \xi_H \mu_T^C. \tag{14}$$

This compares municipal holding returns with the household-scaled taxable holding return. The portfolio conditions separate the gap into three channels. Section 4.5 then translates holding returns into yields, including the convexity adjustment.

**Proposition 1 (Three channels of the one-period return gap).** In the two-investor affine allocation just described, the lower-tax investor holds risky taxable bonds. For municipal bonds, use the household's condition ($i = H$) below $T^*$ and the lower-tax investor's condition ($i = I$) from $T^*$ onward. Subtracting $\xi_H$ times the taxable holder's condition from the municipal holder's condition in (11) gives

$$\begin{aligned} g_T = \; & \underbrace{(\xi_i - \xi_H) r}_{\text{exemption benefit}} + \underbrace{\sigma_r^2 A_{T-1}^M (\gamma_i \widehat{D}_i - \gamma_I \widehat{D}_I)}_{\text{duration cost}} \\ & + \underbrace{\gamma_I \sigma_r^2 \widehat{D}_I (A_{T-1}^M - \xi_H A_{T-1}^C)}_{\text{rate loading}}. \end{aligned} \tag{15}$$

The three channels compare exemption benefits, costs per unit of rate exposure, and the exposures themselves. The exemption-benefit term is zero below the threshold, the duration-cost term is zero from the threshold onward, and the loading term applies to both ranges. A zero duration-cost term does not mean that the lower-tax investor bears no rate risk. It means that the same investor's unit risk cost applies to both bonds in this comparison. The appendix states and proves the decomposition.

*Exemption benefit.* At long maturities, the lower-tax investor holds both municipal and taxable bonds. Because it pays less tax on taxable interest, tax exemption saves it less than it saves the high-tax household. It therefore requires a higher municipal return than the household benchmark predicts. This difference in exemption benefits, $(\xi_I - \xi_H)r$, contributes to this gap.

*Duration cost.* At shorter risky maturities, the household alone holds municipal bonds, while the lower-tax investor holds taxable bonds. The household needs more return to bear the same rate exposure because rate risk is more costly for it. This difference contributes $\sigma_r^2 A_{T-1}^M(\gamma_H \widehat{D}_H - \gamma_I \widehat{D}_I)$ to municipal returns above the household benchmark, even where the lower-tax investor holds no municipal bonds. It is the duration benefit in (13), expressed as the household's higher risk cost at these maturities.

*Rate loading.* When taxes scale returns, they also scale risk. The household benchmark scales taxable yields and rate exposure by $\xi_H$, predicting $A_{T-1}^M = \xi_H A_{T-1}^C$. In the illustration, lower-tax participation raises municipal rate loadings to $A_{T-1}^M \geq \xi_H A_{T-1}^C$. The extra exposure contributes positively to the gap when the lower-tax investor's net rate exposure is positive, as in the illustration. The contribution equals the extra exposure times $\gamma_I \sigma_r^2 \widehat{D}_I$, the lower-tax investor's unit risk cost. Section 4.5 derives the extra rate loading $A_{T-1}^M - \xi_H A_{T-1}^C$.

All three channels vanish at the riskless maturity, and $g_1 = 0$. Since $A_0^M = A_0^C = 0$, the duration-cost and rate-loading terms vanish. The exemption-benefit term is also zero because the household holds the riskless municipal bond. Figure 5 shows how the channels combine across maturities. At a jointly held maturity, either holder's condition gives the same total gap. The choice changes how the gap is divided between the exemption-benefit and duration-cost terms.

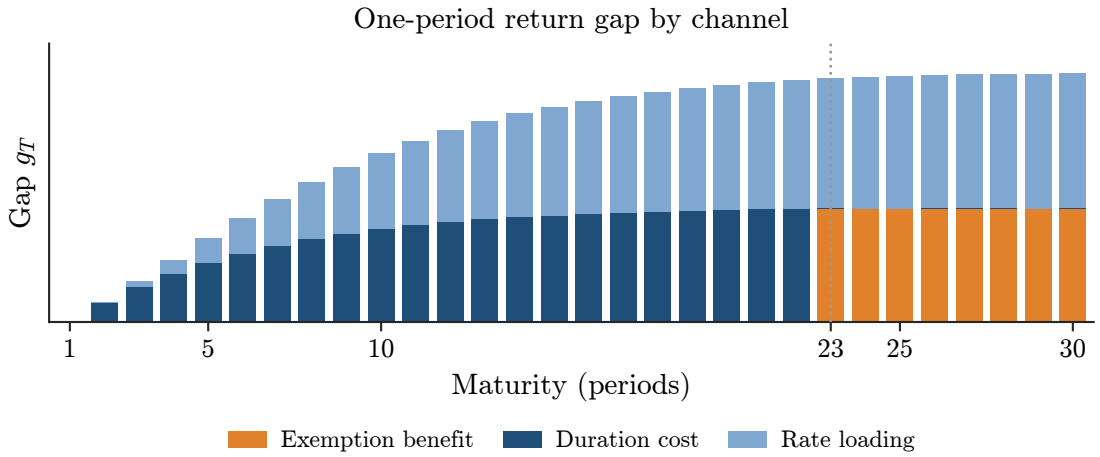



Figure 5: One-period return gap by channel. The allocation is that of Figure 4. Bars decompose $g_T = \mu_T^M - \xi_H \mu_T^C$ into the exemption-benefit, duration-cost and total rate-loading terms of (15). The vertical scale is omitted because levels depend on the illustrative parameters.

### 4.5 Rate loadings and the yield gap

A bond's rate loading depends on how its expected holding return responds to rates. Let $s_T^a = \partial \mu_T^a / \partial z$ denote this response. The cash terms in (11) contribute the holder's after-tax share $\xi_i$ for municipal bonds and 1 for taxable bonds. Changes in the holder's portfolio exposure also change the risk compensation it requires. Both the holder's tax treatment and its portfolio response therefore enter the bond's return response.

As rates rise, the difference in exemption benefits, $(\xi_I - \xi_H)r$, grows, making municipal bonds relatively less attractive to the lower-tax investor. Trading at the shared maturity transfers rate exposure to the household. The household then requires more compensation on all its risky municipal bonds, including those it holds alone. The lower-tax investor sheds exposure and requires less risk compensation on its municipal and taxable bonds. Portfolio substitution thus transmits the effect beyond the bond that is traded.

**Proposition 2 (Participation and rate responses).** On a feasible affine participation branch of the two-investor benchmark, suppose the lower-tax investor alone holds risky taxable bonds. Let $\bar{\xi}_T$ denote the municipal holders' tax mix. It equals $\xi_H$ when the household holds the bond alone, $\xi_I$ when the lower-tax investor holds it alone, and $\xi^* = \omega_H \xi_H + \omega_I \xi_I$ under joint holdings. Here $\omega_i$ is each investor's share of total risk tolerance.[10] Differentiating (11) gives

$$s_T^M = \begin{cases} \xi_H + \gamma_H \sigma_r^2 A_{T-1}^M \frac{\partial \widehat{D}_H}{\partial z}, & \text{household alone,} \\ \xi^*, & \text{both investors,} \\ \xi_I + \gamma_I \sigma_r^2 A_{T-1}^M \frac{\partial \widehat{D}_I}{\partial z}, & \text{lower-tax investor alone,} \end{cases} \tag{16}$$

$$s_T^C = 1 + \gamma_I \sigma_r^2 A_{T-1}^C \frac{\partial \widehat{D}_I}{\partial z}.$$

[10] The weights are $\omega_i = (m_i/\gamma_i)/(m_H/\gamma_H + m_I/\gamma_I)$. They sum to one.

Here $\partial \widehat{D}_i/\partial z$ measures how the investor's net rate exposure changes as rates rise. On the illustrated allocation, $\partial \widehat{D}_H/\partial z > 0 > \partial \widehat{D}_I/\partial z$. The municipal portfolio response is therefore positive at household-only maturities, zero at the shared maturity and negative where the lower-tax investor holds municipals alone. A negative portfolio response means that risk compensation falls as the investor sheds exposure, not that the compensation itself is negative. At the longest illustrated maturity, the total municipal response is 0.664, equal to the lower-tax investor's after-tax share, $\xi_I = 0.780$, less a portfolio offset of 0.116. It remains above the household benchmark, $\xi_H = 0.550$.

Combining the tax mix and municipal portfolio response separates the net effects in the two markets:

$$s_T^M - \xi_H s_T^C = \underbrace{s_T^M - \xi_H}_{\text{municipal side}} + \underbrace{\xi_H(1 - s_T^C)}_{\text{taxable side}}.$$

The municipal term measures how much more the municipal return responds than the household's after-tax cash alternative. The taxable term reflects the lower-tax investor's falling risk compensation in the taxable market. Both are positive at risky maturities in the illustration. Appendix A.2.7 proves the response formulas and increasing relative sensitivity for this illustration's rate-risk-only branch. Figure 6 shows the net effect in each market, combining the municipal tax and portfolio terms.

Because $\mu_T^a(z)$ is the log of the expected gross holding return, it links expected holding returns to prices. Rearranging its definition gives the log-price recursion

$$\log P_T^a(z) = -\mu_T^a(z) + \log \mathbb{E}_z\left[P_{T-1}^a(z')\right].$$

Today's log price equals the log expected resale value less the required holding return. With affine log prices and $\mathbb{E}_z[z'] = \rho z$, the log expected resale value has loading $\rho A_{T-1}^a$. Taking the negative state derivative gives

$$A_T^a = \rho A_{T-1}^a + s_T^a, \qquad A_0^a = 0. \tag{17}$$

The first term carries forward the shorter bond's loading, weighted by rate persistence, and the second adds the current return response. Starting from the riskless holding return, this relation carries the response differences into $A_T^M > \xi_H A_T^C$ for $T \geq 2$. Municipal bonds consequently carry more rate exposure than the household benchmark implies. Figure 6(b) accumulates the municipal and taxable contributions into this loading gap.

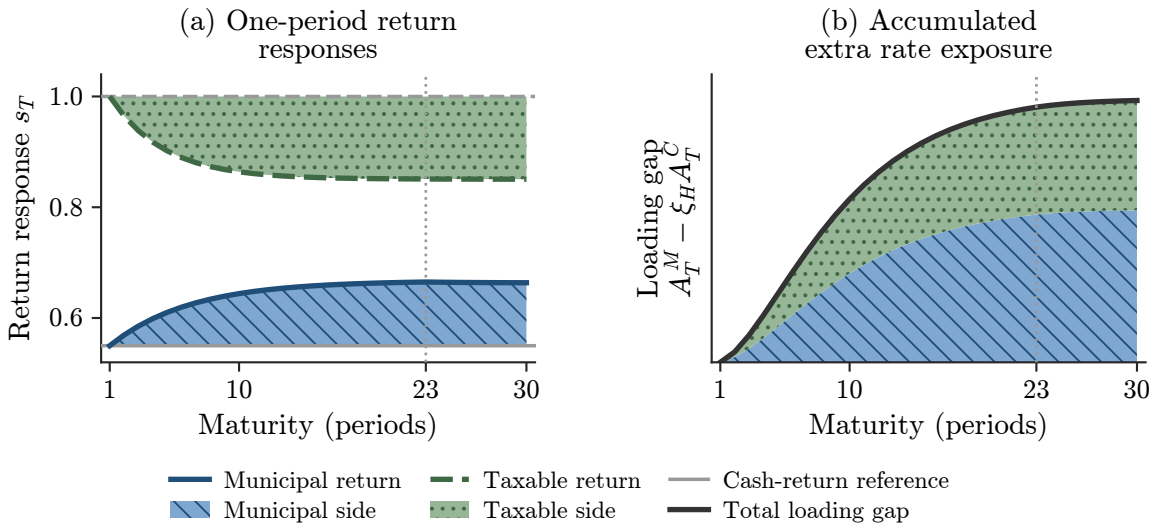



Figure 6: Rate responses in the illustration of Figure 4. Panel (a) shows municipal and taxable holding-return responses. Gray lines show the household's after-tax cash response, $\xi_H$, and the pretax cash response, 1. Blue shading combines the municipal tax mix and portfolio response into $s_T^M - \xi_H$. Green shading shows $1 - s_T^C$. Panel (b) accumulates $s_T^M - \xi_H$ and $\xi_H(1 - s_T^C)$ into the loading gap. The dotted line marks the shared municipal maturity $T^* = 23$.

*From return gaps to yield gaps.* A long bond earns a sequence of holding returns before it matures. Its yield averages the expected returns along this path, less convexity (Appendix equation (56)). Applying this relation to both markets gives

$$y_T^M - \xi_H y_T^C = \frac{1}{T} \sum_{j=0}^{T-1} \mathbb{E}_t[g_{T-j}(z_{t+j})] - \text{convexity difference}. \tag{18}$$

The yield gap accumulates the three return-gap channels over the bond's remaining life. In the illustration, the equilibrium municipal yield exceeds the household benchmark at every maturity beyond one period.

### 4.6 Empirical predictions

The representative-tax benchmark predicts full pass-through of household tax changes at every maturity and no response to tax benefits specific to lower-tax institutions. In the model, household taxes determine the riskless short municipal yield, while longer yields also reflect lower-tax participants' tax treatment and portfolio responses. I first use tax variation specific to households to test whether household tax pass-through declines with maturity (P1). I then use tax variation specific to banks and insurers to test whether lower-tax institutions' benefits affect longer yields while leaving the riskless short yield unchanged (P2). Together, these tests distinguish the participation mechanism from a single household tax rate applied across the yield curve. P3 tests the

associated prediction that municipal rate sensitivity rises relative to the taxable benchmark with maturity.

The outcome of interest is the yield-implied tax $\tau_T^{Y,\text{impl}} = 1 - y_T^M/y_T^C$ from (2). Investor $i$'s pass-through measures how much this equilibrium implied tax changes when investor $i$'s exemption benefit changes:

$$\mathcal{P}_i^Y(T) = \frac{d\tau_T^{Y,\text{impl}}}{d(1-\xi_i)}, \qquad i \in \{H, I\}. \tag{19}$$

Here, pass-through of one means full transmission and zero means no effect. The representative-tax benchmark predicts pass-through of one for household taxes by definition.

In the affine benchmark, the household alone holds the riskless short municipal bond, whose yield equals its after-tax return on taxable cash. Since $y_1^C = r$, this gives $y_1^M = \xi_H r$, so the implied tax at the riskless shortest maturity is $1 - \xi_H$. Therefore

$$y_1^M = \xi_H y_1^C \quad \Longrightarrow \quad \mathcal{P}_H^Y(1) = 1, \qquad \mathcal{P}_I^Y(1) = 0. \tag{20}$$

Only the household's benefit enters this pricing relation.

At longer maturities, lower-tax participation changes whose exemption benefit affects returns. A larger benefit makes municipal bonds more attractive to its recipient and changes portfolios in both markets. The household's benefit therefore no longer has the exclusive role it has at the riskless short end. An institution's tax treatment affects returns where it holds bonds, while portfolio substitution transmits the effect to other risky maturities. A long bond's yield also reflects the returns earned as its remaining maturity shortens. Its yield can therefore respond to household taxes even when an institution holds it today.

I compute tax pass-through by changing one investor's after-tax share and re-solving loadings and portfolios, keeping non-tax primitives fixed. In the two-investor illustration, household pass-through declines from 1 at the shortest maturity to 0.230 at maturity 30, while institutional pass-through rises from 0 to 0.736. These are full yield-ratio derivatives, including the response of the taxable curve. The long-end continuation of the illustrated price recursions has the same ordering:

$$\mathcal{P}_H^Y(\infty) < 1, \qquad \mathcal{P}_I^Y(\infty) > 0. \tag{21}$$

Appendix A.2.7 derives the loading result and reports the re-solved tax derivatives. Appendix Proposition A.9 gives sufficient conditions for the yield-pass-through inequalities in the extension with joint holdings across long maturities. The empirical analysis tests the following three predictions, P1 to P3.

**Household taxes (P1).** A home-state exemption gives resident households an additional benefit on their state's bonds. That benefit is fully reflected in implied state taxes at the riskless short end and less strongly at longer maturities in the solved model. The representative-tax benchmark predicts full pass-through at every maturity.

**Institutional taxes (P2).** A benefit specific to a lower-tax institution has no effect at the riskless short end and matters more at longer maturities. It affects returns where the institution holds bonds and spreads to other risky maturities through portfolio substitution. Bank qualification and the 2017 corporate tax change for insurers test this prediction. For banks, the measure is the share of the bank-only benefit reflected in yields (Section 6.2). The representative-tax benchmark gives these benefits no role.

**Interest-rate responses (P3).** Section 4.5 shows why municipal bonds carry more rate exposure than the household benchmark predicts. Since $y_T^a = -\log P_T^a/T$, yield sensitivity is $A_T^a/T$. Comparing bonds at the same maturity cancels this division by $T$. For the illustration's affine branch, this gives

$$\frac{\partial y_T^M/\partial z}{\partial y_T^C/\partial z} = \frac{A_T^M}{A_T^C}, \qquad \frac{A_1^M}{A_1^C} = \xi_H, \qquad \lim_{T\to\infty} \frac{A_T^M}{A_T^C} = \xi_I > \xi_H.$$

The lower-tax investor prices both markets at long maturities. Its changing rate exposure affects both loading recursions through the same risk-compensation coefficient in (16). This common adjustment cancels in the limiting ratio, leaving $\xi_I$ (Appendix A.2.7). It still affects each bond's loading and return response. On the illustrated branch, the ratio rises with maturity from $\xi_H$ toward $\xi_I$. P3 tests whether municipal yield sensitivity relative to same-maturity taxable bonds increases with maturity. The representative-tax benchmark predicts the constant ratio $\xi_H$.

### 4.7 Liability hedging and other motives

Liability hedging is one source of the advantage that leads lower-tax investors to hold long municipal bonds. Lower trading costs or better default diversification can also make municipal bonds less costly for a lower-tax investor to hold. An investor may trade less often or at lower cost than households, or diversify default risk across markets more effectively than a household concentrated in its home state's bonds. When these advantages grow with maturity, they can outweigh the maturity-invariant difference in exemption benefits and sustain participation in long municipal bonds.

The pricing mechanism extends to these motives when the relevant advantage is neutral across taxable and tax-exempt bonds. Neutrality means that the advantage itself creates no return gap relative to the household benchmark when the household holds both bonds. Lower-tax participation then brings the participant's smaller exemption benefit into the return comparison. Appendix A.3 states the result, and Appendix A.3.4 gives solved trading-cost and default-diversification examples, with the default comparison net of default compensation. A preference for municipal bonds as such, including warm glow from financing local projects, lowers the return an investor requires on them. Such a preference alone does not explain yields above the household benchmark.

I use liability hedging in the quantitative application because rate risk is common to both markets and balance sheets provide a measurable source of differences in investors' risk-bearing costs. Observed differences in maturity holdings help discipline the outside exposures. The trading-

cost and diversification examples show that liability hedging is not the only motive that can sustain lower-tax holdings and bring a smaller exemption benefit into relative returns.

# 5 Quantitative Results

## 5.1 The quantitative model

The quantitative model asks how much of the municipal puzzle measured taxes, balance sheets and ownership can explain. It extends Section 4.1 to municipal markets across the 50 U.S. states and the District of Columbia, a taxable market and six investor groups. Tax rules, quantities and household balance sheets are measured directly. Seventeen common coefficients governing risk aversion, bond-specific risk, outside exposures and within-group asset shares are estimated from twenty-two moment conditions on ownership, maturities, portfolios and the taxable yield curve. Municipal yields are untargeted.

States with a personal income tax have separate municipal markets. The seven states with no state income tax are combined into one market because investors treat their bonds alike in the model. The model applies a common home-state exemption rule to resident households, while institutions have no modeled state-income-tax wedge. A model period is two years, giving fifteen maturities from two to thirty years in each of the 45 municipal markets. The taxable market pools treasury and corporate supply. Cash and the two-year municipal claim are riskless. Prices in both markets are determined in equilibrium.

For calibration, I add a small independent bond-specific return shock, which makes bonds imperfect substitutes and permits overlapping participation regions in which several investor types hold the same maturities. It vanishes at the riskless claim, and changing its variance by 20% moves the thirty-year gap by less than 0.1 basis points (Appendix B.7). Investors solve (7) with market-specific after-tax shares, and prices satisfy the covariance-based portfolio conditions (8) and clearing (9) in every market and rate state. Banks' measured financing-tax cost is subtracted from municipal returns according to their tax treatment (Appendix B.1). The model retains the rate floor in (5) and allows the shorting constraint on cash to bind and participation to change across rate states. I solve $P_T^a(z)$ numerically through the general resale relation in Appendix equation (39), integrating over next-period states. Returns, portfolio risks, holdings and both yield curves are determined jointly. Prices can be nonlinear in the rate state because future participation, binding shorting constraints on cash and the rate floor affect resale values.

The measured inputs cover six three-year blocks, from 2000–2002 through 2015–2017, centered on Survey of Consumer Finances (SCF) waves. Financial Accounts stocks give each sector's portfolio assets and bond quantities. Mergent gives maturity profiles and municipal state shares. The SCF supplies household balance sheets and the income and filing status needed to apply statutory taxes. Historical short rates determine the rate process. Appendix B.1 details the construction, and Table 9 lists the sources.

Table 3: Investor groups in the quantitative model, 2015–2017

| Group | Types | Types differ by | Gross exemption benefit $1-\xi_i$ | Portfolio assets $m_i$ |
|---|---|---|---|---|
| High-tax households | $7 \times 45 = 315$ | Rate exposure, home state | 0.43–0.50 | 6.97 |
| Lower-tax households | $7 \times 45 = 315$ | Rate exposure, home state | 0.21–0.27 | 8.09 |
| P&C insurers | 2 | Rate exposure | 0.31 | 1.11 |
| Life insurers | 3 | Rate exposure | 0.21 | 3.05 |
| Banks | 1 | None | 0.35 | 4.83 |
| Residual group | 7 | Rate exposure | 0 | 26.47 |
| Total | 643 | | | 50.52 |

*Notes.* Each household group has seven rate-exposure types in each of the 45 markets and in the one-market economy used to estimate the coefficients (Section 5.2). Estimated coefficients map each household type's measured outside assets and debt into exposure. Institutional exposures are estimated directly (Appendix Table 12). The gross exemption benefit $1-\xi_i$, with $\xi_i = (1-\tau_i^C)/(1-\tau_i^M)$, follows from tax law. Households multiply the federal and home-state after-tax shares of each SCF tax unit, with the range running across states, and institutions face the 35% corporate rate, with proration for insurers. Banks also pay a financing-tax cost of 0.11% a year on municipal holdings, which lowers their net benefit below this gross value (Table 2). Portfolio assets $m_i$, in trillions of dollars, are all of a group's debt securities and cash in the Financial Accounts, and they scale the group's risk-bearing capacity.

Table 3 lists the groups. For type $i$, $m_i$ is its baseline portfolio assets (debt securities and cash), and $W_i$ is the share available for the asset categories traded in the model. Its investment budget is therefore $m_i W_i$. Positions $x^a_{i,T}$ are measured per dollar of portfolio assets, so aggregate bond holdings are $m_i x^a_{i,T}$. Households differ by federal tax bracket and outside assets and debt. Outside assets add rate risk and liabilities offset it. Their mapping into exposures is estimated. Insurers differ in exposures and shares of sector assets, banks face municipal financing-tax costs, and the residual group is tax-neutral. Section 3.4 and Appendix D.1 distinguish gross exemption benefits from these financing costs.

## 5.2 Estimation and fit

I measure directly the inputs available in the data: investor tax treatment, each sector's portfolio assets, bond quantities and maturity profiles, household assets and debt, and the short-rate process. These inputs vary across the six blocks of 2000–2017. The remaining seventeen coefficients govern risk aversion, bond-specific risk, outside rate exposures and within-group asset shares. I estimate one common set of coefficients across all six time blocks.

Observed ownership provides the main discipline on outside rate exposure, which balance-sheet amounts alone cannot measure. Given each group's tax treatment and portfolio assets, its municipal and taxable holdings reveal how much rate risk it is willing to bear. For households, three coefficients translate measured outside assets and debt into rate exposure. For institutions, the estimated exposures and within-sector asset shares account for both the amounts held and their maturity composition. The taxable yield curve and long-bond allocations discipline risk aversion and bond-specific risk. The coefficients are estimated jointly because portfolio choices and yields adjust together in equilibrium.

The estimation uses twenty-two targets. Nine are ownership stocks: municipal holdings of the two household groups, P&C insurers, life insurers and banks, and taxable holdings of households and the same three institutional sectors. Four are the mean maturities of insurers' municipal and taxable portfolios. Six describe the distribution of insurer holdings: two measure whether municipal holdings are spread across insurers or concentrated in portfolios with high municipal shares, giving larger portfolios more weight, and four measure the fractions of each sector's municipal and taxable portfolios maturing beyond twenty years. The remaining three are taxable term spreads over the two-year yield at ten, twenty and thirty years. I evaluate these conditions at observed rate states, average them across the six blocks and minimize their weighted squared deviations from the data.[11]

For computational tractability, I estimate the coefficients in an economy with one national municipal market, using municipal supply weights to average each household type's home-state after-tax shares. I then restore the 45 municipal markets and their measured tax differences, holding every coefficient fixed. No estimated coefficient is state-specific.[12]

The one-market and full economies give similar relative municipal yields. Their thirty-year gaps average 102.6 and 102.4 basis points across the six blocks, differing by at most 2.5 basis points in any block. Thirty-year taxable yields differ by at most 0.5 basis points. Changing each coefficient group in turn (Appendix Table 15) moves the one-market gap by at most about 4 basis points, except for the common risk aversion.[13] Section 5.4 uses the 2000–2020 bond sample's weights and expresses the model's relative municipal yields on the observed corporate curve for the headline comparison.

Municipal yields, the municipal puzzle and empirical tax-pass-through coefficients do not enter the estimation criterion. Ownership disciplines outside rate exposure, which combines with observed taxes to determine municipal yields in equilibrium. The relative municipal yield curve and cross-state pass-through therefore test the mechanism without being targeted (Sections 5.3–5.4 and 6.1).[14]

### 5.3 Model and data yield curves

Figure 7 sets the calibrated model curve (top) against the data (bottom). To cover the full data sample, I keep the estimated coefficients fixed and add a seventh economy for 2018–2020 after the 2017 Tax Cuts and Jobs Act (TCJA), built from that period's taxes, rates and quantities. Both rows show a municipal yield near the household benchmark at the short end and a spread over it that widens with maturity. In nominal semiannual yields, the model's implied tax $1 - y_T^M / y_T^C$ is 41.0% at two years, 34.7% at ten and 24.8% at thirty, a decline of 16.2 percentage points. The

[11] Each averaged deviation is scaled by 10% of its observed normalizer. Maturity and long-portfolio conditions include model-portfolio weights. Appendix equation (111) gives the criterion, and Tables 12 and 13 report the coefficients and target definitions.

[12] Fourteen of the twenty-two targets are matched to within 10% of their observed scale. Appendix B.3 reports the estimation procedure and fit. Small moves from the estimated coefficients change the thirty-year gap by less than 0.01 basis points.

[13] Changing risk aversion alone scales both curves and the spread, moving the taxable yield level away from its calibration targets.

[14] Appendix B describes the solution method, and Appendix B.6.3 reports accuracy.

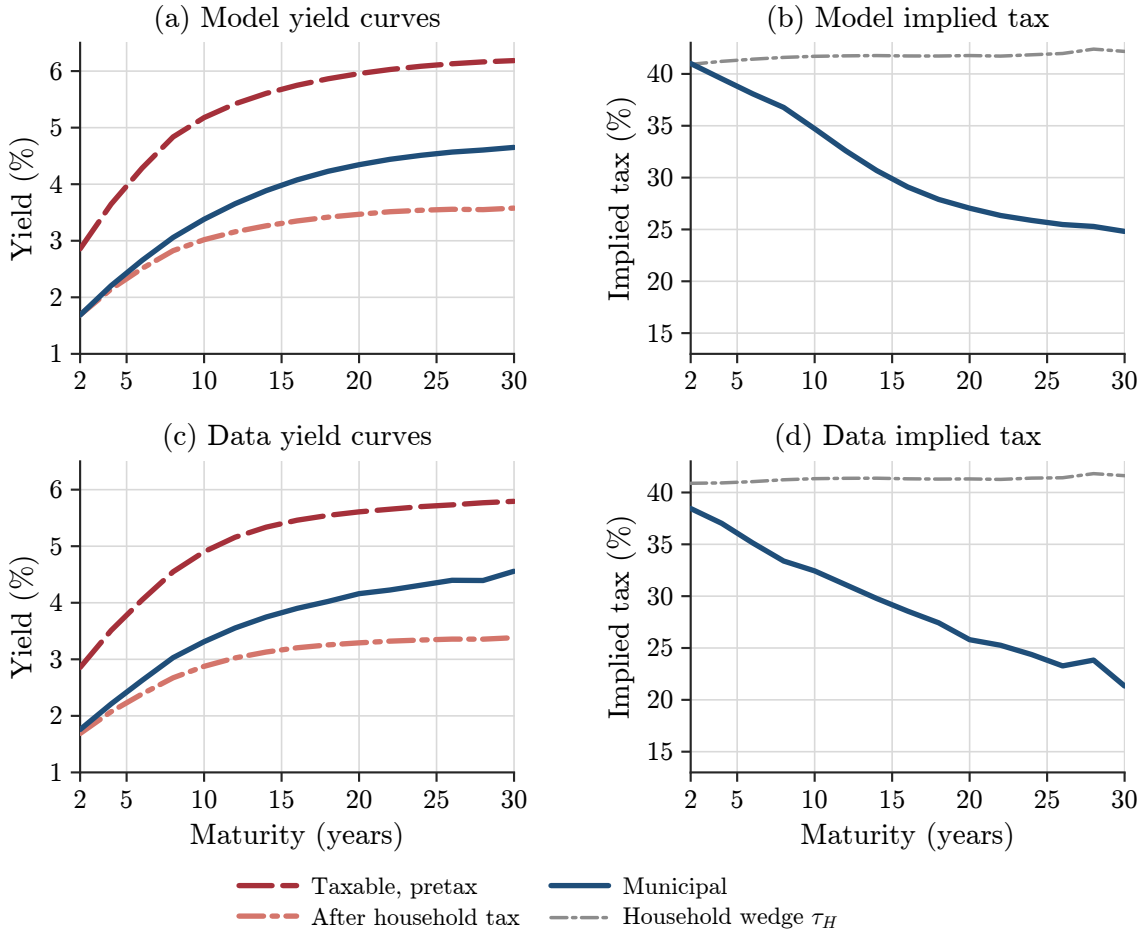



Figure 7: Model (top) and adjusted data (bottom), 2000–2020. Panels (a) and (c) plot the pretax taxable yield, the taxable yield after household tax (the household benchmark) and the municipal yield. Panels (b) and (d) plot the implied tax $1 - y_T^M / y_T^C$ and the household wedge, denoted $\tau_H$ in the legend. Corresponding panels share axis scales. The taxable curve is the model's pooled taxable market (top) and the observed corporate benchmark (bottom). The municipal data curve removes callability (Appendix C.1.2) and the issuance premium, and panel (d) uses this same adjusted curve. The model evaluates the estimated coefficients at observed issue-quarter rate states in the six calibration blocks and a 2018–2020 economy. Both rows use the empirical bond, rating and quarter weights and nominal semiannual yields. The household benchmark multiplies each observation's taxable yield by the high-tax household's relative after-tax share for that bond's state and date before averaging. Implied taxes and the household wedge are ratios of averaged yields.

model's thirty-year municipal and taxable yields are 4.651% and 6.185%, against 4.558% and 5.794% in the adjusted data.

## 5.4 Accounting for the puzzle

The accounting asks how much of the municipal puzzle the model explains, alongside the measured capital-gains, liquidity and default differentials. To compare these contributions, I express them on the same observed taxable yield curve. Otherwise, a higher taxable yield in the model would mechanically enlarge its gap in basis points, even with the same implied tax.

At each maturity, I keep the model's municipal-to-taxable yield ratio and multiply it by the observed taxable yield. I then subtract the household after-tax benchmark. The resulting gap, denoted $\widetilde{\mathcal{G}}_T$, measures the model's contribution on the data's taxable curve:

$$\begin{aligned} \text{Normalized municipal yield} &= \frac{\text{Model municipal yield}}{\text{Model taxable yield}} \times \text{Observed taxable yield}, \\ \widetilde{\mathcal{G}}_T &= \text{Normalized municipal yield} - \text{Household benchmark}. \end{aligned} \tag{22}$$

I call this the corporate normalization. It preserves the model's implied tax and changes only the scale of this comparison. The solved equilibrium and the curves in Figure 7 are unchanged. Appendix C.1.7 gives the averaging and normalization details.

The accounting has three steps at each maturity $T$:

$$\begin{aligned} \text{Spread}(T) - \underbrace{\big[\text{Call}(T) + \text{Prem}(T)\big]}_{\text{comparability adjustments}} &= \underbrace{\text{Puzzle}(T)}_{\text{municipal puzzle}}, \\ \text{Puzzle}(T) - \underbrace{\big[\text{CG}(T) + \text{Liq}(T) + \text{Def}(T)\big]}_{\text{bond-side adjustments}} &= \underbrace{\text{Puzzle}^{\text{adj}}(T)}_{\text{residual before the model}}, \\ \text{Puzzle}^{\text{adj}}(T) - \underbrace{\widetilde{\mathcal{G}}_T}_{\text{model}} &= \underbrace{\text{Residual}(T)}_{\text{unexplained}}. \end{aligned} \tag{23}$$

The spread, call adjustment and issuance premium follow Section 3.7. The terms CG, Liq and Def denote capital-gains, liquidity and default differentials. Since the model excludes these features, its gap is compared with $\text{Puzzle}^{\text{adj}}(T)$. The final residual equals the observed corporate yield times the implied-tax difference (model minus data), with the data reflecting all five adjustments.

At thirty years (Figure 8 and Table 4), the model accounts for 97 basis points of the 118-basis-point puzzle, or 83%, and the bond-side adjustments account for 15, or 12%, leaving about 6 basis points unexplained. The model's contribution grows with maturity in basis points. Its share of the puzzle is 74% at ten years, 97% at eighteen and 83% at thirty. The unexplained residual is negative at every maturity except thirty years. The model gap exceeds $\text{Puzzle}^{\text{adj}}(T)$ by 7 to 11 basis points from twelve to twenty-four years, and the largest absolute residual is 11 basis points at sixteen years. This negative residual is consistent with the convenience premium interpretation in Fleckenstein and Longstaff (2025).

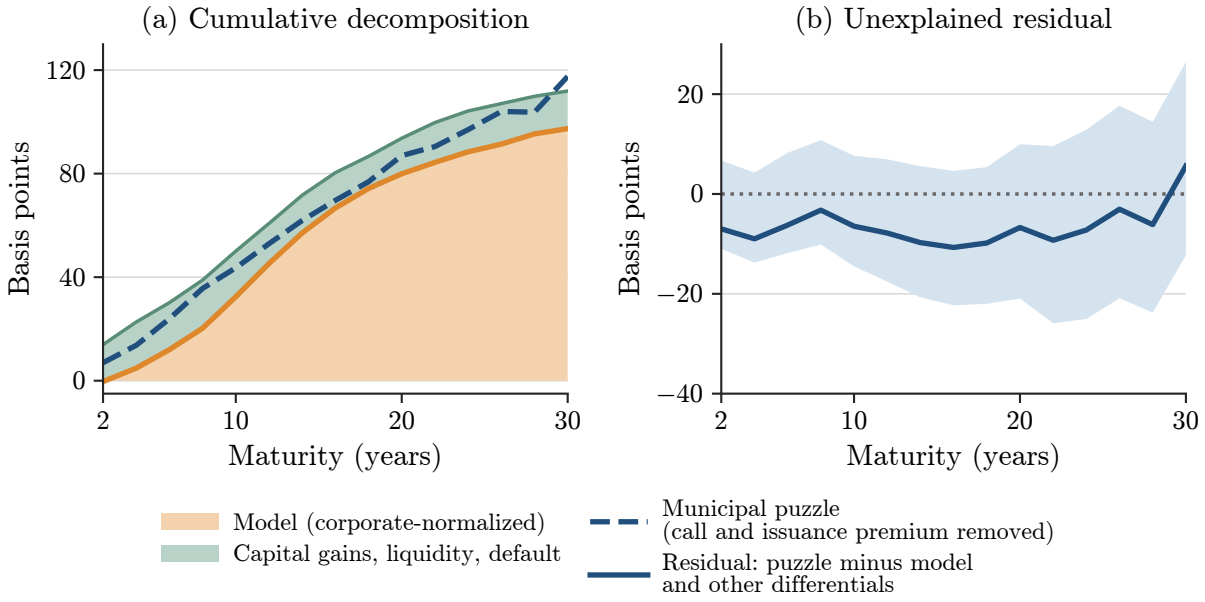



Figure 8: Cumulative decomposition of the municipal puzzle, 2000–2020, as in (23). The puzzle is the observed spread after removing callability and the issuance premium. Panel (a) stacks the model gap $\widetilde{\mathcal{G}}_T$ of (22) and the capital-gains, liquidity and default differentials against that puzzle. Panel (b) shows Residual($T$), the puzzle less their sum. Shading spans the ranges of the five empirical adjustments (Appendix Table 31) and excludes model uncertainty. Weights and yield conventions match Figure 7.

Table 4: Accounting for the municipal puzzle. Removing callability and the issuance premium from the raw spread gives the puzzle. The model contribution applies its implied tax to observed corporate yields, as in equation (22).

| | | Comparability | | | Accounting for the puzzle | | | |
|---|---|---|---|---|---|---|---|---|
| Maturity (years) | Raw spread | Call | Issuance premium | Puzzle | Model | Model share | Bond-side adjustments | Unexplained |
| 2 | 9.9 | 0.1 | 2.9 | 6.9 | −0.3 | | 14.3 | −7.0 |
| 6 | 27.9 | 0.7 | 3.3 | 23.9 | 12.0 | 50% | 18.2 | −6.2 |
| 10 | 52.2 | 4.6 | 4.0 | 43.6 | 32.4 | 74% | 17.6 | −6.5 |
| 14 | 87.0 | 17.9 | 7.2 | 61.9 | 57.1 | 92% | 14.6 | −9.7 |
| 18 | 110.3 | 23.9 | 9.6 | 76.9 | 74.3 | 97% | 12.4 | −9.8 |
| 22 | 128.8 | 27.0 | 11.4 | 90.5 | 84.3 | 93% | 15.4 | −9.3 |
| 26 | 145.0 | 28.5 | 12.5 | 104.0 | 91.4 | 88% | 15.7 | −3.1 |
| 30 | 159.8 | 29.5 | 12.8 | 117.5 | 97.4 | 83% | 14.5 | 5.6 |

*Notes.* Basis points, except maturity and shares. The accounting is puzzle = model + bond-side + unexplained, as in (23). The model column uses the corporate normalization, and the bond-side adjustments sum the capital-gains, liquidity and default differentials. The model share divides its contribution by the puzzle, using unrounded values, and is omitted where the puzzle is below 20 bp. Appendix C.1.7 reports the unnormalized gaps. Model and data both use the 2000–2020 sample, the averaging of Appendix C.1.1 and the nominal semiannual convention. The separately measured adjustments are not jointly estimated. Rows show every second model maturity. Figure 8 shows all fifteen.

# 6 Empirical Tests

I test the model's participation mechanism using variation in household taxes, institutional tax benefits and interest rates (Section 4.6). None of the estimates is a calibration target. P1 uses differences in residents' state income-tax rates. P2 uses bank qualification and the 2017 corporate tax cut for insurers. These designs compare household and institutional tax effects across maturities. P3 compares changes in municipal and treasury yields at the same maturity. For traded coupon bonds, I account for the timing of their payments when constructing the benchmark response.

The P1, TCJA and P3 tests use the AA-and-above municipal sample (agency score three or better) to limit credit differences that could affect both yields and trading. The curve accounting of Sections 3 and 5 uses the wider A-or-above sample (score seven or better), which matches the corporate benchmark's credit band, and Appendix C.11.2 repeats each test on it with the same maturity ordering. The bank-qualified design uses all new issues with rating-by-year fixed effects. Throughout, $n_b$ denotes a bond's remaining maturity in years.

## 6.1 P1. State-tax pass-through

P1 asks whether resident households' state-tax benefit is reflected less in municipal yields at longer maturities. Most states exempt residents' interest on in-state municipal bonds and tax interest on out-of-state bonds. A household in state $s$ has a state after-tax share of one on in-state municipal interest and $1-\tau_s^{\text{st}}$ on out-of-state municipal and taxable interest.

State rates move resident households' exemption benefits, a wedge institutions largely do not face. Pensions and tax-deferred accounts owe no current state tax on interest, and insurer premium taxes generally do not distinguish municipal from taxable interest. Funds pass the benefit through according to shareholder residence and portfolio eligibility: single-state funds preserve the home-state exemption, while national funds pass through only the eligible home-state share (Appendix D.2). Under its stated conditions, Appendix Proposition A.9 gives household pass-through of one at the riskless boundary and below one in the long-end limit, and Appendix Proposition A.16 gives the corresponding responses in state markets.

The yield ratio $y^M/y^C$ contains both federal and state taxes. Under the pre-2018 federal deduction convention, the after-tax share of taxable interest is $\xi_s^{\text{st}}\xi_t^{\text{fed}}$, where $\xi_s^{\text{st}} = 1-\tau_s^{\text{st}}$ is the state share and $\xi_t^{\text{fed}}$ the federal share, while home-state municipal interest is exempt at both levels. Household indifference under the representative-tax benchmark gives $y_{s,t}^M = \xi_t^{\text{fed}}\xi_s^{\text{st}}\, y_t^C$. Dividing out the federal share defines the implied state tax

$$\mathit{ImpliedStateTax}_{s,t}(T) \;\equiv\; 1 - \frac{y_{s,t}^M(T)/y_t^C(T)}{\xi_t^{\text{fed}}}. \tag{24}$$

It equals the statutory rate $\tau_s^{\text{st}}$ under the benchmark, giving pass-through of one. I set $\xi_t^{\text{fed}}$ to one minus the dated top federal ordinary-income rate, excluding the net investment income tax: 0.604 in 2000, 0.609 in 2001, 0.614 in 2002, 0.65 in 2003–2012 and 0.604 in 2013–2017. The sample ends

before the Tax Cuts and Jobs Act capped state and local tax deductions at $10,000 from 2018, and Appendix C.6 reports the later period under its tax mapping.

I estimate pass-through separately in five maturity bins ($[0,5)$, $[5,10)$, $[10,15)$, $[15,20)$ and $[20,30]$ years):

$$\begin{aligned} ImpliedStateTax_{b,t} = &\sum_{g=1}^{5} \beta_g \big( StateTax_{s(b),t} \times \mathbf{1}\{n_b \in g\} \big) + \sum_{g=1}^{5} \delta_g' \big( Z_{b,t} \times \mathbf{1}\{n_b \in g\} \big) \\ &+ \lambda' V_b + \alpha_{t \times n \times callable} + \varepsilon_{b,t}, \end{aligned} \tag{25}$$

where $StateTax_{s,t}$ is the state's top statutory rate and the outcome (24) uses call-adjusted zero-coupon yields. The coefficient $\beta_g$ is the change in the implied state tax per percentage point of statutory rate in maturity bin $g$, and the benchmark gives $\beta_g = 1$ in every bin. The controls $Z_{b,t}$, also interacted with the bins, are the rounded average rating and six lagged state variables: par-weighted outstanding maturity, the maturity of the preceding year's issuance, outstanding par relative to personal income, log income, log preceding-year issuance and the outstanding share in the bond's maturity bin. The stock and issuance variables cover federally exempt bonds, and the stock and income variables are measured at the preceding quarter-end. $V_b$ contains log offering par and security-type indicators, and $\alpha_{t \times n \times callable}$ denotes issue-year–month–maturity–callability fixed effects.

The empirical sample excludes Illinois, Iowa, Wisconsin and Oklahoma and retains bonds recorded as exempt from own-state tax, or issued in a state with a zero rate. The 2000–2017 sample has 681,838 observations with complete controls. Standard errors are clustered by state, allowing dependence among a state's bonds across years and maturities. As an instrumental-variables (IV) comparison, I instrument the current statutory rate with its 1990 historical counterpart, following the policy-history approach of Babina et al. (2021). The results are robust to the choice of historical year, and Appendix C.6 reports other vintages with the same controls and sample. Cross-sectional variation in state tax rates identifies pass-through into yields at different maturities. Historical rates address feedback from current municipal yields to current tax policy. The IV interpretation requires that, conditional on the controls, historical rates affect current yields only through current tax rates.

Figure 9 compares the ordinary least squares (OLS) estimates of (25) with the full-state model's cross-state slope. For the model, I compute (24) using its high-tax household's federal share $\xi^{\text{fed}}$ and nominal semiannual yields. At each maturity, I regress this implied tax on the state's top rate averaged within the three-year block, with an intercept and equal weights across the 43 income-tax states and the District of Columbia. All states and economies are evaluated at the same fixed rate state, and the slopes are averaged equally across the six pre-2018 blocks.

Estimated pass-through declines across all five bins, from 1.069 below five years to 0.305 beyond twenty (Appendix Table 33). A state top rate ten percentage points higher raises the implied tax by about 10.7 points at the short end and 3.0 points at the long end. The short-to-long decline

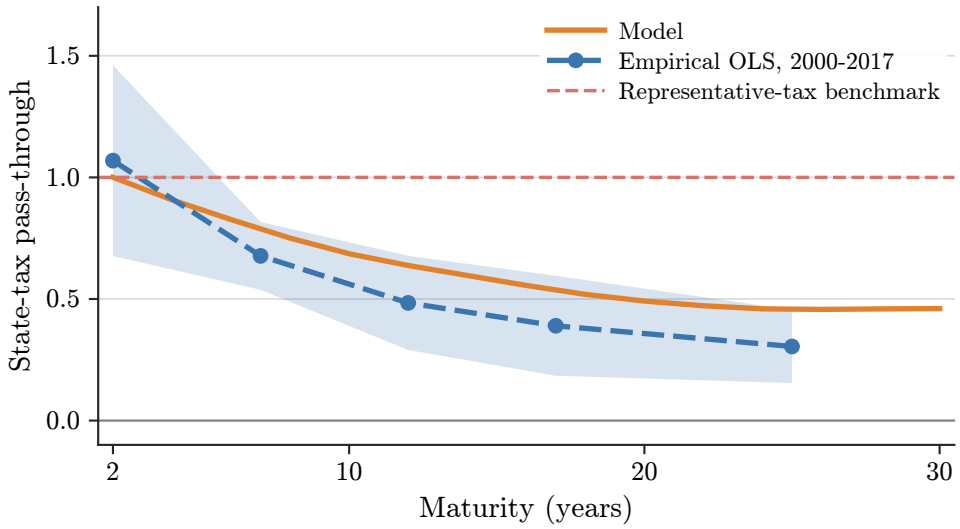



Figure 9: State-tax pass-through by maturity. Blue: OLS estimates for 2000–2017 call-adjusted zeros, with state-clustered 95% intervals using $t(46)$. The regression is unweighted and includes the controls in (25). Orange: the model's cross-state slope, using its own federal after-tax share and nominal semiannual yields at a common fixed rate state and averaged equally across the six pre-2018 economies. Both use (24). The red dashed line is the representative-tax benchmark.

is 0.764 (s.e. 0.235, $p = 0.002$). Unit pass-through cannot be rejected in the shortest bin and is strongly rejected in the longest, and every bin coefficient is positive and significant at one percent. The IV estimates also decline, from 1.088 to 0.229, with an endpoint-equality $p$-value below 0.001.

The model's untargeted profile declines from 1 at two years to 0.69 at ten and 0.46 at thirty. The empirical decline survives alternative yield constructions, credit controls, weights, instruments, state fixed effects and single-state omissions (Appendices C.5 and C.6). The appendix also reports the post-2018 period separately under its changed tax mapping.

## 6.2 P2. Bank qualification

P2 uses a tax benefit that only banks receive. Banks ordinarily cannot deduct interest expense allocated to tax-exempt holdings (Section 265(b) of the Internal Revenue Code). For bank-qualified (BQ) bonds, Sections 265(b)(3) and 291 restore 80% of that deduction. An issuer can designate its bonds as bank-qualified if its reasonably anticipated tax-exempt issuance for the year does not exceed $10 million ($30 million in 2009–10 under the American Recovery and Reinvestment Act). Designation changes neither the bonds' payments nor any other investor's taxes. Dagostino (2025) studies how the expanded cap affected issuance and employment.

A bank allocates interest expense across assets in proportion to tax basis, so each dollar of municipal basis carries the same annual funding cost $c$ at every maturity. Designation saves $X = 0.8\,\tau_B c$ a year, where $\tau_B$ is the federal bank tax rate. Measured from Call Report interest expense, $X$ is about 65 basis points in 2005–08 and 9 to 14 from 2011 to 2019 (Appendix C.7.3). Because a bond's benefit accrues over its life, I also scale premia by $\bar{X}_b$, the expected annual benefit over bond $b$'s life (Appendix C.7.3). Its median is 45 basis points.

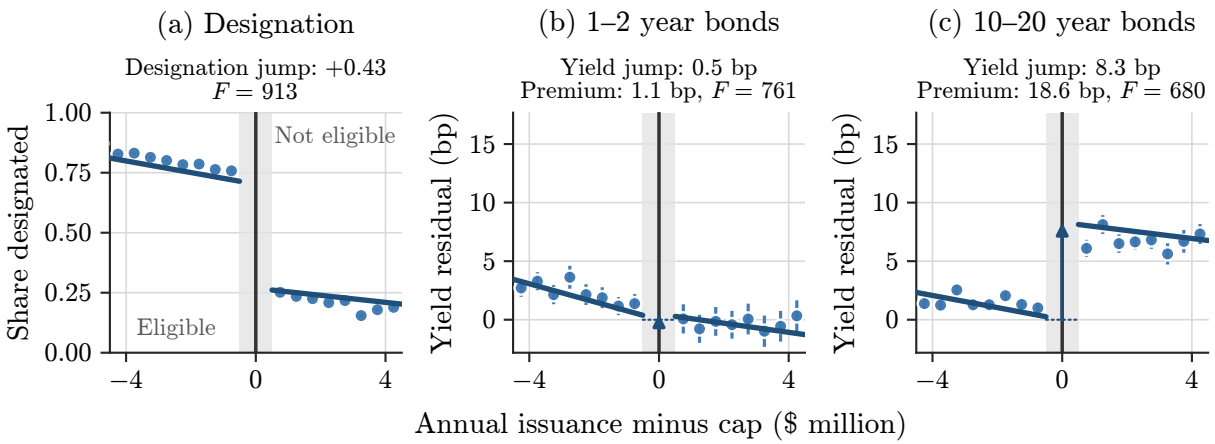



Figure 10: The bank-qualified discontinuity. The horizontal axis is an issuer's measured annual tax-exempt issuance minus that year's cap. Designation is more frequent to the left of zero. Dots are means in \$0.5 million bins with 95% intervals: the share of bonds designated in panel (a), and offering-yield residuals after the fixed effects and log par of (27) for 1–2 and 10–20 year bonds in panels (b) and (c). Lines are the estimated local linear fits, which exclude the shaded donut. Yield fits are shifted to zero on the below-cap side at the cap, and the arrow marks the yield jump. Each premium divides the yield jump by its bin's designation jump. $F$ is the first-stage Wald statistic with issuer-clustered standard errors. Appendix Figure 24 shows the other maturities.

Three hypotheses differ in how much of $X$ is reflected in the designation premium $\varpi_k = y_k^N - y_k^Q$, the non-designated minus the designated yield at maturity $k$. Under the representative-tax benchmark, households hold both bonds at the same yield, $y_k^N = y_k^Q = (1-\tau_H)y_k^C$. A bank then prefers a taxable bond, which pays $(1-\tau_B)y_k^C$ after tax, to a designated bond, which pays the household benchmark yield less the tax value of the disallowed fifth of its funding cost, whenever

$$(1-\tau_B)\,y_k^C > (1-\tau_H)\,y_k^C - 0.2\,\tau_B c \quad \Longleftrightarrow \quad \tau_H\,y_k^C > \tau_B\,y_k^C - 0.2\,\tau_B c,$$

that is, whenever the household's exemption benefit exceeds the bank's net exemption benefit on a designated bond, as it does when household rates are at least as high as banks'. Banks then hold neither municipal bond, and $\varpi_k = 0$. Under a bank-only benchmark, banks' indifference to taxable bonds gives full capitalization, $\varpi_k = X$, flat in maturity. Holding outside investment opportunities and other bond characteristics fixed, the premium is the benefit times the bank's capitalization share $\theta_k$, the average bank share of the total demand response to the bond's yield (Appendix C.7):

$$\varpi_k = \theta_k\,X, \qquad \underbrace{\theta_k = 0}_{\text{representative-tax}}, \qquad \underbrace{\theta_k = 1}_{\text{bank-only}}, \qquad \underbrace{\theta_{\text{short}} \simeq 0 < \theta_{\text{long}}}_{\text{model}}. \tag{26}$$

The capitalization share is the bank counterpart of $\mathcal{P}_I^Y(T)$ in Section 4.6. Household after-tax cash returns anchor short yields and banks hold few short bonds, so $\theta_k$ is near zero at short maturities. In the calibrated model, banks hold intermediate maturities, where their duration benefit offsets their smaller exemption benefit (Section 4.3), and yields there reflect more of the bank benefit.

The cap provides a comparison among issuers of similar size. Designation averages 0.79 below the cap and 0.21 above it across the estimation window. Measured annual issuance need not equal the issuance anticipated at each sale or counted under the statute, so some designated bonds appear above the measured cap. Below-cap status raises the probability of designation without determining it, giving a fuzzy discontinuity. It instruments designation, and the premium is the local yield jump divided by the local designation jump.

The variable $R_{jt}$ is issuer $j$'s annual tax-exempt issuance minus that year's cap, in millions of dollars. The indicator $E_{jt} = \mathbf{1}\{R_{jt} \leq 0\}$ marks below-cap issuer-years, and $B_b$ indicates whether bond $b$ is bank-qualified. The Mergent sample covers tax-exempt new issues from 2000 to early 2021 with $0.5 \leq |R_{jt}| \leq 4$. The half-million exclusion around the cap, the donut, removes managed bunching and issuer-years exactly at the cap. With bins $g$ of maturity at issuance (1–2, 2–5, 5–10, 10–20 and more than 20 years), I estimate

$$y_b = \sum_g \beta_g \, B_b \mathbf{1}\{n_b \in g\} + W_b'\delta + u_b, \tag{27}$$

by two-stage least squares, instrumenting each designation interaction $B_b\mathbf{1}\{n_b \in g\}$ with the eligibility interaction $E_{jt}\mathbf{1}\{n_b \in g\}$ of the same bin. The first stage is the local designation jump conditional on controls, 0.42 to 0.45 across bins (Appendix Table 39). Here $W_b$ contains local linear controls in $R_{jt}$ on each side of the cap, log tranche par, and fixed effects for maturity, state, offering month, rating, security type, use of proceeds, coupon and callability, with the interactions listed in Appendix C.7.2. The outcome is the offering yield in basis points, or that yield divided by $\bar{X}_b$, so $-\beta_g$ is the premium in basis points or in expected-benefit units. Standard errors are clustered by issuer.

Identification requires that, given these controls, crossing the cap moves offering yields only through designation. Callability and the number of maturities per issue are balanced at the cap. Placebo cutoffs at \$5 million and \$15 million show essentially no designation effects, and their 10–20 year yield jumps exceed the 1–2 year jumps by only 0.4 and 2.5 basis points, against 7.8 at the cap (Appendix C.7).

Figure 10 illustrates the discontinuity with separate local fits for each maturity bin. The 1–2 year yield barely moves at the cap, by 0.5 basis points, while the 10–20 year yield jumps by 8.3, implying bank-qualification premia of 1.1 and 18.6 basis points after dividing by the designation jumps. The main estimates instead come from the joint two-stage least squares (2SLS) regression (27), which estimates all five bins together. Figure 11 reports these estimates (Appendix Table 39). The premium is 1.8 basis points at 1–2 years, indistinguishable from zero, and 8.7, 13.8, 18.4 and 18.2 at 2–5, 5–10, 10–20 and more than twenty years. The 10–20 year premium exceeds the 1–2 year premium by 16.5 basis points (s.e. 2.7). In expected-benefit units, the premium rises from 0.17 at 1–2 years to 0.52 at 10–20 years: the bank-only benefit is barely priced in short municipal yields, and about half of it is priced at 10–20 years.

The positive premia at 5–20 years reject the representative-tax benchmark, with $t$ statistics of 6 to 8, and the rise in expected-benefit units, 0.34 (s.e. 0.11), rejects the flat premium of full capitalization. The gradient appears in every benefit regime and survives all robustness checks, including alternative donuts and bandwidths, credit-enhancement controls, a general-obligation sample, issuer effects, year clustering and the binned estimator of Dagostino (2025) (Appendix C.7).

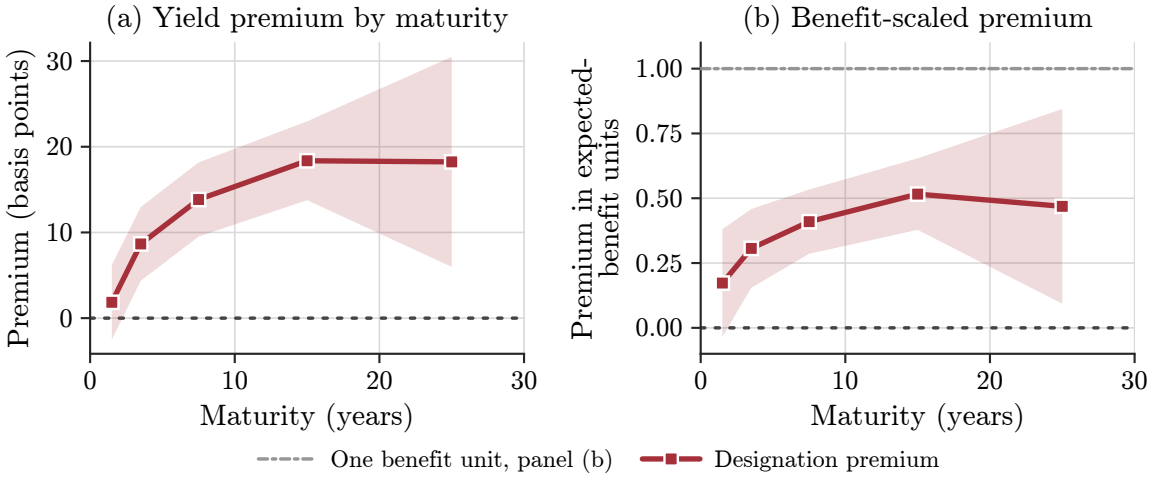



Figure 11: The bank-only benefit and municipal yields. Panel (a): 2SLS yield premium in basis points. Panel (b): the premium in expected-benefit units, with each yield divided by the expected annual benefit $\bar{X}_b$ and the same controls. Funding costs are municipal-holdings-weighted. Expected rates use the Kim–Wright expected-short-rate component, capped at ten years (Appendix C.7). Shading is the issuer-clustered 95% interval. The dashed line marks one benefit unit.

The rising bank profile is the reverse of household pass-through in P1, which falls from 1.07 to 0.30 (Figure 9), as the model predicts for institutional and household tax benefits. The 2017 corporate tax cut provides a second institutional test: it reduced P&C insurers' exemption benefit, and implied taxes then fell more at longer maturities in states where these insurers held more municipal debt. Appendix C.9 presents this second test.

### 6.3 P3. Rate responses by maturity

P3 asks whether municipal yields respond more to a given treasury yield change at longer maturities. Section 4.6 derives this maturity pattern relative to the model's jointly priced taxable market $C$. The empirical test uses treasury yields to measure rate movements. The same prediction holds for this reference in the analytical illustration.[15] The representative-tax benchmark instead makes municipal zero yields a constant fraction $\xi_H$ of treasury zero yields, implying a constant relative response. The observed response is about 0.35 at one year, below $\xi_H$. The test concerns how this relative response varies with maturity.

[15] In the two-investor illustration, municipal sensitivity relative to a taxable reference whose one-period return responds one for one to the rate state increases over the thirty illustrated maturities. Appendix Proposition A.9(iii) establishes the maturity ordering for that reference in the extension with joint long-maturity holdings.

I compare consecutive trades in the same municipal bond, using its yield to maturity and the maturity-matched treasury par yield. With $\Delta y^{\mathrm{M}}_{b,t}$ the change in its yield between the two trades and $\Delta y^{\mathrm{Tsy}}_{n_b,t}$ the change in the treasury yield matched to its remaining maturity $n_b$, I estimate

$$\Delta y^{\mathrm{M}}_{b,t} = \beta_0 \, \Delta y^{\mathrm{Tsy}}_{n_b,t} + \beta_1 \big( \Delta y^{\mathrm{Tsy}}_{n_b,t} \times n_b \big) + \beta_2 \, n_b + \varepsilon_{b,t}. \tag{28}$$

The response to a treasury yield change is $\beta_0 + \beta_1 n_b$. A positive $\beta_1$ means that the same treasury yield change is associated with a larger municipal yield change at longer maturities. Comparing trades within a bond removes its time-invariant characteristics. The sample covers January 2005 to August 2020 and includes non-callable, federally tax-exempt municipal bonds rated AA and above. Trades are at most 90 days apart, and the treasury yield must move by at least 10 basis points in either direction to limit transaction noise. The treasury comparison follows the bond's remaining maturity at each trade, and standard errors are clustered by the second trade's date. Appendix C.10 details the construction.

The estimates are $\hat{\beta}_0 = 0.346$ and $\hat{\beta}_1 = 0.007$ (Appendix Table 45). A 100-basis-point rise in the treasury yield goes with a municipal increase of about 35 basis points at one year and 42 at ten years. Splitting bonds at five years gives the same ordering, with responses of 0.351 and 0.426.

All thirty combinations of treasury-move cutoffs and trade windows give a positive maturity interaction, from 0.0047 to 0.0104, significant at five percent (Appendix Table 47). On the wider rating sample, the gradient is larger, 0.013 against 0.007 in the continuous specification and 0.113 against 0.075 with the five-year split (Appendix Tables 50 and 51).

# 7 Policy Counterfactuals

Tax reforms that change investors' exemption benefits change who holds municipal bonds. The Tax Reform Act of 1986 generally removed banks' interest deduction for new municipal purchases and was followed by a sharp decline in their holdings. The 2017 Tax Cuts and Jobs Act again changed banks' and insurers' relative exemption benefits. I study two reforms with existing policy precedents. The first removes a state's exemption for residents' interest on its own municipal bonds, which Illinois, Iowa and Wisconsin already generally tax. The second taxes municipal interest like taxable interest for every investor, as the 2009–2010 Build America Bonds program did for its taxable municipal bonds while subsidizing issuers' interest.

The key distinction is whether investors can retain the federal benefit by switching municipal markets. Home-state repeal leaves that benefit in place, so high-tax households mainly switch between municipal markets while their replacement investments generate state-taxable interest. Equal treatment removes the federal benefit as well. Investors who previously benefited most shorten their bond portfolios, while tax-neutral investors expand their long-bond holdings. The financing and revenue effects depend on both the interest these investors earn and the compensation they require for bearing rate risk.

### 7.1 Design and fiscal accounting

I solve each reform in the six calibrated three-year economies covering 2000–2017 and report calendar-year-weighted averages. Each economy uses the pre-2018 tax regime and period-specific investor and bond-market inputs (Section 5.2). I change the relevant tax rules and re-solve holdings and both yield curves, keeping bond supplies, investors' total funding, risk preferences and outside exposures at their dated baseline values.[16]

Annual financing costs equal debt outstanding at each maturity times its yield, summed across maturities. Their increase is the replacement-financing cost, the additional annual cost of financing the same debt at the new yields.[17] Each group's pretax interest on bonds and taxable cash divides into tax receipts and after-tax interest income, with households taxed by residence and institutions under their own tax treatment. Because aggregate funding, bond supplies and the short rate are fixed, total cash interest does not change. Every additional dollar issuers pay therefore goes either to governments or to investors. Summing across both markets and all governments gives

$$\Delta(\text{Financing costs}) = \Delta(\text{Tax receipts}) + \Delta(\text{After-tax interest income}). \tag{29}$$

I report financing costs per dollar of added receipts. For home-state repeal, I compare both own-state and all-issuer costs with the repealing state's added revenue, excluding changes in other governments' receipts. For equal treatment, I compare total municipal and taxable issuer costs with the combined increase in federal and state revenue. When costs and receipts cover all issuers and governments, equation (29) gives a ratio of one if aggregate after-tax interest income is unchanged and a ratio above one if that income rises.[18]

### 7.2 The home-state exemption

Most income-tax states exempt residents' interest on their own municipal bonds and tax interest on other states' bonds, which encourages residents to hold local debt. Illinois, Iowa and Wisconsin generally tax interest on their own bonds as well, so the policy I study already exists in some states.[19]

[16] Tax changes are applied in each investor's baseline payoff units, preserving preferences and liability hedges. Yield changes average each economy's 36 monthly rate states. Dollar amounts, California's implied state taxes and Figure 12(a) use $z = 0$. Ratios for 2000–2017 divide average dollar financing costs by average dollar receipts.

[17] The calculation uses continuously compounded yields to construct coupon-equivalent flows. Existing contractual payments and original-issue-discount schedules remain unchanged. The taxable market pools treasury and corporate debt. Other liability flows, issuer deductions, refinancing timing and fiscal feedback are outside this interest-income accounting (Appendix D.4).

[18] These ratios compare financing and interest-income flows. Welfare would also depend on changes in risk bearing, investment opportunities and public spending.

[19] Each of these three states exempts particular issues. The Supreme Court upheld the home-state preference in *Department of Revenue of Kentucky v. Davis* (2008). Home-state preferences also changed within the sample: Oklahoma exempts qualifying local obligations issued after July 1, 2001, and Utah has taxed out-of-state municipal interest on bonds acquired after January 1, 2003, except from states that do not tax Utah's bonds. Since 2018, the federal state-and-local-tax deduction cap has also increased the exemption's value for residents whose deductions it limits.

Panel (a) of Figure 12 gradually removes California households' exemption on local municipal interest in the 2015–2017 economy, keeping statutory state and federal tax rates and household composition fixed. The horizontal axis measures the remaining tax benefit of a local bond relative to an out-of-state bond for the high-tax household. The implied state tax is the percentage by which California's yield lies below the pooled municipal yield of the states without an income tax.[20] At the riskless two-year maturity, the implied state tax moves one for one with this benefit. Removing the exemption entirely lowers the implied state tax by 13.3 percentage points at two years but only 5.0 percentage points at thirty years. As in P1, the household's tax benefit is fully reflected in short yields and only partly in long yields, where lower-tax residents and institutions also hold the bonds.

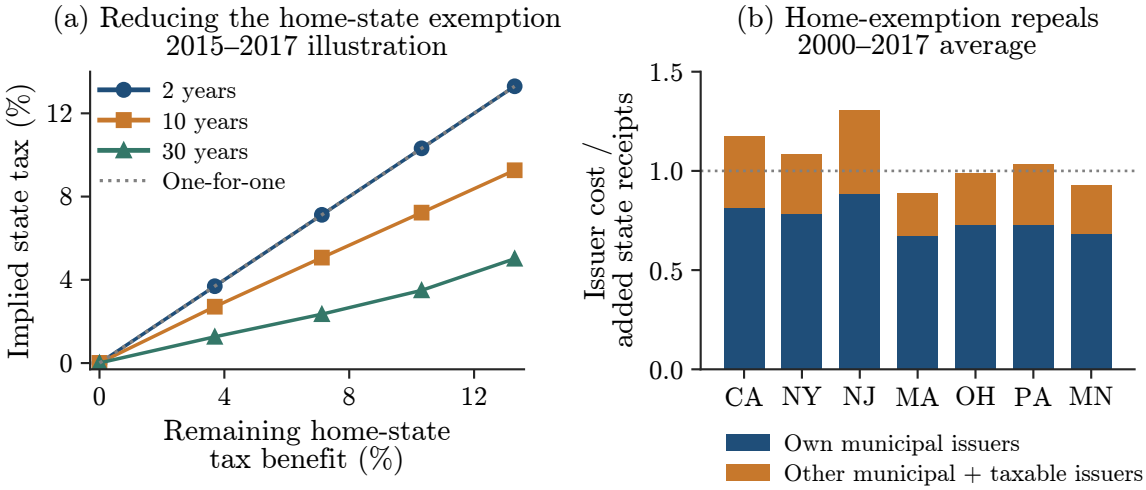



Figure 12: Home-state exemptions, municipal yields and financing costs. Panel (a) makes 0, 25, 50, 75 and 100% of local municipal interest state-taxable for both California household groups in the 2015–2017 economy, keeping all statutory tax rates fixed. It plots $1 - y^M_{\mathrm{CA},T}/y^0_T$ at two, ten and thirty years against the high-tax household's remaining relative home-state tax benefit. The comparison yield $y^0_T$ pools the seven states without an income tax. Markers are solved equilibria at $z = 0$, and the dotted line is one-for-one. Panel (b) removes each listed state's home-state exemption at fixed tax rates in each of the six 2000–2017 economies. It divides average replacement-financing costs for own-state municipal issuers and other municipal and pooled taxable issuers by average added state receipts from residents' bond and taxable cash interest at $z = 0$.

Panel (b) removes residents' state exemption on home-state municipal bonds, separately in California, New York, New Jersey, Massachusetts, Ohio, Pennsylvania and Minnesota, seven states with high supply-weighted tax benefits. Residents then pay the same state tax on home-state as on out-of-state municipal interest, while the tax treatment of institutions, nonresidents, cash and

[20] The implied state tax is $1 - y^M_{\mathrm{CA},T}/y^0_T$, where $y^0_T$ is the pooled no-income-tax-state municipal yield. It is the model counterpart of the P1 measure in (24), with $y^0_T$ in place of $\xi^{\mathrm{fed}}_t y^C_T$. Both yields are solved at rate state $z = 0$. If a fraction $f$ of local interest becomes state-taxable, the high-tax household's relative home-state tax benefit is $(1 - f)\tau^{\mathrm{st}}_H/(1 - f\tau^{\mathrm{st}}_H)$. This compares its after-tax shares on local and out-of-state interest and equals its state rate when the local exemption is full.

taxable bonds is unchanged. Added state receipts generally exceed local issuers' added financing costs, which rise by \$0.68–\$0.88 per receipt dollar, below one in all seven six-period averages (Appendix Table 55). Including other municipal and taxable issuers raises the ratio to \$0.89–\$1.30, above one in California, New York, New Jersey and Pennsylvania. Both ratios use the same receipts, so their difference is the cost transmitted to other markets, of which taxable issuers bear \$0.11–\$0.31 per receipt dollar.

To see the mechanism, consider one dollar of state $s$'s maturity-$T$ debt initially held by a resident. Let $y_T^0$ be the pooled no-income-tax-state municipal yield and $\tau_{s,T}^{\text{impl}} = 1 - y_{s,T}^M / y_T^0$ the implied state tax. Holding $y_T^0$ fixed, the issuer's yield rises in proportion to the fall in the implied state tax. The state, however, collects its full rate $\tau_s^{\text{st}}$ on the resident's interest after repeal. For a dollar that continues to earn $y_T^0$, the annual comparison is

$$\underbrace{\Delta y_{s,T}^M = -y_T^0 \, \Delta \tau_{s,T}^{\text{impl}}}_{\text{added cost to the issuer}} \qquad \text{against} \qquad \underbrace{\tau_s^{\text{st}} \, y_T^0}_{\text{tax on the formerly exempt dollar}} .$$

Under the representative-tax benchmark, residents determine relative yields at every maturity. Repeal lowers the implied state tax by the full rate, $-\Delta \tau_{s,T}^{\text{impl}} = \tau_s^{\text{st}}$, and the two sides are equal. In the model, this holds at the riskless short end. Long municipal bonds are also priced by lower-tax residents and investors who do not receive the home-state exemption, so their yields reflect only part of the top-bracket resident's benefit. Removing the exemption therefore raises long yields by less than the household benchmark predicts, consistent with the partial pass-through documented in P1. Across the six economies, California's thirty-year implied state tax falls by 3.8 percentage points, against a mean high-tax household rate of 10.8%. Per dollar of debt, the thirty-year response is thus about a third of the household benchmark. The aggregate calculation must also account for resident ownership, because higher yields affect all local debt while the state taxes only its residents.

The federal exemption remains after home-state repeal, so residents can retain its benefit by buying other states' municipal bonds. High-tax households therefore mainly switch between municipal markets, with relatively modest changes in the maturities they hold. Their sales of local bonds do not by themselves erode the state's tax base, because interest on the replacement investments is also taxable. Across the seven experiments, residents' aggregate risky-bond share changes by less than two percentage points (Appendix Table 54). Reinvestment at the same-maturity comparison yield $y_T^0$ preserves the benchmark receipts above. Lower-yielding investments reduce those receipts, while lower-tax households' substitution into taxable bonds with higher gross yields increases them. The equilibrium calculation includes interest on purchases and interest forgone on sales, together with changes in yields and maturities. Appendix D.4 separates these effects.

Reallocation also spreads financing costs to other issuers because investors choose jointly across municipal and taxable markets. In the California experiment, other states' municipal issuers bear an additional \$0.11 and taxable issuers \$0.25 per dollar of added California receipts. These spillovers raise the cost from \$0.81 for California's own issuers to \$1.17 for all issuers. The state's receipts

can therefore exceed its own borrowers' costs even when financing costs across all markets exceed those receipts.

### 7.3 Equal municipal and taxable treatment

The second reform taxes each investor's municipal interest as it taxes that investor's taxable interest. It removes federal and state exemption benefits for every investor group.[21] The Build America Bonds program of 2009–2010 tested a version of it: issuers could sell taxable bonds and receive a federal subsidy of 35% of their interest. That program changed individual issues while the rest of the market stayed exempt. Equal treatment changes all municipal debt at once, so it also changes who holds long bonds and the yields they require.

I implement the reform by moving each municipal after-tax share 99.9% of the way to its taxable share and removing the same share of banks' interest-deduction disallowance, so that banks' municipal interest is taxed like their taxable interest. With exact equality, the riskless municipal bond and cash would be perfect substitutes, leaving investors' split between them indeterminate. The remaining 0.1% pins down that split and keeps yields and the allocation of bonds across maturities close to those of full removal (Appendix B.6.5). The tax treatment of cash and taxable interest does not change.

Equal treatment raises financing costs by more than it raises tax receipts. All amounts below are average annual changes across the six 2000–2017 economies. Per dollar of added federal and state receipts, municipal issuers bear about $1.06 in added financing costs and taxable issuers bear another $0.13. Total costs rise by about $1.20, leaving roughly $0.20 per receipt dollar as additional after-tax interest income for investors (Table 5 and Figure 13(a)). In annual dollars, municipal issuers pay an additional $41.2 billion and taxable issuers $5.1 billion, for a total of $46.3 billion. Governments collect $38.7 billion and investors retain the remaining $7.6 billion after tax, by (29). Appendix Table 56 reports each period separately.

Unlike home-state repeal, equal treatment leaves investors no municipal market in which they can retain the federal exemption. High-tax households, banks and P&C insurers shorten the average maturity of their risky bonds by 1.8 to 3.7 years, and high-tax households also increase their riskless holdings (Appendix Table 54). Shorter bonds and cash generate less interest than long bonds, limiting the revenue collected from the affected investors. Tax-neutral investors' share of bonds beyond ten years rises from 39% to 57%, while lower-tax households' and life insurers' shares decline slightly to 27% and 14%. Long-bond interest thus shifts toward the tax-neutral group, on whom the reform imposes no additional tax, limiting receipts.

At the same time, the tax-neutral group's larger long-bond holdings increase its rate exposure and the returns it requires in both markets. Municipal yields rise nearly to taxable yields, increasing by 170.6 basis points at thirty years, while thirty-year taxable yields rise by 23.1 basis points. The transfer of rate risk therefore raises borrowing costs in both markets, while the new holders' smaller additional tax payments limit receipts.

[21] The reform is broader than the investor-specific changes in the 2017 Tax Cuts and Jobs Act (Appendix C.9).

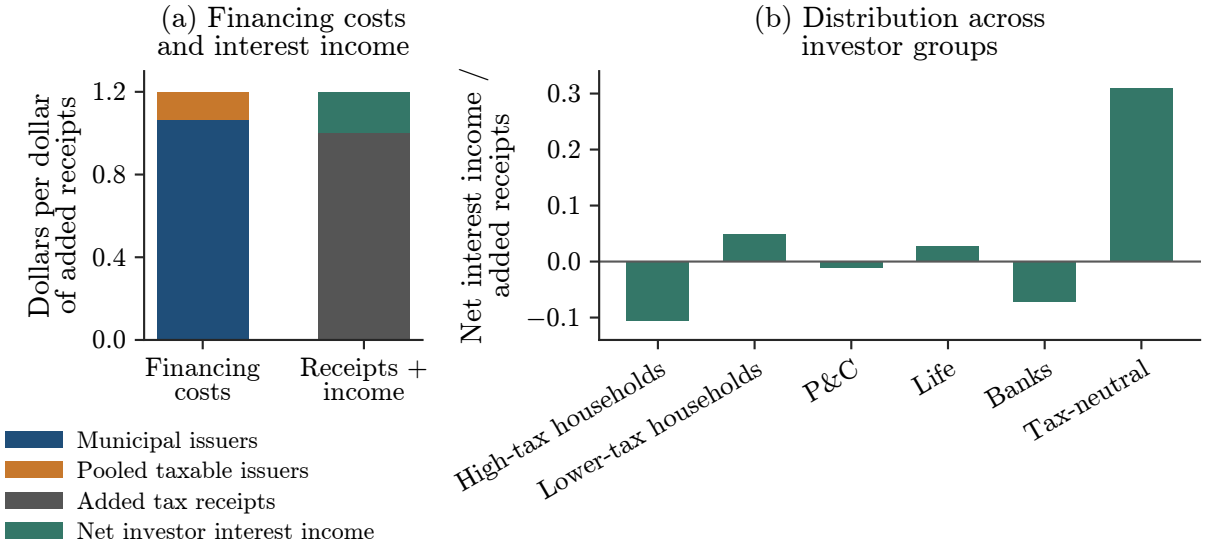



Figure 13: Average annual flows after removing 99.9% of each tax distinction in the six 2000–2017 economies. Panel (a) reports identity (29), dividing average financing costs and net investor interest income by average added federal and state receipts. Panel (b) allocates the income change across investor groups. Receipts and income include taxable cash interest. Each economy uses continuous yields and equilibrium holdings at $z = 0$. Taxable issuers pool treasury and corporate quantities.

Table 5: Average annual fiscal effects of equal municipal and taxable treatment, 2000–2017

| Δ(Financing costs) | | | = | Δ(Tax receipts) + Δ(After-tax interest income) | | |
|---|---|---|---|---|---|---|
| | $bn/ year | Per receipt dollar | | | $bn/ year | Per receipt dollar |
| Municipal issuers | 41.16 | 1.063 | | Federal receipts | 34.29 | 0.886 |
| Pooled taxable issuers | 5.15 | 0.133 | | Residence-state receipts | 4.42 | 0.114 |
| | | | | Total receipts | 38.71 | 1.000 |
| | | | + | Net investor interest income | 7.59 | 0.196 |
| Total financing costs | 46.30 | 1.196 | = | Total receipts and income | 46.30 | 1.196 |

*Notes.* Calendar-year-weighted changes across six three-year economies, each relative to its own baseline, with 99.9% of the tax distinction removed. Dollar amounts use dated-period quantities without inflation adjustment and continuous yields at $z = 0$. Per-receipt-dollar columns divide each average dollar change by average total receipts. Receipts and income include taxable cash interest, and net investor interest income is the change in after-tax interest income. Before rounding, total issuer costs equal receipts plus net investor interest income. Totals may differ because of rounding.

The $7.6 billion net increase in annual after-tax interest income is distributed unevenly (Figure 13(b)). The tax-neutral group gains $12.0 billion a year, lower-tax households gain $1.9 billion and life insurers gain $1.1 billion. High-tax households lose $4.1 billion, banks lose $2.8 billion and P&C insurers lose $0.5 billion. These gains exceed the aggregate because investors who benefited more from the exemption lose income.

A direct subsidy, as under Build America Bonds, can return the added receipts to municipal issuers. Returning all $38.7 billion of added annual receipts still leaves municipal issuers paying $2.4 billion more per year and taxable issuers paying $5.1 billion more. The subsidy covers most of municipal issuers' added costs but not the excess across both markets. With bond supplies and investors' total funding fixed, payments to issuers leave gross yields and investors' portfolios unchanged.

# 8 Conclusion

I develop an equilibrium term-structure model in which heterogeneous investor taxes create endogenous segmentation across bond markets and maturities. Investors jointly choose portfolios, balancing tax benefits against interest-rate risk. With short-sale constraints, these choices determine bond ownership and whose taxes are reflected in yields. Portfolio substitution links markets and maturities, so an investor's tax treatment can affect even bonds it never holds.

The municipal application shows why one representative investor's taxes cannot explain an entire yield curve. Top-bracket households value the exemption most, yet investors with smaller exemption benefits hold substantial municipal portfolios. The model connects this ownership to relative yields: household taxes anchor the riskless short end, while longer yields also reflect lower-tax investors' valuations and the allocation of rate risk. The quantitative and empirical evidence supports this explanation of the municipal bond puzzle.

The counterfactuals show why tax-policy analysis must account for financing costs in both markets and the redistribution of investor income. The framework applies wherever investors with different tax treatments choose portfolios under arbitrage frictions. Tax heterogeneity is a first-order determinant of bond yields and should be considered in term-structure models.

## References


Adelino, M., Cheong, S. C., Choi, J., and Oh, J. Y. J. (2026). Mutual fund flows and the supply of capital in municipal financing. *Review of Financial Studies*. Forthcoming.

Adrian, T., Crump, R. K., and Moench, E. (2013). Pricing the term structure with linear regressions. *Journal of Financial Economics*, 110(1):110–138.

Ang, A., Bhansali, V., and Xing, Y. (2010). Taxes on tax-exempt bonds. *Journal of Finance*, 65(2):565–601.

Asquith, P., Au, A. S., Covert, T., and Pathak, P. A. (2013). The market for borrowing corporate bonds. *Journal of Financial Economics*, 107(1):155–182.

Auerbach, A. J. and King, M. A. (1983). Taxation, portfolio choice, and debt-equity ratios: A general equilibrium model. *Quarterly Journal of Economics*, 98(4):587–609.

Babina, T., Jotikasthira, C., Lundblad, C. T., and Ramadorai, T. (2021). Heterogeneous taxes and limited risk sharing: Evidence from municipal bonds. *Review of Financial Studies*, 34(1):509–568.

Bagley, J., Gissler, S., Hiteshew, K., and Ivanov, I. T. (2023). Pushing bonds over the edge: Investor demand and municipal bond liquidity. Federal Reserve Bank of Chicago Working Paper 2023-05, revised July 2023.

Benson, E. D. and Marks, B. R. (2024). The impact of the Tax Cuts and Jobs Act on the true interest cost of bank-qualified tax-exempt bond issues. *Journal of Public Budgeting, Accounting & Financial Management*, 36(4):471–489.

BlackRock Municipal Bond Fund, Inc. (2025). Statement of additional information: BlackRock National Municipal Fund and BlackRock Short Duration Muni Fund. SEC Form 485BPOS, https://www.sec.gov/Archives/edgar/data/225635/000119312525248402/d56358d485bpos.htm. Filed October 23, 2025.

Brown, D. T. (2011). Callable bonds: What is special about munis? Working paper, SSRN 1925445.

Brown, D. T. and Sirmans, S. (2011). Maturity clienteles in the municipal bond market: Term premiums and the muni puzzle. Working paper, SSRN 1928129.

Cestau, D., Hollifield, B., Li, D., and Schürhoff, N. (2019). Municipal bond markets. *Annual Review of Financial Economics*, 11:65–84.

Chalmers, J. M. R. (1998). Default risk cannot explain the muni puzzle: Evidence from municipal bonds that are secured by U.S. Treasury obligations. *Review of Financial Studies*, 11(2):281–308.

Coppola, A. (2025). In safe hands: The financial and real impact of investor composition over the credit cycle. *Review of Financial Studies*, 38(8):2275–2325.

Dagostino, R. (2025). The impact of bank financing on municipalities' bond issuance and the real economy. *Journal of Financial Economics*, 166:104022.

Dammon, R. M. and Green, R. C. (1987). Tax arbitrage and the existence of equilibrium prices for financial assets. *Journal of Finance*, 42(5):1143–1166.

Domanski, D., Shin, H. S., and Sushko, V. (2017). The hunt for duration: Not waving but drowning? *IMF Economic Review*, 65(1):113–153.

Dybvig, P. H. and Ross, S. A. (1986). Tax clienteles and asset pricing. *Journal of Finance*, 41(3):751–762.

Fama, E. F. (1977). A pricing model for the municipal bond market. Unpublished manuscript, University of Chicago.

Feenberg, D. R. and Poterba, J. M. (1991). Which households own municipal bonds? Evidence from tax returns. *National Tax Journal*, 44(4):93–103.

Fleckenstein, M. and Longstaff, F. A. (2025). Do municipal bond investors pay a convenience premium to avoid taxes? *Review of Financial Studies*. Advance Access, hhaf100.

Gabaix, X., Koijen, R. S. J., Mainardi, F., Oh, S. S., and Yogo, M. (2026). Asset demand of U.S. households. Working paper, SSRN 4251972, May 9, 2026.

Garrett, D., Ordin, A., Roberts, J. W., and Suárez Serrato, J. C. (2023). Tax advantages and imperfect competition in auctions for municipal bonds. *Review of Economic Studies*, 90(2):815–851.

Green, R. C. (1993). A simple model of the taxable and tax-exempt yield curves. *Review of Financial Studies*, 6(2):233–264.

Green, R. C., Hollifield, B., and Schürhoff, N. (2007). Financial intermediation and the costs of trading in an opaque market. *Review of Financial Studies*, 20(2):275–314.

Greenwood, R., Hanson, S., and Vayanos, D. (2024). Supply and demand and the term structure of interest rates. *Annual Review of Financial Economics*, 16(1):115–151.

Greenwood, R., Hanson, S. G., and Liao, G. Y. (2018). Asset price dynamics in partially segmented markets. *Review of Financial Studies*, 31(9):3307–3343.

Greenwood, R., Hanson, S. G., Stein, J. C., and Sunderam, A. (2023). A quantity-driven theory of term premia and exchange rates. *Quarterly Journal of Economics*, 138(4):2327–2389.

Greenwood, R. and Vayanos, D. (2014). Bond supply and excess bond returns. *Review of Financial Studies*, 27(3):663–713.

Gromb, D. and Vayanos, D. (2018). The dynamics of financially constrained arbitrage. *Journal of Finance*, 73(4):1713–1750.

Guibaud, S., Nosbusch, Y., and Vayanos, D. (2013). Bond market clienteles, the yield curve, and the optimal maturity structure of government debt. *Review of Financial Studies*, 26(8):1914–1961.

Gürkaynak, R. S., Sack, B., and Wright, J. H. (2007). The U.S. Treasury yield curve: 1961 to the present. *Journal of Monetary Economics*, 54(8):2291–2304.

Haddad, V. and Sraer, D. A. (2020). The banking view of bond risk premia. *Journal of Finance*, 75(5):2465–2502.

Hanson, S. G. and Stein, J. C. (2015). Monetary policy and long-term real rates. *Journal of Financial Economics*, 115(3):429–448.

Harris, L. E. and Piwowar, M. S. (2006). Secondary trading costs in the municipal bond market. *Journal of Finance*, 61(3):1361–1397.

He, Z. and Krishnamurthy, A. (2013). Intermediary asset pricing. *American Economic Review*, 103(2):732–770.

Hendershott, P. H. and Kidwell, D. S. (1978). The impact of relative security supplies: A test with data from a regional tax-exempt bond market. *Journal of Money, Credit and Banking*, 10(3):337–347.

Hendershott, P. H. and Koch, T. W. (1980). The demand for tax-exempt securities by financial institutions. *Journal of Finance*, 35(3):717–727.

Hines, Jr., J. R. and Horton, E. (2024). Banks and tax-exempt debt arbitrage. Working Paper 32647, National Bureau of Economic Research.

Hull, J. and White, A. (1994). Numerical procedures for implementing term structure models I: Single-factor models. *Journal of Derivatives*, 2(1):7–16.

iShares Trust (2026). Statement of additional information, iShares National Muni Bond ETF and other series. SEC Form 485BPOS, https://www.sec.gov/Archives/edgar/data/1100663/000119312526275730/d124972d485bpos.htm. Filed June 18, 2026.

Kekre, R., Lenel, M., and Mainardi, F. (2024). Monetary policy, segmentation, and the term structure. Working Paper 32324, National Bureau of Economic Research.

Kidwell, D. S. and Koch, T. W. (1983). Market segmentation and the term structure of municipal yields. *Journal of Money, Credit and Banking*, 15(1):40–55.

Kim, D. H. and Wright, J. H. (2005). An arbitrage-free three-factor term structure model and the recent behavior of long-term yields and distant-horizon forward rates. Finance and Economics Discussion Series 2005-33, Board of Governors of the Federal Reserve System.

Koijen, R. S. J. and Yogo, M. (2019). A demand system approach to asset pricing. *Journal of Political Economy*, 127(4):1475–1515.

Krishnamurthy, A. and Vissing-Jorgensen, A. (2012). The aggregate demand for Treasury debt. *Journal of Political Economy*, 120(2):233–267.

Li, Y., O'Hara, M., and Zhou, X. A. (2024). Mutual fund fragility, dealer liquidity provision, and the pricing of municipal bonds. *Management Science*, 70(7):4802–4823.

Li, Z. (2026). Long rates, life insurers, and credit spreads. *Review of Financial Studies*. Forthcoming.

Livingston, M. (1982). The pricing of municipal bonds. *Journal of Financial and Quantitative Analysis*, 17(2):179–193.

Longstaff, F. A. (2011). Municipal debt and marginal tax rates: Is there a tax premium in asset prices? *Journal of Finance*, 66(3):721–751.

Mankiw, N. G. and Poterba, J. M. (1996). Stock market yields and the pricing of municipal bonds. Working Paper 5607, National Bureau of Economic Research.

Miller, M. H. (1977). Debt and taxes. *Journal of Finance*, 32(2):261–275.

Moussawi, H. (2022). Asset purchases and demand in municipal bonds. Working paper, SSRN 4017698, revised April 2025.

Moussawi, H. (2024). The efficiency of municipal bond tax shields. Working paper, SSRN 5065844, revised August 2026.

Mussa, M. L. and Kormendi, R. C. (1979). *The Taxation of Municipal Bonds: An Economic Appraisal*. Number 229 in AEI Studies in Tax Policy. American Enterprise Institute for Public Policy Research, Washington, DC.

Nuveen Municipal Trust (2026). Statement of additional information, Nuveen All-American Municipal Bond Fund and other series. SEC Form 485BPOS, https://www.sec.gov/Archives/edgar/data/1020660/000119312526322012/d113980d485bpos.htm. Filed July 29, 2026.

Pati, A. (2023). Heterogeneous investors and risk premiums in the municipal bond market. Essay in Ph.D. dissertation, Emory University.

Pirinsky, C. A. and Wang, Q. (2011). Market segmentation and the cost of capital in a domestic market: Evidence from municipal bonds. *Financial Management*, 40(2):455–481.

Poterba, J. M. (1986). Explaining the yield spread between taxable and tax-exempt bonds: The role of expected tax policy. In Rosen, H. S., editor, *Studies in State and Local Public Finance*, pages 5–52. University of Chicago Press, Chicago.

Poterba, J. M. and Ramírez Verdugo, A. (2011). Portfolio substitution and the revenue cost of the federal income tax exemption for state and local government bonds. *National Tax Journal*, 64(2, Part 2):591–613.

Schwert, M. (2017). Municipal bond liquidity and default risk. *Journal of Finance*, 72(4):1683–1722.

Sialm, C. (2009). Tax changes and asset pricing. *American Economic Review*, 99(4):1356–1383.

Siani, K. Y. (2022). Raising bond capital in segmented markets. Working paper, MIT Sloan School of Management. November 7 version.

Siriwardane, E. N., Sunderam, A., and Wallen, J. (2025). Segmented arbitrage. *Journal of Finance*, 80(5):2543–2590.

Skelton, J. L. (1983). Banks, firms and the relative pricing of tax-exempt and taxable bonds. *Journal of Financial Economics*, 12(3):343–355.

Spreen, T. L. and Gerrish, E. (2022). Taxes and tax-exempt bonds: A literature review. *Journal of Economic Surveys*, 36(4):767–808.

St. Clair, T. (2024). The supply elasticity of municipal debt: Evidence from bank-qualified bonds. *National Tax Journal*, 77(1):111–139.

Vayanos, D. and Vila, J.-L. (2021). A preferred-habitat model of the term structure of interest rates. *Econometrica*, 89(1):77–112.

# A The model: formal results and proofs

This appendix establishes the model results used in Section 4. Appendix A.1 states portfolio choice, market clearing and the price recursion for the general economy. Appendix A.2 derives participation and pricing in the two-investor affine benchmark, then verifies the illustrations with and without bond-specific risk. Appendices A.3–A.6 consider other participation motives, multiple states and investor groups, and other bond markets.

The analytical solutions allow borrowing through the cash account unless stated otherwise. Wherever cash stays nonnegative, they also solve the problem with borrowing ruled out. Appendix B solves the quantitative model with this constraint, Appendix C constructs the empirical adjustments and tests, and Appendix D supplies the tax rules and institutional evidence. Table 6 collects the model notation.

Table 6: Notation for the model and the proofs

| Symbol | Meaning |
|---|---|
| $z$, $r(z) = \max\{\bar{r} + z, 0\}$ | Latent rate state and the one-period log gross cash return |
| $\tau_i^a$, $\xi_i^a$, $\xi_i = \xi_i^C/\xi_i^M$ | Investor $i$'s tax rate and after-tax share in market $a$. In the two-market benchmark, $\xi_i$ is its relative after-tax share and $1-\xi_i$ its exemption benefit |
| $\xi_H < \xi_I$, $\xi^*$ | Benchmark after-tax shares of the high-tax household and lower-tax investor, and their tax mix $\omega_H\xi_H + \omega_I\xi_I$ with risk-tolerance shares $\omega_i$ |
| $p_T^a(z)$, $A_T^a = -\partial_z p_T^a$, $y_T^a = -p_T^a/T$ | Log price of the maturity-$T$ bond in market $a$, its rate loading (the model's duration), and its yield, with $A_0^a = 0$ |
| $C_T^a(z) = -p_T^a(z) - A_T^a(z)z$ | Level term, the constant intercept in the affine benchmark |
| $\mu_T^a(z) = \log \mathbb{E}_z[1 + R_T^a]$, $s_T^a = \partial_z \mu_T^a$ | Mean adjusted holding return, equal to its holders' required return in the Karush–Kuhn–Tucker (KKT) conditions, and its response to the state |
| $R_{T,t+1}^a$, $\mathcal{R}_{T,t+1}^a = \mu_T^a + u_{T,t+1}^a$ | Simple holding-period return and the return used in portfolio choice, with mean $\mu_T^a$ and centered log-return innovation $u_{T,t+1}^a$ |
| $\sigma_{\eta,T}^a$, $v_T^a = (\sigma_{\eta,T}^a)^2$ | Loading on bond-specific risk and its variance, zero for the one-period claim |
| $\nu_{i,T}^a$ | KKT slack: required return minus offered return, zero for a positive holding |
| $x_{i,T}^a$, $Q_T^a$, $m_i$, $\gamma_i$ | Holdings per unit of mass, total supply, investor mass, risk aversion. In the calibration, $m_i$ measures baseline portfolio assets |
| $W_i$, $W_i'$, $\widehat{W}_i'$, $\kappa_i$ | Funding per unit of mass, next-period net wealth, its discrete-time approximation, and the multiplier on the shorting constraint on cash |
| $\widehat{D}_i$, $\ell_i$ | Net rate exposure and the signed outside rate exposure. Appendix B.1 writes $h_i = -\ell_i$ |
| $\Lambda$, $\widehat{\Lambda}$ | Net-exposure difference weighted by risk aversion, $\gamma_H\widehat{D}_H - \gamma_I\widehat{D}_I$, at the household-only and equilibrium allocations |
| $\omega_{i,T}^a$, $\mathcal{U}_T^a$ | Risk-tolerance share of an active holder and the set of active holders of the maturity-$T$ bond in market $a$ |
| $g_T = \mu_T^M - \xi_H\mu_T^C$ | Gap between the municipal one-period return and the high-tax household's after-tax taxable return |
| $\tau_T^{Y,\text{impl}} = 1 - y_T^M/y_T^C$ | Model yield-implied tax, the counterpart of *ImpliedTax* in (2) |

## A.1 The general model

The quantitative model of Section 5 solves this economy with several investor groups, state municipal markets, binding shorting constraints on cash and the rate floor.

### A.1.1 Assets and portfolio choice

**Assets and investors.** Time is discrete and indexed by $t$. One aggregate rate state $z_t$ follows a stationary first-order autoregression, AR(1), and the short rate, the pretax log gross return on cash, cannot fall below zero:

$$z_{t+1} \;=\; \rho z_t \;+\; \sigma_r \varepsilon_{t+1}, \qquad r_t = r(z_t) := \max\{\bar{r} + z_t, 0\}, \qquad \rho \in (0,1), \quad \varepsilon_{t+1} \sim \mathcal{N}(0,1). \tag{30}$$

Here $\rho$ measures the persistence of the rate state and $\sigma_r$ the size of its shocks. The parameter $\bar{r}$ is the cash rate before the floor at the stationary mean $z = 0$. The floor applies to the cash return, and the state $z_{t+1}$ stays Gaussian. Appendix A.1.2 describes how the floor affects continuation prices.

Let $\mathcal{A}$ be a finite set of bond markets, indexed by $a$ and distinguished by asset class or issuer. In the two-investor benchmark, $\mathcal{A}$ contains a taxable market $C$ and a municipal market $M$. The quantitative model has 45 municipal markets, pooling the seven states without an income tax. Each market contains zero-coupon bonds indexed by their remaining maturity $T \in \{1, \ldots, T_{\max}\}$, the number of periods until the bond pays one unit. At maturity, $P^a_0 = 1$. The one-period claim of the taxable market $C$ is the cash account, with log gross return $r(z)$, so risky taxable bonds begin at $T = 2$. A bond with $T$ periods remaining has price $P^a_T(z_t)$ and holding-period return

$$1 + R^a_{T,t+1} \;=\; \frac{P^a_{T-1}(z_{t+1})}{P^a_T(z_t)} \exp\!\big(\sigma^a_{\eta,T}\, \eta^a_{T,t+1}\big). \tag{31}$$

The standard normal shocks $\eta^a_{T,t+1}$ are independent of each other and of $\varepsilon_{t+1}$, and they add bond-specific log-return variance $v^a_T := (\sigma^a_{\eta,T})^2$. Setting $v^a_T = 0$ gives the main-text return. Every market has $\sigma^a_{\eta,1} = 0$, so a one-period claim is fully riskless. Bond-specific risk represents residual variation in holding returns, such as an issue-specific change in trading demand. It makes longer bonds imperfect substitutes without adding a state variable that affects future prices. Appendix B.7 shows that normalizing the shock's mean to one leaves holding returns unchanged and barely moves the calibrated gap.

Investors belong to finitely many types, indexed by $i$. Type $i$ has mass $m_i > 0$, risk aversion $\gamma_i > 0$, funding $W_i$ and outside exposure $\ell_i$ to the rate shock from assets and liabilities beyond its bonds. It keeps the absolute after-tax share $\xi^a_i = 1 - \tau^a_i > 0$ of returns in market $a$, where $\tau^a_i$ is its tax rate on those returns. Because after-tax shares differ across both investors and markets, the same structure covers municipal and taxable bonds, state taxes on treasury and corporate interest, tax-favored accounts and cross-border withholding (Appendix A.6).

**Portfolio choice.** Investor $i$ holds bonds $x^a_{i,T,t}$ and cash $x^0_{i,t} = W_i - \sum_{a,T} x^a_{i,T,t}$, in dollars per unit of mass. Aggregate holdings are $m_i x^a_{i,T,t}$, in the same units as supply. As in Section 4.1, portfolio choice uses $\mathcal{R}^a_{T,t+1} = \mu^a_T(z_t) + u^a_{T,t+1}$, where $\mu^a_T(z_t) := \log \mathbb{E}_t[1 + R^a_{T,t+1}]$ and $u^a_{T,t+1} := \log(1 + R^a_{T,t+1}) - \mathbb{E}_t[\log(1 + R^a_{T,t+1})]$. Thus $\mu$ is the conditional mean of the adjusted return, including the Jensen adjustment, and $u$ is the centered log-return innovation. Substituting $R \approx \mathcal{R}$ for bonds and $e^r - 1 \approx r$ for cash gives approximated next-period wealth $\widehat{W}'_i$ and the mean–variance problem:[22]

$$\begin{aligned} \max_{x_i \geq 0,\, x^0_{i,t} \geq 0} \quad & \mathbb{E}_t[\widehat{W}'_i] - \frac{\gamma_i}{2} \operatorname{Var}_t(\widehat{W}'_i), \\ \text{subject to} \quad & \widehat{W}'_i = W_i + \xi^C_i r(z_t) x^0_{i,t} \\ & \qquad + \sum_{a,T} \xi^a_i x^a_{i,T,t} \mathcal{R}^a_{T,t+1} + \ell_i \sigma_r \varepsilon_{t+1}. \end{aligned} \tag{32}$$

Positions and cash are nonnegative, which rules out short sales and borrowing (Section 3.5). Funding and masses are fixed across states. A positive $\ell_i$ raises net wealth when rates rise, offsetting the rate exposure of the investor's bonds. The exposure $\ell_i$ can take either sign.

[22] Appendices C.1.4 and D.3 give the capital-income tax correction.

**Portfolio optimality and market clearing.** Collecting the cash and bond terms in approximated wealth (32) gives each bond's expected after-tax excess return over cash

$$\widetilde{\mu}^a_{i,T}(z_t) := \underbrace{\xi^a_i \mu^a_T(z_t)}_{\text{after-tax bond return}} - \underbrace{\xi^C_i r(z_t)}_{\text{after-tax cash return}} . \tag{33}$$

The bond and outside innovations have conditional mean zero, so expected approximated wealth is the return on initial funding plus the portfolio's excess returns:

$$\mathbb{E}_t[\widehat{W}'_i] = W_i\big(1 + \xi^C_i r(z_t)\big) + \sum_{a,T} x^a_{i,T,t}\widetilde{\mu}^a_{i,T}(z_t).$$

The variance of wealth is that of the after-tax portfolio innovation, the random part of (32):

$$\widetilde{\Pi}_i(t+1) = \sum_{a\in\mathcal{A}}\sum_T \xi^a_i x^a_{i,T,t} u^a_{T,t+1} + \underbrace{\ell_i \sigma_r \varepsilon_{t+1}}_{\text{liability hedge}} . \tag{34}$$

Dropping the portfolio-independent term $W_i(1+\xi^C_i r(z_t))$ and eliminating cash gives the equivalent reduced problem:

$$\max_{\substack{x_i\geq 0 \\ \sum_{a,T} x^a_{i,T,t}\leq W_i}} \underbrace{\sum_{a\in\mathcal{A}}\sum_T x^a_{i,T,t}\widetilde{\mu}^a_{i,T}(z_t)}_{\text{after-tax excess return}} - \underbrace{\frac{\gamma_i}{2}\operatorname{Var}_t\Big(\widetilde{\Pi}_i(t+1)\Big)}_{\text{whole-portfolio risk cost}} . \tag{35}$$

Let $\kappa_i(z) \geq 0$ denote the multiplier on the shorting constraint on cash, $\sum_{a,T} x^a_{i,T} \leq W_i$. The problem's first-order, or KKT, conditions are

$$\begin{aligned} x^a_{i,T}(z) > 0 \quad &\Longrightarrow \quad \widetilde{\mu}^a_{i,T}(z) = \underbrace{\gamma_i \xi^a_i \operatorname{Cov}_z\Big(u^a_{T,t+1}, \widetilde{\Pi}_i(t+1)\Big)}_{\text{marginal portfolio-risk cost}} + \underbrace{\kappa_i(z)}_{\text{shorting constraint on cash}} , \\ x^a_{i,T}(z) = 0 \quad &\Longrightarrow \quad \widetilde{\mu}^a_{i,T}(z) \leq \gamma_i \xi^a_i \operatorname{Cov}_z\Big(u^a_{T,t+1}, \widetilde{\Pi}_i(t+1)\Big) + \kappa_i(z). \end{aligned} \tag{36}$$

The multiplier satisfies complementary slackness, $\kappa_i(z)[W_i - \sum_{a,T} x^a_{i,T}(z)] = 0$.

Rate risk comes from the next-period resale price. With log price $p^a_T := \log P^a_T$ and rate loading $A^a_T := -\partial_z p^a_T$, a bond with $T$ periods remaining has rate innovation $-A^a_{T-1}(\rho z_t)\sigma_r \varepsilon_{t+1}$ to first order around the expected resale state. A one-period bond carries no rate innovation, since $A^a_0 = 0$, although its price can respond to today's state.

Every bond other than cash clears at every rate state:

$$x^a_{i,T}(z) \geq 0, \qquad \sum_i m_i\, x^a_{i,T}(z) = Q^a_T, \tag{37}$$

where $Q^a_T$ is the fixed dollar supply. An equilibrium consists of prices and holdings satisfying (31), (36) and (37). These conditions determine participation, that is, which investors hold each bond.

**The discrete-time return approximation.** With simple holding-period returns $R^a_{T,t+1}$, exact next-period net wealth is

$$W'_i = W_i + \xi^C_i x^0_i\big(e^{r(z_t)} - 1\big) + \sum_{a,T} \xi^a_i x^a_{i,T} R^a_{T,t+1} + \ell_i \sigma_r \varepsilon_{t+1}.$$

Taking logs of (31) gives the change in log price plus the bond-specific shock. Following Greenwood et al. (2024), I approximate the simple bond return by this log return plus a Jensen adjustment, and $e^r - 1$ by $r$ for cash. These substitutions give approximated wealth $\widehat{W}_i'$ in (32), with adjusted bond return

$$\mathcal{R}^a_{T,t+1} = \log(1 + R^a_{T,t+1}) + \underbrace{\log \mathbb{E}_t[1 + R^a_{T,t+1}] - \mathbb{E}_t[\log(1 + R^a_{T,t+1})]}_{\text{Jensen adjustment}}.$$

The Jensen adjustment is known at time $t$, so it shifts the mean but leaves log-return variances and covariances unchanged. Taking the conditional expectation gives $\mathbb{E}_t[\mathcal{R}^a_{T,t+1}] = \mu^a_T(z_t) = \log \mathbb{E}_t[1 + R^a_{T,t+1}]$. Subtracting this mean cancels the adjustment:

$$u^a_{T,t+1} = \mathcal{R}^a_{T,t+1} - \mu^a_T(z_t) = \log(1 + R^a_{T,t+1}) - \mathbb{E}_t[\log(1 + R^a_{T,t+1})].$$

In the affine benchmark, log returns are conditionally Gaussian, so

$$\mu^a_T(z_t) = \mathbb{E}_t[\log(1 + R^a_{T,t+1})] + \tfrac{1}{2} \operatorname{Var}_t[\log(1 + R^a_{T,t+1})].$$

This mean and the centered innovation $u$ match equations (7)–(8) of Greenwood et al. (2024). The nonlinear solver uses $\mu = \log \mathbb{E}_t[1 + R]$ directly, retaining the full Jensen adjustment even when log returns are not Gaussian.

The price recursion is exact for the chosen mean $\mu = \log \mathbb{E}_t[1 + R]$. Using this mean in the linear wealth equation approximates simple returns, whose conditional mean is $e^\mu - 1$. Appendix D.3 compares the two conventions for a deterministic return. I re-solve the one-market economy used to estimate the coefficients in Section 5, now with exactly taxed simple returns. I refit risk aversion to restore the targeted taxable yields and hold the other coefficients fixed. The thirty-year gap in that economy rises from 102.6 to 107.7 basis points.

**The cash constraint and riskless claims.** The shorting constraint on cash in (35) keeps cash nonnegative, so investors cannot borrow. Its multiplier $\kappa_i(z)$ in (36) is zero whenever the investor holds cash, since the investor can then buy more bonds without borrowing. A positive multiplier requires the investor to be fully invested, so that buying more of one bond means selling another. The multiplier raises the required pretax return on each held bond by $\kappa_i/\xi^a_i$.

Consider a riskless one-period claim in any market $a \neq C$. Its bond-specific variance $v^a_1 = (\sigma^a_{\eta,1})^2$ is zero and $p^a_0 = 0$, so its return has no covariance with portfolios or liabilities. Retaining $\kappa_i$ in (36), each eligible investor's KKT condition becomes

$$\mu^a_1 \leq \frac{\xi^C_i r + \kappa_i}{\xi^a_i}, \qquad \text{with equality for a positive holding.}$$

With positive supply, at least one investor holds the claim, so its return equals the smallest of these bounds:

$$\mu^a_1(z) = \min_i \frac{\xi^C_i r(z) + \kappa_i(z)}{\xi^a_i}. \tag{38}$$

The price recursion (41) then gives $p^a_1 = -\mu^a_1$, hence $y^a_1 = \mu^a_1$.

### A.1.2 Bond prices and required returns

Taking expectations of the gross holding return (31), including bond-specific risk with variance $v^a_T = (\sigma^a_{\eta,T})^2$, gives

$$P^a_T(z) = e^{-\mu^a_T(z) + v^a_T/2}\, \mathbb{E}_z\big[P^a_{T-1}(z')\big], \qquad P^a_0 = 1. \tag{39}$$

A higher required return lowers today's price for a given expected resale value. Iterating from the maturity payment links yields to the one-period returns along the bond's life. The factor $e^{v^a_T/2}$ accounts for the mean of the bond-specific shock (see the exact recursion below). Its yield effect is small (Appendix B.7).

With fixed supplies and funding, I seek an equilibrium log-price function for each market and maturity:

$$p^a_{T,t} = p^a_T(z_t), \qquad p^a_0(z) \equiv 0, \qquad a \in \mathcal{A}. \tag{40}$$

Let $Z' = \rho z + \sigma_r \varepsilon'$. Since $\mu^a_T$ is the log of the bond's conditional expected gross return (31), solving for today's log price gives

$$p^a_T(z) = \underbrace{\log \mathbb{E}_z\left[\exp\{p^a_{T-1}(Z')\}\right]}_{\text{log expected resale price}} + \underbrace{\tfrac{1}{2}(\sigma^a_{\eta,T})^2}_{\text{bond-specific payoff term}} - \underbrace{\mu^a_T(z)}_{\text{one-period return}}. \tag{41}$$

Equation (41) discounts the expected resale payoff at the equilibrium holding-period return, so a higher return lowers today's price. This recursion links adjacent maturities, while the portfolio conditions (36) link markets through investors' holdings. At every state, prices and portfolios jointly solve these equations and clearing (37). Appendix B gives the numerical implementation. Equations (45)–(46) in Appendix A.1.3 state the portfolio system for a slack shorting constraint on cash.

The rate floor, changes in participation and binding shorting constraints on cash make prices nonlinear in the state. Resale carries these nonlinearities across maturities, because today's continuation value averages future prices over states with different participation and cash constraints.

Iterating (41) from $p^a_0 = 0$ expresses the price through the one-period returns the bond earns as it ages through maturities $T, T-1, \ldots, 1$:

$$P^a_T(z_t) = e^{p^a_T(z_t)} = \mathbb{E}_t\Big[\exp\Big\{-\sum_{j=0}^{T-1}\Big(\mu^a_{T-j}(z_{t+j}) - \tfrac{1}{2}\big(\sigma^a_{\eta,T-j}\big)^2\Big)\Big\}\Big]. \tag{42}$$

Taking the negative log of (42) and dividing by $T$ gives the yield $y^a_T = -p^a_T/T$. In the affine benchmark, this becomes the explicit average of expected one-period returns less convexity in (56). In the general economy, I summarize the yield's dependence on the state by writing the log price in loading form,

$$p^a_T(z) = -A^a_T(z)\, z - C^a_T(z), \qquad A^a_T(z) := -\frac{\partial p^a_T(z)}{\partial z}, \tag{43}$$

so that $P^a_T = \exp\{-A^a_T z - C^a_T\}$ and, at differentiability points,

$$y^a_T(z) = \frac{A^a_T(z) z + C^a_T(z)}{T}, \qquad \frac{\partial y^a_T(z)}{\partial z} = \frac{A^a_T(z)}{T}.$$

The level term $C^a_T$ includes accumulated one-period returns and convexity, and the mean-state yield is $C^a_T(0)/T$. Outside the affine benchmark, both $A^a_T$ and $C^a_T$ can vary with the state. The identity $\partial_z C^a_T = -z\,\partial_z A^a_T$ preserves the yield derivative above.

The loading recursion links the loading $A^a_T$ to the one-period return response $s^a_T = \partial_z \mu^a_T$ of equation (54). Above the rate floor, cash has response one. By (38), a riskless one-period claim in another market has response $\xi^C_i/\xi^a_i$ for its active holder $i$ when the shorting constraint on cash is slack. At longer maturities the response also includes the change in the holders' risk premia. If continuation prices are locally absolutely continuous and the Gaussian exponential moment and its state derivative have an integrable envelope, differentiating under the expectation in (41) gives:

$$A^a_T(z) = \underbrace{s^a_T(z)}_{\text{current return response}} + \underbrace{\rho\,\frac{\mathbb{E}_z\left[\exp\{p^a_{T-1}(Z')\}\, A^a_{T-1}(Z')\right]}{\mathbb{E}_z\left[\exp\{p^a_{T-1}(Z')\}\right]}}_{\text{continuation price response}}. \tag{44}$$

The fraction averages the continuation loading, giving more weight to states where the resale bond is more valuable. In the affine benchmark, (44) reduces to the loading recursion (55).

### A.1.3 Market clearing and the riskless short yield

The one-period taxable claim is riskless cash. An allocation's support is the set of investor–bond pairs with positive holdings. A participation branch fixes the support, and a bond's active set is its holders. Affine results use floorless rates and a fixed participation branch, whereas the numerical model retains the floor and the shorting constraint on cash.

**One KKT system for all markets.** Index risky market–maturity coordinates by $j = (a, T)$ and write $v_j = (\sigma^a_{\eta,T})^2$, $K_i = \operatorname{diag}(\xi^a_i)$ and the relative after-tax share $\xi_{ij} = \xi^C_i/\xi^a_i$, collected in the vector $\boldsymbol{\xi}_i$. In the normalized two-market benchmark, $\xi_{ij} = \xi_i$ at municipal coordinates and $\xi_{ij} = 1$ at taxable coordinates. With resale log price $p_{j,-} = p^a_{T-1}$, the exact covariance and liability credit are

$$\Sigma_{jk}(z) = \operatorname{Cov}_z(p_{j,-}(Z'), p_{k,-}(Z')) + \mathbf{1}_{\{j=k\}} v_j, \qquad c^L_{ij}(z) = -\operatorname{Cov}_z(p_{j,-}(Z'), \ell_i \sigma_r \varepsilon'). \tag{45}$$

With a slack shorting constraint on cash, dividing each first-order condition by the positive after-tax share of its market gives

$$\begin{gathered} \nu_{ij} := \xi_{ij} r + \gamma_i [\Sigma K_i x_i - c^L_i]_j - \mu_j \geq 0, \qquad x_{ij} \geq 0, \qquad x_{ij}\nu_{ij} = 0, \\ \sum_i m_i x_{ij} = Q_j. \end{gathered} \tag{46}$$

Concavity of (35) makes these conditions sufficient for investor optimality. With future resale prices held fixed, return covariances and liability credits are given, and current required returns must adjust until investors' optimal holdings match bond supplies. The next proposition establishes existence for this clearing step and gives sufficient conditions for uniqueness.

**Proposition A.1** (Market clearing with given continuation prices)**.** *Fix the state and continuation prices, with finite $\Sigma \succeq 0$, finite liability credits and at least one eligible investor for every positive-supply coordinate.*

*(i) The KKT-and-clearing system (46) has a current-state solution.*

*(ii) If, in addition, $\Sigma = V + \sigma_r^2 AA'$ with $V$ positive diagonal and a loading vector $A \gg 0$ on the risky coordinates, after-tax shares are positive and supported supplies are positive, the current-state clearing returns and risky holdings are unique, even with endogenous support changes. Deterministic one-period claims are treated separately. Payoff-equivalent holdings in tied boundary claims need not be unique.*

*Neither part establishes the dynamic fixed point of continuation prices.*

**Proof.** (i) Let $X = \{x \geq 0 : \sum_i m_i x_i = Q\}$ be the set of nonnegative allocations that clear, with inner product $\langle x, y\rangle_m = \sum_i m_i x_i' y_i$. Positive masses and finite supplies make $X$ nonempty, compact and convex. For $\Upsilon_i(x_i) = \boldsymbol{\xi}_i r + \gamma_i(\Sigma K_i x_i - c^L_i)$, investor $i$'s required return on each coordinate, Brouwer gives a fixed point of the continuous map $x \mapsto \operatorname{Proj}_X(x - \Upsilon(x))$. The projection characterization implies

$$\sum_i m_i \langle \Upsilon_i(x_i), y_i - x_i\rangle \geq 0 \qquad (y \in X).$$

Transfers from positive holders to other eligible investors show that each positive holder minimizes $\Upsilon_{ij}$. Thus $\mu_j = \min_i \Upsilon_{ij}$ and $\nu_{ij} = \Upsilon_{ij} - \mu_j$ satisfy (46). Conversely, complementarity implies the displayed inequality for every feasible transfer. Predetermined trading exclusions restrict eligibility. (ii) Strict concavity makes individual demand $x_i(\mu)$ continuous, Lipschitz and piecewise affine. On a region where investor $i$ holds exactly the coordinates in a set $B_i$, its derivative is

$$J_i = \gamma_i^{-1} E_{B_i} K^{-1}_{i,B_i} \Sigma^{-1}_{B_i B_i} E'_{B_i},$$

where $E_B$ embeds the held coordinates. Sherman–Morrison gives

$$\Sigma^{-1}_{BB} = V_B^{-1} - \frac{\sigma_r^2 V_B^{-1} A_B A_B' V_B^{-1}}{1 + \sigma_r^2 A_B' V_B^{-1} A_B}, \qquad \Sigma^{-1}_{BB} A_B = \frac{V_B^{-1} A_B}{1 + \sigma_r^2 A_B' V_B^{-1} A_B} \gg 0.$$

Every selection of $J = \sum_i m_i J_i$ therefore has nonpositive off-diagonal entries and $JA \geq 0$, strictly on held coordinates. If $\mu^0, \mu^1$ both cleared, piecewise-affine continuity would give $0 = \overline{J}(\mu^1 - \mu^0)$, with $\overline{J}$ the segment-average selection matrix. Positive supply at $\mu^0$ and demand continuity imply positive holdings on an initial segment for each asset. Finitely many selections then give $(\overline{J}A)_j > 0$ everywhere. Hence $\overline{J}\,\mathrm{diag}(A)$ is strictly diagonally dominant and nonsingular, so returns coincide, and strict concavity gives identical risky holdings. On support-boundary segments, adjacent selections have the same directional derivative. □

**Regularity for derivatives.** Derivatives below are taken at regular states, at which the cash rate and price moments are continuously differentiable, complementarity is strict on risky coordinates and the active KKT-and-clearing Jacobian is nonsingular. The implicit-function theorem then gives a locally fixed participation branch and differentiable holdings and returns. Fixed supplies imply, for each bond's active set $\mathcal{U}_j$,

$$\sum_{i \in \mathcal{U}_j} m_i \frac{\partial x_{ij}}{\partial z} = 0. \tag{47}$$

Each investor's strict concavity does not imply nonsingularity, because with heterogeneous after-tax shares the normalized operator contains $\Sigma K_i$, which need not be symmetric or strongly monotone.

**The one-period boundary.** Specializing the riskless-claim return (38) to the normalized municipal–taxable benchmark gives the one-period yields (12) and tax pass-throughs (20).

**Proposition A.2** (Deterministic one-period boundary)**.** *In the normalized municipal–taxable benchmark with slack shorting constraints on cash, suppose the one-period municipal claim has $v_1^M = 0$ and positive supply, and $r(z) > 0$. Suppose the high-tax household has the minimal after-tax share, $\xi_H = \min_i \xi_i$. Then only investors with the minimal after-tax share can hold it, and*

$$y_1^M = \mu_1^M = \xi_H r(z), \qquad y_1^C = r(z), \qquad \tau_1^{Y,\mathrm{impl}} = 1 - \xi_H. \tag{48}$$

*If the high-tax household is the unique after-tax share minimizer, its tax passes through one-for-one and the tax of an investor with a strictly higher after-tax share has zero effect at this boundary. At a tie the household derivative requires perturbing all tied minimal after-tax shares together. At $r = 0$ the yield ratio is undefined.*

**Proof.** With $\xi_i^M = 1$, $\xi_i^C = \xi_i$ and $\kappa_i = 0$, the riskless-claim KKT condition of Appendix A.1.1 is $\mu_1^M \leq \xi_i r$, with equality for positive holdings. Since $r > 0$, only investors with the minimal after-tax share can satisfy it with equality. Positive supply gives $\mu_1^M = \min_i \xi_i r = \xi_H r$ and $p_1^M = -\mu_1^M$. Taxable cash has $p_1^C = -r$. Differentiating the yield ratio proves the tax statements. □

## A.2 The two-investor affine benchmark

This section specializes the general model to two investors and two jointly priced bond markets. It derives the participation and pricing results behind Sections 4.3–4.6, then verifies the analytical illustrations.

### A.2.1 Assumptions and portfolio conditions

The two-investor results assume $0 < \rho < 1$, $\sigma_r > 0$, $m_i, \gamma_i > 0$, $0 < \xi_H < \xi_I$ and positive traded supplies.

The high-tax household $H$ and the lower-tax investor $I$ choose among municipal maturities $T = 1, \ldots, T_{\max}$, taxable maturities $T = 2, \ldots, T_{\max}$ and cash, the one-period taxable claim. Bond supplies $Q_T^a$ are fixed. The benchmark allows bond-specific risk and different risk aversions $\gamma_i$, and keeps the relative after-tax share $\xi_i = \xi_i^C / \xi_i^M$ of Section 4.2, with $\xi_H < \xi_I$. Payoffs are in units of after-tax municipal income, so taxable returns and cash income are multiplied by $\xi_i$, with $\gamma_i$ and $\ell_i$ normalized to match. I fix a participation branch, which specifies which investors hold each bond, solve its portfolio conditions and market clearing, and check that no investor benefits from changing its holdings. Closed forms use three restrictions:

(i) the shorting constraint on cash is dropped;

(ii) the floored cash rate is replaced by the linear rule $r(z) = \bar{r} + z$; and

(iii) prices are computed as if the same participation branch applied at every state, which makes log prices affine, and the branch is verified at the states studied.

On a fixed branch, holdings, one-period returns and log prices are affine in $z$. The price coefficients solve a fixed-point problem. The construction below states the feasibility and uniqueness conditions. The quantitative model lets participation change across states and keeps the shorting constraint on cash and the rate floor.

Log prices then have constant coefficients,

$$p_T^{a,\mathrm{aff}}(z) = -A_T^{a,\mathrm{aff}} z - C_T^{a,\mathrm{aff}}, \tag{49}$$

and one-period returns are affine as well, $\mu_T^{a,\mathrm{aff}}(z) = \mu_T^{a,\mathrm{aff}}(0) + s_T^{a,\mathrm{aff}} z$, so $s_T^{a,\mathrm{aff}}$ is the state coefficient of the one-period return and $A_T^{a,\mathrm{aff}}$ that of the negative log price. The affine recursions (55) and the yield average (56) of Appendix A.2.2 hold exactly. Below and in the proofs, loadings, level terms and responses refer to this affine branch, marked by the superscript "aff" where shown.

Investor $i$'s net rate exposure, the after-tax exposure of its whole balance sheet to the rate shock, is

$$\widehat{D}_i(z) := \underbrace{\sum_T x_{i,T}^M(z) A_{T-1}^M + \xi_i \sum_T x_{i,T}^C(z) A_{T-1}^C}_{\text{tax-weighted asset exposure}} - \underbrace{\ell_i}_{\text{outside rate exposure}}. \tag{50}$$

In these payoff units, municipal positions count at their loadings and taxable positions at $\xi_i$ times their loadings. With today's positions fixed, $\widehat{D}_i = -\partial \widehat{W}_i'/\partial z_{t+1}$, as in (10). It is dollar exposure before division by wealth. The rate component of the wealth innovation is $-\widehat{D}_i \sigma_r \varepsilon_{t+1}$, so a larger liability hedge $\ell_i$ lowers this exposure. A bond's rate-risk compensation equals its loading times the investor's unit cost of rate risk, $\gamma_i \sigma_r^2 \widehat{D}_i$. Loadings are constant on an affine branch, so changes in $\widehat{D}_i$ across rate states reflect trading alone.

Without bond-specific risk, joint holdings require the investors' difference in rate-risk compensation to equal their difference in exemption benefits. A small bond-specific variance adds a position-dependent cost and can support joint holdings over several maturities.

For the riskless one-period claim, a slack shorting constraint on cash gives $\mu_1^M = \xi_H r$ whenever $r > 0$ (Proposition A.2). The household, which keeps the smallest share of taxable interest, then holds the riskless municipal claim. At longer maturities, participation determines whose taxes and risk costs shape equilibrium returns. Proposition A.1 gives the uniqueness conditions for current-state clearing, and Appendix A.2.4 those for the loading solution.

The main-text illustration shares only municipal maturity 23 (Appendix A.2.7). Adding small bond-specific risk gives overlapping holdings from maturity 17 (Appendix A.2.8). Both appendices report parameters and check holdings and pricing. These are mechanism illustrations, so their level panels omit the vertical scale.

**Constructing affine solutions.** The construction fixes a participation branch, sets inactive positions to zero, sets $r(z) = \bar{r} + z$ and uses affine continuation prices on $\mathbb{R}$. It solves the active equalities jointly across maturities, retaining whole-portfolio risk, and then checks nonnegative holdings, inactive slacks and, for finite budgets, $W_i - \sum_j x_{ij}(z) \geq 0$.

For affine candidates $p_T^a(z) = -A_T^a z - C_T^a$, Gaussian integration gives exactly

$$\log \mathbb{E}_z e^{p_{T-1}^a(Z')} = -\rho A_{T-1}^a z - C_{T-1}^a + \tfrac{1}{2}\sigma_r^2 (A_{T-1}^a)^2. \tag{51}$$

The current-state covariance $\Sigma = \sigma_r^2 AA' + \mathrm{diag}(v)$ and liability credit $c_i^L = \ell_i \sigma_r^2 A$ are constant, where $A$ collects the predecessor loadings. If the active KKT-and-clearing matrix is nonsingular, holdings and $\mu$ are affine in $z$, and the price recursion returns affine prices with updated loadings $\rho A_{T-1}^a + s_T^a$. At a coefficient

fixed point, all active equalities, clearing equations and exact continuation equations hold identically, and, with $A_0^a = 0$,

$$A_T^a = \rho A_{T-1}^a + s_T^a = \sum_{j=1}^{T} \rho^{T-j} s_j^a, \qquad s_T^a = \frac{\partial \mu_T^a}{\partial z}. \tag{52}$$

With affine holdings and slacks, the states where the KKT conditions hold form an interval, possibly empty, given globally affine continuation prices. Gaussian transitions leave any proper interval, so interval feasibility is not global equilibrium feasibility. For example, at the household-only allocation the lower-tax investor's slack on the one-period municipal claim, $(\xi_I - \xi_H)(\bar{r} + z)$, turns negative at negative rates. The cash floor changes continuation values.

For a fixed participation branch, a continuous coefficient update on a nonempty compact convex invariant set has a fixed point by Brouwer. A contraction gives uniqueness within that set. Nonsingularity, invariance and KKT feasibility with the cash floor restored must be checked separately. Appendices A.2.7 and A.2.8 verify the reported analytical illustrations. These checks supply no global existence argument.

### A.2.2 Affine price loadings and yields

Write $p_T^a = \log P_T^a$ for the log price and $y_T^a = -p_T^a/T$ for the yield. Under the affine assumptions above,

$$\log P_T^a(z) = -A_T^a z - C_T^a, \qquad y_T^a(z) = \frac{A_T^a z + C_T^a}{T}. \tag{53}$$

The coefficient $A_T^a$, the bond's rate loading, is the proportional fall in its price per unit rise in the rate state and the model's counterpart of duration. Thus $A_T^a/T$ is the yield's response to the rate state and $C_T^a/T$ is its level at the mean state. The response of the one-period return to the rate state is

$$s_T^a(z) := \frac{\partial \mu_T^a(z)}{\partial z}. \tag{54}$$

Substituting affine prices into the resale equation and matching its state and constant terms gives

$$A_T^a = \underbrace{\rho A_{T-1}^a}_{\text{inherited loading}} + \underbrace{s_T^a}_{\text{current response}}, \qquad C_T^a = C_{T-1}^a + \mu_T^a(0) - \tfrac{1}{2}\left[\sigma_r^2 \left(A_{T-1}^a\right)^2 + \left(\sigma_{\eta,T}^a\right)^2\right]. \tag{55}$$

The first relation adds today's return response to the loading inherited through the expected resale price. The second accumulates the mean-state one-period returns less half-variance terms. Starting from $A_0^a = C_0^a = 0$, the loading $A_T^a = \sum_{j=1}^{T} \rho^{T-j} s_j^a$ accumulates the responses, weighted by rate persistence. With a riskless one-period claim, $\sigma_{\eta,1}^a = 0$, the same recursion gives the yield as the average of the one-period returns along the bond's life less convexity:

$$y_T^a(z_t) = \underbrace{\frac{1}{T}\sum_{j=0}^{T-1} \mathbb{E}_t\left[\mu_{T-j}^a(z_{t+j})\right]}_{\text{one-period returns as the bond ages}} - \underbrace{\frac{1}{2T}\sum_{k=2}^{T}\left[\sigma_r^2\left(A_{k-1}^a\right)^2 + \left(\sigma_{\eta,k}^a\right)^2\right]}_{\text{convexity}}. \tag{56}$$

The formulas allow zero or positive bond-specific variance. The main-text analytical illustration sets both markets' bond-specific variances to zero. The extension in Appendix A.2.8 gives risky municipal bonds positive variance and sets taxable bond-specific variance to zero.

### A.2.3 Participation across maturities

The comparison below asks whether a small municipal purchase would benefit the lower-tax investor when the household initially holds all municipal bonds. This household-only allocation, marked by the superscript hh, assigns all risky taxable bonds to the lower-tax investor. The lower-tax investor's net rate exposure already includes its taxable portfolio and the offset from any positive outside liability exposure.

Proposition A.3 sets bond-specific risk to zero, $v_T^M = 0$, as in the equilibrium comparison of Section 4.3. Here, I evaluate the comparison at household-only holdings to identify a first-order purchase incentive. Theorem A.1 adds positive bond-specific risk at the equilibrium allocation. In the zero-variance case, the high-tax household's active KKT equality and the lower-tax investor's condition for a nonnegative first-order purchase incentive are

$$\underbrace{\mu_T^M - \xi_H r}_{\text{offered to the household}} = \underbrace{\gamma_H \sigma_r^2 A_{T-1}^M \widehat{D}_H^{\text{hh}}}_{\text{household marginal risk cost}}, \tag{57}$$

$$\underbrace{\mu_T^M - \xi_I r}_{\text{offered to the lower-tax investor}} \geq \underbrace{\gamma_I \sigma_r^2 A_{T-1}^M \widehat{D}_I^{\text{hh}}}_{\text{lower-tax investor's marginal risk cost}}. \tag{58}$$

If the second inequality fails, an initial purchase is unprofitable at this allocation. Substituting the household's equality cancels the common return $\mu_T^M$ and leaves one condition:

$$\underbrace{\sigma_r^2 \, A_{T-1}^M \, \Lambda}_{\text{duration benefit (grows with maturity)}} \geq \underbrace{(\xi_I - \xi_H)\, r}_{\text{difference in exemption benefits (flat in maturity)}}, \quad \Lambda := \gamma_H \widehat{D}_H^{\text{hh}} - \gamma_I \widehat{D}_I^{\text{hh}}. \tag{59}$$

The duration benefit on the left is the lower-tax investor's saving in rate-risk cost on the bond. It equals the bond's loading times $\sigma_r^2 \Lambda$, the difference between the two investors' unit costs of rate risk. Assumption A.1 below requires $\Lambda > 0$: risk-weighted net rate exposure is larger for the household at this allocation. At the household-only allocation, clearing fixes every holding, so $s_T^M = \xi_H$ and (55) gives $A_{T-1}^M = \xi_H \sum_{j=0}^{T-2} \rho^j$. Starting from $A_0^M = 0$, the loading rises strictly with maturity, so the duration benefit grows from zero and crosses the difference in exemption benefits at most once.

**Assumption A.1** (Household duration burden at the household-only allocation)**.** *At the household-only allocation, where the high-tax household holds all municipal bonds and the lower-tax investor all taxable bonds,*

$$\widehat{D}_H^{\text{hh}} = \sum_T A_{T-1}^M Q_T^M / m_H - \ell_H, \quad \widehat{D}_I^{\text{hh}} = \xi_I \sum_T A_{T-1}^C Q_T^C / m_I - \ell_I, \quad \Lambda := \gamma_H \widehat{D}_H^{\text{hh}} - \gamma_I \widehat{D}_I^{\text{hh}} > 0.$$

Under the assumption, the duration benefit $\sigma_r^2 A_{T-1}^M \Lambda$ is nonnegative and grows with the bond's loading.

**Proposition A.3** (Participation from the long end)**.** *In the two-investor affine benchmark without bond-specific risk, suppose $\Lambda > 0$ at the household-only allocation and $r_t > 0$. At that allocation, the maturities with a nonnegative first-order purchase incentive for the lower-tax investor form the upper interval $\{T : T \geq T^*(z_t)\}$, where*

$$T^*(z_t) \;=\; \min\Big\{T \geq 1 :\; A_{T-1}^M \geq \frac{(\xi_I - \xi_H)\, r_t}{\sigma_r^2 \, \Lambda}\Big\}, \tag{60}$$

*with $\min \varnothing = +\infty$. A strict inequality makes participation strictly profitable, and equality need not imply a positive holding.*

**Proof.** Under the proposition's zero-variance assumption, the household holds $Q_T^M / m_H$ of each municipal bond, and its active municipal equality sets

$$\mu_T^{M,\text{hh}} = \xi_H r + \gamma_H \sigma_r^2 A_{T-1}^M \widehat{D}_H^{\text{hh}}.$$

At zero holdings, the lower-tax investor requires $\xi_I r + \gamma_I \sigma_r^2 A_{T-1}^M \widehat{D}_I^{\text{hh}}$. Subtracting the offered return gives its KKT slack:

$$\nu_{I,T}^{\text{hh}} = (\xi_I - \xi_H) r - \sigma_r^2 A_{T-1}^M \Lambda. \tag{61}$$

The slack is nonpositive exactly when (59) holds. Sole-holder quantities are state invariant, giving $s_T^M = \xi_H$ and strictly increasing $A_T^M = \xi_H \sum_{q=0}^{T-1} \rho^q$. Since $\Lambda > 0$, the slack decreases strictly with maturity, proving the upper-interval result. Zero slack gives indifference at a zero position, while negative slack makes an initial purchase strictly profitable. □

Taxes, balance sheets and loadings set the threshold $T^*$, and no investor is assigned a maturity habitat. At fixed benchmark loadings, a larger difference in exemption benefits raises the threshold, and a larger $\Lambda$ lowers it. With positive bond-specific risk, the household's required return also includes compensation for the bond-specific risk of its whole holding. A new buyer bears no bond-specific risk cost at a zero position. This additional participation incentive appears explicitly in the equilibrium index (62) of Theorem A.1.

In equilibrium, prices and both portfolios are re-solved. With positive bond-specific risk, Theorem A.1 gives conditions under which the lower-tax investor's municipal holdings form an upper maturity interval. It uses the equilibrium duration gap $\widehat{\Lambda} := \gamma_H \widehat{D}_H - \gamma_I \widehat{D}_I$. Section 4.3 states the comparison without bond-specific risk in (13). The theorem adds the bond-specific risk cost to that comparison. From here on, $T^*$ denotes the start of the lower-tax investor's holdings.

The joint-ownership formulas below consider a branch with positive bond-specific risk: the household alone holds municipal bonds below $T^*$, and both investors hold them from $T^*$ on. This holding pattern is an additional branch assumption, not a consequence of the participation cutoff alone. The main-text illustration instead sets bond-specific risk to zero and has one shared maturity, above which only the lower-tax investor holds municipal bonds (Appendix A.2.7). The calibrated economy allows overlapping holdings by several investor groups (Section 5).

**Theorem A.1** (Upper-interval participation by the lower-tax investor on an affine branch)**.** *At a regular state of a two-investor affine solution, suppose $v_T^M > 0$ on the risky maturities considered. Define the duration gap at the equilibrium allocation, $\widehat{\Lambda} = \gamma_H \widehat{D}_H - \gamma_I \widehat{D}_I$, and*

$$\Psi_T = (\xi_I - \xi_H) r - \sigma_r^2 A_{T-1}^M \widehat{\Lambda} - \gamma_H v_T^M Q_T^M / m_H. \tag{62}$$

*Then $x_{I,T}^M > 0$ if and only if $\Psi_T < 0$. If $\widehat{\Lambda} > 0$ and the predecessor loadings and $v_T^M Q_T^M$ are nondecreasing, the lower-tax investor's positive support is an upper interval, possibly empty. This identifies the cutoff within a given equilibrium and does not establish uniqueness across equilibria.*

**Proof.** If the lower-tax investor is inactive, $x_{H,T}^M = Q_T^M / m_H$. Substituting the household equality into the lower-tax investor's inequality gives $\Psi_T = \nu_{I,T} \geq 0$. If both investors hold the bond, subtracting their equalities and using clearing gives

$$\Psi_T = -v_T^M \left(\gamma_I + \gamma_H m_I / m_H\right) x_{I,T}^M < 0.$$

If the lower-tax investor alone is active, its equality and the household inequality imply $(\xi_I - \xi_H) r - \sigma_r^2 A_{T-1}^M \widehat{\Lambda} \leq -\gamma_I v_T^M Q_T^M / m_I < 0$, which also gives $\Psi_T < 0$. These cases prove the equivalence. Finally, adjacent indices differ by

$$\Psi_{T+1} - \Psi_T = -\sigma_r^2 (A_T^M - A_{T-1}^M) \widehat{\Lambda} - \frac{\gamma_H}{m_H} (v_{T+1}^M Q_{T+1}^M - v_T^M Q_T^M) \leq 0.$$

With $\widehat{\Lambda} > 0$ and nondecreasing predecessor loadings and $v_T^M Q_T^M$, the index is nonincreasing in maturity, and the first maturity with a strictly negative index is the participation threshold $T^*$. Equality is a zero-holding boundary. □

**Joint holdings.** The participation cutoff marks the start of lower-tax holdings, not necessarily the end of household holdings. Bond-specific risk adds a position-dependent cost, allowing both investors' portfolio conditions to bind over a range of maturities. For candidate net rate exposures, define $N_T = (\xi_I - \xi_H) r - \sigma_r^2 A_{T-1}^M \widehat{\Lambda}$, the difference in exemption benefits less the duration benefit. For $v_T := v_T^M > 0$, subtracting the two active KKT equalities and imposing clearing gives

$$x_{H,T}^M = \frac{\gamma_I Q_T^M + m_I N_T / v_T}{\gamma_H m_I + \gamma_I m_H}, \qquad x_{I,T}^M = \frac{\gamma_H Q_T^M - m_H N_T / v_T}{\gamma_H m_I + \gamma_I m_H}. \tag{63}$$

Both holdings are nonnegative exactly when $-\gamma_I Q_T^M / m_I \leq N_T / v_T \leq \gamma_H Q_T^M / m_H$. Strict inequalities give positive holdings for both investors. Each maturity's holdings depend on the net rate exposures, which are common to all maturities and must match the resulting portfolios.

### A.2.4 Return-gap decomposition and rate loadings

This subsection gives the formal results behind Sections 4.4 and 4.5. It first splits the one-period return gap into three channels: the exemption benefit, the duration cost and the rate loading. It also shows where compensation for bond-specific risk enters when that risk is present. It then derives the tax-mix contribution and the municipal and taxable portfolio responses within the rate-loading channel.

**The three channels of the one-period return gap.** Fix maturity $T$, and let the active set $\mathcal{U}_T^M \subseteq \{H, I\}$ be the investors who hold that municipal bond. In the affine benchmark, an active holder's portfolio condition (36) becomes the main-text condition (11) with bond-specific risk added:

$$\underbrace{\mu_T^M - \xi_i r}_{\text{after-tax excess return offered to } i} = \gamma_i \Bigg[ \underbrace{\sigma_r^2 A_{T-1}^M \widehat{D}_i}_{\text{whole-portfolio rate risk}} + \underbrace{(\sigma_{\eta,T}^M)^2 x_{i,T}^M}_{\text{bond-specific risk}} \Bigg]. \tag{64}$$

On the branches studied here the lower-tax investor alone holds the risky taxable bonds. Dividing its taxable condition by its after-tax share gives, for $T \geq 2$,

$$\mu_T^C - r = \gamma_I \big[ \sigma_r^2 A_{T-1}^C \widehat{D}_I + \xi_I (\sigma_{\eta,T}^C)^2 x_{I,T}^C \big], \tag{65}$$

and $\mu_1^C = r$ for cash. Subtracting $\xi_H$ times the taxable equality from a municipal holder's equality separates the one-period return gap $g_T = \mu_T^M - \xi_H \mu_T^C$ into the three channels plus bond-specific compensation. Proposition 1 states this decomposition, the formal counterpart of (15) in Section 4.4.

**Proposition 1** (Three channels of the one-period return gap)**.** *In the two-investor affine benchmark, consider a state at which the lower-tax investor alone holds the risky taxable bonds. Let $i$ denote a holder of the maturity-$T$ municipal bond, with $i = I$ whenever the lower-tax investor holds it. Then*

$$\begin{aligned} g_T = \; & \underbrace{(\xi_i - \xi_H) r}_{\textit{exemption benefit}} + \underbrace{\sigma_r^2 A_{T-1}^M \big( \gamma_i \widehat{D}_i - \gamma_I \widehat{D}_I \big)}_{\textit{duration cost}} + \underbrace{\gamma_I \sigma_r^2 \widehat{D}_I \big( A_{T-1}^M - \xi_H A_{T-1}^C \big)}_{\textit{rate loading}} \\ & + \underbrace{\gamma_i (\sigma_{\eta,T}^M)^2 x_{i,T}^M - \xi_H \gamma_I \xi_I (\sigma_{\eta,T}^C)^2 x_{I,T}^C}_{\textit{bond-specific compensation}}. \end{aligned} \tag{66}$$

*At $T = 1$ the duration-cost, rate-loading and bond-specific contributions are zero, and so is the exemption-benefit contribution when $r > 0$. Where the household holds the bond alone, the exemption-benefit contribution is zero and the duration-cost contribution is at most $(\xi_I - \xi_H) r$. Where the lower-tax investor holds it, the exemption-benefit contribution is $(\xi_I - \xi_H) r$ and the duration-cost contribution is zero.*

**Proof.** Investor $i$'s active municipal equality (64) gives $\mu_T^M = \xi_i r + \gamma_i [\sigma_r^2 A_{T-1}^M \widehat{D}_i + (\sigma_{\eta,T}^M)^2 x_{i,T}^M]$. The lower-tax investor's taxable equality (65) gives $\mu_T^C = r + \gamma_I [\sigma_r^2 A_{T-1}^C \widehat{D}_I + \xi_I (\sigma_{\eta,T}^C)^2 x_{I,T}^C]$, which also covers cash at $T = 1$ because $A_0^C = 0$ and $\sigma_{\eta,1}^C = 0$. Subtracting $\xi_H$ times the second equality from the first, then adding and subtracting $\gamma_I \sigma_r^2 A_{T-1}^M \widehat{D}_I$, gives the three-channel identity (66).

At $T = 1$, $A_0^a = 0$ and there is no bond-specific risk, so the duration-cost, rate-loading and bond-specific terms vanish. When $r > 0$, the household alone holds the claim (Proposition A.2), so $i = H$ and the exemption-benefit term also vanishes.

If the household alone holds maturity $T$, then $i = H$ and $x_{H,T}^M = Q_T^M / m_H$. The lower-tax investor's inactive condition gives $\Psi_T \geq 0$ in (62), with $v_T^M = (\sigma_{\eta,T}^M)^2$, or $\sigma_r^2 A_{T-1}^M \widehat{\Lambda} + \gamma_H (\sigma_{\eta,T}^M)^2 Q_T^M / m_H \leq (\xi_I - \xi_H) r$. The duration-cost contribution is the first term on the left and is therefore at most $(\xi_I - \xi_H) r$. If the lower-tax investor holds the bond, the identity with $i = I$ has a zero duration-cost contribution and an exemption-benefit contribution of $(\xi_I - \xi_H) r$. □

Below the participation threshold, where $i = H$, the duration-cost contribution equals $\sigma_r^2 A_{T-1}^M \widehat{\Lambda}$. At jointly held maturities, either investor's equality gives the same total gap. The tax mix and portfolio responses derived below determine the loadings in the rate-loading channel.

The rate-loading channel links maturities. Iterating the loading recursion (55) from $A_0^a = 0$ gives

$$A_{T-1}^M - \xi_H A_{T-1}^C = \sum_{j=1}^{T-1} \rho^{T-1-j}\left(s_j^M - \xi_H s_j^C\right),$$

so the rate loading accumulates the response gaps $s_j^M - \xi_H s_j^C$ of the shorter maturities the bond passes through. The rate loading is zero at two periods, since the one-period responses are $s_1^M = \xi_H$ and $s_1^C = 1$ when $r > 0$.

Figure 5 shows the main-text decomposition without bond-specific risk. At thirty periods, the rate-loading term accounts for 54% of the return gap and the exemption-benefit term for the remainder. Figure 14 gives the counterpart with bond-specific risk, using the same channel names.

**Tax mix and portfolio responses.** The rate-loading channel depends on how one-period returns respond to the rate state. They respond for two reasons. Each holder's after-tax return on cash rises by its after-tax share per unit of the rate state. Trading also changes how much net rate exposure each holder carries, and with it the compensation each requires. Aggregating the holders' equalities separates the tax-mix and portfolio responses. Multiplying investor $i$'s equality (64) by $m_i/\gamma_i$, summing over the active holders and differentiating with respect to $z$ removes the bond-specific terms, because fixed supply implies $\sum_i m_i\, \partial x_{i,T}^M(z)/\partial z = 0$. What remains is

$$\left(\sum_{i\in\mathcal{U}_T^M} \frac{m_i}{\gamma_i}\right) s_T^M - \sum_{i\in\mathcal{U}_T^M} \frac{m_i\xi_i}{\gamma_i} = \sigma_r^2 A_{T-1}^M \sum_{i\in\mathcal{U}_T^M} m_i \frac{\partial \widehat{D}_i(z)}{\partial z}. \tag{67}$$

Define the active holders' risk-tolerance shares and aggregate exposure response as

$$\omega_{i,T}^M := \frac{m_i/\gamma_i}{\sum_{j\in\mathcal{U}_T^M} m_j/\gamma_j}, \qquad E_T^M := \frac{\sum_{i\in\mathcal{U}_T^M} m_i\, \partial \widehat{D}_i(z)/\partial z}{\sum_{i\in\mathcal{U}_T^M} m_i/\gamma_i}.$$

The weight $\omega_{i,T}^M$ is investor $i$'s risk tolerance, $m_i/\gamma_i$, as a share of the active holders' total. The term $E_T^M$ is the response of their aggregate net rate exposure to the rate state, divided by that total risk tolerance. Likewise, $E_T^C$ is defined from the holders of the taxable bond. Dividing (67) by total risk tolerance gives Proposition A.4, the formal counterpart of the response formulas (16) in Section 4.5.

**Proposition A.4** (Holding-return response to the rate state)**.** *Consider a regular state of the two-investor affine benchmark: price moments are differentiable, active risky holdings and inactive KKT slacks are strictly positive, and the active KKT-and-clearing Jacobian is nonsingular. The active set $\mathcal{U}_T^M \subseteq \{H, I\}$ is therefore locally fixed. Then*

$$s_T^M = \underbrace{\sum_{i\in\mathcal{U}_T^M} \omega_{i,T}^M \xi_i}_{\text{holders' tax mix}} + \underbrace{\sigma_r^2 A_{T-1}^M E_T^M}_{\text{portfolio-risk response}} \tag{68}$$

*exactly. The tax-mix term depends on who holds maturity $T$. The portfolio-risk response depends on how equilibrium trading changes those holders' exposure across their whole portfolios. This decomposes one market's return response. Comparing the two markets below separates the portfolio response into its municipal and taxable contributions to the rate-loading channel.*

**Proof of Proposition A.4.** Divide (67) by the active holders' total risk tolerance. Fixed supply cancels the bond-specific derivative by (47). The definitions of $\omega_{i,T}^M$ and $E_T^M$ then give (68). □

Proposition A.5 extends the same aggregation to nonlinear prices. It collects the active holders' tax treatment and their covariance-based risk compensation before taking the state derivative.

**Proposition A.5** (Required returns aggregated across a bond's holders)**.** *At a regular state with slack shorting constraints on cash for the active holders, let $\Gamma_j = \sum_{i\in\mathcal{U}_j} m_i/\gamma_i$ be their total risk tolerance,*

$\bar{\xi}_j = \Gamma_j^{-1} \sum_{i \in \mathcal{U}_j} m_i \xi_{ij}/\gamma_i$ *and* $\Omega_j = \Gamma_j^{-1} \sum_{i \in \mathcal{U}_j} m_i [\Sigma K_i x_i - c_i^L]_j$. *Here* $\bar{\xi}_j$ *is their tax mix and* $\Omega_j$ *their average risk compensation, weighted by risk tolerance. Then*

$$\begin{aligned} \mu_j &= \bar{\xi}_j r + \Omega_j, \\ s_j &= \bar{\xi}_j \, \partial_z r + \frac{1}{\Gamma_j} \sum_{i \in \mathcal{U}_j} m_i \left[ (\partial_z \Sigma) K_i x_i + \Sigma K_i \partial_z x_i - \partial_z c_i^L \right]_j . \end{aligned} \tag{69}$$

*In the normalized municipal–taxable benchmark, above the cash floor the response difference is*

$$s_T^M - \xi_H s_T^C = (\bar{\xi}_T^M - \xi_H) + \partial_z \Omega_T^M - \xi_H \partial_z \Omega_T^C. \tag{70}$$

**Proof.** Weighting the active equalities $\mu_j - \xi_{ij} r = \gamma_i [\Sigma K_i x_i - c_i^L]_j$ by $m_i/\gamma_i$, summing and dividing by $\Gamma_j$ gives the first line. Differentiating on the locally fixed participation branch gives the second line. In the two-market normalization, $\bar{\xi}_j = \bar{\xi}_T^M$ at municipal and $\bar{\xi}_j = 1$ at taxable coordinates. Subtraction gives (70). □

When the shorting constraint on cash binds, each active equality also contains $\kappa_i/\xi_i^a$ by (36). The return identity in (69) then gains $\Gamma_j^{-1} \sum_{i \in \mathcal{U}_j} m_i \kappa_i / (\gamma_i \xi_i^a)$, and the response identity gains its state derivative. A per-dollar charge, such as the bank financing-tax cost, appears in the return identity in the same units. When it does not vary with the rate state, it leaves the response identity unchanged on a given participation branch.

In the affine benchmark, substituting the covariance and liability credits into the active KKT equalities gives each bond's response

$$s_j = \bar{\xi}_j + \frac{\sigma_r^2 A_j}{\Gamma_j} \sum_{i \in \mathcal{U}_j} m_i \frac{\partial \widehat{D}_i}{\partial z} + \frac{v_j}{\Gamma_j} \sum_{i \in \mathcal{U}_j} m_i \xi_i^a \frac{\partial x_{ij}}{\partial z}, \tag{71}$$

where $A_j = A_{T-1}^a$ is the predecessor loading, the bond's loading one period later. For municipal bonds, whose after-tax shares are normalized to one, (47) cancels the final term.

The tax-mix term is $\xi_H$ under household-only holdings, $\xi_I$ under lower-tax-only holdings and, under joint holdings, the tax mix

$$\xi^* := \omega_H \xi_H + \omega_I \xi_I, \qquad \omega_i := \frac{m_i/\gamma_i}{m_H/\gamma_H + m_I/\gamma_I},$$

where $\omega_i$ is investor $i$'s share of total risk tolerance. Thus the lower-tax investor's after-tax share enters this term only at maturities it holds.

To sign the portfolio response, consider the following branch. The lower-tax investor alone holds risky taxable bonds, both investors hold municipal maturities in a nonempty set $J \subseteq \{2, \ldots, T_{\max}\}$, and the household alone holds every other municipal maturity. Proposition 2 shows that at a regular state with $r(z) > 0$, positive loadings and positive bond-specific variance on $J$,

$$d := \frac{\partial \widehat{D}_H(z)}{\partial z} > 0, \qquad \frac{\partial \widehat{D}_I(z)}{\partial z} = -\frac{m_H}{m_I} d < 0.$$

A higher rate state raises the difference in exemption benefits, $(\xi_I - \xi_H) r$, and trading in jointly held bonds shifts net rate exposure to the household. The household's unit cost of rate risk rises, raising its required compensation on all its risky municipal holdings, including maturities the lower-tax investor never holds. The sign holds for net rate exposure. Individual jointly held positions can move either way. At a household-only risky maturity,

$$s_T^M - \xi_H = \gamma_H \sigma_r^2 A_{T-1}^M \frac{\partial \widehat{D}_H(z)}{\partial z} > 0. \tag{72}$$

The taxable response moves the other way. Because the lower-tax investor's net rate exposure falls as the rate state rises, the response of its taxable return is

$$s_T^C = 1 - bA_{T-1}^C, \qquad b := \gamma_I \sigma_r^2 \frac{m_H}{m_I} d > 0,$$

which lies below one at every risky maturity. The coefficient $b$ measures how portfolio reallocation feeds back into taxable returns.

If joint municipal holding starts at $T^*$ and continues to the longest maturity, the municipal responses are

$$s_T^M = \begin{cases} \xi_H + \gamma_H \sigma_r^2 A_{T-1}^M d, & T < T^*, \\ \xi^* = \omega_H \xi_H + \omega_I \xi_I, & T \geq T^*. \end{cases} \tag{73}$$

Below the threshold, the lower-tax investor's after-tax share affects the municipal return response through the household's portfolio adjustment. From the threshold on, both investors hold the bond. Clearing cancels their aggregate exposure response, so the tax mix $\xi^*$ remains. These are the household-only and joint-ownership cases of (16). Without bond-specific risk, the illustration also has lower-tax-only maturities, treated in Appendix A.2.7.

Figure 6(a) plots the main-text illustration's responses. Figure 14(c) shows the joint-ownership responses just derived.

**Taxable counterpart.** Proposition 2 also uses the taxable version of (71). With $\xi_{ij} = 1$ at taxable coordinates, its tax term is one. Its bond-specific term vanishes if $v_T^C = 0$ and in two other cases. If one investor holds the taxable bond, clearing fixes its holding. If all holders have equal after-tax shares, these factor out of (47). Otherwise, unweighted clearing does not cancel the term. With the lower-tax investor the sole holder of the taxable bond,

$$s_T^C = 1 + \gamma_I \sigma_r^2 A_{T-1}^C \partial_z \widehat{D}_I. \tag{74}$$

Proposition 2 supplies the signs in (16) and the positive response gap in (80).

**Proposition 2** (Joint-ownership case with bond-specific risk)**.** *At a regular state with $r(z) > 0$ in the two-investor affine benchmark, suppose the lower-tax investor alone holds risky taxable bonds. Let $\varnothing \neq J \subseteq \{2, \ldots, T_{\max}\}$ be the fixed set of municipal maturities held by both investors. The high-tax household alone holds every municipal maturity outside $J$, including the riskless one-period claim. Assume $A_{T-1}^M, A_{T-1}^C > 0$ for every $T \geq 2$ and $v_T^M > 0$ on $J$. Define*

$$P_J = \sum_{T \in J} \frac{A_{T-1}^M}{v_T^M}, \quad S_J = \sum_{T \in J} \frac{(A_{T-1}^M)^2}{v_T^M}, \quad d := \frac{\partial \widehat{D}_H}{\partial z}.$$

*Then*

$$d = \frac{(\xi_I - \xi_H) m_I P_J}{(\gamma_H m_I + \gamma_I m_H)(1 + \sigma_r^2 S_J)} > 0, \qquad \frac{\partial \widehat{D}_I}{\partial z} = -\frac{m_H}{m_I} d. \tag{75}$$

*Writing $\omega_i = (m_i/\gamma_i)/(m_H/\gamma_H + m_I/\gamma_I)$, $\xi^* = \omega_H \xi_H + \omega_I \xi_I$ and $b = \gamma_I \sigma_r^2 (m_H/m_I) d$, the responses are*

$$s_T^M = \begin{cases} \xi_H + \gamma_H \sigma_r^2 A_{T-1}^M d, & T \notin J, \\ \xi^*, & T \in J, \end{cases} \qquad s_T^C = 1 - bA_{T-1}^C. \tag{76}$$

*Hence $s_T^M - \xi_H s_T^C > 0$ and $A_T^M - \xi_H A_T^C > 0$ at every risky maturity $T \geq 2$, with equality in both comparisons at $T = 1$.*

**Proof.** Clearing fixes each sole holder's position: $Q_T^M/m_H$ for the household at maturities outside $J$ and $Q_T^C/m_I$ for the lower-tax investor's risky taxable bonds. Only jointly held municipal positions vary. With constant liabilities and fixed taxable holdings, differentiating clearing, weighting by municipal predecessor loadings and summing gives

$$m_H \partial_z \widehat{D}_H + m_I \partial_z \widehat{D}_I = 0. \tag{77}$$

At each $T \in J$, subtracting the lower-tax investor's differentiated municipal equality from the household's gives

$$\xi_I - \xi_H = \sigma_r^2 A_{T-1}^M(\gamma_H \partial_z \widehat{D}_H - \gamma_I \partial_z \widehat{D}_I) + v_T^M(\gamma_H \partial_z x_{H,T}^M - \gamma_I \partial_z x_{I,T}^M).$$

With $k = (\gamma_H m_I + \gamma_I m_H)/m_H$, clearing reduces this to $\partial_z x_{I,T}^M = -(\xi_I - \xi_H)/(k v_T^M) - \sigma_r^2 A_{T-1}^M \partial_z \widehat{D}_I / v_T^M$. Since $\partial_z \widehat{D}_I = \sum_{T \in J} A_{T-1}^M \partial_z x_{I,T}^M$, multiplying by $A_{T-1}^M$ and summing over $J$ gives

$$(1 + \sigma_r^2 S_J)\partial_z \widehat{D}_I = -\frac{\xi_I - \xi_H}{k} P_J < 0.$$

Equation (77) gives (75). The responses follow from (71). Household-only quantities have zero derivatives, giving the first municipal response. At jointly held maturities, risk-tolerance aggregation cancels the aggregate exposure and bond-specific derivatives and leaves $\xi^*$, and (74) gives the taxable formula.

Equation (76) gives the response gap $\partial_z g_T = s_T^M - \xi_H s_T^C$ as

$$\partial_z g_T = \begin{cases} \sigma_r^2 d \left[ \gamma_H A_{T-1}^M + \xi_H \gamma_I \dfrac{m_H}{m_I} A_{T-1}^C \right], & T \notin J, \\ \omega_I(\xi_I - \xi_H) + \xi_H b A_{T-1}^C, & T \in J. \end{cases}$$

Thus $\partial_z g_T > 0$ for every $T \geq 2$, while $\partial_z g_1 = 0$ because both predecessor loadings vanish. Iterating (52), with $0 < \rho < 1$, yields

$$A_T^M - \xi_H A_T^C = \sum_{j=2}^{T} \rho^{T-j} \partial_z g_j > 0, \qquad T \geq 2.$$

□

Combining the two markets gives the response gap $s_j^M - \xi_H s_j^C = \partial_z g_j$. Its persistence-weighted sum over $j < T$ is the loading difference $A_{T-1}^M - \xi_H A_{T-1}^C$ in the rate-loading channel of (66). Applying the decomposition of Proposition A.4 in both markets, with the lower-tax investor the sole holder of risky taxable bonds, gives

$$s_T^M - \xi_H s_T^C = \underbrace{(\bar{\xi}_T^M - \xi_H)}_{\text{difference in tax mix}} + \underbrace{\sigma_r^2 \left( A_{T-1}^M E_T^M - \xi_H A_{T-1}^C E_T^C \right)}_{\text{difference in portfolio-risk responses}}, \tag{78}$$

where $\bar{\xi}_T^M := \sum_i \omega_{i,T}^M \xi_i$ is the active holders' tax mix: $\xi_H$ under household-only holdings, $\xi^*$ under joint holdings and $\xi_I$ under lower-tax-only holdings. The second term combines both markets' portfolio responses. Below the participation threshold, trading raises the household's municipal response and lowers the taxable response. Joint municipal holders' exposure responses cancel in aggregate. A sole lower-tax holder instead has a negative municipal portfolio response when its net exposure falls. The joint-ownership branch with bond-specific risk therefore adds $\xi^* - \xi_H$ from the threshold onward. In the main-text illustration, lower-tax-only maturities add $\xi_I - \xi_H$ and a portfolio offset. Figures 6 and 14 combine the municipal tax and portfolio terms and show the taxable contribution separately. All contributions vanish at the riskless one-period claim. Equation (70) gives the exact counterpart with nonlinear covariances.

On an upper-interval branch, the rate-loading channel of Proposition 1 incorporates portfolio-risk responses from three periods on and the change in tax mix from $T^* + 1$ on.

The empirical treasury comparison in Section 6.3 divides municipal loadings by $G_T := \sum_{j=1}^{T} \rho^{T-j} = A_T^{\text{Tsy}}$, the loading of a reference whose one-period return responds one-for-one to the rate state. Since $A_T^M / G_T = \sum_{j=1}^{T} (\rho^{T-j}/G_T) s_j^M$, the relative yield sensitivity averages municipal responses with positive weights. It equals $\xi_H$ at the household-only allocation. Its difference from $\xi_H$ reflects municipal holders' tax treatment and portfolio responses, both of which depend on participation. Figure 6(b) instead uses the equilibrium taxable market $C$, whose own portfolio response also contributes to $A_T^M - \xi_H A_T^C$.

Consider a per-dollar benefit or cost that does not vary with the rate state, such as a trading-cost advantage or the calibration's fixed bank financing-tax cost. On the same participation branch, it drops out

of these response formulas. Such a benefit or cost still changes which maturities an investor holds and the levels of one-period returns, as shown next.

**State-independent holding benefits.** Suppose an investor's objective includes $\sum_{a,T} \vartheta^a_{i,T} x^a_{i,T}$, with each $\vartheta^a_{i,T}$ independent of the rate state. Its active municipal condition becomes

$$\mu^M_T - \xi_i r + \vartheta^M_{i,T} = \gamma_i \left[\sigma_r^2 A^M_{T-1} \widehat{D}_i + v^M_T x^M_{i,T}\right],$$

and the taxable condition similarly gains $\vartheta^C_{i,T}/\xi_i$. With covariance, financing and constraints fixed, differentiation removes state-independent benefits and costs. The preceding proof and loading equations hold for the same feasible set $J$ of jointly held maturities, so a given loading solution is unchanged on the same participation branch. Benefits and costs affect return intercepts, gap levels (Appendix A.3) and the feasible support. State-dependent benefits add their derivatives. Fixed participation fees affect global participation without a per-unit KKT term.

**Uniqueness of the loading fixed point.** Appendix A.2.8 verifies the following sufficient condition for a unique loading fixed point in the illustration. For given $d$, let $\mathcal{F}_J(d)$ be the right side of (75) at the municipal loadings that (76) and the loading recursion (52) generate. If $\mathcal{F}_J$ maps $[0,\infty)$ into $[0, d_{\max}]$ and $\sup_{[0,d_{\max}]} |\mathcal{F}'_J| < 1$, every nonnegative exposure fixed point lies in $[0, d_{\max}]$, and Banach's fixed-point theorem gives exactly one, with its loading path on this branch. Holdings and non-holders' inequalities still need checking. Equilibria with other feasible participation thresholds remain possible. The derivative of $\mathcal{F}_J$ is

$$\mathcal{F}'_J(d) = c\frac{P'_J(1+\sigma_r^2 S_J) - \sigma_r^2 P_J S'_J}{(1+\sigma_r^2 S_J)^2}, \qquad c = \frac{(\xi_I - \xi_H) m_I}{\gamma_H m_I + \gamma_I m_H},$$

with the derivatives of $P_J, S_J$ obtained by differentiating the same loading recursions.

The last result in this subsection gives the household-only benchmark. Without lower-tax participation, municipal loadings are exactly $\xi_H$ times taxable loadings.

**Proposition A.6** (Household-only loadings)**.** *In the two-investor affine benchmark, if the high-tax household alone holds every municipal bond and the lower-tax investor alone holds every taxable bond, then $s^M_T = \xi_H$, $s^C_T = 1$ and $A^M_T = \xi_H A^C_T$ at every maturity.*

**Proof.** Clearing fixes every holding, so both net rate exposures have zero derivatives. Equations (71) and (52) give the result. The identity concerns loadings, and different intercept risk premia can still affect the yield ratio. □

### A.2.5 Yield gaps and implied taxes

The channels act on one-period returns, while the municipal puzzle concerns yields to maturity. This subsection carries the channels into yields, formalizing the yield relation (18) in Section 4.5.

Applying (56) to both markets turns the one-period gap $g_T$ of Proposition 1 into the yield gap:

$$\begin{aligned} y^M_T - \xi_H y^C_T = \quad & \underbrace{\frac{1}{T}\sum_{j=0}^{T-1} \mathbb{E}_t\left[g_{T-j}(z_{t+j})\right]}_{\text{average gap along the bond's life}} \\ & \underbrace{- \frac{1}{2T}\sum_{k=2}^{T}\left[\sigma_r^2\left(\left(A^M_{k-1}\right)^2 - \xi_H\left(A^C_{k-1}\right)^2\right) + \left(\sigma^M_{\eta,k}\right)^2 - \xi_H\left(\sigma^C_{\eta,k}\right)^2\right]}_{\text{convexity difference}}. \end{aligned} \tag{79}$$

The riskless one-period gap is zero. The yield-implied tax is

$$\tau^{Y,\text{impl}}_T := 1 - \frac{y^M_T}{y^C_T}, \qquad \text{the model counterpart of } \textit{ImpliedTax} \text{ in (2).}$$

It equals the household wedge less the yield gap per unit of taxable yield, $\tau_T^{Y,\text{impl}} = (1-\xi_H) - (y_T^M - \xi_H y_T^C)/y_T^C$, as in the identity (3) of Section 3.6. A positive gap accumulated in (79) pulls the implied tax below the household wedge, more so at longer maturities when the gap grows proportionally faster than the taxable yield (Proposition A.7(ii)).

The relative yield response sums the response gaps in (78). On a joint-participation branch, the response gap is positive at every risky maturity (Proposition 2). Below the participation threshold it reflects portfolio responses, and from that threshold on both the tax mix and portfolio responses. Hence, for $T \geq 2$

$$\underbrace{\frac{\partial y_T^M}{\partial z} - \xi_H \frac{\partial y_T^C}{\partial z}}_{\text{relative yield response}} = \frac{1}{T}\sum_{j=2}^{T} \rho^{T-j} \underbrace{\left(\frac{\partial \mu_j^M}{\partial z} - \xi_H \frac{\partial \mu_j^C}{\partial z}\right)}_{\text{holding-return response gap}} > 0, \tag{80}$$

in model-period units, and both sides are zero for the riskless one-period bond. Long municipal yields therefore move more with the short rate than the household benchmark implies. The lower-tax investor's tax affects these responses even at maturities it does not hold.

Proposition A.7 uses one restriction on taxes, risk-tolerance shares and rate persistence:

$$\text{(S)} \qquad \omega_H \frac{\xi_I - \xi_H}{\xi_H} \leq \rho. \tag{81}$$

Equation (86) writes the same condition in terms of masses and risk aversions. This condition keeps the feedback $b$ below $\rho$, so taxable loadings rise with maturity. Municipal loadings also rise beyond the participation threshold, because the threshold loading lies below its long-end value, $A_{T^*-1}^M < \xi^*/(1-\rho)$ (equation (87)). The long-end taxable response $s_\infty^{C,\text{aff}} := \lim_{T\to\infty} s_T^{C,\text{aff}}$ lies in $(0,1)$, because the lower-tax investor's net rate exposure falls as the rate state rises.

For the yield levels in Proposition A.7(ii), define the forward rate implied by today's prices for the period from $t+T-1$ to $t+T$:

$$f_T^a(z) = \log \frac{P_{T-1}^a(z)}{P_T^a(z)}, \qquad y_T^a(z) = \frac{1}{T}\sum_{j=1}^{T} f_j^a(z). \tag{82}$$

With positive taxable forwards, the yield ratio is a weighted average of municipal-to-taxable forward-rate ratios. The forward rate compares adjacent-maturity prices at the same date, whereas the holding return uses next period's resale price. The affine recursions (55) link the two:

$$f_T^a(z) = \mu_T^a(z) - (1-\rho)A_{T-1}^a z - \tfrac{1}{2}\left[\sigma_r^2 (A_{T-1}^a)^2 + (\sigma_{\eta,T}^a)^2\right]. \tag{83}$$

For taxable bonds, the term premium is the yield less the average expected future short rate. Equation (56) expresses it as average expected bond risk premia less convexity. The municipal–taxable yield gap also reflects differences in holders' tax benefits.

**Proposition A.7** (Maturity ordering of the loading ratio and the yield-implied tax)**.** *In the two-investor affine benchmark, consider the following equilibrium of Proposition 2. The high-tax household alone holds municipal bonds shorter than $T^*$, both investors hold municipal bonds from $T^*$ to $T_{\max}$, and the lower-tax investor holds the taxable bonds. Assume condition (S) in (81).*

***(i) Rate sensitivity.*** *The municipal-to-taxable loading ratio is nondecreasing in maturity. It starts at the household's after-tax share. When the recursions are extended to infinite maturity with the household's exposure response $d$ held at its solved value, the ratio converges to the long-end tax mix relative to the taxable response:*

$$\frac{A_1^{M,\text{aff}}}{A_1^{C,\text{aff}}} = \xi_H \qquad \textit{and} \qquad \lim_{T\to\infty} \frac{A_T^{M,\text{aff}}}{A_T^{C,\text{aff}}} = \frac{\xi^*}{s_\infty^{C,\text{aff}}}. \tag{84}$$

***(ii) Yield levels.*** *Suppose also that taxable forward rates are positive and the municipal-to-taxable forward-rate ratio is nondecreasing in maturity at the stationary mean $z = 0$. Then the yield ratio $y_T^M/y_T^C$ is nondecreasing at $z = 0$ (Proposition A.8), so the yield-implied tax $\tau_T^{Y,\text{impl,aff}}$ is nonincreasing.*

***(iii) Long end.*** *With the recursions extended to infinite maturity, suppose the loadings, the risk-premium coefficients $\pi_T^a$ and the bond-specific compensation terms $\iota_T^a$ converge, with limits marked by the subscript $\infty$. Assume the limiting taxable forward is positive. Then*

$$\lim_{T\to\infty} \tau_T^{Y,\text{impl,aff}} \;=\; 1 - \frac{\xi^*\bar{r} + A_\infty^M\pi_\infty^M - \frac{1}{2}\sigma_r^2(A_\infty^M)^2 + \iota_\infty^M}{\bar{r} + A_\infty^C\pi_\infty^C - \frac{1}{2}\sigma_r^2(A_\infty^C)^2 + \iota_\infty^C}, \tag{85}$$

*where $\pi_T^a = \sum_i \omega_{i,T}^a \gamma_i \sigma_r^2 \widehat{D}_i(0)$ is the holders' cost of rate risk per unit of loading, weighted by risk tolerance, and $\iota_T^a$ is their compensation for bond-specific risk net of its convexity term (equation (88)).*

The proof below also gives the treasury-reference ordering in Proposition A.9(iii).

**The sufficient condition and the threshold loading.** For Proposition 2 with the upper interval $J = \{T^*, \dots, T_{\max}\}$, the maturity ordering uses condition (S) and the bound (M) on the threshold loading:

$$\textbf{(S)} \quad \frac{\gamma_I m_H(\xi_I - \xi_H)}{\xi_H(\gamma_H m_I + \gamma_I m_H)} \le \rho, \tag{86}$$

$$\textbf{(M)} \quad A_{T^*-1}^M < \frac{\xi^*}{1-\rho}. \tag{87}$$

Condition (S) is assumed. It restates condition (81) with $\omega_H$ written out. The bound (M) holds at every solution of (75) on this branch. To prove it, put $n = T^* - 1$ and $\lambda = \gamma_H\sigma_r^2 d$, so that (75) reads $\lambda = \omega_I(\xi_I - \xi_H)\sigma_r^2 P_J/(1 + \sigma_r^2 S_J)$. The recursions at household-only and jointly held maturities give $A_n^M = \xi_H \sum_{q=0}^{n-1}(\rho+\lambda)^q$ and $A_{n+l}^M = \rho^l A_n^M + (1-\rho^l)\xi^*/(1-\rho)$ for $l \ge 0$. Suppose $A_n^M \ge \xi^*/(1-\rho)$. Every loading on $J$ is then at least $\xi^*/(1-\rho)$, so $S_J/P_J \ge \xi^*/(1-\rho)$ and $\lambda < (1-\rho)(\xi^* - \xi_H)/\xi^* < 1-\rho$, using $\xi^* - \xi_H = \omega_I(\xi_I - \xi_H)$. Then $A_n^M < \xi_H/(1-\rho-\lambda) < \xi^*/(1-\rho)$, a contradiction. Both statements concern the upper-interval branch of the two-investor affine benchmark and do not prove that this branch exists.

**Proof of Proposition A.7(i).** Step 1 concerns municipal loadings. At household-only maturities, $A_T^M = (\rho + \gamma_H\sigma_r^2 d)A_{T-1}^M + \xi_H$, so $A_T^M$ is increasing from zero. At jointly held maturities, $A_T^M = \rho A_{T-1}^M + \xi^*$, whose increments have the sign of $\xi^* - (1-\rho)A_{T^*-1}^M$. Condition (M) makes them positive. Step 2 concerns the taxable leg. The bound $A_{T-1}^M \ge \xi_H$ on $J$ gives $P_J/S_J \le 1/\xi_H$, and (75) and (S) give $0 \le b \le \rho$. Therefore

$$A_T^C = \sum_{q=0}^{T-1}(\rho - b)^q, \qquad s_T^C = 1 - bA_{T-1}^C \downarrow \frac{1-\rho}{1-\rho+b} > 0.$$

Step 3 concerns the municipal response, which rises at household-only maturities, is constant at jointly held maturities and does not jump down when the lower-tax investor begins to participate. For the last claim, the response at $T^* - 1$ must not exceed $\xi^*$. Substituting (75) shows that the required inequality $\omega_I(\xi_I - \xi_H) \ge \gamma_H\sigma_r^2 A_{T^*-2}^M d$ is equivalent to $\sigma_r^2 A_{T^*-2}^M P_J \le 1 + \sigma_r^2 S_J$. This inequality follows from $A_{T-1}^M > A_{T^*-2}^M$ for every $T \in J$, established above using (M). Since Step 2 makes $s_j^C$ positive and nonincreasing, the response ratio $s_j^M/s_j^C$ is nondecreasing in maturity.

Step 4 is an averaging argument. For any positive $e_j$ and nondecreasing ratios $q_j = a_j/e_j$, the ratio weighted by rate persistence,

$$\bar{q}_T = \frac{\sum_{j\le T}\rho^{T-j}a_j}{\sum_{j\le T}\rho^{T-j}e_j}$$

is nondecreasing: as a positive-weight average of $q_1, \dots, q_T$, it is at most $q_T \le q_{T+1}$, and

$$\bar{q}_{T+1} - \bar{q}_T = \frac{e_{T+1}(q_{T+1} - \bar{q}_T)}{\rho\sum_{j\le T}\rho^{T-j}e_j + e_{T+1}} \ge 0.$$

By (52), $A_T^M/A_T^C$ has this form. Applying the argument with $a_j = s_j^M$ and $e_j = s_j^C$ proves the loading-ratio ordering, starting from $\xi_H$.

Step 5 gives the limits and the treasury reference. Extending the same branch recursions to arbitrarily long maturities with fixed $d$ and eventual joint participation gives the limits $A^M_\infty = \xi^*/(1-\rho)$ and $A^C_\infty = 1/(1-\rho+b)$, hence the endpoint in (84). With $e_j = 1$, for which (M) suffices, the same argument proves that $A^M_T/G_T$ is nondecreasing. Here $G_T = \sum_{j\le T}\rho^{T-j}$ is the treasury-reference loading of Appendix A.2.4. Thus the same ordering holds for the empirical treasury comparison of Section 6.3. □

Part (ii) of Proposition A.7 rests on the following bridge.

**Proposition A.8** (Exact forward-ratio bridge)**.** *For any price functions, affine or nonlinear, let* $f^a_T(z) = -p^a_T(z) + p^a_{T-1}(z)$*. If* $f^C_T(z) > 0$ *and* $f^M_T(z)/f^C_T(z)$ *is nondecreasing over the maturities considered, then* $y^M_T(z)/y^C_T(z)$ *is nondecreasing and* $\tau^{Y,\mathrm{impl}}_T(z)$ *is nonincreasing.*

**Proof.** Telescoping gives $Ty^a_T = -p^a_T = \sum_{j\le T} f^a_j$. The yield ratio averages forward ratios with positive weights $f^C_j > 0$. Adding a ratio at least as large as its predecessors cannot lower the average. On a two-year grid, $T$ counts steps and annualization cancels. □

The premise concerns full forward rates, including bond-specific and convexity terms, so monotone loadings alone do not establish it. Section 5 checks the nonlinear model's maturity profiles numerically.

Part (iii) adds relative term premia, convexity and bond-specific compensation to the long-end tax mix. With $\widehat{\Lambda} > 0$, small taxable feedback $b$ and negligible convexity and bond-specific terms, the long-end implied tax is, to first order, below $1-\xi^*$ and hence below $1-\xi_H$ (equations (90) and (91)).

**Affine forward levels and Proposition A.7(iii).** For coordinate $j = (a,T)$, put $\Gamma^a_T := \sum_{i\in\mathcal{U}_j} m_i/\gamma_i$ and

$$\bar\xi^a_T = \bar\xi_j = \frac{\sum_i (m_i/\gamma_i)\xi_{ij}}{\Gamma^a_T}, \quad \pi^a_T = \frac{\sigma_r^2}{\Gamma^a_T}\sum_i m_i \widehat{D}_i(0), \quad \iota^a_T = v_j\left[\frac{\sum_i m_i \xi^a_i x^a_{i,T}(0)}{\Gamma^a_T} - \tfrac{1}{2}\right],$$

where sums run over active holders. Here $\pi^a_T$ is the holders' cost of rate risk per unit of loading and $\iota^a_T$ their compensation for bond-specific risk net of its convexity term. Aggregating the holders' level KKT equalities at $z=0$ and subtracting the two Jensen terms in the exact Gaussian mean gives

$$C^a_T - C^a_{T-1} = \bar\xi^a_T \bar r + A^a_{T-1}\pi^a_T - \tfrac{1}{2}\sigma_r^2 (A^a_{T-1})^2 + \iota^a_T. \tag{88}$$

For municipal bonds the weighted holding sum is $Q^M_T$. For taxable bonds heterogeneous after-tax shares remain inside it. The increment $C^a_T - C^a_{T-1}$ is the forward rate at $z=0$. If the displayed coefficients and loadings converge and the limiting taxable forward $f^C_\infty$ is positive, Cesàro convergence gives

$$y^a_T(z) \longrightarrow f^a_\infty = \bar\xi^a_\infty \bar r + A^a_\infty \pi^a_\infty - \tfrac{1}{2}\sigma_r^2 (A^a_\infty)^2 + \iota^a_\infty \tag{89}$$

at fixed $z$, because $C^a_T/T$ converges to the forward limit and $A^a_T z/T \to 0$. At jointly held long maturities, $\bar\xi^M_\infty = \xi^*$. Relative after-tax shares on taxable bonds equal one, so $\bar\xi^C_\infty = 1$. The implied-tax limit $1 - f^M_\infty/f^C_\infty$ is therefore (85). This proves part (iii).

For the first-order approximation, aggregating by risk tolerance over both investors for long municipal bonds and over the lower-tax investor alone for long taxable bonds gives

$$\pi^M_\infty - \pi^C_\infty = \omega_H\, \sigma_r^2\, \widehat{\Lambda}, \qquad \widehat{\Lambda} = \gamma_H \widehat{D}_H - \gamma_I \widehat{D}_I, \tag{90}$$

where $\widehat{D}_i = \widehat{D}_i(0)$ is the net rate exposure at the stationary mean. The premium difference does not bound the convexity and bond-specific terms. With small taxable loading feedback ($b \ll 1-\rho$), negligible convexity and bond-specific terms and a positive denominator, substituting $A^M_\infty = \xi^*/(1-\rho)$, $A^C_\infty \approx 1/(1-\rho)$ and (90) into (89) gives the first-order endpoint

$$\tau^{Y,\mathrm{impl,aff}}_\infty \approx \underbrace{(1-\xi^*)}_{\text{one minus tax mix}} - \underbrace{\frac{\xi^*\,\omega_H\,\sigma_r^2\,\widehat{\Lambda}}{(1-\rho)\bar r + \pi^C_\infty}}_{\text{municipal term-premium excess}}. \tag{91}$$

With $\widehat{\Lambda} > 0$, this first-order endpoint lies below one minus the tax mix, $1-\xi^*$, and hence below the one-period implied tax $1-\xi_H$.

### A.2.6 Tax pass-through and relative rate sensitivity

The mechanism predicts whose taxes are priced at each maturity. This subsection gives sufficient conditions on the branch with bond-specific risk and joint long-maturity holdings. Appendix A.2.7 treats the main-text illustration separately. P1 and P2 concern yield levels, while P3 concerns rate sensitivities. To keep these objects distinct, I pair the yield-implied tax of Appendix A.2.5 with a loading-implied tax:

$$\tau_T^{Y,\text{impl}} = 1 - \frac{y_T^M}{y_T^C}, \qquad \tau_T^A = 1 - \frac{A_T^M}{A_T^C}.$$

To measure pass-through, I vary one investor's after-tax share at a fixed rate state and re-solve the equilibrium. Holdings, risk costs and both yield curves adjust while non-tax primitives stay fixed. The pass-throughs of investor $i$'s exemption benefit $1-\xi_i$ into the two implied taxes are

$$\mathcal{P}_i^Y(T) := \frac{d\tau_T^{Y,\text{impl}}}{d(1-\xi_i)}, \qquad \mathcal{P}_i^A(T) := \frac{d\tau_T^A}{d(1-\xi_i)}, \qquad i \in \{H, I\},$$

where an argument $\infty$ denotes the limit as $T \to \infty$. Under the representative-tax benchmark of Section 3.3, $y_T^M = \xi_H y_T^C$ and $A_T^M = \xi_H A_T^C$ at every maturity. Hence, in both measures, household pass-through is one and the lower-tax investor's is zero at every maturity. With heterogeneous taxes, the household alone holds the riskless short bond, whose return is $\xi_H r$. Both investors hold long bonds, whose returns depend on both after-tax shares through the tax mix $\xi^*$. For the empirical treasury comparison of Section 6.3, I also use a taxable reference whose one-period return moves one-for-one with the rate state, giving loading $A_T^{\text{Tsy}} = G_T$ (Appendix A.2.4).

The long-end yield inequalities in (21) hold in the main-text illustration. Proposition A.9 gives sufficient conditions for them on the joint-ownership branch with positive bond-specific risk.

**Proposition A.9** (Tax pass-through at the short and long ends)**.** *Consider the joint-ownership branch of Proposition 2 in Appendix A.2.4, under condition (S) of (81), with $0 < \rho < 1$, $\sigma_r > 0$, a positive cash rate and a slack household shorting constraint on cash. With its recursions extended to infinite maturity, suppose long yields and their tax derivatives converge locally uniformly, with a positive long taxable yield.*

*(i) (Short end) Both implied taxes equal the household wedge, $\tau_1^{Y,\text{impl}} = \tau_1^A = 1-\xi_H$, so $\mathcal{P}_H^Y(1) = \mathcal{P}_H^A(1) = 1$ and $\mathcal{P}_I^Y(1) = \mathcal{P}_I^A(1) = 0$.*

*(ii) (Long end, P1 and P2) Suppose the after-tax shares satisfy the tax-difference condition*

$$2\omega_H e + \omega_I e^2 \le 1, \qquad e := \frac{\xi_I - \xi_H}{\xi_H}.$$

*Then*

$$\mathcal{P}_H^A(\infty) < \omega_H, \qquad \mathcal{P}_I^A(\infty) > \omega_I.$$

*Along the algebraic continuation of a fixed participation branch, the difference $\mathcal{P}_i^Y(\infty) - \mathcal{P}_i^A(\infty)$ shrinks to zero as the mean rate $\bar{r}$ rises. Consequently, for mean rates above a bound set by the other primitives, the following inequalities hold whenever the branch and its local tax perturbations remain feasible:*

$$\mathcal{P}_H^Y(\infty) < \omega_H < 1, \qquad \mathcal{P}_I^Y(\infty) > \omega_I > 0.$$

*(iii) (Response ratios, P3) The municipal-to-taxable zero-coupon response ratio $A_T^M/A_T^C$ is nondecreasing in maturity, equals $\xi_H$ at the short end and converges to $\xi^*/s_\infty^C$, with $0 < s_\infty^C < 1$. For the unit-response reference, $A_T^M/A_T^{\text{Tsy}}$ has the same ordering and short endpoint, and converges to $\xi^* > \xi_H$.*

Part (ii) first separates the holders' tax mix from portfolio adjustment. With the tax mix alone, each investor's pass-through equals its risk-tolerance share, $\omega_H$ or $\omega_I$. Under the tax-difference condition, portfolio

adjustment offsets the household's tax-mix effect and reinforces the lower-tax investor's, giving $\mathcal{P}_H^A(\infty) < \omega_H$ and $\mathcal{P}_I^A(\infty) > \omega_I$. Lemma A.1 proves these bounds while allowing loadings and portfolios to adjust.

Raising $\xi_H$ narrows the tax difference, so less net rate exposure shifts to the household as rates rise. It also raises the tax mix and bond loadings, strengthening the risk feedback. The first effect dominates when the tax difference is not too large relative to $\xi_H$. The condition bounds this relative difference, $e$, and holds whenever $\xi_I < \frac{3}{2}\xi_H$, whatever the risk-tolerance shares or the participation threshold. Lemma A.1 proves the derivative bounds with both effects included.

For yield levels, write $R_\infty = A_\infty^M / A_\infty^C$. The proof shows that $\mathcal{P}_i^Y(\infty)$ equals $\mathcal{P}_i^A(\infty)$ plus two corrections divided by the long taxable yield. The first is the tax response of $y_\infty^M - R_\infty y_\infty^C$, which contains only relative risk costs, convexity and bond-specific compensation and does not depend on $\bar{r}$. The second reflects the tax response of the taxable yield. The taxable yield rises with $\bar{r}$, so both corrections shrink as $1/\bar{r}$ along the fixed-branch continuation. The yield-level inequalities in part (ii) apply where the mean rate is high enough for these corrections to be small and the participation branch remains feasible.

The tax-difference condition holds for a top-bracket household and a corporate investor. The illustration with bond-specific risk satisfies it and the yield-level inequalities with wide margins (Appendix A.2.8). In that illustration, convexity explains most of the difference between loading and yield pass-throughs, and the yield-level inequalities hold for every nonnegative taxable risk cost $\pi_\infty^C$. Lower mean rates or stronger hedging demand can enlarge the corrections by reducing the long taxable yield in their denominator.

The proposition establishes the short-end and long-end contrasts that the empirical tests use. At intermediate maturities, the tax-mix part of $\mathcal{P}_i^A(T)$ shifts weight from the household to the lower-tax investor as more of a bond's remaining life falls at jointly held maturities (Proposition A.10). In the illustration with bond-specific risk, $\mathcal{P}_H^Y(T)$ falls and $\mathcal{P}_I^Y(T)$ rises strictly with maturity. Appendix A.2.7 verifies the same finite-maturity ordering for the main-text illustration. With more clienteles, each has its own participation threshold and risk-tolerance share (Appendix A.5). Section 4.6 states the resulting P1–P3 predictions, and Section 6 tests them.

Pass-throughs are the derivatives of $y_T^M / y_T^C$ and $A_T^M / A_T^C$ with respect to $\xi_i$.

**Proof of Proposition A.9.**

**(i) The riskless short end.** The one-period claims have zero variance. With a slack household shorting constraint on cash, Proposition A.2 gives $y_1^M = \xi_H r$ and $y_1^C = r$ at every rate state with $r > 0$. The one-period loadings equal the responses $s_1^M = \xi_H$ and $s_1^C = 1$ by (55), so $y_1^M / y_1^C = A_1^M / A_1^C = \xi_H$. Differentiating the two ratios gives $\mathcal{P}_H^Y(1) = \mathcal{P}_H^A(1) = 1$ and $\mathcal{P}_I^Y(1) = \mathcal{P}_I^A(1) = 0$.

**(ii) The long end.** Let $d$ be the household's exposure response in (75) and $b = \gamma_I \sigma_r^2 (m_H / m_I) d$ the feedback into taxable responses, as in (76). The solved long responses and loadings satisfy

$$s_\infty^M = \xi^*, \qquad s_\infty^C = \frac{1-\rho}{1-\rho+b}, \qquad A_\infty^a = \frac{s_\infty^a}{1-\rho}.$$

Thus $R_\infty = A_\infty^M / A_\infty^C = \xi^* / s_\infty^C$. At the mean state, the solved yield formula (89) gives

$$\begin{aligned} y_\infty^M &= \xi^* \bar{r} + A_\infty^M \pi_\infty^M - \tfrac{1}{2}\sigma_r^2 (A_\infty^M)^2 + \iota_\infty^M, \\ y_\infty^C &= \bar{r} + A_\infty^C \pi_\infty^C - \tfrac{1}{2}\sigma_r^2 (A_\infty^C)^2 + \iota_\infty^C. \end{aligned}$$

The four terms are interest income, portfolio rate-risk compensation, convexity and bond-specific compensation. Dividing the two yields separates the loading ratio from the remaining yield-level difference:

$$\frac{y_\infty^M}{y_\infty^C} = R_\infty + \frac{y_\infty^M - R_\infty y_\infty^C}{y_\infty^C}.$$

All derivatives below are with respect to $\xi_i$, with holdings, loadings and both yields adjusting. The assumed convergence permits differentiation of the limiting ratio. The quotient rule gives

$$\begin{aligned}\mathcal{P}_i^Y(\infty) = \mathcal{P}_i^A(\infty) + \frac{1}{y_\infty^C}\frac{d(y_\infty^M - R_\infty y_\infty^C)}{d\xi_i} \\ - \frac{y_\infty^M - R_\infty y_\infty^C}{(y_\infty^C)^2}\frac{dy_\infty^C}{d\xi_i}.\end{aligned} \tag{92}$$

The three terms are the responses of the loading ratio, the relative yield compensation and the taxable denominator. The first retains the endogenous taxable response:

$$\mathcal{P}_i^A(\infty) = \frac{\omega_i}{s_\infty^C} - \frac{\xi^*}{(s_\infty^C)^2}\frac{ds_\infty^C}{d\xi_i}.$$

Lemma A.1 below proves $\mathcal{P}_H^A(\infty) < \omega_H$ and $\mathcal{P}_I^A(\infty) > \omega_I$ under the tax-difference condition, with uniform positive margins.

The calculation after the lemma shows that $y_\infty^M - R_\infty y_\infty^C$ and its tax derivatives are independent of $\bar{r}$. Also, $y_\infty^C = A_\infty^C(1-\rho)\bar{r} + O(1)$ and $dy_\infty^C/d\xi_i = O(\bar{r})$. Both correction terms in (92) therefore vanish at rate $1/\bar{r}$. The loading and derivative bounds in the lemma are uniform over the finitely many cutoffs. On each cutoff's algebraic continuation, a lower bound on $\bar{r}$ that depends only on the other primitives makes both yields positive and keeps the corrections smaller than the loading margins. Applying this comparison to an equilibrium also requires nonnegative holdings and the inactive KKT inequalities to remain valid locally. This proves $\mathcal{P}_H^Y(\infty) < \omega_H < 1$ and $\mathcal{P}_I^Y(\infty) > \omega_I > 0$.

**(iii) The yield-response ratios.** At a common maturity, dividing both price loadings by $T$ leaves their ratio unchanged. Proposition A.7(i) gives the ordering and endpoints for the equilibrium taxable market $C$: the ratio starts at $\xi_H$ and converges to $\xi^*/s_\infty^C$. For the unit-response reference of Appendix A.2.4, $A_T^{\mathrm{Tsy}} = G_T$, with $A_1^{\mathrm{Tsy}} = 1$ and $A_\infty^{\mathrm{Tsy}} = 1/(1-\rho)$. Its ratio therefore starts at $\xi_H$ and converges to $\xi^*$. Appendix A.2.5 proves monotonicity for both comparisons. □

**Lemma A.1** (The loading response)**.** *Under the tax-difference condition in Proposition A.9, $\mathcal{P}_H^A(\infty) < \omega_H$ and $\mathcal{P}_I^A(\infty) > \omega_I$. With other primitives fixed, both strict margins have positive lower bounds over all feasible upper-interval cutoffs.*

**Proof.** The proof differentiates the exposure fixed point and bounds its feedback into the loading ratio. Put $n = T^* - 1$ and $\lambda = \gamma_H \sigma_r^2 d$, the response of the household's unit risk cost to the rate state. Equation (75) gives

$$\lambda = \omega_I(\xi_I - \xi_H)\sigma_r^2 f, \qquad f = \frac{P_J}{1+\sigma_r^2 S_J}, \qquad R_\infty = \xi^*\left[1 + \frac{\omega_H \lambda}{\omega_I(1-\rho)}\right].$$

The sums $P_J$ and $S_J$ are those of Proposition 2. With $G_n(p) = \sum_{q=0}^{n-1} p^q$, the loading recursion gives

$$A_n^M = \xi_H G_n(\rho+\lambda), \qquad A_{n+l}^M = \rho^l A_n^M + (1-\rho^l)A_\infty^M, \quad l \geq 0.$$

By (87), $A_n^M < A_\infty^M$, so loadings rise along the jointly held block. Thus $S_J/P_J \geq A_n^M$ and

$$\lambda G_n(\rho+\lambda) < \omega_I e, \qquad \lambda\frac{G_n'(\rho+\lambda)}{G_n(\rho+\lambda)} < \omega_I e,$$

where $e$ is the relative tax difference of Proposition A.9. The first bound uses $\lambda < \omega_I(\xi_I - \xi_H)P_J/S_J$, and the second follows from $G_n' \leq G_n^2$. The tax-difference condition implies $e < 1$ and hence $b = (\omega_H/\omega_I)\lambda < \omega_H e < 1$, ensuring convergence of the taxable loadings.

To differentiate the loading sums, let $t_j = \rho^l A_n^M / A_{j-1}^M$ for $j = T^* + l$ and

$$X = \bar{t}_P - \frac{2\sigma_r^2 S_J}{1+\sigma_r^2 S_J}\bar{t}_S.$$

Here $\bar{t}_P$ is the mean of $t_j$ with weights proportional to $A^M_{j-1}/v^M_j$, and $\bar{t}_S$ uses weights proportional to $(A^M_{j-1})^2/v^M_j$. Since $0 \le t_j \le 1$, $-2 < X \le 1$. Let $\partial_i$ denote differentiation with respect to $\xi_i$ holding $\lambda$ fixed. Differentiating the loadings gives

$$\partial_H \ln f = \left(\frac{1}{\xi_H} - \frac{\omega_H}{\xi^*}\right) X + \frac{\omega_H}{\xi^*}\frac{1-\sigma_r^2 S_J}{1+\sigma_r^2 S_J},$$
$$\partial_I \ln f = \frac{\omega_I}{\xi^*}\left(\frac{1-\sigma_r^2 S_J}{1+\sigma_r^2 S_J} - X\right).$$

The partial derivatives above hold $\lambda$ fixed. To include its equilibrium response, differentiate its fixed point. With $G_n$ and $G_n'$ evaluated at $\rho+\lambda$,

$$\frac{d\lambda}{d\xi_i} = \frac{\lambda\,\partial_i \ln[(\xi_I - \xi_H)f]}{D}, \qquad D = 1 - \lambda\frac{G_n'}{G_n}X > 1 - \omega_I e > 0.$$

Substitute into $dR_\infty/d\xi_i = \mathcal{P}^A_i(\infty)$. The coefficient of $X$ is positive in the household derivative and negative in the lower-tax derivative. Using $X \le 1$ and $\lambda G_n'/G_n < \omega_I e$ gives

$$\mathcal{P}^A_H(\infty) - \omega_H \le \frac{\omega_H\lambda}{\omega_I(1-\rho)D}\left(\frac{2\omega_H}{1+\sigma_r^2 S_J} - \frac{1}{e} + \omega_I e\right) < 0,$$
$$\mathcal{P}^A_I(\infty) - \omega_I > \frac{\omega_H\lambda}{\omega_I(1-\rho)D}(1/e - \omega_I^2 e) > 0.$$

The tax-difference condition makes the household's parenthesized term at most $-2\omega_H\sigma_r^2 S_J/(1+\sigma_r^2 S_J)$, giving a positive margin for each investor.

For uniform margins, consider $n \in \{1,\dots,T_{\max}-1\}$ in the finite loading system and trial feedbacks $x \ge 0$ with $xG_n(\rho+x) \le \omega_I e$. This compact domain, defined by primitives, contains every equilibrium pair. Its positive loading sums give positive minima for $f$ and $\sigma_r^2 S_J/(1+\sigma_r^2 S_J)$, while $D \ge 1 - \omega_I e$ keeps full tax derivatives bounded. At a fixed point, $\lambda \ge \omega_I(\xi_I - \xi_H)\sigma_r^2 \min f > 0$ and $D < 1 + 2\omega_I e$. The displayed margins consequently have uniform positive lower bounds. □

**The relative yield compensation.** Subtract the two solved long yields, weighting the taxable yield by $R_\infty = A^M_\infty/A^C_\infty$. Equation (90) gives $\pi^M_\infty - \pi^C_\infty = \omega_H\sigma_r^2\widehat{\Lambda}$, while $\xi^* - R_\infty = -bA^M_\infty$. Therefore

$$y^M_\infty - R_\infty y^C_\infty = A^M_\infty(\omega_H\sigma_r^2\widehat{\Lambda} - b\bar{r}) + \tfrac{1}{2}\sigma_r^2 A^M_\infty(A^C_\infty - A^M_\infty) + \iota^M_\infty - R_\infty\iota^C_\infty.$$

The three terms are the difference in risk costs on the bonds each investor holds alone, relative convexity and bond-specific compensation. On the jointly held block, subtract the lower-tax investor's municipal equality from the household's and use clearing:

$$\omega_H\sigma_r^2\widehat{\Lambda} - b\bar{r} = \frac{\omega_H\sigma_r^2}{1+\sigma_r^2 S_J}\left[\gamma_H\left(\sum_{T\notin J}\frac{A^M_{T-1}Q^M_T}{m_H} - \ell_H\right) - \gamma_I\left(\frac{\xi_I}{m_I}\sum_T A^C_{T-1}Q^C_T - \ell_I\right)\right].$$

Every term, and hence its tax derivative, is independent of $\bar{r}$ on the fixed branch. Common duration compensation per unit of loading cancels from $y^M_\infty - R_\infty y^C_\infty$.

The joint holders' risk-premium coefficient is also independent of $\bar{r}$. With $\Gamma := m_H/\gamma_H + m_I/\gamma_I$, the value of $\Gamma_T^M$ on the jointly held block, clearing gives

$$\pi_\infty^M = \frac{\sigma_r^2}{\Gamma}\left(\sum_T A_{T-1}^M Q_T^M + \xi_I \sum_T A_{T-1}^C Q_T^C - m_H \ell_H - m_I \ell_I\right).$$

By (90) and the two previous displays, the mean-rate term in $\pi_\infty^C$ is $-b\bar{r}$. Since $1/A_\infty^C = 1-\rho+b$, dividing the taxable yield by its loading cancels $b\bar{r}$ against this term. This establishes $y_\infty^C = A_\infty^C(1-\rho)\bar{r} + O(1)$. The bond-specific terms are $\iota_\infty^M = v_\infty^M(Q_\infty^M/\Gamma - \frac{1}{2})$ and $\iota_\infty^C = v_\infty^C(\gamma_I \xi_I Q_\infty^C/m_I - \frac{1}{2})$.

**Convexity in the long-end correction.** For the illustration (Appendix A.2.8), dividing by $y_\infty^C$ isolates convexity:

$$\Delta\tau_\infty := \frac{y_\infty^M}{y_\infty^C} - R_\infty = R_\infty(1-R_\infty)\chi_\infty + \psi_\infty,$$
$$\chi_\infty = \frac{\frac{1}{2}\sigma_r^2 (A_\infty^C)^2}{y_\infty^C}, \qquad \psi_\infty = \frac{A_\infty^M(\omega_H \sigma_r^2 \widehat{\Lambda} - b\bar{r}) + \iota_\infty^M - R_\infty \iota_\infty^C}{y_\infty^C}.$$

The signed convexity contribution is at most $\chi_\infty/4$, since $R_\infty(1-R_\infty) \le 1/4$. The derivatives of the two terms sum to $\mathcal{P}_i^Y(\infty) - \mathcal{P}_i^A(\infty)$ in (92).

**The tax-mix pass-through across maturities.** The next result supplies the intermediate-maturity profile used above and in Appendix A.4.

**Proposition A.10** (Tax-mix pass-through)**.** *Switch off exposure-response terms in the affine pricing recursions, fixing each maturity's holders and risk-tolerance weights. Put*

$$G_T = \sum_{j=1}^{T} \rho^{T-j}, \qquad \Theta_S(T) = \frac{\sum_{j=1}^{\min(T,T^*-1)} \rho^{T-j}}{G_T}, \qquad \Theta_L(T) = 1 - \Theta_S(T).$$

*The tax-mix pass-throughs $\mathcal{D}_i(T) := -\partial_{\xi_i}\tau_T^A$ are the part of $\mathcal{P}_i^A(T)$ that remains with exposure responses switched off. For the loading-implied tax $\tau_T^A$ above, they are*

$$\mathcal{D}_H(T) = \Theta_S(T) + \omega_H \Theta_L(T), \qquad \mathcal{D}_I(T) = \omega_I \Theta_L(T). \tag{93}$$

*They equal $(1,0)$ before participation and converge monotonically to $(\omega_H, \omega_I)$.*

**Proof.** With exposure responses switched off, the recursions give $A_T^C = G_T$ and

$$A_T^M = \xi_H \sum_{j=1}^{\min(T,T^*-1)} \rho^{T-j} + \xi^* \sum_{j=T^*}^{T} \rho^{T-j}.$$

Differentiating $A_T^M/G_T$ gives (93). For $T \ge T^*$, $\Theta_S(T+1) = \Theta_S(T)\rho G_T/(\rho G_T + 1) < \Theta_S(T)$ and $\Theta_S(T) \to 0$. □

The pass-throughs average $(1,0)$ and $(\omega_H, \omega_I)$, using the shares of household-only and jointly held maturities in the bond's remaining life weighted by rate persistence.

**Price derivatives in the nonlinear model.** On a differentiable nonlinear equilibrium branch, let $\theta$ be a tax parameter, such as an after-tax share $\xi_i$, and put $q_T^a = \partial_\theta p_T^a$. The return derivative $\partial_\theta \mu_T^a$ includes portfolio and covariance responses. Hold the state process, cash rule and bond-specific variances fixed. If the exponential moments and their derivatives share an integrable envelope, differentiating the price recursion (41) gives

$$q_T^a = \mathcal{K}_T^a q_{T-1}^a - \partial_\theta \mu_T^a, \qquad q_0^a = 0, \qquad (\mathcal{K}_T^a q_{T-1}^a)(z) = \mathbb{E}^{\mathbb{Q}_{T,z}^a} q_{T-1}^a(Z'), \tag{94}$$

where $F_z$ is the conditional law of $Z'$ and $d\mathbb{Q}^a_{T,z}/dF_z = e^{p^a_{T-1}(Z')}/\mathbb{E}_z e^{p^a_{T-1}(Z')}$. The identity follows from $\partial_\theta \log \mathbb{E}_z e^{p^a_{T-1}} = \mathbb{E}^{\mathbb{Q}^a_{T,z}} q^a_{T-1}$. The full implied-tax response is

$$\partial_\theta \tau^{Y,\text{impl}}_T = \frac{q^M_T - (y^M_T/y^C_T) q^C_T}{T y^C_T}, \qquad y^C_T > 0. \tag{95}$$

Both price responses therefore enter the implied tax measured in the data.

### A.2.7 Illustration without bond-specific risk

The main-text figures have rate risk alone, common risk aversion $\gamma_H = \gamma_I = 5$, and different liability exposures, $\ell_H \approx 1.053$ and $\ell_I \approx 1.452$. Values below are rounded, but calculations use unrounded inputs. The remaining parameters are $\rho = 0.8$, $\sigma_r = 0.025$, $(\xi_H, \xi_I) = (0.55, 0.78)$, equal unit masses, $\bar{r} = 0.04$, municipal supply 0.1 per maturity and taxable supply 0.03 at each $T \geq 2$. With common risk aversion, the lower-tax investor's balance-sheet advantage comes entirely from its greater liability hedge. The solved allocation has household-only municipal holdings at $T = 1, \ldots, 22$, joint holdings at maturity $k = 23$, and lower-tax-only holdings at $T = 24, \ldots, 30$. The lower-tax investor holds about 47% of maturity 23 and all risky taxable bonds. The analytical price recursions use an affine continuation on this allocation branch. Positive holdings and the inactive KKT inequalities are checked locally around the reference state. The quantitative model retains bond-specific risk and solves state-dependent participation.

**Proposition 2** (Rate-risk-only case)**.** *Consider a feasible affine branch with equal masses, zero bond-specific variance, household-only municipal holdings below maturity $k$, both investors holding maturity $k$, and lower-tax-only holdings above $k$. The lower-tax investor holds risky taxable bonds. Put $d = \partial_z \widehat{D}_H$. If outside exposures and supplies are fixed, then*

$$d = \frac{\xi_I - \xi_H}{(\gamma_H + \gamma_I)\sigma_r^2 A^M_{k-1}} > 0, \qquad \partial_z \widehat{D}_I = -d. \tag{96}$$

*The three municipal responses in (16) follow. Write $b = \gamma_I \sigma_r^2 d$ and suppose $0 < \rho - b < 1$. If $A^M_T$ is increasing through $k$, then $A^M_T/A^C_T$ increases from $\xi_H$ toward $\xi_I$.*

**Proof.** Only the shared municipal quantity changes on the branch. Clearing gives $\partial_z \widehat{D}_H + \partial_z \widehat{D}_I = 0$. Differentiating the two active municipal equalities at $k$ gives $\xi_I - \xi_H = (\gamma_H + \gamma_I)\sigma_r^2 A^M_{k-1} d$, proving (96). Sole-holder conditions give $s^M_T = \xi_H + \gamma_H \sigma_r^2 A^M_{T-1} d$ below $k$ and $s^M_T = \xi_I - b A^M_{T-1}$ above it. At $k$, risk-tolerance weighting cancels the exposure responses and gives $s^M_k = \xi^*$. The taxable response is $s^C_T = 1 - b A^C_{T-1}$.

Let $\alpha = \rho - b$ and $\beta = \rho + \gamma_H \sigma_r^2 d$. The loading recursions become

$$A^C_T = \alpha A^C_{T-1} + 1, \qquad A^M_T = \begin{cases} \beta A^M_{T-1} + \xi_H, & T \leq k, \\ \alpha A^M_{T-1} + \xi_I, & T > k. \end{cases}$$

The first municipal recursion holds at $k$ as well because (96) makes its response equal $\xi^*$. For $T \leq k$, $A^M_T/\xi_H = \sum_{j=0}^{T-1} \beta^j$ and $A^C_T = \sum_{j=0}^{T-1} \alpha^j$. Since $\beta > \alpha > 0$, their ratio increases with $T$. To bound it above, put $F_T = A^M_T - \xi_I A^C_T$. For $T \leq k$,

$$F_T = \alpha F_{T-1} + (\xi_I - \xi_H)\left(\frac{A^M_{T-1}}{A^M_{k-1}} - 1\right).$$

The term in parentheses is nonpositive, so $F_T < 0$ from $T = 1$. For $T > k$, $F_T = \alpha F_{T-1}$. Moreover, $A^M_T/A^C_T$ is a positive weighted average of its preceding value and $\xi_I$, so it continues to increase. As $T \to \infty$, $A^C_T \to 1/(1-\alpha) > 0$ and $F_T \to 0$, proving the limit $\xi_I$. □

**Numerical and tax checks.** The solved branch has $d \approx 11.24$, $\widehat{D}_H \approx 4.47$, $\widehat{D}_I \approx 3.57$ and $\rho - b \approx 0.765$. With common risk aversion, the lower-tax investor's smaller net exposure gives it the lower unit cost of

rate risk. The active portfolio equalities and market clearing hold to numerical precision, and inactive inequalities are strict. Independent portfolio optimization reproduces the return responses, and direct finite sums reproduce the loading recursion. At $T = 30$, the return gap is about 201 basis points, comprising 92 of exemption benefit and 109 of rate loading. The duration-cost contribution is about 92 basis points at $T = 22$ and zero from $T = 23$, where the lower-tax investor's conditions are used for both bonds.

For P1 and P2, I differentiate the exposure fixed point, loading recursions, portfolio conditions and yield ratio with respect to one after-tax share at a time, fixing the remaining primitives. Re-solving after small tax changes confirms these derivatives, with the shared maturity and portfolio feasibility preserved. Household pass-through falls from 1 at $T = 1$ to 0.581 at $T = 10$, 0.301 at 23 and 0.230 at 30. Lower-tax pass-through rises from 0 to 0.402, 0.668 and 0.736 at those maturities. These monotone profiles are results for the solved illustration.

For the long-end price continuation, the risk-premium coefficients remain those of the solved finite-supply economy. Since the lower-tax investor's conditions bind in both markets, its exposure response enters both loading recursions through $-b = \gamma_I \sigma_r^2 \partial \widehat{D}_I / \partial z$. Thus $A_\infty^C = 1/(1-\rho+b)$ and $A_\infty^M = \xi_I/(1-\rho+b)$. The duration-risk response changes both loadings, but the common denominator cancels in their ratio. The individual return responses are $s_\infty^C = (1-\rho)/(1-\rho+b) = 0.851$ and $s_\infty^M = \xi_I s_\infty^C = 0.663$, so both responses include the portfolio adjustment. With $\pi_I = \gamma_I \sigma_r^2 \widehat{D}_I$, the limiting yields are

$$y_\infty^C = \bar{r} + \pi_I A_\infty^C - \tfrac{1}{2}\sigma_r^2 (A_\infty^C)^2, \qquad y_\infty^M = \xi_I \bar{r} + \pi_I A_\infty^M - \tfrac{1}{2}\sigma_r^2 (A_\infty^M)^2.$$

Re-solving the finite economy under the same tax perturbations and differentiating these limits gives $\mathcal{P}_H^Y(\infty) = 0.001$ and $\mathcal{P}_I^Y(\infty) = 0.948$, compared with loading-implied limits of 0 and 1. For the empirical unit-response reference, the ratio $A_T^M/G_T$, where $G_T = \sum_{j=0}^{T-1} \rho^j$, increases over the thirty illustrated maturities. Its infinite-maturity limit is $(1-\rho)A_\infty^M$, not $\xi_I$. This reference is distinct from the jointly priced taxable market $C$.

### A.2.8 Small bond-specific risk and overlapping holdings

Adding bond-specific risk to the main-text illustration broadens joint participation without materially changing the return gap. I keep its other primitives fixed and add municipal log-return standard deviation $0.004[1 + 7\sqrt{(T-1)/29}]$ for $2 \le T \le 30$. Taxable bonds have no bond-specific shock, and the one-period claims remain riskless. Both investors now hold municipal maturities 17 through 30, while without the shock they share only maturity 23.

With rate risk alone, joint holdings require both investors to demand exactly the same return. Bond-specific risk adds a separate cost that rises with each investor's holding of that bond, beyond the rate risk of its whole portfolio. This cost can offset a small difference in the returns the investors require for taxes and rate risk. In this illustration, those return requirements are close across long maturities, where rate loadings flatten. A small bond-specific cost therefore supports joint holdings over a broad maturity range while adding little to required returns.

Figure 14 uses the main text's three pricing channels and its municipal and taxable rate-loading contributions. Bond-specific compensation is at most 2.7% of the return gap. At thirty periods, the gap is about 205 basis points, comprising 92 of exemption benefit, 109 of rate loading and 4 of bond-specific compensation. At the last household-only maturity, $T = 16$, the 182-basis-point gap comprises 88 of duration cost, 91 of rate loading and 3 of bond-specific compensation. Removing the shock changes the return gap by at most 3.91 basis points over all thirty maturities, municipal yields by at most 0.33 basis points and taxable yields by at most 0.07 basis points. The shock's main role here is to support overlapping holdings.

**A unique loading solution.** For $a_{\min} = \xi_H \sum_{q=0}^{15} \rho^q$, predecessor loadings on the jointly held block satisfy $A_{T-1}^M \ge a_{\min}$ and $P_J/S_J \le 1/a_{\min}$. The exposure map $\mathcal{F}_J$ of Appendix A.2.4 therefore maps $[0,\infty)$ into $[0, d_{\max}]$, where

$$d_{\max} = \frac{\xi_I - \xi_H}{(\gamma_H + \gamma_I)\sigma_r^2 a_{\min}} < 13.77.$$

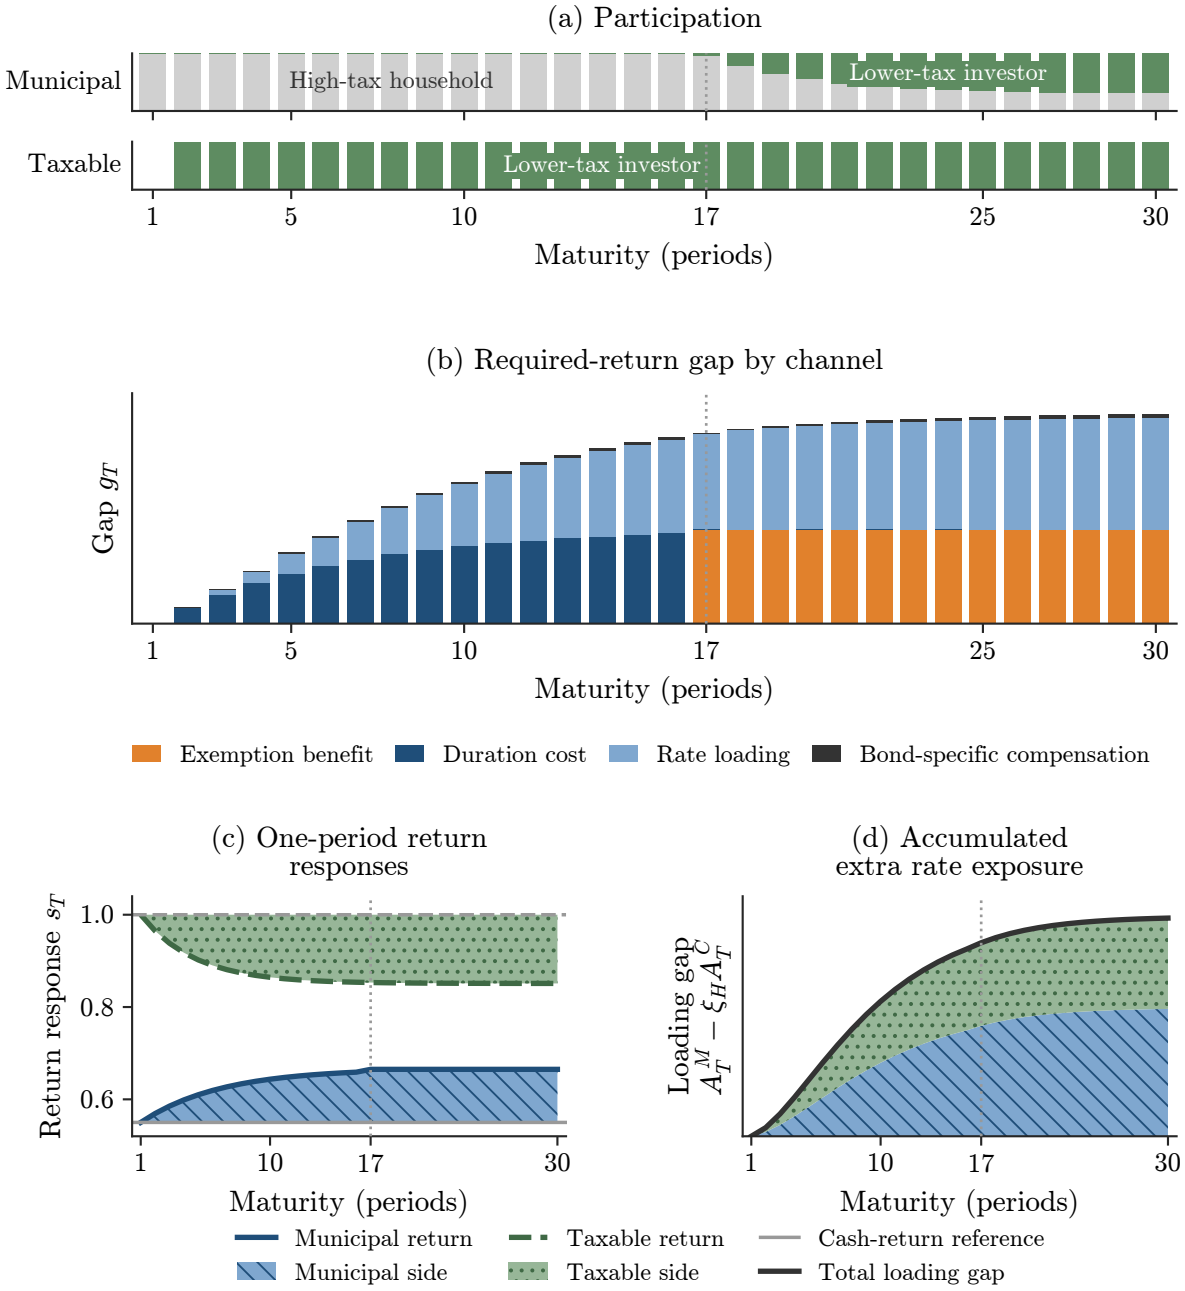



Figure 14: The analytical illustration with bond-specific risk. All other primitives match Appendix A.2.7. Panel (a) shows ownership. Panel (b) gives the three pricing channels of Figure 5, with small bond-specific compensation in gray. Panels (c) and (d) use the grouping of Figure 6: the municipal side combines the tax mix and municipal portfolio response, while the taxable side reflects the taxable portfolio response. Panel (c) shades $s_T^M - \xi_H$ and $1 - s_T^C$. Panel (d) accumulates $s_T^M - \xi_H$ and $\xi_H(1 - s_T^C)$ into the total loading gap. The dotted line marks joint municipal holdings from $T^* = 17$.

Interval bounds give $\sup|\mathcal{F}_J'| < 0.066 < 1$, so the map is a contraction and has a unique fixed point, $d^* \approx 11.20$. The resulting $A_{16}^M \approx 3.147$ is below $\xi^*/(1-\rho) = 3.325$, verifying condition (M) in (87). The left side of condition (S) in (81) is approximately 0.209, below $\rho = 0.8$.

**Levels and the KKT conditions.** The liability exposures above give $\widehat{\Lambda} \approx 4.53$, which satisfies the participation inequalities selecting $T^* = 17$ under Theorem A.1. At the mean rate state, the active KKT conditions, market clearing and the return-gap identity hold to numerical precision, and all inactive portfolio inequalities are strict.

**Pass-through across maturities.** Re-solving the exposure fixed point, loadings and holdings for tax changes gives household pass-through declining from 1 to 0.230 and lower-tax pass-through rising from 0 to 0.739. Lower-tax pass-through is 0.06 at $T = 2$, a maturity the lower-tax investor does not hold. Interval bounds confirm that household pass-through falls and lower-tax pass-through rises at every successive maturity. Their sum ranges from about 0.969 to 1, and the participation cutoff remains $T^* = 17$.

**Long-end continuation.** Continuing the price recursions with the solved economy's exposure responses and risk-premium coefficients, and with $Q_\infty^M = 0.1$ and $v_\infty^M = 0.001024$, gives $A_\infty^M = 3.325$, $A_\infty^C = 4.255$ and $s_\infty^C = 0.851$. The risk-premium coefficients are $(\pi_\infty^M, \pi_\infty^C) = (0.0126, 0.0112)$ and $y_\infty^C = 0.0819$. Loading-implied pass-throughs are $(-0.012, 1.013)$ and yield-implied pass-throughs $(-0.006, 0.959)$. The tax-difference condition of Proposition A.9 has left side $0.506 < 1$. These results include both markets' portfolio responses and convexity, so yield pass-throughs differ from the risk-tolerance weights.

**Equilibrium selection.** I keep an enumerated affine configuration only when its level KKT conditions and clearing, a linear system, have a unique feasible solution. This check applies to each configuration separately, and other feasible cutoffs remain possible (Appendix A.2.4). Full-state nonlinear solutions pass the independent grid checks in Appendix B.6.3, and their maturity patterns are numerical results.

## A.3 Alternative motives for municipal participation

Lower trading costs or better default diversification can induce long-maturity participation without a liability-hedging advantage (Section 4.7). The pricing comparison depends on neutrality: the motive must create no gap relative to the household benchmark when the household holds both bonds. Definition A.1 states this condition for common risks and wealth costs. Separate default shocks require subtracting each market's holder-specific default compensation.

Proposition A.11 separates the participation incentive from its pricing effect. Proposition A.14 gives solved trading-cost and diversification examples, and Proposition A.15 shows why a municipal-specific preference has the opposite pricing implication. Throughout, $g_T = \mu_T^M - \xi_H \mu_T^C$ and $\xi_H < \xi_I$.

**Proposition A.11** (Participation creates the gap)**.** *In the two-investor affine benchmark without bond-specific risk, suppose every non-tax characteristic other than rate risk is neutral and zero at the one-period claim. Each is a risk common to both markets, independent of rate risk, with municipal exposure $\xi_H$ times the taxable exposure, or a cost that reduces after-tax wealth by the same amount in both markets.*

*(i) If the household held every municipal and taxable bond, $g_T = 0$ at every maturity.*

*(ii) Wherever the lower-tax investor holds the maturity-$T$ municipal bond,*

$$g_T \;\geq\; \underbrace{(\xi_I - \xi_H) r}_{\text{exemption benefit}} \;+\; \underbrace{\gamma_I \sigma_r^2 \widehat{D}_I \left(A_{T-1}^M - \xi_H A_{T-1}^C\right)}_{\text{rate loading}}.$$

Part (i) is proved in Appendix A.3.2, and part (ii) follows from Proposition A.12. The rate-loading term is positive on the branch of Proposition 2 when the lower-tax investor's net rate exposure is positive.

### A.3.1 Non-tax required returns and neutral characteristics

Add a differentiable concave term $\Phi_i(x_i)$, in the same after-tax units, to investor $i$'s objective in (35). Dividing each KKT condition (36) by its after-tax share gives, under the normalized municipal–taxable benchmark,

$$\mu_T^M \le \xi_i r + K_{i,T}^M, \qquad \mu_T^C \le r + K_{i,T}^C, \qquad \text{with equality for a positive holding,} \tag{97}$$

where

$$K_{i,T}^a := \gamma_i \operatorname{Cov}_z\big(u_{T,t+1}^a, \widetilde{\Pi}_i(t+1)\big) + \frac{\kappa_i}{\xi_i^a} - \frac{1}{\xi_i^a}\frac{\partial \Phi_i}{\partial x_{i,T}^a}.$$

With additive shocks, $K_{i,T}^a$ splits into components $k_{ij,T}^a$ for rate risk, bond-specific risk, each additional shock or wealth characteristic, and the shorting constraint on cash. Additional shocks are independent of the rate and bond-specific shocks.

A common return characteristic is a unit-variance shock $q$ with nonnegative maturity-$T$ exposures $e_T^M$, $e_T^C$ in both markets, taxed with returns. Its component is $k_{iq,T}^a = c_{iq} e_T^a$, where the cost per unit of exposure, $c_{iq} := \gamma_i \operatorname{Cov}_z(q_{t+1}, \widetilde{\Pi}_i(t+1))$, is the same in both markets. Affine rate risk has $q = -\varepsilon$, $e_T^a = \sigma_r A_{T-1}^a$, $c_{iq} = \gamma_i \sigma_r \widehat{D}_i$ and component $\gamma_i \sigma_r^2 A_{T-1}^a \widehat{D}_i$.

A wealth characteristic changes after-tax wealth by $w_T^a$ per dollar, scaled by investor $i$'s intensity $\lambda_i \ge 0$, as with trading costs. This intensity is unrelated to the feedback $\lambda$ in the proof of Lemma A.1. For $\Phi_i = -\lambda_i \sum_{a,T} w_T^a x_{i,T}^a$, the component is $k_{iw,T}^a = \lambda_i w_T^a / \xi_i^a$. The shorting constraint on cash similarly gives $\kappa_i / \xi_i^a$.

**Definition A.1** (Neutral characteristic)**.** *A common return characteristic is neutral if $e_T^M = \xi_H e_T^C$, and a wealth characteristic if $w_T^M = w_T^C$, at every maturity. For a household that bears it, a neutral characteristic satisfies $k_{H,T}^M = \xi_H k_{H,T}^C$, so it creates no gap when the household holds both bonds.*

**Neutrality for each investor.** A neutral common return characteristic gives $k_{iq,T}^M = c_{iq} e_T^M = c_{iq} \xi_H e_T^C = \xi_H k_{iq,T}^C$ for every investor $i$. A neutral wealth characteristic gives $k_{iw,T}^M = \lambda_i w_T = \xi_i (\lambda_i w_T / \xi_i) = \xi_i k_{iw,T}^C$, so

$$k_{Iw,T}^M - \xi_H k_{Iw,T}^C = (\xi_I - \xi_H) k_{Iw,T}^C. \tag{98}$$

In the affine benchmark where the household alone holds both bonds, rate risk is neutral. Clearing fixes the household's exposure, so differentiation gives $s_T^M = \xi_H$ and $s_T^C = 1$. Starting at $A_1^M = \xi_H$, $A_1^C = 1$, recursion (55) then gives $A_{T-1}^M = \xi_H A_{T-1}^C$ at every maturity by induction, as in Proposition A.6.

Neutrality is stated for the household because the puzzle is defined against the household benchmark. A characteristic that weighs more on municipal bonds would add a premium even there.

### A.3.2 The gap without and with participation

**No one-period return gap when the household holds both bonds.** Suppose the household alone holds both bonds, there is no bond-specific risk and every additional characteristic is neutral and zero at the one-period claim. Both household conditions in (97) bind, so $g_T = \sum_j (k_{Hj,T}^M - \xi_H k_{Hj,T}^C)$. The terms $\kappa_H$ and $\kappa_H / \xi_H$ from the household's shorting constraint on cash cancel. Definition A.1 removes the additional components. In the affine benchmark, the neutrality of rate risk makes the rate contribution $\gamma_H \sigma_r^2 \widehat{D}_H (A_{T-1}^M - \xi_H A_{T-1}^C)$ zero. With a constant cash rate, there is no rate component. At riskless $T = 1$, the same equalities give $\mu_1^M = \xi_H \mu_1^C$, reducing to $\mu_1^M = \xi_H r$ and $\mu_1^C = r$ when the shorting constraint on cash is slack. Hence

$$g_T = 0 \qquad \text{at every maturity.} \tag{99}$$

This is a statement about holding returns. The yield gap also includes the convexity difference in (79), so a zero one-period return gap need not imply a zero yield gap.

**Participation incentive for any motive.** Evaluate components at the household-only allocation of Appendix A.2.3. There the household's binding municipal condition gives $\mu_T^M = \xi_H r + K_{H,T}^M$. The lower-tax

investor's first-order gain from buying the maturity-$T$ municipal bond, $\mu_T^M - \xi_I r - K_{I,T}^M$, is then nonnegative if and only if

$$\Delta_T := K_{H,T}^M - K_{I,T}^M \;\geq\; (\xi_I - \xi_H)\, r. \tag{100}$$

If $r > 0$, $\Delta_1 = 0$ and $\Delta_T$ is nondecreasing in $T$, the maturities satisfying (100) form an upper interval, possibly empty, that excludes the one-period claim. A nondecreasing $\Delta_T$ crosses the maturity-invariant tax difference at most once, and $\Delta_1 = 0 < (\xi_I - \xi_H)r$.

Rate risk gives $\Delta_T = \sigma_r^2 A_{T-1}^M \Lambda$ (Proposition A.3). At that allocation, the advantages from several motives add. This comparison holds prices and portfolios fixed. Equilibrium upper-interval participation follows from Proposition A.13(i) for the default and trading-cost motives and under the conditions of Theorem A.1 for rate risk.

**Proposition A.12** (The gap when the lower-tax investor participates)**.** *In the economy of Appendix A.1.1 with one municipal market, suppose every additional characteristic is a neutral common return or wealth characteristic. If the lower-tax investor holds the maturity-$T$ municipal bond at state $z$, then*

$$g_T \;\geq\; (\xi_I - \xi_H)\, r + g_T^{\text{risk}} + (\xi_I - \xi_H)\Big(\sum_w k_{Iw,T}^C + \frac{\kappa_I}{\xi_I}\Big),$$

$$g_T^{\text{risk}} := \gamma_I \Big[\operatorname{Cov}_z\big(u_{T,t+1}^M, \widetilde{\Pi}_I(t+1)\big) - \xi_H \operatorname{Cov}_z\big(u_{T,t+1}^C, \widetilde{\Pi}_I(t+1)\big)\Big],$$

*where the covariances run over the rate and bond-specific shocks, with equality if this investor also holds the maturity-$T$ taxable bond. In the affine benchmark without bond-specific risk, $g_T^{\text{risk}} = \gamma_I \sigma_r^2 \widehat{D}_I (A_{T-1}^M - \xi_H A_{T-1}^C)$ and the shorting constraint on cash is slack. With nonnegative costs, this gives Proposition A.11(ii).*

**Proof.** The active municipal condition gives $\mu_T^M = \xi_I r + K_{I,T}^M$. Taxable KKT gives $\mu_T^C \leq r + K_{I,T}^C$, with equality when taxable holdings are positive. Subtract $\xi_H$ times this inequality:

$$g_T \;\geq\; (\xi_I - \xi_H) r + \sum_j \big(k_{Ij,T}^M - \xi_H k_{Ij,T}^C\big).$$

Rate and bond-specific components give $g_T^{\text{risk}}$. Neutral common-return components cancel. Neutral wealth costs give $(\xi_I - \xi_H) k_{Iw,T}^C$ by (98). The terms $\kappa_I$ and $\kappa_I/\xi_I$ from the shorting constraint on cash give $(\xi_I - \xi_H)\kappa_I/\xi_I$. □

The third term prices wealth costs and a binding shorting constraint on cash at the lower-tax investor's after-tax share. Nonnegative $g_T^{\text{risk}}$ and nonnegative costs suffice for $g_T \geq (\xi_I - \xi_H) r$.

### A.3.3 Three motives that produce the long-end gap

**Proposition A.13** (Three motives that produce the long-end gap)**.** *Let $r > 0$ and consider one of the three cases below. Cases (b) and (c) are set in the two-investor economy with a constant cash rate, no bond-specific risk and a slack shorting constraint on cash, so the stated characteristic is the only non-tax component. In each case the characteristic is zero at the one-period claim. In (b), investors have no outside exposure to the default shocks.*

- *(a) Liability hedging and risk-bearing capacity: the two-investor affine benchmark with $\Lambda > 0$ at the household-only allocation, without bond-specific risk for parts (i) and (iv).*
- *(b) Default diversification: municipal default risk loads on one unit-variance shock and taxable default risk on another, with correlation $\varrho \in [0,1]$ and exposures $\xi_H e_T$ and $e_T$, where $e_T$ is strictly increasing for $T \geq 1$. Write $X_i^M := \xi_H \sum_T e_T x_{i,T}^M$ and $X_i^C := \xi_i \sum_T e_T x_{i,T}^C$ for investor $i$'s municipal and taxable default exposures, and $k_{id,T}^M := \gamma_i \xi_H e_T (X_i^M + \varrho X_i^C)$ and $k_{id,T}^C := \gamma_i e_T (X_i^C + \varrho X_i^M)$ for its default components. Assume $\gamma_H X_H^{M,\text{hh}} > \varrho\, \gamma_I X_I^{C,\text{hh}}$.*
- *(c) Trading costs: the same round-trip cost $w_T \geq 0$ in both markets, strictly increasing for $T \geq 1$, and trading intensities $\lambda_I < \lambda_H$.*

*Then:*

(i) *the participation advantage* $\Delta_T$ *in* (100) *is zero at* $T = 1$ *and nondecreasing in* $T$*. In (b) and (c), the municipal maturities the lower-tax investor holds in any equilibrium form an upper interval;*

(ii) $g_1 = 0$*: the household anchors the one-period claim;*

(iii) *at every maturity at which the lower-tax investor holds the municipal bond,* $g_T \geq (\xi_I - \xi_H)r + g_T^{\text{risk}}$ *in (a), with* $g_T^{\text{risk}}$ *of Proposition A.12,* $g_T \geq (\xi_I - \xi_H)r$ *in (c), and in (b)*

$$g_T = (\xi_I - \xi_H)r + k^M_{Id,T} - \xi_H k^C_{d,T},$$

*where* $k^C_{d,T}$ *is the default component of any investor who holds the taxable bond, so the gap net of these default components is* $(\xi_I - \xi_H)r$*;*

(iv) *with the household alone holding both bonds,* $g_T = 0$ *in (a) and (c), and* $g_T = k^M_{Hd,T} - \xi_H k^C_{Hd,T}$ *in (b).*

**Proof.** (i) At household-only holdings, (a) gives $\Delta_T = \sigma_r^2 A^M_{T-1}\Lambda$ with $A^M_{T-1} = \xi_H \sum_{j=0}^{T-2} \rho^j$ (Proposition A.6). Case (b) gives $X^C_H = X^M_I = 0$ and $\Delta_T = \xi_H e_T(\gamma_H X^{M,\text{hh}}_H - \varrho\,\gamma_I X^{C,\text{hh}}_I)$. Case (c) gives $\Delta_T = (\lambda_H - \lambda_I)w_T$. All vanish at $T = 1$ and increase weakly. For equilibria in case (b), define investor $i$'s unit municipal-default cost $c_i := \gamma_i(X^M_i + \varrho X^C_i)$. The lower-tax investor's required municipal return is no greater than the household's exactly when $\xi_H e_T(c_H - c_I) \geq (\xi_I - \xi_H)r$. If the lower-tax investor holds maturity $T$, this comparison holds there and $c_H > c_I$. Strictly increasing $e_T$ makes it strict beyond $T$, precluding household holdings and forcing lower-tax-investor holdings. Case (c) uses $(\lambda_H - \lambda_I)w_T$ identically.

(ii) Every non-tax component vanishes at the one-period claim, so the municipal conditions reduce to $\mu^M_1 \leq \xi_i r$. Positive supply gives $\mu^M_1 = \xi_H r$, while $\mu^C_1 = r$, as in Proposition A.2.

(iii) Apply Proposition A.12 to (a), with no additional characteristic, and (c), with $g_T^{\text{risk}} = 0$ and cost $(\xi_I - \xi_H)\lambda_I w_T/\xi_I \geq 0$. In (b), active conditions give $\mu^M_T = \xi_I r + k^M_{Id,T}$ and $\mu^C_T = r + k^C_{d,T}$. With a constant rate, every taxable holder has the same default component. Subtraction proves the expression.

(iv) In (a) and (c) this is (99). In (b), the household's conditions give $\mu^M_T = \xi_H r + k^M_{Hd,T}$ and $\mu^C_T = r + k^C_{Hd,T}$. □

In (a), $\Lambda$ is positive when liability hedges or lower risk aversion reduce the lower-tax investor's duration costs. In (b), households bear all municipal default risk at the household-only allocation. The lower-tax investor's taxable holdings are exposed to it only through $\varrho$, so this investor can add municipal default risk more cheaply. Independent shocks guarantee this advantage. Raw gaps also contain holder-specific default-premium differences driven by relative risk supplies. Appendix C.1.6 adjusts these with credit default swap (CDS) spreads. In (c), trading costs reduce after-tax wealth (Appendix C.1.5). For matched long bonds at 26 to 30 years, institutional round-trip costs are 0.53% of price in both markets (Table 25), supporting equal costs. With combined motives, the decomposition of Proposition A.12 still applies. Its tax-difference bound requires the remaining risk term $g_T^{\text{risk}}$ to be nonnegative at the resulting equilibrium. Separate default shocks are treated as in (b).

### A.3.4 Two solved examples without a duration motive

Consider two investors and three maturities with constant $r > 0$ and no bond-specific risk, eliminating rate-risk advantages. Both investors have risk aversion $\gamma$, masses $m_H$, $m_I$, after-tax shares $\xi_H < \xi_I$ and slack shorting constraints on cash. Supplies are $Q^M_T$ for $T = 1, 2, 3$ and $Q^C_T$ for $T = 2, 3$. Write $M1$, $M2$, $M3$ for the municipal bonds and $C2$, $C3$ for the longer taxable bonds, indexed by maturity. Each example adds only its stated non-tax characteristic.

**Proposition A.14** (Participation for diversification or lower trading costs)**.** *Consider the economy just described.*

(a) *Default diversification. Municipal and taxable default risk load on independent unit-variance shocks, with exposures* $\xi_H e_T$ *and* $e_T$ *and* $0 = e_1 < e_2 < e_3$*. Let* $X^M_H := \xi_H e_2 Q^M_2/m_H$ *and* $X^M_I := \xi_H e_3 Q^M_3/m_I$ *be the investors' municipal default exposures when the household holds* $M1$ *and* $M2$ *and the lower-tax investor holds* $M3$*, and let*

$$X^C := \frac{e_2 Q^C_2 + e_3 Q^C_3}{m_H/\xi_H + m_I/\xi_I}.$$

*There is an equilibrium with this municipal allocation if and only if*

$$\gamma\,\xi_H e_2(X_H^M - X_I^M) \;\le\; (\xi_I - \xi_H)r \;\le\; \gamma\,\xi_H e_3(X_H^M - X_I^M).$$

*Both investors then hold taxable bonds, each with taxable default exposure $X^C$, and*

$$g_1 = 0, \qquad g_2 = \gamma\,\xi_H e_2(X_H^M - X^C), \qquad g_3 = (\xi_I - \xi_H)r + \gamma\,\xi_H e_3(X_I^M - X^C).$$

*Net of the default components of the investors who hold each bond, the gaps are 0, 0 and $(\xi_I - \xi_H)r$.*

*(b) Trading costs. Each investor pays the after-tax cost $\lambda_i w_T$ per period on either bond of maturity $T$, with $0 = w_1 < w_2 < w_3$ and $0 \le \lambda_I < \lambda_H$. If $(\lambda_H - \lambda_I)w_2 < (\xi_I - \xi_H)r \le (\lambda_H - \lambda_I)w_3$, there is an equilibrium in which the household holds $M1$ and $M2$, the lower-tax investor holds $M3$, $C2$ and $C3$, and*

$$g_1 = 0, \qquad g_2 = \Big(\lambda_H - \frac{\xi_H}{\xi_I}\lambda_I\Big)w_2, \qquad g_3 = (\xi_I - \xi_H)\Big(r + \frac{\lambda_I w_3}{\xi_I}\Big),$$

*with $0 < g_2 < g_3$. The ratio $\mu_T^M/\mu_T^C$ equals $\xi_H$ at $T = 1$, lies strictly between $\xi_H$ and $\xi_I$ at $T = 2$ and equals $\xi_I$ at $T = 3$.*

**Proof.** (a) Independent default shocks separate market risks. Investor $i$'s KKT conditions are $\mu_T^M \le \xi_i r + \gamma\xi_H e_T X_i^M$ and $\mu_T^C \le r + \gamma e_T X_i^C$, for municipal and taxable exposures $X_i^M$, $X_i^C$. Lower taxable exposure implies lower required return and profitable purchases, so both investors hold taxable bonds with equal exposures. Clearing $\sum_i m_i X_i^C/\xi_i = e_2 Q_2^C + e_3 Q_3^C$ gives $X_H^C = X_I^C = X^C$ and $\mu_T^C = r + \gamma e_T X^C$. Each investor's split of its taxable exposure across the two maturities need not be unique. Supply-proportional holdings implement these exposures and clear both maturities. Active municipal conditions give $\mu_1^M = \xi_H r$, $\mu_2^M = \xi_H r + \gamma\xi_H e_2 X_H^M$ and $\mu_3^M = \xi_I r + \gamma\xi_H e_3 X_I^M$. The inactive KKT conditions require $\xi_H r < \xi_I r$ at $M1$, the left threshold inequality at $M2$ and the right at $M3$. Subtracting returns gives the gaps. Remove the holder-specific municipal component $\gamma\xi_H e_T X_i^M$ and the taxable component $\gamma e_T X^C$ to leave $\xi_H r - \xi_H r$ at $T = 2$ and $\xi_I r - \xi_H r$ at $T = 3$.

(b) Investor $i$ requires $\xi_i r + \lambda_i w_T$ on municipal bonds and $r + \lambda_i w_T/\xi_i$ on taxable bonds. With slack shorting constraints on cash, each bond's return equals the minimum eligible requirement. Since $\lambda_I/\xi_I < \lambda_H/\xi_H$, the lower-tax investor holds all taxable bonds. It holds the maturity-$T$ municipal bond when $\xi_I r + \lambda_I w_T \le \xi_H r + \lambda_H w_T$, here only at $T = 3$. Subtract returns to obtain the gaps and

$$g_2 = (\lambda_H - \lambda_I)w_2 + (\xi_I - \xi_H)\frac{\lambda_I w_2}{\xi_I} < (\xi_I - \xi_H)\Big(r + \frac{\lambda_I w_2}{\xi_I}\Big) \le g_3.$$

At $T = 2$ the ratio is below $\xi_I$ because the household's required municipal return is below the lower-tax investor's. At $T = 3$ it is $(\xi_I r + \lambda_I w_3)/(r + \lambda_I w_3/\xi_I) = \xi_I$. □

Case (a) implements default diversification (Proposition A.13), requiring $X_H^M > X_I^M$. The intermediate raw gap $\gamma\xi_H e_2(X_H^M - X^C)$ is a default-premium difference. With $\varrho \in (0,1)$, the same allocation, where it is an equilibrium, has default-net gaps 0, 0, $(\xi_I - \xi_H)r$ (Proposition A.13(b)). Case (b)'s intermediate gap reflects higher household trading costs, independent of other holdings.

**Examples with rate risk.** With the affine rate process $z' = \rho z + \sigma_r\varepsilon$ added, the examples use $\xi_H = 0.6$, $\xi_I = 0.8$, $\rho = 0.8$, $\bar{r} = 4\%$ and small municipal bond-specific variance. They set $\Lambda = 0$ at household-only holdings to remove the initial rate-risk advantage. For each motive, I solve 120 economies with randomly drawn other parameters, including $\varrho \in [0, 0.9]$, and screen them for a participation incentive at household-only holdings. In all 120 trading-cost economies, the lower-tax investor holds only the longest municipal bond, with $0 = g_1 < g_2 \le g_3$ and $g_3 \ge (\xi_I - \xi_H)r$. With default diversification, it does so in 72 economies and holds no municipal bond in 48. In the 72, the default-net gap is the tax difference plus Proposition A.12's rate and bond-specific terms. The bond-specific term is below 0.01 bp. The default-net long gap is at least the tax difference in 68 economies. In the other four, negative rate terms lower it below the tax difference by at most 0.09 bp. Rate terms across both motives span −0.09 to 1.7 bp.

### A.3.5 A municipal-specific benefit

A municipal-only benefit $\zeta_T > 0$, from a mandate, issuer relationship or favorable regulation at matched maturity and credit, lowers the required municipal return without changing the taxable alternative. It therefore violates neutrality. For the lower-tax investor, it gives the participation advantage $\Delta_T = \zeta_T$ but reduces the return gap by $\zeta_T$. A benefit that grows with maturity can thus induce long-end participation while lowering municipal returns. Unlike default compensation, it has no taxable counterpart to subtract.

**Proposition A.15** (Same participation, opposite pricing)**.** *Consider the two-investor economy without rate risk or bond-specific risk, with a constant cash rate $r > 0$, a slack shorting constraint on cash and no additional household characteristic. A fixed per-dollar benefit lowers the lower-tax investor's required municipal return by $\zeta_T \geq 0$, with $\zeta_1 = 0$, $\zeta_T$ nondecreasing and $\zeta_T < \xi_H r$. Compare two benefits. A neutral benefit is paid and taxed with returns: $\zeta_T$ on the municipal bond and $\zeta_T/\xi_H$ before tax on the taxable bond. It also lowers the investor's required taxable return by the household-equivalent $\zeta_T/\xi_H$. A municipal-specific benefit leaves the required taxable return unchanged.*

*(i) In both cases the lower-tax investor holds the maturity-$T$ municipal bond only if $\zeta_T \geq (\xi_I - \xi_H)r$, so participation begins at the same maturity.*

*(ii) With the neutral benefit, $\mu_T^C = r - \zeta_T/\xi_H$ and*

$$g_T = \min\{\zeta_T,\ (\xi_I - \xi_H)r\} \geq 0,$$

*so the gap equals the full difference in exemption benefits wherever the lower-tax investor participates. The ratio $\mu_T^M/\mu_T^C$ rises with $\zeta_T$.*

*(iii) With the municipal-specific benefit, $\mu_T^C = r$ and*

$$g_T = \min\{0,\ (\xi_I - \xi_H)r - \zeta_T\} \leq 0,$$

*and the ratio $\mu_T^M/\mu_T^C = \min\{\xi_H,\ \xi_I - \zeta_T/r\}$ is weakly below $\xi_H$ where the lower-tax investor participates and strictly below it wherever $\zeta_T > (\xi_I - \xi_H)r$.*

*Because the rate is constant, each yield averages the one-period returns $\mu_1^a, \ldots, \mu_T^a$. The yield gap $y_T^M - \xi_H y_T^C$ therefore has the sign of the $g_j$, and the yield-implied tax is nonincreasing in maturity in case (ii) and nondecreasing in case (iii).*

**Proof.** As in Proposition A.14(b), minimum required returns determine holdings. The household requires $\xi_H r$ and $r$. The lower-tax investor requires $\xi_I r - \zeta_T$ and either $r - \zeta_T/\xi_H$ or $r$. This proves (i). With neutrality, $g_T = \mu_T^M - \xi_H r + \zeta_T$ equals $\zeta_T$ when $\xi_H r$ is minimal, otherwise $(\xi_I - \xi_H)r$. The ratios $\xi_H r/(r - \zeta_T/\xi_H)$ below the threshold and $(\xi_I r - \zeta_T)/(r - \zeta_T/\xi_H)$ above increase in $\zeta_T$ since $\xi_I > \xi_H$. A municipal-only benefit gives $g_T = \mu_T^M - \xi_H r$. Since $\zeta_T$ is nondecreasing, $\mu_j^M/\mu_j^C$ is nondecreasing in $j$ in (ii) and nonincreasing in (iii). Adding a maturity whose ratio is at least (in (iii), at most) the running yield ratio cannot lower (raise) that ratio, which proves the yield statements. □

## A.4 State-specific exemptions across municipal markets

A household resident in state $s'$ pays its home state's income tax on interest from a municipal bond issued by another state $s \neq s'$ but not on a home-state municipal bond. Institutional types in the quantitative model have no state-income-tax wedge on either bond, so state taxation is a household-specific channel (Appendix D.2). Statutory top state rates vary across states, with a supply-weighted mean of 5.8%. This cross-state variation is both a calibration input and P1's source of identification.

**Multi-state aggregation.** The multi-state model replaces $\xi_i$ with $\xi_{i,s} = \xi_i^C/\xi_{i,s}^M$, which depends on the issuing state $s$ and on the investor, whose type includes a household's state of residence. For state $s$'s market $M_s$, put $\Gamma_T^{M_s} := \sum_{i \in \mathcal{U}_T^{M_s}} m_i/\gamma_i$ and $\omega_{i,T}^s := (m_i/\gamma_i)/\Gamma_T^{M_s}$, where $\mathcal{U}_T^{M_s}$ is the set of holders of the maturity-$T$

bond. The exact KKT aggregation (69) applies state by state, with $\xi_{ij} = \xi_{i,s}$. Applying (71) coordinate by coordinate gives, in the affine construction,

$$\begin{aligned} s_T^{M_s} = & \sum_{i \in \mathcal{U}_T^{M_s}} \omega_{i,T}^s \xi_{i,s} + \frac{\sigma_r^2 A_{T-1}^{M_s}}{\Gamma_T^{M_s}} \sum_{i \in \mathcal{U}_T^{M_s}} m_i \partial_z \widehat{D}_i \\ & + \frac{v_{M_s,T}}{\Gamma_T^{M_s}} \sum_{i \in \mathcal{U}_T^{M_s}} m_i \xi_{i,s}^M \partial_z x_{i,M_s}(T). \end{aligned} \tag{101}$$

Here each investor uses one payoff normalization across all markets, with after-tax shares $\xi_i^a$ as in (32). Its net rate exposure in these units is $\widehat{D}_i = \sum_{a,T} \xi_i^a x_{i,T}^a A_{T-1}^a - \ell_i$. When the municipal after-tax shares all equal one, this reduces to (50). The final term in (101) vanishes under common active municipal after-tax shares, a single active holder or zero bond-specific variance, as for the taxable counterpart before (74). Unweighted clearing alone does not remove it.

Proposition A.16 is the state-tax counterpart of Proposition A.10. Its implied state tax is the model counterpart of (24), and $T_s^*$ is the first maturity at which investors other than the resident household hold state $s$'s bonds.

**Proposition A.16** (State-tax anchor and tax-mix pass-through)**.** *If the household $H_s$ resident in state $s$ uniquely minimizes $\xi_{i,s}$, the shorting constraint on cash is slack and the one-period municipal boundary is deterministic with $r > 0$, then $y_1^{M_s}/y_1^C = \xi_{H_s,s}$. Under the multiplicative mapping $\xi_{H_s,s} = \xi^{\text{fed}}(1 - \tau_s^{\text{st}})$, with federal after-tax share $\xi^{\text{fed}}$ and state income-tax rate $\tau_s^{\text{st}}$, define*

$$\tau_{s,T}^{\text{st,impl}} = 1 - \frac{y_T^{M_s}/y_T^C}{\xi^{\text{fed}}}, \qquad \tau_{s,T}^{A,\text{st}} = 1 - \frac{A_T^{M_s}/A_T^C}{\xi^{\text{fed}}}. \tag{102}$$

*Then $\tau_{s,1}^{\text{st,impl}} = \tau_s^{\text{st}}$, with unit pass-through. Switch off exposure responses in the affine recursions, and suppose only the resident household's effective after-tax share changes, with its active weight one below $T_s^*$ and constant at $\omega_{H_s}^s$ thereafter. Then the tax-mix derivative of the loading-implied state tax $\tau_{s,T}^{A,\text{st}}$ with respect to $\tau_s^{\text{st}}$ is $\Theta_S^s(T) + \omega_{H_s}^s \Theta_L^s(T)$, where $\Theta_S^s, \Theta_L^s$ are the weights in (93) with $T_s^*$ in place of $T^*$.*

**Proof.** The boundary return is $\min_i \xi_{i,s} r$ (Proposition A.2). Equation (102) gives the anchor. Applying (93) with $T_s^*$ in place of $T^*$ gives the tax-mix derivative, because $\partial \xi_{H_s,s}/\partial \tau_s^{\text{st}} = -\xi^{\text{fed}}$ cancels the factor $1/\xi^{\text{fed}}$ in (102). □

Changes to out-of-state after-tax shares or other clienteles add their derivatives. Other statutory mappings require a corresponding implied-tax transformation. The primitive boundary remains $y_1^{M_s}/y_1^C = \xi_{H_s,s}$. Reported state-tax responses use the full multi-state model, which supplies P1's full-yield pass-through profile, including level and exposure corrections.

## A.5 Participation and tax pass-through with multiple investors

Extend the two-investor affine benchmark of Appendix A.2 to $N$ investor types with slack shorting constraints on cash. Order after-tax shares $\xi_1 < \cdots < \xi_N$, so type 1 is the household with the largest exemption benefit. Without bond-specific risk, subtracting type 1's active municipal equality from type $j > 1$'s municipal condition gives its first-order purchase incentive (Appendix A.2.3). Type $j$ has a nonnegative incentive at maturity $T$ when

$$\underbrace{\sigma_r^2 A_{T-1}^M \widehat{\Lambda}_j}_{\text{duration benefit}} \geq \underbrace{(\xi_j - \xi_1) r}_{\text{difference in exemption benefits}}, \qquad \widehat{\Lambda}_j := \gamma_1 \widehat{D}_1 - \gamma_j \widehat{D}_j. \tag{103}$$

At fixed net rate exposures the duration benefit grows with the bond's rate loading, while the difference in exemption benefits is flat in maturity. Positive bond-specific risk adds the incumbent's bond-specific compensation to the comparison, as in the KKT system (46) and the participation index (62).

**Assumption A.2** (Long-end participation, clientele by clientele)**.** *Type 1 is active at every maturity, so* $T_1^* = 1$*. For each type* $j > 1$*,* $\widehat{\Lambda}_j > 0$*, and its active municipal set is an upper maturity interval* $[T_j^*, T_{\max}]$ *with* $T_j^* \geq 2$*. The thresholds need not be ordered by tax advantage.*

**Many-clientele tax-mix pass-through.** To isolate the tax mix under the nested holding intervals of Assumption A.2, switch off exposure-response terms and fix each maturity's holders and risk-tolerance weights, as in Proposition A.10. Let $\omega_{k,j}$ be type $k$'s share of the holders' risk tolerance at municipal maturity $j$, zero when it does not hold that bond. Then $A_T^C = G_T$ and $A_T^M = \sum_{j \leq T} \rho^{T-j} \sum_k \omega_{k,j} \xi_k$. Differentiating $\tau_T^A = 1 - A_T^M / G_T$ in $\xi_i$ gives its tax-mix pass-through $-\partial \tau_T^A / \partial \xi_i$,

$$\mathcal{D}_i(T) = \frac{\sum_{j=T_i^*}^{T} \rho^{T-j} \omega_{i,j}}{\sum_{j=1}^{T} \rho^{T-j}}, \tag{104}$$

with an empty numerator equal to zero. It is positive after participation begins and converges to $\omega_i$ if the active weight is eventually constant at $\omega_i$, because the finite pre-participation block receives vanishing weight at the long end. Participation in order of increasing after-tax shares makes the tax mix, the clientele-weighted after-tax share, nondecreasing in maturity: adding an after-tax share no smaller than the incumbents' does not lower their risk-tolerance-weighted mean. Such ordered participation does not by itself order total slopes, loading ratios or yields, since other terms in (101) and the taxable leg remain unrestricted. Clientele-specific pass-through and endpoint comparisons need no tax ordering.

Under Assumption A.2, the household anchors the riskless boundary, and (104) gives each clientele's tax-mix derivative. Full yield responses also include relative compensation and convexity (Appendix A.2.5). The portfolio-risk response can transmit a later clientele's tax to shorter risky bonds it never holds. P1 concerns the always-active household, and P2 a later-participating clientele. Assumption A.2 restricts only this affine extension. The calibration does not impose it (Section 5).

### A.6 Applications to other bond markets

The mechanism requires common portfolio risk, heterogeneous relative taxes and a shorting friction. For a tax-favored bond and a comparable fully taxed bond, $\xi_i = (1-\tau_i^{\text{tax}})/(1-\tau_i^{\text{fav}})$, with investor $i$'s tax rates on the taxed and favored bonds. Symmetric taxation gives $\xi_i = 1$ and removes the relative tax incentive. The response decomposition (69) applies when after-tax shares differ. Positive taxable forwards and an increasing relative forward ratio imply a declining yield-implied tax (Proposition A.8). Condition (S) orders the affine loading ratio (Proposition A.7).

Applications include state-exempt treasury versus corporate interest; state-exempt Federal Home Loan Bank and Farm Credit issues versus Fannie Mae and Freddie Mac issues; treaty-dependent withholding on cross-border sovereign debt; and the inflation accrual on treasury inflation-protected securities (TIPS), taxed before payment in taxable accounts. Taxable accounts, individual retirement accounts (IRAs) and pensions also assign different after-tax payoffs to the same security, which retains one market price. In each application, the maturity profile depends on which clienteles hold the bonds. Municipal bonds combine a large federal wedge with tight shorting restrictions.

## B Quantitative model, solution and calibration

This appendix explains how the quantitative economy is constructed, solved and calibrated. I first describe the investors, markets and inputs measured from data. I then solve for portfolios and prices at given coefficients, before explaining which coefficients are estimated, the moments used to estimate them and the resulting fit. Numerical implementation and accuracy checks come last. Throughout, one model period is two years, and calendar maturities are stated in years.

## B.1 The economy, investors and measured inputs

**Markets and dated economies.** The application has 45 municipal markets and one taxable market. Each of the 43 states with a personal income tax and the District of Columbia has a separate municipal market. The seven states without one, Alaska, Florida, Nevada, South Dakota, Texas, Washington and Wyoming, share a market because investors treat their bonds alike in the model. Each municipal market has fifteen maturities in two-year steps from two to thirty years. The taxable market has fourteen risky maturities, four through thirty years, and cash is the one-period taxable claim. There are therefore $45 \times 15 + 14 = 689$ bond positions to clear at each rate state. Cash and two-year municipal bonds are riskless. The taxable market pools treasury and corporate supply, is taxed like corporate interest and has its yields compared with the corporate curve.

The measured inputs vary across six three-year blocks, 2000–2002 through 2015–2017, centered on the 2001, 2004, 2007, 2010, 2013 and 2016 Survey of Consumer Finances (SCF) waves. One set of estimated coefficients applies to all six economies. To estimate these coefficients, I first combine the municipal markets into one national market, using supply weights to average household tax treatment across states. This smaller economy has 27 investor types and 29 bond positions. After estimation, I restore all 45 municipal markets and solve again with the coefficients fixed. Municipal yields, the puzzle and the empirical tax-pass-through coefficients are not calibration targets.

**Investors, portfolio assets and budgets.** The six investor groups are high-tax households (federal marginal rates of at least 35%), lower-tax households, property-and-casualty (P&C) insurers, life insurers, banks and a tax-neutral residual group. Households have seven balance-sheet types within each tax group in each municipal market. P&C insurers have two exposure types, life insurers three, banks one and the residual group seven. Thus the full economy has $45 \times 2 \times 7 + 2 + 3 + 1 + 7 = 643$ types. Table 3 reports their portfolio assets and tax benefits in 2015–2017.

For type $i$, baseline portfolio assets $m_i$ are its debt securities and cash in the Financial Accounts. A measured fraction $W_i$ is available for the modeled investments, giving a total budget $m_i W_i$. These investments are municipal bonds, corporate bonds excluding asset-backed securities, treasury notes and bonds, and cash excluding bank cash. Positions $x_i$ are measured per dollar of baseline portfolio assets, so aggregate bond holdings are $m_i x_i$ and cash is $m_i x_i^0$, where $x_i^0 = W_i - \mathbf{1}' x_i$. The amount $m_i$ scales the investor's risk-bearing capacity, while $W_i$ limits the fraction invested in these markets. Neither is chosen to fit yields.

Each type also has an outside rate exposure $h_i = -\ell_i$. Positive $h_i$ adds exposure to losses when rates rise. Negative $h_i$ offsets those losses, as when a rise in rates lowers the value of liabilities. It is measured in the same units as a bond's rate loading: an exposure $h$ carries the rate risk of investing the type's entire baseline portfolio assets in a bond whose log price falls by $h$ per unit rise in the rate state. Balance-sheet amounts are measured directly, while the coefficients translating them into these exposures are estimated below.

**Which parameters belong to which investors.** Table 7 gives the parameterization of each group. Within a state and household tax group, and within each institutional sector, types share tax treatment and differ in outside exposure. For households, measured balance sheets determine the differences, using three common estimated coefficients. For insurers, the estimated exposure levels and portfolio-asset shares describe heterogeneity within the sector. All types in a sector have the same measured budget fraction $W_i$. Table 8 separates the common rate parameters measured from data from the common risk coefficients fitted jointly with investor exposures.

**Measured inputs and after-tax shares.** Tax law fixes after-tax shares, Financial Accounts stocks fix portfolio assets, budgets and national supplies, and the SCF supplies household balance sheets. Municipal state shares and maturity profiles come from outstanding bonds. Table 9 gives the sources, and Tables 11 and 10 report the dated quantities and balance-sheet ratios. These inputs remain fixed during estimation. The coefficients fitted below determine risk costs and outside exposures, rather than changing measured amounts or tax rules.

Table 7: Investor types and the parameters governing their outside exposure

| Investor group | Outside rate exposure $h_i = -\ell_i$ | Division of measured group portfolio assets |
|---|---|---|
| High-tax households 7 types per market | $h_k = a_H + (b_A Out_k - b_D Debt_k)/2$. The intercept $a_H$ is common exposure, while $b_A$ and $b_D$ translate outside assets and debt into rate risk. | Measured SCF portfolio-asset weights divide the group into seven balance-sheet bins. Each bin is then allocated across states by municipal supply shares. |
| Lower-tax households 7 types per market | The same mapping and coefficients, using the outside-asset and debt ratios measured separately for this tax group. | The same construction, using this group's own SCF bin weights. The two household groups share the measured household budget fraction. |
| P&C insurers 2 types | $a_P$ and $a_P + d_P$. The fitted exposures are $-1.96$ and 1.07. The coefficient $d_P$ is the exposure span. | Shares $w_P$ and $1 - w_P$, fitted at 8.0% and 92.0%. These shares divide the observed sector total, rather than changing it. |
| Life insurers 3 types | $a_L$, $a_L + t_L d_L$ and $a_L + d_L$. The fitted exposures are $-5.39$, 0.04 and 0.91. The coefficient $t_L$ locates the middle type within span $d_L$. | Low share $w_L$, middle share $(1 - w_L)u_L$ and the remainder for the high type. Fitted shares are 26.2%, 1.2% and 72.6%. |
| Banks 1 type | A single fitted exposure $h_B = 1.20$. | All measured bank portfolio assets. Banks also bear the measured financing-tax cost on municipal holdings. |
| Tax-neutral residual 7 types | A normal exposure distribution with fitted mean $\mu_R = 1.07$ and standard deviation $s_R = 1.72$. | Seven fixed Gauss–Hermite weights divide the observed residual-group total. Only the distribution's mean and dispersion are estimated. |

*Notes.* Positive $h_i$ adds rate risk and negative $h_i$ provides a liability hedge. Low, middle and high insurer types refer to this exposure, not their tax rates. The measured household ratios $Out_k$ and $Debt_k$ are outside assets and debt divided by the portfolio base (Table 10). All exposure and mixture coefficients are common across the six economies. Table 12 reports their estimates. Household bin and geographic weights and the residual quadrature weights are fixed before estimation.

Table 8: Common rate and risk parameters

| Symbol | Economic role | How its value is determined | Value |
|---|---|---|---|
| $\bar{r}$ | Central short-rate level | Observed short-yield mean | 0.06041 |
| $\rho$ | Rate-state persistence | Historical short-rate estimate | 0.7292 |
| $\sigma_r$ | Rate-state innovation standard deviation | Historical short-rate estimate | 0.04278 |
| $\gamma$ | Common penalty on portfolio variance | Joint calibration | 9.697 |
| $v_{30}$ | Thirty-year bond's two-year bond-specific variance | Joint calibration | 0.000722 |
| $p_v$ | Maturity exponent of bond-specific variance | Joint calibration | 1.260 |

*Notes.* All values use the model's two-year period. The common state process is $z' = \rho z + \sigma_r \varepsilon$, with $\varepsilon \sim N(0,1)$, and $r(z) = \max\{\bar{r} + z, 0\}$. The central level is twice the 2000–2017 mean annualized two-year corporate zero yield. The annual short-rate estimates are persistence 0.8539 and innovation standard deviation 0.01627, transformed to the two-year state. The first three parameters are fixed before the joint calibration. The last three are common to all investors and both markets, in normalized payoff units: $\gamma_i = \gamma$ and $v_n = v_{30}(n/30)^{p_v}$ for calendar maturities $n = 4, 6, \ldots, 30$, with $v_2 = 0$. A larger $\gamma$ raises the cost of risk. The small $v_n$ limits concentrated holdings of otherwise similar bonds. Table 12 reports all seventeen jointly fitted coefficients.

Table 9: Inputs measured before calibration

| Input | Source and construction |
|---|---|
| Portfolio assets, budgets and ownership | Financial Accounts vintage of September 8, 2023. Fund holdings are assigned to their owners, nonprofits are separated and sector stocks are adjusted for tax-exempt holdings. SCF direct and fund holdings split household municipal ownership by tax group. Tax-exempt interest flows are not used as ownership stocks. |
| Household balance sheets | Each SCF wave supplies the seven equal-portfolio-asset bins within each tax group. Outside assets and debt are divided by the portfolio base before bin averaging. |
| Household taxes | Dated federal tax calculations and annual state-rate schedules, applied to SCF tax units. For 209 of 792 state-years, missing lower-tax filing-status cutoffs or related bracket rules are imputed from nearest-year thresholds or simple proxies. Published rates and top rates are unchanged. |
| Institutional taxes | Statutory P&C and life-insurer proration and bank financing-tax rules. The bank input is interest expense divided by assets, weighted across banks by municipal holdings. |
| Insurer portfolio moments | Annual positive-value P&C and life positions matched to tax status and remaining maturity. Cross-sectional moments are constructed separately each year and equally averaged within a block. |
| Municipal supply | Dated exempt Mergent outstanding amounts, projected onto maturity bins, with measured issuer-state shares. Household geographic weights equal these supply shares. |
| Taxable supply | Financial Accounts treasury and corporate totals, FISD corporate maturity shares and treasury maturity information from the Monthly Statement of the Public Debt. Quantities are not fitted to yields. |
| Rate dynamics and yield targets | Annual TB3MS for persistence and innovations. The A-or-above corporate zero curve sets the rate mean and monthly states and supplies the ten-, twenty- and thirty-year yield conditions in annual log units. |

Portfolio equations use the normalized payoff units of Section 4.2: a household's home-state municipal after-tax share and an institution's municipal share are set to one, with risk aversion and outside exposures transformed accordingly. Let $K_i = \operatorname{diag}(\xi_i^{a(j)})$ collect bond shares for the market–maturity positions $j$, and let $\xi_i^C$ be the cash share. For a household, the normalized share on out-of-state municipal interest is one minus its state rate. Fiscal calculations instead retain absolute federal and state tax rates, as explained in Appendix D.4.

## B.2 Solving equilibrium at given coefficients

An equilibrium consists of municipal and taxable price functions and the portfolios they support at every rate state. The difficulty is that current holdings depend on resale risk, which itself depends on future prices. I solve this consistency problem by iterating on price functions. Within each price iteration, a separate market-clearing problem determines current holdings and required returns at each state. The coefficient search described below lies outside this equilibrium calculation.

**From prices to return moments.** Represent log prices $p_j(z) = \log P_j(z)$ on a grid of rate states $\{z_g\}_{g=1}^G$. Position $j$ identifies a market and a maturity, and $p_{j,-}(Z')$ is the next-period log price of its one-period-shorter resale bond. Given the current price functions, $Z' = \rho z + \sigma_r U$, with $U \sim N(0,1)$, determines

$$\mathcal{M}_j(z) = \log \mathbb{E}_z \exp\{p_{j,-}(Z')\} + \tfrac{1}{2} v_j, \tag{105}$$

$$\Sigma_{jk}(z) = \operatorname{Cov}_z(p_{j,-}(Z'), p_{k,-}(Z')) + \mathbf{1}_{\{j=k\}} v_j. \tag{106}$$

Here $v_j$ is bond-specific return variance. The first expression is the log expected resale payoff, including the bond-specific mean correction. It gives $\mu_j = \mathcal{M}_j - p_j = \log \mathbb{E}_z[1 + R_j]$, the mean adjusted holding return of Appendix A.1.1. The covariance matrix $\Sigma$ measures risk in the centered log-return innovations. The liability credit is $c_{ij}^L = -\operatorname{Cov}_z(p_{j,-}(Z'), \ell_i \sigma_r U)$, as in (45).

**From return moments to portfolios and clearing.** At each state, hold these continuation moments fixed and find the return vector $\mu$ that clears all markets. The cash rate is $r(z) = \max\{\bar{r} + z, 0\}$. For any trial $\mu$, investor $i$ chooses its portfolio by solving

$$\max_{x_i \geq 0,\, \mathbf{1}'x_i \leq W_i} x_i'\big(K_i\mu - \xi_i^C r\mathbf{1} - \tau_i^{\text{fin}}\big) + \gamma_i x_i' K_i c_i^L - \frac{\gamma_i}{2} x_i' K_i \Sigma K_i x_i. \tag{107}$$

The terms are the return above cash, the liability hedge and the cost of portfolio risk. The financing-tax vector $\tau_i^{\text{fin}}$ is zero except on banks' municipal positions, where it is twice their measured annual cost. Cash is the uninvested budget $x_i^0 = W_i - \mathbf{1}'x_i$. With $\kappa_i$ its constraint multiplier, the numerical portfolio conditions are

$$\begin{aligned} g_i &= \gamma_i K_i[\Sigma K_i x_i - c_i^L] - K_i\mu + \xi_i^C r\mathbf{1} + \tau_i^{\text{fin}} + \kappa_i\mathbf{1}, \\ g_{ij} &= \frac{\epsilon w_j}{x_{ij}}, \qquad \kappa_i = \frac{\epsilon}{x_i^0}, \qquad x_{ij} > 0, \quad x_i^0 > 0, \end{aligned} \tag{108}$$

jointly with market clearing,

$$\sum_i m_i x_i(z_g) = Q. \tag{109}$$

Here $Q$ is the fixed vector of bond supplies. The first line of (108) is the marginal net cost of a position. The second uses a small log barrier, $\epsilon = 10^{-9}$, to approximate exact nonnegativity and complementarity. Costly positions become correspondingly small rather than exactly zero. Municipal barrier weights $w_j$ equal issuer-state supply shares, and taxable weights equal one. Splitting the municipal market therefore does not multiply the barrier weight. All investors can adjust every bond position at every iteration.

**From clearing returns back to prices.** The returns required to clear markets imply updated current log prices

$$\widehat{p}_j(z_g) = \mathcal{M}_j(z_g) - \mu_j(z_g). \tag{110}$$

I move the candidate price functions toward these updated prices and repeat. All markets and maturities are updated together because portfolio covariances and budgets link them. At convergence, the same prices determine resale moments, support optimal portfolios and clear supply. The one-period municipal claim pays one, so its continuation log price and return variance are zero. Cash has log price $-r$, supplying the continuation price for the next-longer taxable bond.

**Algorithm 1** Equilibrium prices and portfolios at given coefficients

**Require:** Measured taxes, portfolio assets, budgets, supplies, rate dynamics, state grid and candidate coefficients

1: Construct investor exposures and bond-specific variances from the coefficients
2: Initialize log price functions and portfolios
3: **repeat**
4: Compute resale means, covariances and liability credits from the candidate prices
5: **for all** rate-grid nodes, independently **do**
6: Solve (108) and (109) jointly for returns, bond holdings and cash
7: **end for**
8: Form $\widehat{p} = \mathcal{M} - \mu$ and update all candidate price functions together
9: **until** the maximum price-update error is below $10^{-10}$
10: Recompute the original moments and check portfolio conditions, clearing and budgets
11: Evaluate yields at monthly states and holdings at end-year states

During calibration, solved nearby coefficient vectors provide initial prices and portfolios. For the full-state evaluation, initial municipal prices come from the one-market equilibrium at the same coefficients. Household starting municipal holdings are assigned to their home market, institutional holdings are divided by supply shares and taxable holdings are unchanged. These are starting values only. Holdings and partici-

pation are then solved endogenously. Grid construction, numerical steps and acceptance checks are given in Appendix B.6.

## B.3 Estimated parameters and calibration moments

**What is estimated.** The data measure taxes, amounts invested, supplies and household balance-sheet ratios. They do not directly measure risk aversion or how much rate risk an investor's outside assets and liabilities create. Seventeen coefficients, collected in $\theta$, govern these objects and the small bond-specific risk. All are estimated jointly and are common across blocks and issuer states.

Three coefficients govern portfolio risk: common risk aversion $\gamma_i = \gamma$ in the normalized payoff units, thirty-year bond-specific variance $v_{30}$ and its maturity exponent $p_v$ (Table 8). For $n = 4, 6, \ldots, 30$ calendar years, the two-year return variance is

$$v_n = v_{30}(n/30)^{p_v}, \qquad v_2 = 0.$$

It is the same for municipal and taxable bonds at a given maturity. The numerical model uses this positive variance in both the portfolio problem and resale moments. Appendix B.7 discusses its normalization.

The remaining fourteen coefficients determine investor exposures and their distribution, as shown in Table 7. Three map measured household outside assets and debt into rate exposure, eight describe the insurer mixtures, one sets bank exposure and two set the mean and dispersion of residual-group exposure. Together with the three common risk coefficients, these account for all seventeen coefficients. Their values are reported in Table 12.

**Which moments discipline them.** Twenty-two conditions compare model outcomes with data. Nine concern ownership: municipal stocks held by the two household groups, P&C insurers, life insurers and banks, and taxable stocks held by households and those three institutional sectors. Four concern the mean maturities of P&C and life municipal and taxable portfolios. Six describe insurer portfolio composition: two measure the concentration of municipal shares across insurers, and four the fractions of municipal and taxable holdings beyond twenty years. Three compare taxable yields at ten, twenty and thirty years with the observed curve.

Ownership provides the main information about outside exposure. Given taxes and portfolio assets, the amounts and maturities held help reveal how strongly the rest of an investor's balance sheet offsets or reinforces its bond risk. Insurer portfolio composition also disciplines differences within sectors. Taxable term spreads and long-bond allocations inform risk aversion and bond-specific risk. These links explain the choice of targets. They are not a one-to-one assignment of moments to coefficients, since equilibrium prices and portfolios determine the joint fit. Municipal yields and P1–P3 remain untargeted.

**The distance minimized.** Let $\mathcal{C}_{kj}(\theta)$ be the signed model-minus-data condition for moment $j$ in block $k$, after solving that block's equilibrium. Its observed normalizer $o_{kj} > 0$ puts the difference on a comparable scale. The criterion is

$$f_j(\theta) = \frac{6^{-1}\sum_{k=1}^{6} \mathcal{C}_{kj}(\theta)}{0.1\,(6^{-1}\sum_{k=1}^{6} o_{kj})}, \qquad \mathcal{L}(\theta) = \frac{1}{2}\sum_{j=1}^{22} f_j(\theta)^2. \tag{111}$$

Thus I average each signed condition across the six blocks before squaring it. A deviation equal to 10% of its mean observed normalizer gives $|f_j| = 1$. The weights are fixed using data alone, rather than estimated from a covariance matrix. Estimation repeats three steps: choose a common coefficient vector $\theta$, solve all six dated economies using Algorithm 1, and compare their averaged outcomes with the twenty-two data conditions. The search updates $\theta$ to reduce this distance while keeping measured inputs fixed.

**Constructing the conditions.** For ownership, $\mathcal{C}$ is the model stock less the observed stock and $o$ is the observed stock. For insurer mean maturity, it is $B(\overline{T} - \overline{T}^{\text{obs}})$, with $o = B^{\text{obs}}\overline{T}^{\text{obs}}$. Here $B$ is the model portfolio value and $\overline{T}$ its holdings-weighted maturity. Weighting by the model portfolio prevents a nearly empty book from contributing as much as a large one. These conditions include the riskless two-year municipal stock.

The six composition conditions instead use the risky window above two through thirty years. For insurer $i$, let $M_i$ and $C_i$ be municipal and taxable holdings in that window and $B_i = M_i + C_i$. The concentration statistic is $\sum_i (M_i^2/B_i)/\sum_i B_i$: square each insurer's municipal portfolio share and average with portfolio-size weights. A larger value at a given aggregate municipal share means greater concentration in portfolios with high municipal shares. Its condition is the model statistic less the observed one. For either asset's long fraction, with model book $B$, holdings $L$ beyond twenty years and observed long fraction $f^{\text{obs}}$, the condition is

$$\mathcal{C}^{\text{long}} = \frac{L - f^{\text{obs}} B}{B^{\text{obs}}_{\text{window}}}.$$

The observed window book equals the national sector stock times its measured coverage in the window. Normalizers are the observed concentration statistics and long fractions.

Holdings conditions are computed at the three observed end-year states of each block and then averaged. Products and ratios are formed before this averaging. Yield conditions use annual log yields at all 36 monthly states. Since observed two-year yields set those states, the ten-, twenty- and thirty-year yield errors are also errors in spreads over the two-year yield. Their normalizers are the observed block-mean term spreads.

## B.4 Construction of the calibration inputs

**Household construction and geography.** High-tax households have a federal marginal rate of at least 35% under the law of the SCF year. In 2015–2017 they are about 5% of families and hold 46% of household portfolio assets. Within each tax group, tax units are sorted by outside financial assets less debt relative to their portfolio base of deposits, bonds and bond funds. Seven bins contain approximately equal amounts of portfolio assets. Outside financial assets exclude retirement accounts and this portfolio base. Table 10 reports the measured outside-asset and debt ratios that enter the exposure mapping in Table 7. The coefficients $b_A$ and $b_D$ convert these ratios into rate exposure, and the factor of one half fixes their units.

The same national household composition is repeated across the 45 markets, with portfolio assets allocated by issuer-state municipal supply shares. High-tax households face the home state's top rate, and lower-tax households face tax-unit bracket rates. The model applies a common home-state exemption in all income-tax jurisdictions and taxes out-of-state municipal interest at the resident state's rate. Institutions have no modeled state-income-tax wedge and choose across all markets. Separate municipal bond-specific shocks allow them to diversify geographically.

**Household and institutional tax treatment.** Statutory rules determine after-tax shares before calibration. Household calculations use each SCF tax unit's income and filing status, including the net investment income tax and the additional tax on Social Security benefits induced by interest income. Federal and state after-tax products are formed at the tax-unit level before aggregation. The baseline multiplies federal and state shares for taxable interest and gives no federal deduction credit for state taxes on out-of-state municipal interest. Appendix D.2.1 varies the latter convention. The Social Security component on exempt interest is approximated by the tax response to a small ordinary-income increment less the tax on that increment. It can make the absolute home-municipal after-tax share slightly below one.

Before 2018, institutions face a 35% corporate rate. P&C proration taxes 15% of municipal interest at that rate, giving relative after-tax share $\xi_P = 0.65/0.9475 = 0.686$. Life-insurer proration taxes one half of municipal interest, giving $\xi_L = 0.65/0.825 = 0.788$. Banks have $\xi_B = 0.65$ and separately incur the tax cost of disallowed financing-interest deductions on municipal holdings. This annual cost falls from 1.02% in 2000–2002 to 0.11% in 2015–2017. The calibration uses non-bank-qualified treatment (Appendix D.1). The residual group has $\xi = 1$ and consists mainly of foreign investors, pensions, governments, money market funds and nonprofits.

Table 10: Measured household outside assets and debt by balance-sheet type

| Bin | 2000–02 | 2003–05 | 2006–08 | 2009–11 | 2012–14 | 2015–17 |
|---|---|---|---|---|---|---|
| *High-tax households* | | | | | | |
| 1 | 0.93, 3.17 | 1.17, 4.21 | 0.83, 3.00 | 0.77, 2.85 | 0.63, 3.23 | 0.63, 2.93 |
| 2 | 0.39, 0.23 | 0.10, 0.07 | 0.19, 0.10 | 0.22, 0.18 | 0.26, 0.15 | 0.12, 0.08 |
| 3 | 0.62, 0.12 | 0.43, 0.12 | 0.45, 0.07 | 0.25, 0.06 | 0.57, 0.11 | 0.35, 0.12 |
| 4 | 1.18, 0.24 | 0.84, 0.12 | 0.91, 0.14 | 0.60, 0.11 | 0.87, 0.15 | 0.70, 0.05 |
| 5 | 1.98, 0.26 | 1.53, 0.26 | 1.35, 0.11 | 1.15, 0.19 | 1.46, 0.25 | 1.32, 0.16 |
| 6 | 3.77, 0.29 | 2.55, 0.38 | 2.19, 0.28 | 1.80, 0.16 | 2.50, 0.30 | 2.77, 0.23 |
| 7 | 18.49, 1.55 | 12.06, 1.28 | 10.56, 0.90 | 7.76, 0.81 | 11.08, 0.73 | 12.12, 0.94 |
| *Lower-tax households* | | | | | | |
| 1 | 1.42, 11.55 | 1.03, 12.34 | 1.32, 16.95 | 0.90, 12.88 | 0.75, 14.31 | 0.72, 13.74 |
| 2 | 0.30, 0.65 | 0.35, 0.69 | 0.39, 1.20 | 0.16, 0.26 | 0.31, 0.87 | 0.23, 0.88 |
| 3 | 0.06, 0.04 | 0.04, 0.02 | 0.08, 0.08 | 0.05, 0.02 | 0.01, 0.01 | 0.01, 0.02 |
| 4 | 0.30, 0.10 | 0.28, 0.12 | 0.19, 0.07 | 0.32, 0.12 | 0.17, 0.07 | 0.38, 0.18 |
| 5 | 0.78, 0.15 | 0.71, 0.16 | 0.63, 0.14 | 0.58, 0.08 | 0.71, 0.13 | 0.71, 0.10 |
| 6 | 1.73, 0.29 | 1.40, 0.19 | 1.50, 0.29 | 1.25, 0.19 | 1.57, 0.17 | 1.53, 0.15 |
| 7 | 12.82, 1.26 | 8.93, 0.91 | 9.85, 0.96 | 6.59, 0.71 | 8.70, 0.76 | 10.19, 0.80 |

*Notes.* Each cell gives two ratios, in order: outside financial assets and debt, each divided by the portfolio base. These are $Out_k$ and $Debt_k$ in Table 7. Outside financial assets exclude retirement accounts and the portfolio base of deposits, direct bonds and bond funds. Within each tax group, bins are ordered by outside assets less debt relative to the portfolio base and contain approximately one seventh of the group's portfolio assets. SCF waves 2001, 2004, 2007, 2010, 2013 and 2016 supply the six blocks. These ratios are measured inputs. The coefficients $a_H, b_A, b_D$ translating them into rate exposure are estimated jointly.

**Rate process and bond supplies.** Annual 1952–2024 short rates give persistence 0.854 and an innovation standard deviation of 1.63 percentage points. Temporal aggregation gives two-year persistence $\rho = 0.729151$ and innovation standard deviation $\sigma_r = 0.042778$. The mean $\bar{r} = 0.060412$ is twice the 2000–2017 mean annualized two-year corporate zero yield. The short rate is $r(z) = \max\{\bar{r} + z, 0\}$, and observed monthly states are twice that month's continuously compounded two-year corporate zero yield less $\bar{r}$. Thus long yields do not determine the rate states.[23]

Financial Accounts stocks determine sector portfolio assets, budgets and national supplies. Mergent outstanding amounts determine municipal maturity profiles and state shares. Corporate maturity shares come from the Mergent Fixed Income Securities Database (FISD), and treasury maturities from the Monthly Statement of the Public Debt. Table 9 summarizes the construction.

**The 2018–2020 extension.** For the post-2017 curve comparison, I keep all estimated coefficients fixed and update measured inputs using the 2019 SCF composition, annually inflation-adjusted incomes and mean 2018–2020 end-year stocks. Institutions face the 21% corporate rate, P&C proration of 25% and a life-insurer policyholders' share of 30%. The annual bank financing-tax cost is 0.117%, and insurer portfolio fractions use available 2018–2019 holdings. This seventh economy adds no calibration targets.

[23] Annual persistence is 0.853903 and the innovation standard deviation is 0.016266. The two-year state is twice the annualized rate deviation, so $\rho = 0.853903^2$ and $\sigma_r = 2(0.016266)\sqrt{1 + 0.853903^2}$.

Table 11: Inputs fixed from data in each calibrated economy

| | 2000–02 | 2003–05 | 2006–08 | 2009–11 | 2012–14 | 2015–17 |
|---|---|---|---|---|---|---|
| *Panel A. Group portfolio assets $m_g$ (trillions of current dollars)* | | | | | | |
| High-tax households | 2.02 | 1.97 | 3.47 | 3.82 | 5.72 | 6.97 |
| Lower-tax households | 3.99 | 5.94 | 6.50 | 8.44 | 7.97 | 8.09 |
| P&C insurers | 0.58 | 0.80 | 0.91 | 1.00 | 1.06 | 1.11 |
| Life insurers | 1.56 | 2.11 | 2.10 | 2.52 | 2.84 | 3.05 |
| Banks | 1.50 | 1.89 | 2.38 | 3.28 | 4.21 | 4.83 |
| Tax-neutral residual | 11.54 | 14.94 | 20.04 | 21.42 | 23.29 | 26.47 |
| *Panel B. Investment budget as a share of portfolio assets, $100W_g$* | | | | | | |
| Households (both tax groups) | 86.9 | 88.0 | 86.2 | 90.6 | 91.1 | 91.4 |
| P&C insurers | 73.9 | 71.3 | 71.2 | 73.5 | 73.3 | 72.6 |
| Life insurers | 73.2 | 68.9 | 66.4 | 69.7 | 69.8 | 71.5 |
| Banks | 30.1 | 25.6 | 24.4 | 24.4 | 25.2 | 28.0 |
| Tax-neutral residual | 50.9 | 52.0 | 50.1 | 57.9 | 68.8 | 72.5 |
| *Panel C. Tax inputs and national bond supplies* | | | | | | |
| High-tax household $\xi^C$ | 0.577 | 0.609 | 0.610 | 0.606 | 0.565 | 0.538 |
| Lower-tax household $\xi^C$ | 0.766 | 0.771 | 0.760 | 0.764 | 0.786 | 0.761 |
| Bank financing-tax cost (%/year) | 1.020 | 0.556 | 0.858 | 0.323 | 0.134 | 0.108 |
| Municipal supply ($ trillion) | 1.08 | 1.92 | 2.36 | 2.78 | 2.99 | 3.14 |
| Taxable supply ($ trillion) | 6.18 | 7.98 | 9.74 | 13.72 | 17.39 | 20.02 |

*Notes.* In Panel A, $m_g$ is group $g$'s total of type-level portfolio assets $m_i$, its measured debt securities and cash. Panel B is the measured budget for modeled bonds and cash divided by those portfolio assets. The dollar investment budget is $m_g W_g$. For example, banks' 2015–2017 budget is $4.83 trillion times 0.280, or $1.35 trillion. The same $W_g$ applies to every type within a sector and to both household tax groups. Panel C's household shares are the issuer-supply-weighted averages used in the one-market estimation. The full-state evaluation retains each state's tax shares. Institutional relative taxable shares are constant before 2018: 0.686 for P&C, 0.788 for life insurers, 0.650 for banks and 1 for the residual group. Banks' financing-tax cost is additional to this relative share. Stocks are mean end-year values in each block. Taxable supply includes treasury and corporate debt. Sources and adjustments are in Table 9.

## B.5 Estimated coefficients and fit

The bounded search lowers the distance (111) by re-solving the six economies after each coefficient change. It starts from Table 14's initial vector and reaches $\mathcal{L} = 36.3575$. A restart with finer derivative steps reaches 36.2870, whose coefficients are reported in Table 12. The numerical search and stopping rule are described in Appendix B.6.2.

Table 13 reports observed and fitted values in the one-market economies used for estimation. It shows the economic quantities underlying all twenty-two conditions. Fourteen conditions are within 10% of their observed normalizers. Maturity and long-fraction conditions weight discrepancies by model holdings, so their contribution to the criterion is not the unweighted difference between the averages displayed in the table.

Restoring the 45 municipal markets does not re-estimate any coefficient. Across the six blocks, the full-state minus one-market thirty-year gaps are $(-0.85, -0.72, 2.48, -0.22, -1.15, -0.93)$ basis points, and taxable yields differ by at most 0.52 basis points. The full-state equilibria supply the municipal yield curves, cross-state pass-through and policy calculations. On the observed corporate curve and empirical sample weights, the model explains 83% of the 117.5-basis-point thirty-year puzzle (Figure 8). These municipal outcomes are predictions from the ownership and taxable-yield calibration.

Table 12: Common estimated coefficients, held fixed in the full-state evaluation

| Symbol | Economic role | Value |
|---|---|---|
| *Risk. Main information: taxable term spreads and the allocation of long bonds* | | |
| $\gamma$ | Common risk-aversion curvature, per unit of portfolio assets | 9.6969 |
| $v_{30}$ | Bond-specific variance of a thirty-year bond's two-year return | 0.000722 |
| $p_v$ | Maturity exponent of the bond-specific variance | 1.2597 |
| *Households. Main information: household municipal and taxable stocks* | | |
| $a_H$ | Outside rate-exposure intercept | 1.3045 |
| $b_A$ | Loading on the outside-asset ratio | 2.704 |
| $b_D$ | Loading on the debt ratio | 0.8379 |
| *P&C insurers. Main information: P&C stocks, maturities, composition and long fractions* | | |
| $a_P$ | Low type's outside rate exposure | −1.961 |
| $d_P$ | High-minus-low outside rate exposure | 3.0311 |
| $w_P$ | Low type's share of portfolio assets | 0.0802 |
| *Life insurers. Main information: life-insurer stocks, maturities, composition and long fractions* | | |
| $a_L$ | Low type's outside rate exposure | −5.3927 |
| $d_L$ | High-minus-low outside rate exposure | 6.3007 |
| $w_L$ | Low type's share of portfolio assets | 0.2616 |
| $t_L$ | Middle type's location within the exposure span | 0.8625 |
| $u_L$ | Middle type's share of portfolio assets excluding the low type | 0.0162 |
| *Banks and residual group. Main information: bank stocks and taxable term spreads* | | |
| $h_B$ | Bank outside rate exposure | 1.2045 |
| $\mu_R$ | Residual-group mean outside rate exposure | 1.0683 |
| $s_R$ | Residual-group outside rate-exposure dispersion | 1.7168 |

*Notes.* All seventeen coefficients are fitted jointly to the twenty-two averaged conditions of Table 13. Every coefficient affects every condition through equilibrium prices and clearing, and the headings name the conditions most informative about each block. Outside exposures $h_i = -\ell_i$ use the same units as bond price loadings, in the model's two-year convention. Positive values add rate risk and negative values hedge it (Table 7). Portfolio-asset shares divide each insurer sector's measured debt securities and cash across its types. Appendix Table 14 reports the search intervals, starting values and step scales.

Table 13: Calibration targets and fitted quantities, 2000–2017

| Quantity | Data | Model |
|---|---|---|
| *Ownership stocks (billions of dollars)* | | |
| High-tax households, municipal | 906.8 | 1,135.1 |
| Lower-tax households, municipal | 903.6 | 662.2 |
| P&C insurers, municipal | 295.0 | 289.1 |
| Life insurers, municipal | 21.7 | 22.0 |
| Banks, municipal | 251.6 | 266.5 |
| Households, taxable | 1,635.7 | 1,544.0 |
| P&C insurers, taxable | 335.8 | 244.7 |
| Life insurers, taxable | 1,589.7 | 1,624.4 |
| Banks, taxable | 537.3 | 522.4 |
| *Insurer mean maturities (years)* | | |
| P&C insurers, municipal | 12.1 | 15.1 |
| P&C insurers, taxable | 7.6 | 6.3 |
| Life insurers, municipal | 14.6 | 17.9 |
| Life insurers, taxable | 12.3 | 11.7 |
| *Taxable term spreads (basis points)* | | |
| 10-year taxable spread | 209.0 | 214.7 |
| 20-year taxable spread | 279.9 | 284.6 |
| 30-year taxable spread | 299.3 | 303.6 |
| *Insurer portfolio composition* | | |
| P&C municipal-share concentration | 0.4374 | 0.5613 |
| Life municipal-share concentration | 0.0065 | 0.0063 |
| P&C municipal holdings beyond twenty years | 11.2% | 19.8% |
| P&C taxable holdings beyond twenty years | 7.1% | 4.1% |
| Life municipal holdings beyond twenty years | 24.9% | 28.4% |
| Life taxable holdings beyond twenty years | 25.6% | 26.4% |

*Notes.* Model values are from the one-market economies used to estimate the common coefficients. Entries average the six blocks equally. Stocks and mean maturities include two-year municipal bonds. Concentration and long fractions use the risky window above two through thirty years. Long fractions are percentages held beyond twenty years. The concentration statistic is the portfolio-size-weighted average of squared municipal shares, defined in Appendix B.3. Term spreads are in annual log-yield units over the two-year yield. The table displays average economic quantities. The criterion first constructs dated portfolio-weighted maturity and long-fraction discrepancies and then averages them, as in (111), rather than subtracting these displayed averages.

## B.6 Numerical implementation and checks

### B.6.1 State grid, integration and equilibrium iteration

Each block uses fifteen equally spaced states on $[-6\sigma_z, 6\sigma_z]$, including zero, plus its three December rate states. Here $\sigma_z = \sigma_r/\sqrt{1-\rho^2}$. Log prices are interpolated linearly between the resulting eighteen nodes and extrapolated along the end segments. The baseline has no node exactly at the rate floor $z = -\bar{r}$, whose addition is checked below. Yields are evaluated at the 36 monthly states and holdings at the three included annual nodes. The comparisons in Figures 7 and 8 instead read each bond's issuer-state yield at its issue-quarter-end state and use empirical bond, rating and quarter weights.

All moments integrate the interpolant's segments and tails analytically. For a standard normal $U$ and a log-price segment $a + bU$ on $l < U < u$,

$$\mathbb{E}[e^{a+bU}\mathbf{1}_{\{l<U<u\}}] = e^{a+b^2/2}\{\Phi(u-b) - \Phi(l-b)\}.$$

This removes quadrature truncation. Interpolation and finite-precision errors remain and are assessed below. The local rate loadings $A_T^a = -\partial p_T^a/\partial z$ satisfy Appendix equation (44) under its differentiation conditions.

At each node, Newton's method solves portfolios and clearing jointly. The covariance matrix has a low-rank common component plus diagonal bond-specific variance. Exploiting this structure to eliminate holdings leaves a linear system in market returns, of dimension 689 in the full model. A positivity-preserving line search keeps bond holdings and cash feasible. The first price iteration lowers the barrier from $10^{-5}$ to $10^{-9}$ in factor-of-ten stages. Later iterations use the previous portfolios at $10^{-9}$, and failed inner solves restart from proportional holdings with a larger barrier. Nodes can be solved in parallel.[24]

The price step is damped, with each log-price change capped at 0.025. Anderson acceleration uses at most seven previous iterates and restarts when the price error more than doubles.[25] In the six baseline blocks, the maximum price-update error falls from about $10^{-2}$ to $10^{-6}$ within forty iterations and below $10^{-10}$ within seventy. Later price iterations typically require about two Newton steps per node. All moments and residuals are recomputed before accepting an equilibrium.

### B.6.2 Coefficient search and sensitivity

The coefficient search is a bounded trust-region method. Central differences measure how the averaged conditions respond to coefficients, with feasible one-sided differences at bounds. The search compares damped Gauss–Newton, backtracked and projected-gradient steps, using a secant estimate of the remaining curvature. Table 14 records intervals, initial values and coordinate scales. A scaled move $\delta_n$ changes coefficient $n$ by $c_n\delta_n$. The two searches stop when the trust-region radius falls below $10^{-5}$, with final improvements below 0.001.

Small coordinate and joint perturbations lower the criterion by less than 0.001% at best and change the average thirty-year gap by 0.004 basis points. Larger one-group changes all worsen the fit (Table 15). Lowering risk aversion by 20% has the largest effect: the thirty-year taxable yield falls by 70 basis points and term spreads lie about 20% below their targets. Changing bond-specific variance by 20% moves the gap and taxable yield by less than 0.1 basis points. The other changes hold all remaining coefficients and measured inputs fixed.

[24] Investor Hessians are inverted with the Woodbury identity. Covariance singular values below $10^{-12}$ of the largest are trimmed during iteration, and the final check uses the untrimmed moments. The reduced system uses LU factorization and refinement to relative residual $10^{-12}$. Inner stopping thresholds are $10^{-11}$ for stationarity and clearing and $2.5 \times 10^{-12}$ for holding-weighted stationarity errors.

[25] The damping fraction is $\min\{0.65, 0.025/\|\widehat{p} - p\|_\infty\}$. Acceleration is used only if all least-squares weights are below 50 in absolute value and the step obeys the same price-change cap.

Table 14: Coefficient-search intervals, starting values and step scales

| Coefficient | Search interval | Initial value $\theta_n^0$ | Step scale $c_n$ |
|---|---|---|---|
| $\gamma$ | $[1, 10]$ | 9 | 1 |
| $a_H$ | $[-3, 3]$ | 1 | 0.3 |
| $b_A$ | $[0, 12]$ | 2 | 0.4 |
| $b_D$ | $[0, 20]$ | 0.5 | 0.5 |
| $a_P$ | $[-10, 2]$ | $-1.5$ | 0.5 |
| $d_P$ | $[0.02, 12]$ | 2.4 | 0.6 |
| $w_P$ | $[0.01, 0.99]$ | 0.22 | 0.08 |
| $a_L$ | $[-12, 2]$ | $-4$ | 0.6 |
| $d_L$ | $[0.02, 15]$ | 3.2 | 0.8 |
| $w_L$ | $[0.01, 0.99]$ | 0.1 | 0.08 |
| $h_B$ | $[-3, 3]$ | 0.6 | 0.3 |
| $\mu_R$ | $[-4, 4]$ | 1 | 0.4 |
| $s_R$ | $[0.01, 5]$ | 2.2 | 0.4 |
| $v_{30}$ | $[0.000020, 0.006000]$ | 0.0015 | 0.0004 |
| $p_v$ | $[0, 3]$ | 1.15 | 0.3 |
| $t_L$ | $[0.01, 0.99]$ | 0.75 | 0.08 |
| $u_L$ | $[0.01, 0.99]$ | 0.2 | 0.08 |

Table 15: Sensitivity of the one-market calibration to estimated coefficients

| Coefficient change | 30-year gap (bp) | Change (bp) | Criterion $\mathcal{L}$ |
|---|---|---|---|
| Estimated coefficients | 102.6 | | 36.29 |
| Risk aversion $\gamma \times 0.8$ | 89.0 | $-13.6$ | 55.02 |
| Risk aversion $\gamma = 10$ (upper bound) | 104.7 | $+2.1$ | 36.74 |
| Bond-specific variance $v_{30} \times 0.8$ | 102.7 | 0.0 | 36.61 |
| Bond-specific variance $v_{30} \times 1.2$ | 102.7 | 0.0 | 36.59 |
| Maturity power $p_v - 0.2$ | 102.6 | 0.0 | 36.55 |
| Maturity power $p_v + 0.2$ | 102.7 | 0.0 | 36.47 |
| Household exposure intercept $a_H - 0.25$ | 99.0 | $-3.7$ | 78.82 |
| Household exposure intercept $a_H + 0.25$ | 106.7 | $+4.1$ | 160.84 |
| Household asset loading $b_A \times 0.8$ | 100.9 | $-1.7$ | 62.45 |
| Household asset loading $b_A \times 1.2$ | 105.2 | $+2.6$ | 400.24 |
| Household debt loading $b_D \times 0.8$ | 106.2 | $+3.5$ | 6,720 |
| Household debt loading $b_D \times 1.2$ | 103.9 | $+1.3$ | 82.17 |
| P&C exposures $\times 0.8$ | 102.2 | $-0.4$ | 56.35 |
| P&C exposures $\times 1.2$ | 103.3 | $+0.6$ | 62.51 |
| Life-insurer exposures $\times 0.8$ | 101.9 | $-0.7$ | 59.68 |
| Life-insurer exposures $\times 1.2$ | 103.4 | $+0.8$ | 46.24 |
| Bank exposure $\times 0.8$ | 100.5 | $-2.1$ | 119.11 |
| Bank exposure $\times 1.2$ | 104.5 | $+1.8$ | 241.78 |
| Residual exposure mean $\times 0.8$ | 102.6 | $-0.1$ | 538.79 |
| Residual exposure mean $\times 1.2$ | 101.7 | $-0.9$ | 65.73 |
| Residual exposure s.d. $\times 0.8$ | 99.3 | $-3.3$ | 92.74 |
| Residual exposure s.d. $\times 1.2$ | 105.8 | $+3.1$ | 1,687 |

*Notes.* Results average the six blocks of 2000–2017. Exposure changes scale each type's exposure in the group. The gap is the municipal yield less the household after-tax taxable yield at thirty years. The criterion is (111).

### B.6.3 Equilibrium accuracy

Four diagnostics check prices against resale values, market clearing, portfolio optimality and complementarity:

$$\begin{aligned} R_p &= \|p - \mathcal{M} + \mu\|_\infty, \qquad & R_c &= \|Q - \textstyle\sum_i m_i x_i\|_\infty, \\ R_s &= \|g - \epsilon w/x\|_\infty, & R_{\text{comp}} &= \max\{\|x \odot g\|_\infty, \|\kappa \odot x^0\|_\infty\}. \end{aligned} \tag{112}$$

Independent checks in Table 16 give zero budget and negative-holding violations and zero variance for riskless claims in all six blocks. Acceptance requires $R_p < 3 \times 10^{-7}$, $R_s < 10^{-7}$, $R_c < 10^{-8}$ and a budget violation below $10^{-10}$. Complementarity is reported separately because it equals the barrier weight at an exact barrier solution.

Table 16: Full-state equilibrium diagnostics

| Block | $R_p$ | $R_s$ | $R_c$ | $R_{\text{comp}}$ |
|---|---|---|---|---|
| 2000–2002 | $7.511 \times 10^{-11}$ | $5.884 \times 10^{-12}$ | $5.551 \times 10^{-17}$ | $1.002 \times 10^{-9}$ |
| 2003–2005 | $4.456 \times 10^{-11}$ | $8.340 \times 10^{-12}$ | $1.110 \times 10^{-16}$ | $1.002 \times 10^{-9}$ |
| 2006–2008 | $8.156 \times 10^{-11}$ | $8.547 \times 10^{-12}$ | $1.110 \times 10^{-16}$ | $1.002 \times 10^{-9}$ |
| 2009–2011 | $5.498 \times 10^{-11}$ | $8.687 \times 10^{-12}$ | $2.220 \times 10^{-16}$ | $1.002 \times 10^{-9}$ |
| 2012–2014 | $9.289 \times 10^{-11}$ | $9.864 \times 10^{-12}$ | $2.220 \times 10^{-16}$ | $1.002 \times 10^{-9}$ |
| 2015–2017 | $8.495 \times 10^{-11}$ | $8.352 \times 10^{-12}$ | $2.220 \times 10^{-16}$ | $1.002 \times 10^{-9}$ |

*Notes.* $R_p$, $R_s$, $R_c$ and $R_{\text{comp}}$ are the pricing, stationarity, clearing and complementarity residuals of (112). Maxima over each block's eighteen state nodes and the applicable coordinates, with $\epsilon = 10^{-9}$. Complementarity is the largest unpenalized product of a bond or cash holding and its multiplier, which equals $\epsilon$ at an exact barrier solution. The diagnostics assess only the discretized funded equations.

The fixed-coefficient 2018–2020 extension also passes, with $R_p = 4.52 \times 10^{-11}$, $R_s = 7.78 \times 10^{-12}$, $R_c = 2.22 \times 10^{-16}$ and complementarity $1.00 \times 10^{-9}$. These checks apply at grid nodes. They do not give a uniform error bound between nodes, which motivates the grid comparison below.

The 45 riskless municipal returns are solved within investors' budgets. Their household pricing frontier is

$$\min_{i \in \mathcal{H}} \frac{\xi_i^C r(z) + \kappa_i(z)}{\xi_i^{M_s}},$$

where $\mathcal{H}$ contains all 630 household types and $s$ identifies the issuer. Every issuer's one-period return is within 0.01 basis points of this frontier at every grid node. The minimum can differ from the issuing state's high-tax resident benchmark. In the six calibrated economies, the national one-period municipal yield is also within 0.01 basis points of that household benchmark at every observed month.

### B.6.4 Grid and barrier sensitivity

Three independently checked 2015–2017 re-solves keep coefficients and measured inputs fixed. They refine the equally spaced grid from fifteen to twenty-nine nodes, lower the barrier from $10^{-9}$ to $10^{-10}$, or add a node exactly at the rate floor. Table 17 reports each change relative to the baseline.

Gap changes remain below 0.23 basis points. Holdings can move more: grid refinement lowers lower-tax household municipal holdings by 4.02%, while adding the floor node lowers those holdings by 1.69% and banks' taxable holdings by 2.80%. Lowering the barrier changes risky positions at annual nodes by at most 0.071% of investor budgets. Riskless allocations are less stable because cash and a riskless bond can offer the same after-tax return, leaving their split undetermined by prices. At annual nodes, this split changes by up to 24.84% of a budget, while complementarity falls from about $10^{-9}$ to $10^{-10}$.[26]

[26] Across all grid nodes, the largest riskless and risky position changes are 77.06% and 0.1179% of budgets. The largest riskless change occurs at a positive cash rate. Tiny residual-group municipal and high-tax-household taxable holdings fall about 90% when the barrier is reduced. With fixed prices and controlled stationarity error, concavity bounds the barrier portfolio's utility loss by $\epsilon(\sum_j w_j + 1)$.

Table 17: Numerical sensitivity at fixed coefficients, 2015–2017

| Comparison | Actual nodes | $\epsilon$ | Δ 30-year gap (bp) | Δ lower-tax household municipal book (%) | Δ bank taxable book (%) |
|---|---|---|---|---|---|
| Nested grid | 32 | $10^{-9}$ | +0.199 | -4.02 | -2.59 |
| Smaller barrier | 18 | $10^{-10}$ | +0.0001 | -0.004 | +0.001 |
| Exact floor knot | 19 | $10^{-9}$ | -0.228 | -1.69 | -2.80 |

*Notes.* Each row compares one change with the original eighteen-node, $10^{-9}$-barrier economy. Rows are not a sequence. The yield gap is the national supply-weighted municipal yield minus the original resident-household after-tax taxable benchmark, in continuously compounded yields over 36 actual months. Its baseline is 104.27 bp. Holdings use the three annual measurement nodes, and percentage changes use each group's original book as denominator. The rows are neither confidence intervals nor the seven-block empirical-weight-matched gap. The floor is economically active in every row.

### B.6.5 Policy calculations

Counterfactuals use the same price-and-portfolio algorithm. The forty-two state repeals use the baseline barrier and tolerances. Independent reconstruction gives maximum $R_p = 9.79 \times 10^{-11}$, $R_s = 9.20 \times 10^{-11}$ and $R_c = 3.33 \times 10^{-16}$, zero budget and negative-holding violations and zero variance for riskless claims. The riskless frontier holds within $1.83 \times 10^{-5}$ basis points at every observed month. Separate checks verify the changed tax rules and fiscal identities.

The equal-treatment calculation uses price-update tolerance $2 \times 10^{-7}$, inner tolerance $2 \times 10^{-8}$, holding-weighted tolerance $5 \times 10^{-9}$ and linear-solve tolerance $2 \times 10^{-9}$. Its 2015–2017 endpoint removes 99.9% of tax distinctions, starting from 99% removal, and initializes the other five blocks, each checked independently. Near a binding cash constraint, the cash Newton direction uses the low-rank inverse to avoid cancellation. In a three-municipal-market comparison, moving from 99% to 99.9% removal changes thirty-year municipal and taxable yields by 1.9 and 0.2 basis points and group maturity shares by at most 0.2 percentage points. Municipal and taxable bonds become close substitutes, however, and within-maturity allocations between them can change by 9.5 percentage points of group assets. Table 54, Panel B, therefore reports riskless shares and the maturity composition of risky holdings.

### B.6.6 Bond-supply sensitivity

The preceding checks vary numerical settings. This check instead relaxes the economic assumption of fixed bond quantities. It uses the one-market economies and reports gaps on their own taxable curves, before the corporate normalization used for the full-state comparison.

**Elastic bond supply.** I allow borrowing to respond to its cost in both markets, holding the estimated coefficients fixed. For market $a \in \{M, C\}$ and maturity $T$, supply follows

$$Q_T^a(z) = \overline{Q}_T^a \left( \frac{y_T^a(z)}{\overline{y}_T^a} \right)^{-\chi_{a,T}},$$

where $y_T^a$ is the annual continuously compounded yield. Each schedule passes through the original quantity and fixed-supply yield at the mean rate state. These anchors remain fixed when both yield curves are resolved. Risky maturities from four to thirty years have elastic supply. The two-year riskless claims retain their baseline treatment.

St. Clair (2024) estimates that a one-percent increase in borrowing costs reduces municipal issuance by approximately 0.3–0.6 percent. I use 0.45 as the midpoint of this rounded range and impose the same elasticity in both markets. The estimate concerns annual borrowing near the bank-qualified threshold. Holding financing terms fixed maps proportional borrowing-cost changes locally to proportional yield changes. Applying that response to outstanding model supply defines the comparative-static exercise.

Table 18 reports the one-market economy over the six calibration blocks. At equal elasticities of 0.3, 0.45, and 0.6, the average thirty-year gap lies between 101.0 and 102.0 basis points, compared with 102.6

under fixed supply. At 0.45, municipal yields fall by 14.1 basis points and the household after-tax taxable benchmark by 13.4, so the gap changes by only 0.7 basis points.

I also consider a larger taxable response. The November 2022 estimates of Siani (2022) imply approximately 1.5% less corporate issuance per ten-basis-point increase in credit spreads, holding the benchmark rate fixed. I match that local slope by setting $\chi_{C,T} = 15\overline{y}_T^C$, while keeping $\chi_{M,T} = 0.45$. The thirty-year taxable elasticity is 0.86–0.97 across blocks. Applying this corporate issuance response to the pooled taxable supply is an additional sensitivity assumption. The average gap is 99.2 basis points, retaining 96.7% of its fixed-supply value. Both yields adjust: municipal yields fall by 19.5 basis points and the after-tax taxable benchmark by 16.1. A larger taxable supply elasticity therefore need not widen the gap.

Table 18: Yield gaps with elastic supply in both markets

| Municipal elasticity | Taxable elasticity | 10 years | 20 years | 30 years |
|---|---|---|---|---|
| 0 | 0 | 35.7 | 84.9 | 102.6 |
| 0.30 | 0.30 | 36.4 | 84.1 | 101.9 |
| 0.45 | 0.45 | 36.7 | 84.0 | 102.0 |
| 0.60 | 0.60 | 36.8 | 83.4 | 101.0 |
| 0.45 | $15\overline{y}_T^C$ | 37.2 | 83.2 | 99.2 |

*Notes.* Gaps are in basis points, averaged over the original rate-state observations within each block and equally across the six blocks of 2000–2017. All other coefficients and inputs remain fixed. Elasticities are reported as positive magnitudes: higher borrowing costs reduce supply. The last row converts a corporate spread semi-elasticity to a yield elasticity at each fixed anchor, with yields in annual decimal units.

## B.7 Bond-specific return risk and its normalization

The calibrated model adds the independent shock in (31) to the main-text return. It can represent an issue-specific change in trading demand affecting sale proceeds without moving aggregate rates. Its variance cost grows with position size, making bonds imperfect substitutes and supporting joint holdings by several investors, as Figure 14 illustrates. The shock is zero for the one-period claim. In the one-market calibration, 20% changes in its variance move the thirty-year gap by less than 0.1 basis points (Table 15).

Writing the shock into realized resale proceeds, $P_{T-1}^a(z')e^{\sigma_{\eta,T}^a\eta}$, is algebraically identical to multiplying the holding return by it. This leaves every return, conditional moment and portfolio condition unchanged and adds no persistent price state.

The shock's lognormal mean can also be normalized to one without changing holding returns. Replace the shock factor by $e^{\sigma_{\eta,T}^a\eta-(\sigma_{\eta,T}^a)^2/2}$ and shift log prices to

$$\check{p}_T^a(z) = p_T^a(z) - \frac{1}{2}\sum_{k=1}^{T}(\sigma_{\eta,k}^a)^2, \qquad \check{P}_T^a(z) = e^{\check{p}_T^a(z)}.$$

The adjacent-maturity shifts exactly offset the normalization:

$$\frac{\check{P}_{T-1}^a(z')e^{\sigma_{\eta,T}^a\eta-(\sigma_{\eta,T}^a)^2/2}}{\check{P}_T^a(z)} = \frac{P_{T-1}^a(z')e^{\sigma_{\eta,T}^a\eta}}{P_T^a(z)}.$$

With fixed dollar supplies, portfolios, optimality and clearing are unchanged. For $T$ periods of $\Delta = 2$ years, the annualized yield shift is

$$\check{y}_T^a - y_T^a = \frac{1}{2T\Delta}\sum_{k=1}^{T}(\sigma_{\eta,k}^a)^2.$$

The fitted variance profile is common to both markets, so both thirty-year continuously compounded yields rise by 0.855 basis points and their difference is unchanged. Corporate normalization uses their ratio, so the normalized thirty-year gap rises by 0.20 basis points, from 97.39 to 97.59. The riskless boundary is unchanged.

# C Adjustments and empirical tests

This appendix first constructs the five adjustments used to measure and account for the municipal puzzle. Each adjustment states the economic comparison, the data that identify it, the calculation and the validation or sensitivity checks. Paired issues and high-grade general-obligation bonds then check whether the puzzle depends on the benchmark or security sample. The remaining subsections report the tax and interest-rate tests, their robustness checks and supporting ownership evidence.

## C.1 Comparability and bond-side adjustments

The comparison has two stages. First, municipal offering yields must be made comparable with the option-free, seasoned corporate curve. Removing the value of the issuer's call option and the issuance premium puts the two markets on that basis and defines the municipal puzzle. Second, capital-gains taxation, trading costs and default risk can affect the two markets differently even on this comparable basis. The remaining three adjustments measure these municipal-minus-corporate differences. A positive contribution explains part of the puzzle and is subtracted before comparison with the model in Sections 3.7 and 5.4.

### C.1.1 The comparison, units and averaging

**Common unit and sample.** Each adjustment is a basis-point contribution to

$$\text{Spread}(T) = y_T^{M,\text{raw}} - \bar{\xi}_H(T)\, y_T^C. \tag{113}$$

The raw municipal zero-coupon curve $y_T^{M,\text{raw}}$ includes callable and noncallable issues before either comparability adjustment. The corporate curve $y_T^C$ is defined in Section 3.6. The benchmark $\bar{\xi}_H(T)y_T^C$ averages each bond's household after-tax corporate yield at its dated federal and state tax rates. Dividing this average by the averaged corporate yield gives $\bar{\xi}_H(T)$. These rates include the net investment income tax (NIIT) and, from 2018, the cap on state and local tax (SALT) deductions. The household wedge $1 - \bar{\xi}_H(T)$ averages 41.3%.

Both curves and the benchmark use the three-stage averaging applied throughout: equal weights within maturity × issue-quarter × half-notch rating cells, then across ratings within quarters, then across quarters. Figure 7(d) shows the resulting puzzle in implied-tax units.

Table 31 places all five contributions in basis points of the full-sample, all-bond Spread($T$). Each adjustment uses the observations that identify its particular effect: callable contracts, early secondary trades, first-sale behavior, matched trading-cost cells or CDS quotes. The subsections below explain how each estimate is carried back to the maturity grid of the accounting.

**Additive accounting.** The two comparability adjustments reduce the municipal yield itself. The three bond-side adjustments are differences between markets, so they are contributions to the spread rather than five successive versions of one municipal curve. Each is estimated separately and their reported basis-point contributions are added. Interactions are not estimated, so this accounting does not depend on the order of subtraction.

Appendices C.1.2–C.1.6 give the five estimators in the order of (23), and Appendix C.1.7 combines them.

### C.1.2 Adjustment I: embedded call options

When rates fall, an issuer can redeem a callable bond instead of continuing to pay its higher coupon. This limits the holder's price appreciation and raises the yield required at issuance. The call adjustment asks what the same bond would yield if its promised payments could not be cut short. Removing the option raises its value and lowers its yield. More than 95% of bonds beyond twenty years are callable, so this valuation preserves the long-maturity sample that would be lost by simply dropping callable bonds.

**The call option.** Long municipal bonds are usually callable at par on coupon dates from year ten. Callable bonds make up 51.1% of curve observations and 51.8% of those at two to thirty years. Investment-grade corporate bonds are usually noncallable bullets or have make-whole strikes that discourage exercise

as rates fall. I first fit an option-adjusted spread (OAS) over the model discount curve to reproduce each callable bond's observed price. Keeping that spread and the contractual payments fixed, I value the otherwise identical noncallable bond. The difference between its bullet value $V_b^{\text{bul}}$ and the fitted callable value $V_b^{\text{call}}$ is

$$CV_b \;=\; V_b^{\text{bul}} - V_b^{\text{call}} \;\geq\; 0. \tag{114}$$

Adding $CV_b$ to the observed price gives the option-free price. The unchanged OAS retains compensation for credit, liquidity and other factors. Fitting the observed price anchors the valuation level to the data, while the call value depends on the contract and interest-rate dynamics. The steps below reconstruct contractual payments and call dates, value the bond with and without the call option, and convert the price difference into a yield adjustment.

**What identifies the option value.** The observed coupon and call schedule determine what the issuer can save by redeeming early, while interest-rate volatility determines how often that option becomes valuable. I use measured municipal yield volatility and fit its scale to callable-minus-noncallable yield differences within the same maturity, issue-quarter, rating and coupon cells. Thus the volatility fit uses the observed price of callability in comparable bonds. The fitted scale, its uncertainty and the noncallable validation are reported below.

The equilibrium curve supplies the reference discount prices, while contracts and observed callable–noncallable differences discipline the option value. The call adjustment is not fitted to the municipal gap over the household benchmark.

**Step 1: set the promised payments and call rights.** For each bond, use its annual coupon rate $c$, remaining payment dates $t_1, \ldots, t_n$, accrual fractions $a_k$ and applicable call prices $K_k$. Coupons are $100ca_k$ per 100 face, with principal added at the final date. The coupon interval is $\Delta = 0.5$ years, with first stub $s$. All valuations use dirty prices, which include accrued interest. Recorded call prices apply on coupon dates from the first contractual call. The exact calendar and accrual construction is given after the valuation results.

**Step 2: specify the state-dependent discount curve.** The issue year selects one of seven dated economies: the six calibration blocks or a 2018–2020 economy using 2018–2019 insurer portfolio fractions and 2018–2020 quantities, rates and taxes. Supply-weighted municipal log prices are interpolated over maturities $\{0, 2, \ldots, 30\}$, including $p_0 = 0$:

$$\log P_T(z) = \text{Interp}_T\left\{\sum_s w_s p_T^{M_s}(z)\right\}, \qquad w_s = \frac{\text{municipal supply in state } s}{\text{total municipal supply}}. \tag{115}$$

The OAS is measured relative to this mean curve.

The model's two-year AR(1) in (5) is converted to annual and $h$-year steps:

$$\rho_y = \sqrt{\rho}, \quad \sigma_y = \frac{\sigma_r}{\sqrt{1+\rho_y^2}}, \quad \rho_h = \rho_y^h, \quad \sigma_h = \sigma_y\sqrt{\frac{1-\rho_y^{2h}}{1-\rho_y^2}}. \tag{116}$$

Before the pricing drift is fitted, I rescale log prices around $z = 0$ so that pricing yield volatility equals measured volatility times a scale $m$:

$$\log P_T^{\text{scaled}}(z) = \log P_T(0) + m\,\frac{\sigma_T^{\text{data}}}{\sigma_T^{\text{model}}}\,[\log P_T(z) - \log P_T(0)]. \tag{117}$$

Here $\sigma_T^{\text{model}} = T^{-1}|\partial_z \log P_T|_{z=0}\sigma_y\sqrt{2/(1+\rho_y)}$ is the model's standard deviation of twelve-month changes in the $T$-year yield. Its counterpart $\sigma_T^{\text{data}}$ uses AA-and-above zero-coupon yields, with rating cells averaged to maturity–quarter means. It compares the same quarter of consecutive years, uses fifty quarters per maturity and is smoothed by a count-weighted $\alpha_0 + \alpha_1 e^{-T/\alpha_2}$ fit.

The option model fits the conditional mean $\mu_j$ of the next rate state at each grid node $z_j$. This pricing mean differs from the physical conditional mean and the holding return $\mu_T^a$. The fit matches one-step discounted expected prices to model prices at $T = 1, 1.5, \ldots, 30$, using fifteen-point Gauss–Hermite nodes $x_q$ and weights $w_q$:

$$\mu_j = \arg\min_{\mu} \sum_T \left[ \log P_\Delta(z_j) + \log \sum_q \frac{w_q}{\sqrt{\pi}} P_{T-\Delta}(\mu + \sigma_\Delta \sqrt{2} x_q) - \log P_T(z_j) \right]^2 . \tag{118}$$

Prices and drifts are interpolated to 201 nodes spanning the model grid. Transitions $\Pi_{zj}$ over one coupon interval split quadrature mass between adjacent nodes and assign tails to the end nodes $z_1$ and $z_N$. Over the first stub, $\Pi_{zj}^{(s)}$ uses variance $\sigma_s^2$ and drift $\mu^{(s)}(z) = z + (s/\Delta)[\mu(z) - z]$. Each bond starts at $x_0 = z_t$, the observed rate state of Appendix B.2 in the final month of its issue quarter: $z_t = r_t - \bar{r}$, where $r_t$ is twice the continuously compounded two-year corporate zero yield. Every observed state lies inside $[z_1, z_N]$. Contractual dates remain unchanged.

**Step 3: fit the callable value to the observed price.** The callable value is computed backward from $V_n(z) = 100ca_n + 100$ at a trial flat continuous spread $\theta$:

$$C_k(z) = P_\Delta(z) e^{-\theta\Delta} \sum_j \Pi_{zj} V_{k+1}(z_j), \qquad V_k(z) = 100ca_k + \min\{C_k(z), K_k\}. \tag{119}$$

The issuer calls, after paying that date's coupon, when the continuation value $C_k(z)$ exceeds the strike $K_k$. The first stub is discounted using $P_s(z)e^{-\theta s} \sum_j \Pi_{zj}^{(s)} V_1(z_j)$ and interpolated at $x_0$. Bisection finds the OAS $\theta^*$ that matches the dirty callable price, with a repricing error below $10^{-9}$ per 100 face.

**Step 4: remove the call right.** Removing the strike cap in (119), with the OAS and cash flows held fixed, gives the call value $CV_b = \max\{V^{\text{bul}}(\theta^*) - V^{\text{call}}(\theta^*), 0\}$ and the call-adjusted price $P_b^{\text{adj}} = P_b^{\text{dirty}} + CV_b$.

**Step 5: convert the price difference to a yield difference.** Newton's method solves for a flat continuous spread $\hat{s}$ over the maturity-interpolated mean discount curve $D_0$ at $z = 0$, separately for the raw dirty price treated as a bullet and for the adjusted price:

$$100c \sum_k a_k D_0(t_k) e^{-\hat{s} t_k} + 100 D_0(T_b) e^{-\hat{s} T_b} = P_b^{\text{target}}, \tag{120}$$

$$y_b^{\text{target}} = 200 \left[ \exp\left\{ \frac{y_0(T_b) + \hat{s}}{2} \right\} - 1 \right], \qquad y_0(T_b) = -\frac{\log D_0(T_b)}{T_b}. \tag{121}$$

These nominal semiannual yields, in percent, give a call spread $y_b^{\text{raw}} - y_b^{\text{adj}}$ that is nonnegative because the adjusted price is no lower.

**From bonds to the curve.** The 1,828,379 A-or-above pricing bonds are drawn from 1,990,791 issues (Table 1). Bonds must have $T_b \in (1, 30]$, at least two coupons, a complete offering price, coupon and yield, a fixed-rate or original-issue-discount (OID) structure, an offering price of at least $100 - 0.25\,T_b$ per 100 face and consistent state-tax flags. Calendar and price screens exclude 16,295 bonds for irregular first coupons, unsupported conventions, missing accrual dates or unusable call schedules. Eligible bonds carry 97.6%–99.8% of each maturity's three-stage weight. The three-stage mean of their call spreads is Call($T$) in (23), which is subtracted from the observed curve.[27]

[27] Premium bonds with early par calls can require a negative OAS and produce extreme call spreads. I keep them in the sample. Under full three-stage weights, call spreads above 100 basis points contribute at most 0.34 basis points to any maturity's three-stage mean, and 0.25 at thirty years.

The call table uses the screened A-or-above pricing sample. The main puzzle keeps the full observed curve and subtracts the same call adjustment. Screening changes three-stage municipal yields by at most 1.1 bp and the household benchmark by less than 0.1 bp.

**Identification of the volatility level.** The scale $m$ is fitted to AA-and-above callable and noncallable bonds matched within maturity × issue-quarter × rating × half-point-coupon cells. The market gap $\text{gap}_T^{\text{mkt}}$ is their mean callable-minus-noncallable yield difference, averaged across cells and then quarters, with standard error $\text{se}_T$ across quarters. Model call spreads in the same cells give the model gap $\text{gap}_T(m)$. If callability is the only difference, the correct $m$ equates the two gaps. The fitted $m$ is applied to A-or-above bonds without re-estimation. Repricing at $m = 1$ and 1.15 gives the local elasticity

$$e_T \;=\; \frac{\ln \text{gap}_T(1.15) - \ln \text{gap}_T(1)}{\ln 1.15}, \tag{122}$$

so that $\text{gap}_T(m) \approx \text{gap}_T(1)\, m^{e_T}$ locally. The scale $m$ is fitted by weighted least squares at integer maturities six through twenty with twelve or more quarters and positive market and model gaps:

$$\ln \text{gap}_T^{\text{mkt}} \;=\; \ln \text{gap}_T(1) \;+\; e_T \ln m \;+\; \varepsilon_T, \qquad \text{weight } w_T = \left(\text{gap}_T^{\text{mkt}}/\text{se}_T\right)^2, \tag{123}$$

whose solution is

$$\ln \hat{m} \;=\; \frac{\sum_T w_T\, e_T\, \ln\!\left(\text{gap}_T^{\text{mkt}}/\text{gap}_T(1)\right)}{\sum_T w_T\, e_T^2}. \tag{124}$$

The estimate $\hat{m} = 1.077$ has a conditional 95% band of $[0.964, 1.189]$ from 2,000 bootstrap draws. Draws resample quarters independently, preserving within-quarter dependence across maturities, and fix $\text{gap}_T(1)$, $e_T$ and $w_T$. The band contains measured volatility ($m = 1$).

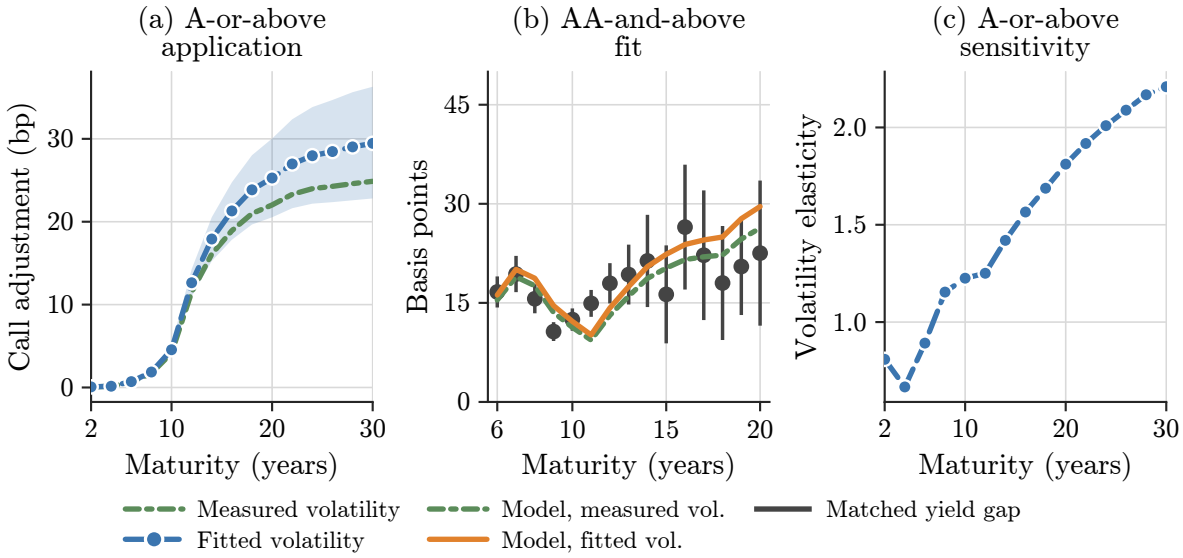



Figure 15: Call adjustment and volatility calibration. (a) A-or-above call adjustments at measured and fitted volatility, with direct-repricing endpoints at the bootstrap band of $\hat{m}$. (b) AA-and-above coupon-matched market gap (±1.96 standard errors) and model gaps at the fitting maturities, 6–20 years. (c) All-bond elasticity $e_T^{\text{all}}$ of the A-or-above call adjustment, distinct from the matched-cell elasticity $e_T$ used in the volatility fit.

**Result.** At $\hat{m}$, the call adjustment for screened A-or-above bonds remains small through ten years, at 4.6 bp, and reaches 29.5 bp at thirty (Table 19 and panel (a) of Figure 15). The directly repriced thirty-year range is $[22.8, 36.3]$ bp, and measured volatility alone gives 24.9 bp.

Table 19: Call adjustment by maturity for screened A-or-above bonds (bp, except the elasticity $e_T^{\text{all}}$). The central estimate is a direct repricing at $\hat{m}$, and the range reprices at the 2.5th and 97.5th percentiles of the bootstrap distribution of $\hat{m}$. The measured-volatility column ($m = 1$) and $e_T^{\text{all}}$ are diagnostics.

| $T$ | measured ($m = 1$) | $e_T^{\text{all}}$ | central ($\hat{m}$) | range |
|---|---|---|---|---|
| 2 | 0.1 | 0.81 | 0.1 | [0.1, 0.1] |
| 4 | 0.1 | 0.67 | 0.2 | [0.1, 0.2] |
| 6 | 0.7 | 0.89 | 0.7 | [0.6, 0.8] |
| 8 | 1.7 | 1.15 | 1.9 | [1.6, 2.1] |
| 10 | 4.2 | 1.23 | 4.6 | [4.0, 5.1] |
| 12 | 11.5 | 1.25 | 12.6 | [11.0, 14.3] |
| 14 | 16.1 | 1.42 | 17.9 | [15.3, 20.5] |
| 16 | 18.9 | 1.57 | 21.3 | [17.8, 24.8] |
| 18 | 21.0 | 1.69 | 23.9 | [19.7, 28.0] |
| 20 | 22.0 | 1.81 | 25.3 | [20.5, 30.0] |
| 22 | 23.3 | 1.92 | 27.0 | [21.6, 32.4] |
| 24 | 24.0 | 2.01 | 28.0 | [22.2, 33.8] |
| 26 | 24.3 | 2.09 | 28.5 | [22.4, 34.7] |
| 28 | 24.6 | 2.17 | 29.0 | [22.6, 35.6] |
| 30 | 24.9 | 2.21 | 29.5 | [22.8, 36.3] |

**Direct-repricing validation of the elasticity.** The all-bond elasticity $e_T^{\text{all}}$ applies the definition of $e_T$ to the A-or-above call adjustment, so the adjustment is approximately its value at $m = 1$ times $m^{e_T^{\text{all}}}$. Direct repricing at $\hat{m}$ and at both band endpoints differs from this approximation by at most 0.156 bp at all fifteen reporting maturities.

**Validation against the noncallable subsample.** If removing callability puts bonds on a comparable basis, the adjusted all-bond curve should track yields on bonds that were noncallable from issuance. Figure 16 makes this comparison in shared maturity, quarter and rating cells. The curves differ by less than three basis points through ten years. On the displayed two-year grid, the largest absolute difference is eight basis points, at eighteen years. At twenty years the adjusted all-bond yield lies 3.5 basis points below the noncallable yield. This checks the adjusted curve against observed noncallable pricing over their common support.

**Robustness to option-pricing dynamics.** I repeat the valuation with a normal one-factor Hull–White model (Hull and White, 1994), whose time-varying drift fits each dated input discount curve exactly, with separately set mean reversion and volatility. With annual mean reversion fixed at 0.10, fitting volatility to the same AA-and-above market gaps at six to twenty years gives diffusion volatility 0.0126 (rates in decimals, time in years). Cash flows, call schedules and the OAS procedure are unchanged. Across the seven dated economies, the Hull–White adjustments for A-or-above bonds are 4.7, 23.4 and 28.5 basis points at ten, twenty and thirty years, against 4.6, 25.3 and 29.5 in the baseline.[28] The root-mean-square error of the Hull–White market-gap fit is 3.7 basis points, against 4.2 for the baseline model. In the same validation cells, the adjusted curve lies 1.9 basis points below the observed noncallable curve at twenty years, compared with 3.5 in the baseline.

[28] The comparison reprices a probability sample of 3,136 draws, covering 2,843 distinct callable contracts, under both models. Weighted mean differences from baseline pricing are added to the baseline population adjustments, which keep the same three-stage weights and zero call values for noncallables. The conditional numerical-integration standard error of the paired thirty-year difference is 0.057 basis points.

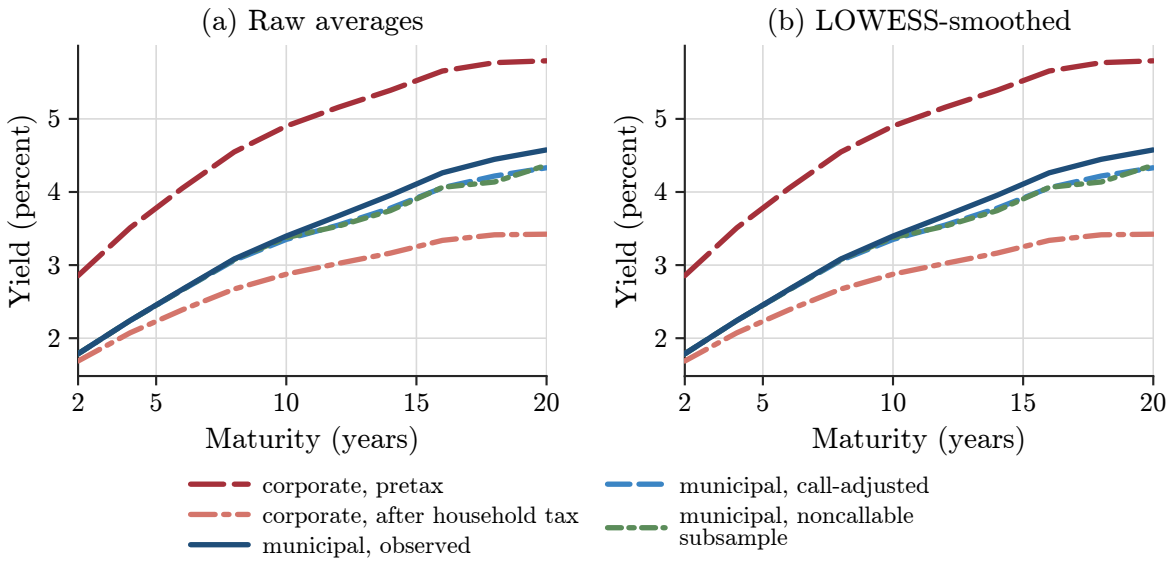



Figure 16: Noncallable validation of Table 19: raw means (a) and locally weighted scatterplot smoothing (LOWESS) (b). Lines show pretax and household after-tax corporate yields alongside observed, call-adjusted and noncallable municipal yields. The call-adjusted and noncallable curves share A-or-above maturity–quarter–rating cells containing noncallables, with fifty or more quarters per maturity and no extrapolation past twenty years. Remaining coupon, credit and other differences prevent identification of individual option values.

**Contract-calendar implementation.** Times are in years under the 30/360 US (NASD) day count, without an end-of-February adjustment or business-day calendar: $\mathrm{DC}(u,v) = [360(Y_v - Y_u) + 30(M_v - M_u) + D_v^* - D_u^*]/360$ for year $Y$, month $M$ and day $D$, where $D_u^* = \min(D_u, 30)$, and $D_v^* = 30$ if $D_v = 31$ and $D_u^* = 30$, otherwise $D_v^* = D_v$. A bond with remaining maturity $T_b$ and coupon interval $\Delta = 0.5$ years has $n = \lceil T_b/\Delta - 10^{-9} \rceil$ coupon dates, a first stub period $s = T_b - \Delta(n-1)$ and coupon nodes $t_k = s + \Delta(k-1)$, $k = 1, \ldots, n$.

With $f = \mathrm{DC}(\text{settlement}, \text{first coupon})$, the first paying node $\underline{k}$ satisfies $\underline{k} - 1 = \mathrm{round}((f-s)/\Delta)$. Each bond must satisfy $|f - s - \Delta(\underline{k}-1)| \le 0.02$ years, $0 \le \underline{k} - 1 \le 2$ and $\underline{k} < n$. Stubs are rounded to 1/360 year to form groups with common $(n, s, \underline{k})$, which keeps $t_n$ within half a day of $T_b$. The accrual fraction $a_k$ is zero before $\underline{k}$, equals $s + \Delta(\underline{k}-1)$ at $\underline{k}$ and equals $\Delta$ afterward. With annual coupon rate $c$, the cash flow at node $k$ is $100ca_k + 100 \cdot \mathbf{1}\{k = n\}$.

Interest accrues from the dated date. If that date precedes settlement, accrued interest $100c\alpha$ is added to the clean target price and to the first paying coupon, with $\alpha = \mathrm{clip}(\mathrm{DC}(\text{dated}, \text{first coupon}) - [s + \Delta(\underline{k} - 1)], 0, 1)$. About 29% of observations require this correction.

**Implementing the recurring call schedule.** Recorded call times $t^{\mathrm{call}} = \mathrm{DC}(\text{settlement}, \text{call date})$ in $(0, T_b + 0.15)$ are kept, with the lowest strike at ties. Each call maps to node $k = \mathrm{round}((t^{\mathrm{call}} - s)/\Delta) + 1$, which must satisfy $\underline{k} \le k \le n$ and $|t_k - t^{\mathrm{call}}| \le 0.02$ years. The strike is $K_k = \infty$ before the first accepted call date. Afterward, the most recent recorded strike applies on every coupon date. Special-call provisions are not modeled. A bond marked callable with no accepted call date is excluded, and other bonds without calls have $CV_b = 0$.

### C.1.3 Adjustment II: the issuance premium

The municipal curve is measured at issuance, whereas the corporate curve represents bonds already trading in the secondary market. Even an otherwise identical bond can be offered at a different price from its subsequent trading price. The issuance-premium adjustment measures the extra yield at the offering relative to seasoned pricing and removes this difference from the municipal curve. Its source of variation is the same

bond's offering yield and its early secondary-market purchase yields, net of market movements between the two dates.

**Comparing the offering with subsequent purchases.** For bond $b$, I compare its offering yield $y_b^{\text{off}}$ with the yield $y_{bj}^{\text{tr}}$ on each customer purchase $j$ within thirty days. A change in market rates during that interval should not be attributed to issuance. I therefore subtract the corresponding yield of seasoned bonds in the same maturity bucket and week, denoted $\bar{y}_k^{\text{seas}}(w)$. The difference in these two relative yields is

$$\text{Premium}_{bj} = [y_b^{\text{off}} - \bar{y}_k^{\text{seas}}(w_b^{\text{off}})] - [y_{bj}^{\text{tr}} - \bar{y}_k^{\text{seas}}(w_{bj})]. \tag{125}$$

A positive value means that the investor received a higher yield at issuance than on an early secondary purchase after removing the market move. Averaging eligible prints gives one estimate per bond. MSRB trades must be customer purchases, because an offering is also an investor purchase. Using customer sales instead would add the bid–ask spread and overlap with the liquidity adjustment. When-issued trades are excluded. The seasoned benchmark is the par-weighted purchase yield of bonds older than three years, with at least 25 prints in the same bucket-week. Both early and seasoned prints must have at least $250,000 par. Comparing the bond with itself removes fixed credit differences, while the seasoned benchmark removes common rate changes. Prints of all sizes give a thirty-year premium of 21.1 bp (Panel B of Table 20).

**Sample and aggregation.** Bonds with an institutional-size print within thirty days of the offering identify the premium. They cover 4% of the target curve, and the estimated premium is applied to the whole curve. The trade tape begins in 2005, so issues from 2000–2004 receive the premium estimated on later issues.

The sample is restricted to the filtered bonds that form the curve. Trade age and benchmark weeks are measured from the offering date, when the primary-market price is set. The dated date, which starts interest accrual, can differ from that date and is not used to measure trade age.

The three-stage averaging of Appendix C.1.1 is applied within issue-quarter and rating cells and the maturity buckets 1–5, 5–8, 8–12, 12–17, 17–23 and 23–31 years. Estimates and bootstrap endpoints are interpolated from the bucket midpoints $(3, 6.5, 10, 14.5, 20, 27)$ to the two-year grid, with ends held flat. In the longest bucket, the mean is 11.9 bp when weighting dollars of par and 12.8 when weighting curve observations.

**Applying the estimate to the zero-coupon curve.** The panel has 1,828,379 A-or-above bonds, each observed once at issuance. The premium is measured in offering yields and subtracted one for one from the zero-coupon curve. The two yields have correlation 0.998 at five years, 0.993 at ten and 0.937 at thirty. Converting coupon to zero-coupon yields raises mean yields at these maturities by 2.8, 12.7 and 58.0 bp.

**Result.** The issuance premium rises steeply with maturity, from 2.9 bp at two years to 12.8 bp at thirty (Table 31 and panel (b) of Figure 17). Its range comes from a 10,000-draw issue-quarter bootstrap that resamples four-quarter blocks with replacement, wrapping at the ends. Lag-one autocorrelations of quarterly bucket means range from $-0.34$ (longest bucket) to 0.02. At 1–5 years, the block bootstrap gives [1.7, 4.2] and a bootstrap of independent quarters gives a similar [1.8, 4.0]. The thirty-year interval is [8.1, 17.3] bp, and the intervals exclude zero at every maturity. Intervals condition on the matched sample and seasoned benchmark.

**Sensitivity to the trade-age window.** In all five trade-age windows of Table 20, premia beyond twelve years exceed those below twelve years. From fourteen days on, the two longest buckets give 9.3–12.8 bp. Three- and seven-day samples have fewer institutional-size prints and are noisier. Panel B shows the same gradient at higher levels.

**Timing, conventions and size.** Across the six maturity buckets, premia average 5.8 bp within three days of the offering, 6.4–7.0 at three to thirty days, 8.5 at thirty to sixty days and 11.6 at sixty to ninety days. In the two longest buckets, par and discount bonds give 10.9 and 13.2 bp and premium bonds give 9.3

Table 20: Issuance premium by maturity bucket (bp) under alternative trade-age windows, in days; the baseline uses thirty. Panel A, the baseline, requires prints of at least $250,000 par in both the early print and the seasoned benchmark. Panel B uses prints of all sizes.

| maturity bucket (yr) | 3 | 7 | 14 | 30 | 60 |
|---|---|---|---|---|---|
| *Panel A: institutional-size prints (baseline)* | | | | | |
| 1–5 | 3.7 | 3.6 | 1.4 | 2.9 | 3.3 |
| 5–8 | 2.8 | 5.2 | 3.0 | 3.3 | 4.0 |
| 8–12 | 3.4 | 5.2 | 3.9 | 4.0 | 4.6 |
| 12–17 | 6.2 | 9.0 | 7.7 | 7.5 | 7.8 |
| 17–23 | 9.1 | 7.3 | 9.3 | 10.8 | 12.2 |
| 23–31 | 9.5 | 7.2 | 11.9 | 12.8 | 12.0 |
| *Panel B: prints of all sizes* | | | | | |
| 1–5 | 4.5 | 6.7 | 5.6 | 3.6 | 2.2 |
| 5–8 | 5.9 | 10.1 | 7.6 | 6.1 | 5.0 |
| 8–12 | 7.2 | 10.5 | 9.0 | 7.9 | 7.4 |
| 12–17 | 15.5 | 16.2 | 13.8 | 12.8 | 12.4 |
| 17–23 | 15.0 | 17.3 | 16.7 | 16.6 | 16.7 |
| 23–31 | 18.1 | 18.7 | 20.0 | 21.1 | 21.0 |

*Notes.* Each cell is the drift-controlled difference-in-differences premium for bonds with at least one print inside the window, averaging each bond's prints inside the window and then averaging bonds with the curve's three-stage weights (Appendix C.1.1). Column heads are windows in days.

and 10.9, so yield-to-worst differences do not explain the gradient. Among all-size prints, long-end premia are positive in every trade-size decile and smallest for the largest trades (panel (c) of Figure 17). Small-trade markups also raise the all-size estimates.

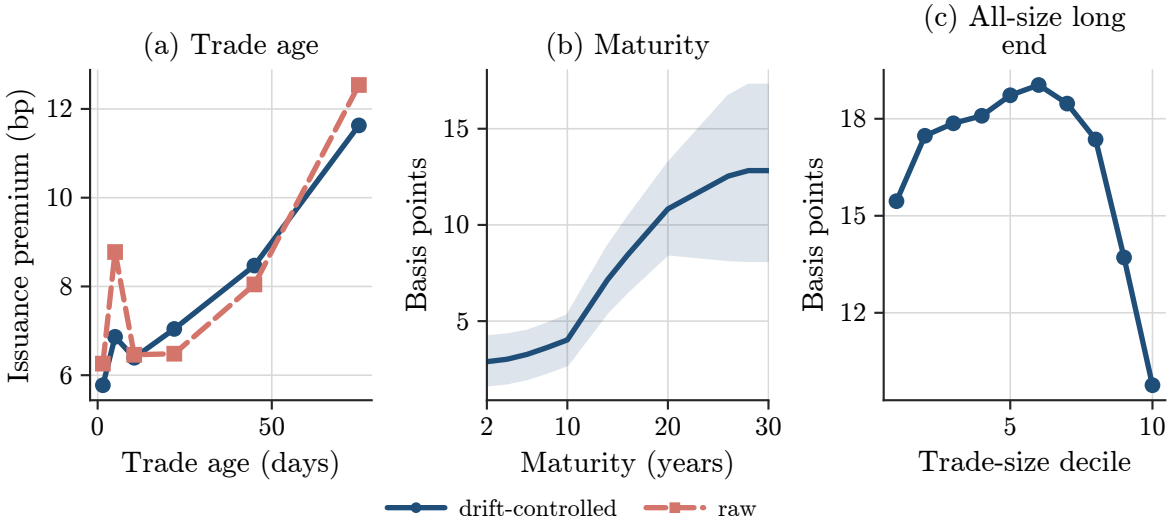



Figure 17: Issuance-premium diagnostics. (a) Institutional-size estimates by trade age, with and without drift controls. (b) Institutional-size maturity profile with block-bootstrap band. (c) All-size long-end estimates by trade-size decile.

**Interpretation.** Net of market drift, the post-offering yield change captures the mismatch between primary and seasoned pricing that matters for the curve comparison. It does so whether the mismatch arises from the issuer's placement cost, dealer markups, ongoing distribution, liquidity seasoning or price discovery. The adjustment removes this mismatch from offering yields.

### C.1.4 Adjustment III: capital-gains taxation

**What the model abstracts from.** The model applies the interest-income after-tax share to the whole holding return in (32). It does not separately track how much of that return is a coupon, accrued discount or a gain on resale. Municipal price gains therefore receive the same exemption as municipal interest, while corporate price gains receive the ordinary-income treatment of corporate interest.

Actual tax treatment distinguishes these sources of return. For example, an original buyer who pays 100 for a par municipal bond and sells it for 105 after a coupon payment has a 5 capital gain even though the coupons are exempt. On a corporate bond, a long-term resale gain can be taxed at the capital-gains rate rather than the higher ordinary-income rate. These differences can raise the municipal yield needed to match the corporate bond's after-tax return. The adjustment measures this additional yield requirement.

**The comparison and what is held fixed.** For a bond bought at issuance and first sold after $h$ years, I make two otherwise identical calculations. Tax treatment $B$, the benchmark, applies the model's interest-income tax classification to both interest and price changes. Tax treatment $A$ separates interest accrual, market discount and capital gains under the rules and rates below. Under each treatment, I find the municipal issuance-yield change that makes the original buyer's expected after-tax return equal to the corporate buyer's. Subtracting the benchmark change from the detailed-tax change removes any pre-existing return difference between the two saved curves. It leaves the yield correction due to the different tax treatment.

The calculation keeps the dated corporate prices, rate process and reinvestment curve fixed. It includes the next buyer's tax as well as the original seller's: a buyer who will owe tax on a discounted municipal bond pays less for it today, reducing the seller's proceeds. The next buyer holds to maturity. This gives a tax-adjusted resale value for each possible sale-date rate state without re-solving investors' portfolios.

**Tax rules, basis and rates.** The taxable resale gain is the sale price less adjusted basis, not simply the sale price less the original purchase price. Original-issue discount (OID) raises basis as interest accrues, so that accrual is not taxed again at sale. Original-issue premium lowers basis without a municipal deduction, and acquisition premium reduces OID. Municipal coupons and OID are exempt, while corporate coupons and OID are ordinary income. Taxable premium amortization assumes the Section 171 election. The original buyer's basis is the adjusted issue price $\text{AIP}_h$ constructed in Step 1, and a secondary buyer's basis starts at the purchase price. Appendix D.3 gives the underlying accrual rules.

Market discount below a quarter point per full remaining year is de minimis and treated as capital gain. Larger discount is ordinary income at disposition. With $n$ remaining years and face 100, the de minimis threshold is

$$\bar{D}_n = 0.0025 \times 100 \times n. \tag{126}$$

Both treatments use the same fixed federal rates in all seven economies, $\tau_{oi} = 0.37$ on ordinary income and $\tau_{cg} = 0.20$ on long-term capital gains. This calculation excludes NIIT and state tax, rather than using the dated household rates in (113). Under treatment $A$, residual gains and losses are treated symmetrically,

$$\mathcal{T}(G) = \tau_{cg} G, \tag{127}$$

so losses are fully deductible in this calculation.

Table 21 summarizes the distinctions that enter the calculation. A bond held from issuance to maturity has no residual gain because its adjusted basis reaches face value. Both treatments therefore give the same after-tax return, so the adjustment is zero.

**Step 1: specify the bonds and sale horizon.** I construct representative par and zero-coupon instruments from each dated economy's discount factors $d_n$. These come from the model's supply-weighted municipal mean log prices or corporate prices at $z = 0$, interpolated annually from two-year knots. At even $T = 4, \ldots, 30$ and sale horizons $h = 2, 4, \ldots, T-2$, the two instruments are

$$\text{PAR:}\quad c = \frac{1 - d_T}{\sum_{n=1}^{T} d_n},\quad \text{offer} = 100,\quad \text{AIP}_j \equiv 100; \qquad \text{ZERO:}\quad \text{offer} = 100\, d_T,\quad \text{AIP}_j = \text{offer} \cdot e^{y_0 j}, \tag{128}$$

Table 21: Tax treatments compared in the capital-gains adjustment

| Return component | Model-based treatment $B$ | Detailed treatment $A$ |
|---|---|---|
| Municipal coupons and OID | Exempt | Exempt |
| Corporate coupons and OID | Ordinary income | Ordinary income |
| Original municipal buyer's residual resale gain | Exempt | Capital gain |
| Original corporate buyer's residual resale gain | Ordinary income | Capital gain |
| Municipal secondary buyer's market discount | Exempt | Capital gain if de minimis, otherwise ordinary income |
| Corporate secondary buyer's market discount | Ordinary income | Capital gain if de minimis, otherwise ordinary income |

*Notes.* Residual gains are measured relative to adjusted basis. Ordinary-income and capital-gains rates are 37% and 20%. Corporate interest and OID are taxed as they accrue. A secondary buyer holds to maturity, with market-discount tax due at redemption. Under treatment $A$, residual capital losses receive full offsets at the capital-gains rate.

with $y_0 = -\log(d_T)/T$, where the adjusted issue price $\mathrm{AIP}_j$ is the original buyer's accreted basis after $j$ years.

For a zero purchased at issue, the interest accruing during year $j+1$ is

$$\mathrm{AIP}_{j+1} - \mathrm{AIP}_j = \mathrm{AIP}_j\big(e^{y_0} - 1\big), \qquad \mathrm{AIP}_T = 100.$$

At a sale in year $h$, the original buyer's residual capital gain is $P_h - \mathrm{AIP}_h$, so accrued interest is not taxed again as a gain.[29]

**Step 2: value resale under each tax treatment.** At sale date $h$, the original investor sells to a new buyer who holds for the remaining $T-h$ years. That buyer compares after-tax municipal cash flows with the return available on the corporate bond. I first infer this corporate return at the saved corporate price, then value municipal resale at the same after-tax return. This is done separately under $B$ and $A$ and at every sale-date rate state.

The seasoned corporate equilibrium price $P_h^{\mathrm{corp}}(z)$ at the sale date is held fixed. For each treatment $\mathcal{M} \in \{B, A\}$, the remaining annual after-tax cash flows $X_j^{\mathcal{M}}$ are contractual payments less the tax due under $\mathcal{M}$. At the corporate price, they determine a buyer's continuous after-tax return:

$$\sum_{j=1}^{T-h} X_j^{\mathcal{M}}\, e^{-r^*_{\mathcal{M}}(h,z)\, j} \;=\; P_h^{\mathrm{corp}}(z). \tag{129}$$

Flows are annual, with continuous discounting and no day-count or semiannual convention. For a corporate buyer holding to maturity, the two treatments differ only in their de minimis treatment of market discount $D$:

$$\mathrm{rate}_{\mathcal{M}}(D) = \begin{cases} 0, & D \le 0, \\ \tau_{cg}, & \mathcal{M} = A \text{ and } D < \bar{D}_n, \\ \tau_{oi}, & \text{otherwise.} \end{cases} \tag{130}$$

[29] IRS Publication 1212, *Guide to Original Issue Discount (OID) Instruments*, and Publication 550, *Investment Income and Expenses*, describe accrual and basis. Annual effective yield $e^{y_0} - 1$ corresponds to continuous issuance yield.

**The municipal buyer's price.** A municipal buyer holding to maturity under $\mathcal{M}$ must also earn $r^*_{\mathcal{M}}$. Benchmark municipal flows $CF_k$ are exempt, giving

$$P^B(h,z) = \sum_{k=1}^{T-h} CF_k\, e^{-r^*_B(h,z)\,k} \;\equiv\; G_B(h,z). \tag{131}$$

The purchase price determines market discount $\mathrm{AIP}_h - P$, while the resulting tax affects how much the buyer will pay. Let $G_A$ be the present value before market-discount tax at the after-tax corporate return under treatment $A$, and $q = e^{-r^*_A(h,z)(T-h)}$ the terminal discount factor. Taxing the discount once at maturity at rate $\tau$ gives $P = G_A - q\,\tau\,(\mathrm{AIP}_h - P)$. Solving this linear equation accounts for the feedback between price and tax:

$$P(\tau) = \frac{G_A(h,z) - q\,\tau\,\mathrm{AIP}_h}{1 - q\,\tau}, \tag{132}$$

and, with $n = T - h$ remaining years, the regime is selected in this order:

$$P^A(h,z) = \begin{cases} G_A, & G_A \geq \mathrm{AIP}_h \quad \text{(no market discount)}, \\ P(\tau_{cg}), & \mathrm{AIP}_h - P(\tau_{cg}) < \bar{D}_n \quad \text{(de minimis)}, \\ P(\tau_{oi}), & \text{otherwise.} \end{cases} \tag{133}$$

When both regime-specific prices satisfy their conditions, this order selects the de minimis price, the higher of the two.

**Step 3: find the municipal issuance-yield correction.** The original investor's after-tax proceeds include coupons received before sale, sale proceeds and taxes due during the holding period. I average these proceeds over sale-date rate states using the rate process, then vary the municipal issuance yield until its expected after-tax holding return equals the corporate return. Repeating this under the two tax treatments gives the two required yield changes.

A scalar $\delta$ shifts the municipal log-discount curve in parallel:

$$\hat{p}^{\,\delta}_n(z) = \hat{p}_n(z) - \delta n \quad \Longleftrightarrow \quad y^{\delta}_n = y_n + \delta. \tag{134}$$

Let $z_i$ denote a possible sale-date rate state and $w_i$ its probability under the rate process. Both sellers are compared using the annualized log of the expected after-tax gross holding return:

$$R \;=\; \frac{1}{h}\,\log \frac{\sum_i w_i\, \mathcal{W}(z_i)}{P_0}, \tag{135}$$

where $\mathcal{W}(z_i)$ is terminal wealth at rate state $z_i$ and $P_0$ is the offering price. Both sellers reinvest interim coupons to $h$ on the unshifted municipal curve. Corporate wealth carries forward coupons net of $\tau_{oi}$ and the annual OID tax. The residual gain $P^{\mathrm{corp}}_h - \mathrm{AIP}^{\mathrm{corp}}_h$ is taxed at $\tau_{oi}$ under $B$ and at $\tau_{cg}$ under $A$. Municipal wealth equals the resale price plus exempt coupons, less $\tau_{cg}$ on its residual gain under $A$. Bisection solves for each treatment's shift,

$$R^{\mathrm{muni}}(\delta_B;\,\text{benchmark}) = R^{\mathrm{corp}}_B, \qquad R^{\mathrm{muni}}(\delta_A;\,\text{detailed tax rules}) = R^{\mathrm{corp}}_A, \tag{136}$$

and the additional yield required solely by the change in tax treatment is

$$\Delta_{\mathrm{cg}}(h,T) = 10^4\big(\delta_A - \delta_B\big)\ \text{basis points}. \tag{137}$$

For a par bond, a higher issuance yield raises the coupon at a 100 issue price. For a zero, it lowers the issue price and changes the accreted basis. The municipal resale price is recalculated for these changed payments and basis using Step 2. When solving for each issuance-yield change, corporate prices, corporate cash flows, the corporate return target and the next buyer's required return stay fixed within that tax treatment. Both treatments use the same rate states, probabilities and reinvestment curve. The correction is computed in

continuously compounded annual yields and enters the additive accounting as a basis-point approximation to the corresponding nominal semiannual shift.

**Step 4: average over bond types and first-sale dates.** The conditional correction depends on both cash-flow structure and when the original buyer sells. I first weight the four corporate–municipal instrument pairs by their observed issuance shares, then weight sale horizons by the probability of a first sale. Finally, I average each calculation equally across the seven dated economies covering 2000–2020.

**Bond-type weights.** Issues with positive coupons and prices of at least 90, including premium issues above 110, are PAR-like. Issues priced below 90 or with zero coupons are ZERO-like. In the Mergent Fixed Income Securities Database (FISD), nonpositive prices denote missing values.

Pooled 1998–2021 amount-weighted ZERO shares are 1.3% municipal and 2.2% corporate (Table 22). The baseline shares use original maturities within three years of each node, pooling when offering amounts are insufficient. At thirty years they are 0.97% municipal and 0.31% corporate. Each of the four pairs in $\{\text{PAR}, \text{ZERO}\}^2$ receives the product of its corporate and municipal shares. The PAR/PAR pair receives 98.7% of the weight at thirty years and drives the adjustment.

Table 22: Four instrument pairs at $T = 30$: pooled issuance weights and each pair's maximum and mean conditional adjustment $\Delta_{\text{cg}}(h, 30)$ over sale horizons $h$, averaged across the seven dated economies. The baseline uses maturity-matched shares, which put PAR/PAR at 98.7% at thirty years.

| Corporate / Municipal | weight (%) | max over $h$ (bp) | mean over $h$ (bp) |
|---|---|---|---|
| PAR / PAR | 96.5 | 6.51 | 5.49 |
| ZERO / PAR | 2.2 | 20.07 | 14.90 |
| PAR / ZERO | 1.3 | 7.71 | 5.74 |
| ZERO / ZERO | 0.0 | 15.75 | 10.91 |

**First-sale weights.** After the issuance-weighted average over instrument pairs, write $\Delta_{\text{cg}}(h, T)$ for the correction at sale horizon $h$. It is multiplied by the probability that the original buyer first sells then. Total turnover cannot be used as that probability: the same dollar of bonds can trade repeatedly, while the next buyer's future tax is already included in the resale price in Step 2. Weighting every subsequent trade would count that tax effect repeatedly.

Turnover $t_a$ is seasoned traded par over outstanding par at age $a$, measured from the MSRB dealer-to-customer tape (71.6 million records, 2005–2021) and pooled over original maturities within $\pm 4$ years of $T$. To exclude primary distribution, I omit when-issued trades and the first quarter after the dated date. A trading segment, a share $\phi(T)$ of par, carries all resales, and the remainder is held to maturity. With aggregate turnover fixed, the first-sale hazard at age $a$ is

$$\lambda_a = \min\big(t_a/\phi,\, 0.99\big), \tag{138}$$

where the cap binds only at the first two ages of thirty-year bonds in the central calibration, at which $t_a/\phi$ would otherwise reach 1.003 and 0.991. The first-sale density and held-to-maturity mass are

$$f_\phi(a) = \phi\,\lambda_a \prod_{j<a}\big(1-\lambda_j\big), \qquad S_\phi(T) = 1 - \phi\Big(1 - \prod_{a<T}\big(1-\lambda_a\big)\Big), \qquad \sum_a f_\phi(a) + S_\phi(T) = 1, \tag{139}$$

and the unconditional adjustment is $\text{CG}(T) = \sum_h \pi_h(T)\,\Delta_{\text{cg}}(h, T)$, where $\pi_h(T)$ sums $f_\phi$ over the ages assigned to sale horizon $h$. Bonds held to maturity receive a zero adjustment. Each $h$ receives ages $h-1$ and $h$. The first node also receives age 0, and the last node, with two years remaining, also receives the final pre-maturity age $T-1$.

**Calibrating the trading segment.** Repeated trades in a narrow segment can generate high turnover without nearly universal resale. Cumulative seasoned volume in each CUSIP, capped at its outstanding par,

bounds the par ever sold (Table 23). The parameter $\phi$ is set so that the model's probability of a first sale within the observation window equals this ceiling, and it applies over the bond's full life. If the homogeneous hazard stays below the ceiling, the segment share is one. At thirty years, par weights over original maturities within four years give a ceiling of 0.602 (single-bin 0.617). Almost all of the trading segment sells within the window, so $\phi = 0.602$ and $S_\phi = 0.398$, with first sales concentrated at young ages.

Table 23: The measured ceiling on the fraction of par ever sold, for selected one-year bins of original maturity. Cumulative traded par counts repeat trades of the same dollar, so the ceiling is an upper bound, and the capital-gains range also scales it by 0.5 and 0.75. Never traded is the par share of CUSIPs with no seasoned trade. The observation window is the par-weighted mean of the shorter of original maturity and the years from offering to 2021. Entries are par-weighted over issues offered from 2005 with at least three observable years.

| orig. maturity | ceiling on $P$(ever sold) | never traded | obs. window (y) | CUSIPs |
|---|---|---|---|---|
| 4 | 0.466 | 0.195 | 3.9 | 124520 |
| 6 | 0.571 | 0.155 | 5.6 | 121412 |
| 8 | 0.556 | 0.227 | 7.0 | 118265 |
| 10 | 0.625 | 0.132 | 8.0 | 110693 |
| 14 | 0.558 | 0.171 | 9.4 | 73189 |
| 18 | 0.546 | 0.216 | 9.7 | 46382 |
| 22 | 0.586 | 0.219 | 10.0 | 13056 |
| 26 | 0.612 | 0.193 | 10.6 | 6551 |
| 30 | 0.617 | 0.159 | 10.3 | 22151 |

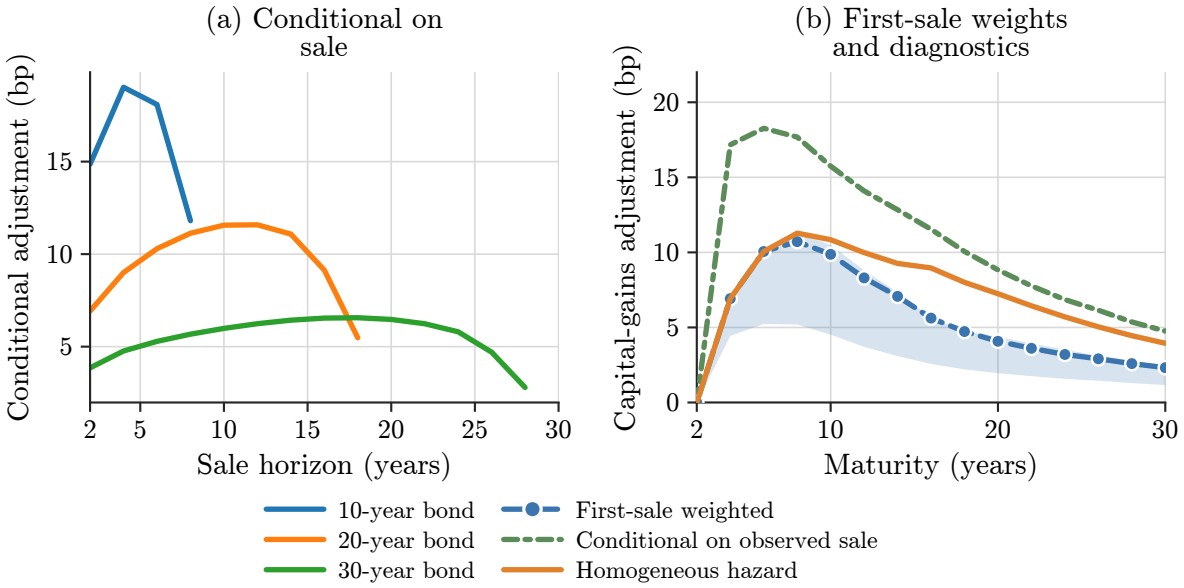



Figure 18: Capital-gains adjustment, averaged over the seven dated economies. (a) Pair-weighted conditional adjustment $\Delta_{\text{cg}}(h, T)$ by sale horizon for ten-, twenty- and thirty-year bonds. (b) Adjustment weighted by first-sale probabilities, with the range of its eighteen variants, and the conditional-on-observed-sale and homogeneous-hazard diagnostics.

**Results and sensitivity.** A positive adjustment is the additional municipal yield attributable to these tax differences and is subtracted from the puzzle. The thirty-year adjustment is 2.3 bp, with range [1.2, 2.5] bp (Table 24 and panel (b) of Figure 18). Because the tax in (127) is linear, the expected tax on the residual gain depends only on its mean.

The range is the envelope of eighteen variants, each averaged equally over the seven 2000–2020 economies. The variants combine three tape windows (the full window and its two halves), three segment multipliers (0.5, 0.75 and 1.0) and two weighting schemes (pooled and maturity-matched). For multipliers below one, the segment share is refitted on the full window before first-sale weights are measured on each half window. The central estimate uses the ceiling. Window and weighting changes move it in both directions.

With the same economies, tax rates and trade tape, the thirty-year adjustment is 4.8 bp conditional on an observed sale and 3.9 bp with a homogeneous hazard ($\phi = 1$). At long maturities, the homogeneous-hazard model sells more par before maturity than the measured resale ceiling allows, and its adjustment lies above the band. At short maturities the fitted segment share is $\phi = 1$, so the reported and homogeneous-hazard estimates coincide.

Table 24: Capital-gains adjustment by maturity (bp, except $\phi$ and $P$(hold)). The reported column uses first-sale probabilities, with the eighteen-variant envelope as its range. The cond. on sale and homogeneous columns are diagnostics. The maximum is over sale horizons of the dated-average, pair-weighted conditional curve. Columns $\phi$ and $P$(hold) give the trading-segment share and the held-to-maturity probability. At two years no sale horizon precedes maturity on the two-year grid, so the adjustment is zero.

| $T$ | cond. on sale | max over $h$ | $\phi$ | $P$(hold) | homogeneous | reported | range |
|---|---|---|---|---|---|---|---|
| 2 | 0.0 | 0.0 | – | 1.00 | 0.0 | 0.0 | [0.0, 0.0] |
| 4 | 17.2 | 17.2 | 1.00 | 0.60 | 6.9 | 6.9 | [4.4, 7.0] |
| 6 | 18.3 | 19.2 | 1.00 | 0.45 | 10.1 | 10.1 | [5.2, 10.3] |
| 8 | 17.7 | 19.9 | 0.88 | 0.40 | 11.3 | 10.7 | [5.2, 11.3] |
| 10 | 15.7 | 19.0 | 0.82 | 0.37 | 10.8 | 9.9 | [4.5, 10.5] |
| 12 | 14.1 | 17.8 | 0.73 | 0.39 | 10.0 | 8.3 | [3.7, 8.8] |
| 14 | 12.8 | 16.2 | 0.67 | 0.40 | 9.3 | 7.1 | [3.1, 7.3] |
| 16 | 11.5 | 14.6 | 0.58 | 0.44 | 9.0 | 5.6 | [2.6, 5.6] |
| 18 | 10.1 | 13.0 | 0.57 | 0.44 | 8.0 | 4.7 | [2.2, 4.8] |
| 20 | 8.8 | 11.6 | 0.57 | 0.43 | 7.2 | 4.1 | [2.0, 4.4] |
| 22 | 7.8 | 10.4 | 0.58 | 0.42 | 6.4 | 3.6 | [1.8, 3.9] |
| 24 | 6.9 | 9.3 | 0.59 | 0.41 | 5.7 | 3.2 | [1.6, 3.5] |
| 26 | 6.1 | 8.5 | 0.60 | 0.40 | 5.0 | 2.9 | [1.5, 3.1] |
| 28 | 5.4 | 7.4 | 0.60 | 0.40 | 4.5 | 2.6 | [1.3, 2.8] |
| 30 | 4.8 | 6.6 | 0.60 | 0.40 | 3.9 | 2.3 | [1.2, 2.5] |

**Integration over sale-date rate states.** Sale-date rate states, $h$ years after $z = 0$, follow the grid-truncated AR(1) density. Integrals use composite Gauss–Legendre nodes $z_i$ and weights $\omega_i$, split where resale prices cross a tax-treatment threshold in either market:

$$w_i \;\propto\; \omega_i \,\frac{1}{\sigma_h}\varphi\Big(\frac{z_i}{\sigma_h}\Big). \tag{140}$$

Here $\varphi$ is the standard normal density and $\sigma_h$ is given by (116). Results with twenty-four and with twelve nodes per interval differ by less than 0.001 bp across 2,940 cells. The residuals of (136), in annualized log return, are at most $10^{-12}$.

### C.1.5 Adjustment IV: liquidity

**The economic comparison.** Investors require compensation for the cost of buying and reselling a bond. For the municipal puzzle, the relevant quantity is the extra trading cost relative to corporate bonds, not the municipal cost alone. I measure costs in comparable municipal and corporate transactions, then spread the difference over a common holding period to express it as an annual yield contribution.

**Measuring comparable trading costs.** The imputed round-trip cost (RTC) is $(\max - \min)/\max$ price within CUSIP × day × size cells that contain both customer buys and customer sells. The extreme

prices need not belong to one dealer's round trip. The comparison uses 2005–2020 MSRB municipal trades and TRACE corporate trades of at least $250,000, made more than three months after issue and more than one year before maturity. Remaining maturity is rounded down to whole years and assigned to the six five-year bins in Table 25. Costs are averaged to bond-months before matching. Trades at the TRACE size caps of 1MM+ and 5MM+ are retained at their lower bounds, and trade size is not otherwise used.

Cells are matched by month × rating family × maturity bin, with ten or more bonds per market. Agency letter grades map AAA to notch one, AA to two–four and A to five–seven on the twenty-two-notch scale. Averages weight ratings equally within months, then months within years, then years.

**Converting a trading cost into an annual yield difference.** Let $c_m$ and $c_c$ be the matched municipal and corporate costs in percent of price. A round-trip cost is incurred once, whereas a yield is paid each year. I annualize one purchase and resale over $h$ years as $c_m/h$ and $c_c/h$, giving the differential in basis points

$$\Delta y^{\text{liq}} = \underbrace{\frac{100\,c_m}{h}}_{\text{municipal cost per year}} - \underbrace{\frac{100\,c_c}{h}}_{\text{corporate cost per year}} = \frac{100(c_m - c_c)}{h}.$$

Both markets use $h = 5$. Trading costs reduce after-tax wealth in both markets, so the differential is not multiplied by an interest-income after-tax share. A positive difference raises the yields municipal holders require and is subtracted from the puzzle.

**Alternative conversions and uncertainty.** A bond-level regression provides a data-based alternative to the holding-period conversion. Regressing A-or-above offering yields on bond-average round-trip costs, with maturity-bin × rating × offer-year effects, gives 15.4 bp per percentage point of price in cost. The $h = 5$ conversion gives 20, so the regression conversion scales it by 0.772. Standard errors for the matched cost difference are the standard deviation of sixteen annual differences divided by $\sqrt{16}$. The reported band spans $h = 2$, $h = 10$, the regression conversion and the central estimate ±1.96 standard errors.

The six bin estimates at $3, 8, 13, 18, 23, 28$ years are interpolated linearly to $2, 4, \ldots, 30$, with end values held flat, to give Liq($T$) in (23). Thirty years uses the 26–30 bin, and the interpolation avoids artificial jumps across maturities.

**Result.** Institutional-size municipal round trips are slightly cheaper than corporate ones before twenty years and costlier afterward (Table 25 and Figure 19). At 26–30 years, municipal and corporate costs of 0.5335% and 0.5310% imply 0.050 bp annually at $h = 5$, reported as 0.1 bp, with a band of $[-2.2, 2.3]$ bp. The higher-cost 21–25 bin produces the peak adjustment of 2.9 bp at twenty-four years.

At $h = 5$, standard errors are 0.8–2.3 bp. The 21–25 estimate is about 1.5 standard errors from zero and depends on the holding period. The issuer-level regressions below provide separate evidence on liquidity and yields within the municipal market.

Table 25: Liquidity by remaining-maturity bin (years), matched trades of 2005–2020. Round-trip costs (RTC) and their municipal-minus-corporate difference in percent of price, and adjustments in bp under three holding periods $h$ and the regression map.

| Bin | Municipal RTC | Corporate RTC | Difference | $h = 2$ | $h = 5$ | $h = 10$ | regr. |
|---|---|---|---|---|---|---|---|
| 1–5 | 0.225 | 0.246 | -0.022 | -1.08 | -0.43 | -0.22 | -0.33 |
| 6–10 | 0.382 | 0.403 | -0.021 | -1.05 | -0.42 | -0.21 | -0.33 |
| 11–15 | 0.582 | 0.588 | -0.006 | -0.30 | -0.12 | -0.06 | -0.09 |
| 16–20 | 0.609 | 0.619 | -0.010 | -0.49 | -0.20 | -0.10 | -0.15 |
| 21–25 | 0.684 | 0.506 | 0.178 | 8.92 | 3.57 | 1.78 | 2.75 |
| 26–30 | 0.533 | 0.531 | 0.003 | 0.13 | 0.05 | 0.03 | 0.04 |

**Sensitivity to the cell rule.** Table 26 varies the minimum number of bonds in each market's cells, using the same underlying trades and size screen. A one-bond minimum admits thin cells and changes

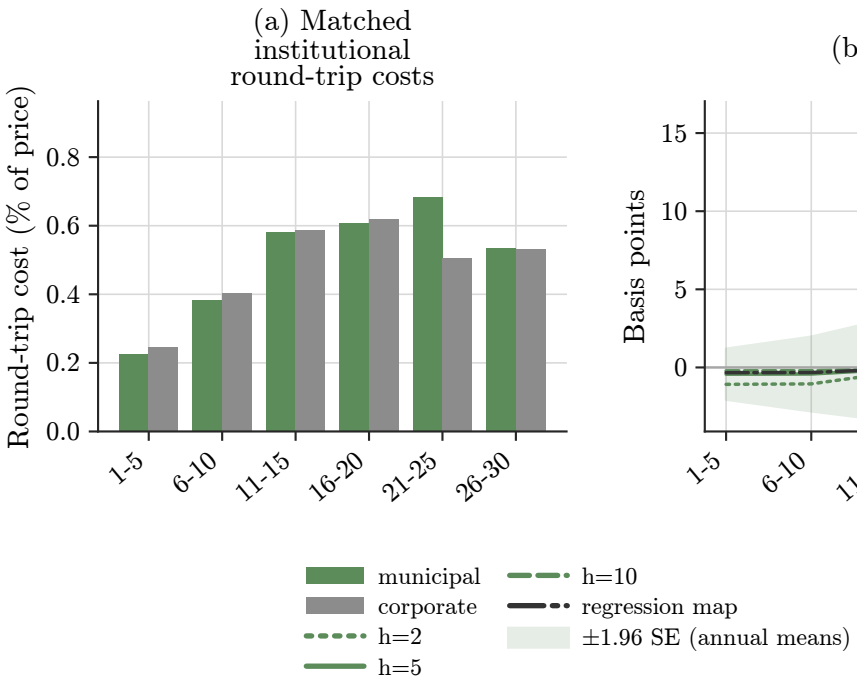



Figure 19: Matched institutional-size customer trades of at least $250,000, 2005–2020. Panel (a): imputed round-trip costs in common month, rating and maturity cells, percent of price. Panel (b): the municipal-minus-corporate cost difference annualized under three holding periods and under the regression map. Shading is ±1.96 standard errors calculated across annual means at $h = 5$.

the 26–30-year estimate from 0.05 to −6.53 bp. A twenty-bond minimum leaves only 19 common months at 11–15 years and changes the estimate there from −0.12 to −2.99 bp. The ten-bond minimum balances coverage against the influence of thin cells.

Table 26: Liquidity adjustment at $h = 5$ (bp) under alternative minimum bond counts per market cell.

| bin | ≥ 1 | ≥ 5 | ≥ 10 (used) | ≥ 20 |
|---|---|---|---|---|
| 1–5 | -0.59 | -0.50 | -0.43 | -0.39 |
| 6–10 | -0.79 | -0.49 | -0.42 | -0.27 |
| 11–15 | -0.12 | 0.42 | -0.12 | -2.99 |
| 16–20 | -0.09 | -0.30 | -0.20 | -0.31 |
| 21–25 | 2.05 | 2.20 | 3.57 | 1.81 |
| 26–30 | -6.53 | -5.81 | 0.05 | 0.30 |

**Issuer-level evidence on liquidity and yields.** The adjustment above compares trading costs across the two markets. Table 27 instead asks whether more liquid municipal issuers have lower offering yields within that market. These regressions provide supporting evidence on liquidity pricing. Their coefficients are not added to the municipal–corporate cost difference and do not enter its uncertainty band.

**Issuer measures and regression specification.** Table 27 relates offering yields to issuer liquidity measured over the bonds' lives. Trading frequency averages, across an issuer's bonds, each bond's mean daily MSRB trades from its dated date to maturity. The spread measure averages each bond's mean imputed round-trip cost (Appendix C.1.5) in the same way. Spreads are in bond-price units, so 0.01 is 1% of price, or 100 basis points.

Table 27: Issuer liquidity and offering yield, by maturity bin.

| | Trading frequency | | Transaction-cost spread | |
|---|---|---|---|---|
| | (1) Above corp median | (2) Top 25% trade | (3) Issuer-avg spread (cont.) | (4) Lowest-spread quartile |
| Maturity 1–5y | $-0.046^{***}$ | $-0.028^{***}$ | $13.446^{***}$ | $-0.057^{***}$ |
| | (0.009) | (0.004) | (1.200) | (0.003) |
| Maturity 6–10y | $-0.007$ | $0.021^{***}$ | $11.624^{***}$ | $-0.048^{***}$ |
| | (0.009) | (0.005) | (1.415) | (0.005) |
| Maturity 11–15y | $-0.011$ | $0.041^{***}$ | $11.860^{***}$ | $-0.040^{***}$ |
| | (0.009) | (0.005) | (0.830) | (0.005) |
| Maturity 16–20y | $-0.035^{***}$ | $0.024^{***}$ | $13.447^{***}$ | $-0.043^{***}$ |
| | (0.010) | (0.006) | (1.084) | (0.006) |
| Maturity 21–25y | $-0.091^{***}$ | $-0.041^{***}$ | $20.427^{***}$ | $-0.104^{***}$ |
| | (0.019) | (0.015) | (2.588) | (0.016) |
| Maturity 26–30y | $-0.093^{***}$ | $-0.082^{***}$ | $27.262^{***}$ | $-0.146^{***}$ |
| | (0.023) | (0.027) | (4.383) | (0.026) |
| Fixed effects | Yes | Yes | Yes | Yes |
| Issuer cluster | Yes | Yes | Yes | Yes |
| Observations | 785,256 | 785,256 | 779,988 | 785,256 |

*Notes.* AA-and-above federally exempt issues, 2000–2021. The outcome is the offering yield in percent, so $-0.046$ is $-4.6$ bp. Each row reports the coefficient on an issuer-average liquidity measure in that maturity bin: (1) an indicator that daily MSRB trades per bond exceed the TRACE median for plain-vanilla corporate debentures and medium-term notes; (2) an indicator for the top quartile of municipal trade frequency; (3) the spread, the imputed round-trip cost in price units, so 13.4 is 13.4 bp per 0.01 of spread; (4) an indicator for the lowest-spread quartile. Fixed effects interact rounded maturity, rating rounded down, state, issue year–month and call status. Standard errors are clustered by issuer (20,346 issuers; 19,806 in column (3)). $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

**Interpreting the coefficients.** At 26–30 years, the indicator coefficients in columns (1), (2) and (4) are about 8–15 bp in absolute value. In column (3), the spread coefficient of 27.262 implies 27.3 bp per 0.01 of spread. At these maturities, more liquid issuers have lower offering yields on all four measures.

### C.1.6 Adjustment V: default

**The economic comparison and credit proxy.** Default risk can raise yields in both markets. The adjustment therefore measures the difference in credit compensation, rather than subtracting municipal credit risk alone. I use credit default swap (CDS) spreads, the annual price of default protection, as proxies for that compensation. State-government Markit CDS supply the municipal proxy, while pooled AAA, AA and A corporate Markit CDS supply the corporate proxy.

**Putting the two credit spreads on a common tax scale.** Following the simple taxable-equivalent convention in Schwert (2017), first express the yield comparison as $y^M/\xi_H - y^C$. Under this convention, the two CDS proxies are credit compensation on that taxable-equivalent scale. The puzzle is instead in household after-tax yield units, $y^M - \xi_H y^C = \xi_H(y^M/\xi_H - y^C)$. The credit difference must therefore receive the same multiplier $\xi_H$. For curve observation $b$, let $s_b^M, s_b^C$ be the matched CDS spreads and $\xi_{H,b}$ its dated household after-tax share. The adjustment is

$$\mathrm{Def}(T) = \mathrm{Avg}_T[\xi_{H,b} s_b^M] - \mathrm{Avg}_T[\xi_{H,b} s_b^C]. \tag{141}$$

A positive differential is the part of the puzzle assigned to higher municipal credit compensation and is subtracted. After-tax shares are applied to individual observations before averaging. They include federal and state rates, NIIT from 2013 and the 2018 change in state-tax deductibility. No further corporate-tax or loss-deduction factor is applied.

**Matching credit data to the yield comparison.** The CDS sample begins in 2004 and contains 1,559,971 curve observations. Its municipal spread is reported in Table 28 under the same dated household tax convention as the full-sample accounting. State general-obligation (GO) CDS are matched to A-or-above bonds by rating, maturity and quarter. The corporate benchmark curve blends bonds rated AAA, AA and A by market weight, so the corporate leg weights pooled quarterly AAA, AA and A corporate CDS by each year's par shares of these ratings among U.S. dollar corporate bonds with at least $250 million outstanding and at most thirty years to maturity. The shares average 8.6%, 22.8% and 68.6% over 2000–2020. After-tax shares are applied before $\mathrm{Avg}_T$ averages within maturity–quarter–whole-notch-rating cells, then across ratings and quarters, retaining each leg's nonmissing support. The AA-and-above sample of Appendix C.11.1 uses the same construction. Beyond quoted tenors, the cumulative spread $T\,\mathrm{CDS}(T)$ is extrapolated linearly, with negative filled spreads floored at zero and quoted spreads left unchanged.

**Coverage and sensitivity.** In the A-or-above CDS panel, exactly observed cells cover 49% at 1–10 years, 47% at 11–20 and 13% at 21–30, and the remaining cells use filled curves (panel (c) of Figure 20). The sensitivity range is the pointwise minimum and maximum of five variants, each with the same household tax mapping: the baseline, cells with at least half of their source observations quoted or interpolated (quality-capped), the 2007-onward sample and CDS–cash basis factors 0.75 and 0.50 (Table 29).

**Result.** The thirty-year default adjustment is 12.2 bp, with a sensitivity range of $[6.1, 16.4]$ bp. It equals 7.3% of the CDS-matched spread and 7.6% of the full-sample denominator $\mathrm{Spread}(T)$, and it remains positive at every maturity (Table 28). A pooled AA corporate leg gives 16.5 bp and a pooled A leg 6.9 bp.

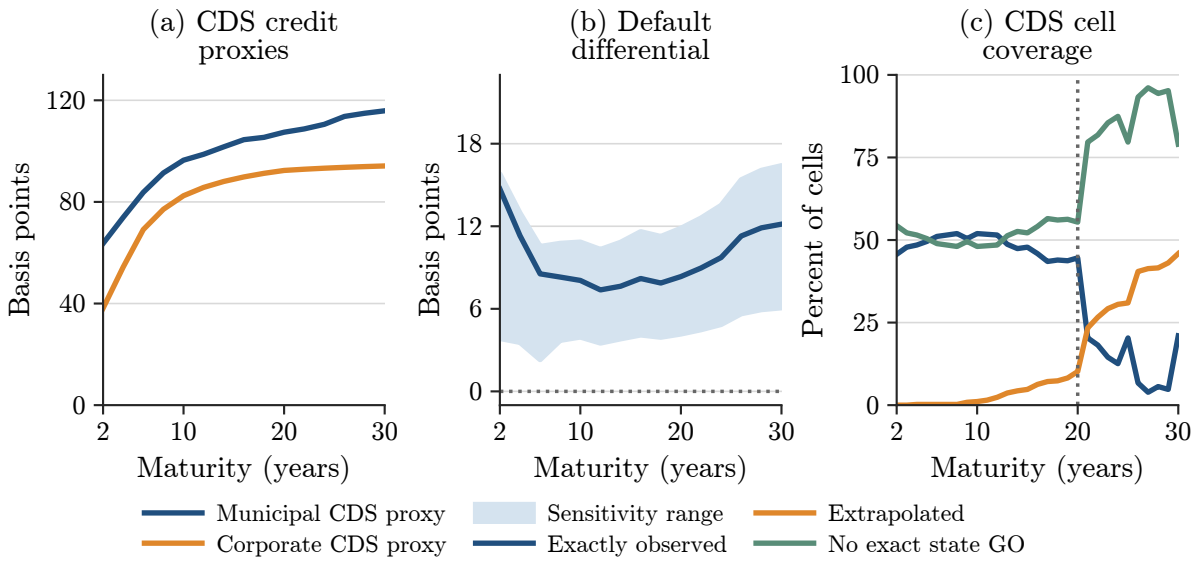



Figure 20: Default adjustment on the A-or-above CDS panel. (a) Municipal and corporate CDS proxies in taxable-equivalent units. (b) Their differential after the dated household tax mapping in (141), with the five-variant sensitivity range. (c) Exactly observed, extrapolated and no-exact-state-GO shares of rating-matched CDS cells. The last two overlap, so the three shares are not a partition. The vertical rule marks twenty years.

Table 28: Default by maturity on the A-or-above CDS-matched sample. Yields in percent, and spread, default adjustment and band in bp. The spread is the municipal yield minus the corporate yield after household tax, with each observation's dated federal and state after-tax share as in the full-sample Spread($T$). The band is the range of Table 29.

| $T$ | municipal yield | corporate after-tax | spread (bp) | adjustment (bp) | band |
|---|---|---|---|---|---|
| 2 | 1.56 | 1.51 | 5.7 | 14.71 | [3.85, 15.79] |
| 4 | 1.99 | 1.85 | 14.6 | 11.31 | [3.56, 12.95] |
| 6 | 2.39 | 2.15 | 24.8 | 8.54 | [2.31, 10.53] |
| 8 | 2.79 | 2.42 | 37.0 | 8.31 | [3.71, 10.75] |
| 10 | 3.12 | 2.63 | 49.5 | 8.06 | [3.95, 10.85] |
| 12 | 3.47 | 2.79 | 68.1 | 7.37 | [3.51, 10.31] |
| 14 | 3.73 | 2.89 | 84.0 | 7.64 | [3.82, 10.83] |
| 16 | 3.93 | 2.96 | 97.1 | 8.20 | [4.10, 11.60] |
| 18 | 4.09 | 3.01 | 108.4 | 7.87 | [3.94, 11.25] |
| 20 | 4.26 | 3.04 | 122.2 | 8.34 | [4.17, 11.85] |
| 22 | 4.36 | 3.07 | 129.0 | 8.97 | [4.48, 12.59] |
| 24 | 4.47 | 3.08 | 138.3 | 9.73 | [4.86, 13.51] |
| 26 | 4.59 | 3.09 | 149.3 | 11.29 | [5.65, 15.37] |
| 28 | 4.58 | 3.09 | 149.5 | 11.88 | [5.94, 16.06] |
| 30 | 4.77 | 3.11 | 165.7 | 12.15 | [6.08, 16.40] |

Table 29: Default sensitivity in basis points. Every column uses the dated household tax mapping of (141). The 2007+ column starts the sample in 2007, the basis columns scale both CDS proxies by the stated factor, and the quality-capped column keeps cells with at least half of their source observations quoted or interpolated. The band in Table 28 spans these five columns.

| $T$ | 2007+ | basis 0.50 | basis 0.75 | baseline | quality-capped |
|---|---|---|---|---|---|
| 2 | 15.79 | 7.35 | 11.03 | 14.71 | 3.85 |
| 6 | 10.53 | 4.27 | 6.41 | 8.54 | 2.31 |
| 10 | 10.85 | 4.03 | 6.05 | 8.06 | 3.95 |
| 14 | 10.83 | 3.82 | 5.73 | 7.64 | 4.00 |
| 18 | 11.25 | 3.94 | 5.91 | 7.87 | 4.53 |
| 22 | 12.59 | 4.48 | 6.72 | 8.97 | 6.84 |
| 26 | 15.37 | 5.65 | 8.47 | 11.29 | 12.64 |
| 30 | 16.40 | 6.08 | 9.11 | 12.15 | 13.85 |

### C.1.7 The combined effect

The tables bring the five calculations back to the same spread. Read Table 30 from left to right: subtract the call adjustment and issuance premium from the raw spread to obtain the puzzle, then subtract the capital-gains, liquidity and default differentials to obtain the residual before the model. Table 31 gives every maturity and each adjustment's band. The combined band adds endpoints of the separate ranges, so it is an accounting envelope rather than a joint confidence interval.

Table 30: What explains the municipal spread over the household benchmark: contributions in basis points to Spread($T$) of equation (113). The two comparability columns remove the call option and the issuance premium from the municipal yield itself, giving the puzzle of Figure 2. The three bond-side differentials are municipal-minus-corporate and do not map to a single adjusted curve.

| | | comparability | | | bond-side differentials | | | | | |
|---|---|---|---|---|---|---|---|---|---|---|
| $T$ | Spread | Call | Issuance premium | Puzzle | Capital gains | Liquidity | Default | Total differential | Share of puzzle | Residual |
| 2 | 9.9 | 0.1 | 2.9 | 6.9 | 0.0 | −0.4 | 14.7 | 14.3 | | −7.3 |
| 6 | 27.9 | 0.7 | 3.3 | 23.9 | 10.1 | −0.4 | 8.5 | 18.2 | 76% | 5.8 |
| 10 | 52.2 | 4.6 | 4.0 | 43.6 | 9.9 | −0.3 | 8.1 | 17.6 | 40% | 26.0 |
| 14 | 87.0 | 17.9 | 7.2 | 61.9 | 7.1 | −0.1 | 7.6 | 14.6 | 24% | 47.3 |
| 18 | 110.3 | 23.9 | 9.6 | 76.9 | 4.7 | −0.2 | 7.9 | 12.4 | 16% | 64.5 |
| 22 | 128.8 | 27.0 | 11.4 | 90.5 | 3.6 | 2.8 | 9.0 | 15.4 | 17% | 75.1 |
| 26 | 145.0 | 28.5 | 12.5 | 104.0 | 2.9 | 1.5 | 11.3 | 15.7 | 15% | 88.3 |
| 30 | 159.8 | 29.5 | 12.8 | 117.5 | 2.3 | 0.1 | 12.2 | 14.5 | 12% | 103.0 |

*Notes.* The puzzle is the spread less the call and issuance-premium adjustments. The total differential sums the three bond-side adjustments, and the share divides that sum by the puzzle (omitted where the puzzle is below 20 basis points). The residual before the model is the puzzle less this total differential. Thirty-year bands, in basis points: call $[22.8, 36.3]$, issuance premium $[8.1, 17.3]$, capital gains $[1.2, 2.5]$, liquidity $[-2.2, 2.3]$, default $[6.1, 16.4]$. Using unrounded band endpoints gives a thirty-year residual before the model of $[85.1, 123.8]$ bp, an envelope with no joint coverage level. Appendix C.1 gives the construction. The issuance-premium estimate uses the 4% of curve bonds with an institutional-size trade within one month of issuance.

The call and issuance-premium adjustments grow with the spread. The three bond-side adjustments together remain near 12–16 bp from twelve years onward. Capital gains peak at intermediate maturities because the resale-tax wedge is spread over the remaining life, and the de minimis threshold rises with that life. At thirty years, the residual before the model is 103.0 bp, and the normalized model leaves an unexplained residual of 5.6 bp (Figure 8). At the shortest maturities, the residual before the model is negative on the full curve, but an adjustment's own matched sample can give a different sign.

Table 31: All five adjustments to Spread($T$) of equation (113) by maturity, in bp. Brackets give each adjustment's band. The total band adds the band endpoints and has no joint coverage level. The final column divides the total by Spread($T$) and is suppressed below 20 bp of observed spread. Table 30 divides the three bond-side adjustments by the puzzle.

| $T$ | call pt band | issuance premium pt band | capital gains pt band | liquidity pt band | default pt band | total pt band | obs. Spread($T$) | % expl. |
|---|---|---|---|---|---|---|---|---|
| 2 | 0.1 [0.1, 0.1] | 2.9 [1.7, 4.2] | 0.0 [0.0, 0.0] | −0.4 [−2.1, 1.2] | 14.7 [3.8, 15.8] | 17.2 [3.5, 21.3] | 9.9 | |
| 4 | 0.2 [0.1, 0.2] | 3.0 [1.8, 4.3] | 6.9 [4.4, 7.0] | −0.4 [−2.2, 1.3] | 11.3 [3.6, 13.0] | 21.0 [7.7, 25.7] | 16.9 | |
| 6 | 0.7 [0.6, 0.8] | 3.3 [2.0, 4.5] | 10.1 [5.2, 10.3] | −0.4 [−2.5, 1.7] | 8.5 [2.3, 10.5] | 22.1 [7.7, 27.8] | 27.9 | 79 |
| 8 | 1.9 [1.6, 2.1] | 3.6 [2.3, 4.9] | 10.7 [5.2, 11.3] | −0.4 [−2.8, 2.0] | 8.3 [3.7, 10.8] | 24.1 [10.1, 31.0] | 41.1 | 59 |
| 10 | 4.6 [4.0, 5.1] | 4.0 [2.7, 5.3] | 9.9 [4.5, 10.5] | −0.3 [−3.1, 2.5] | 8.1 [4.0, 10.8] | 26.2 [12.1, 34.2] | 52.2 | 50 |
| 12 | 12.6 [11.0, 14.3] | 5.6 [4.1, 7.1] | 8.3 [3.7, 8.8] | −0.2 [−3.3, 3.0] | 7.4 [3.5, 10.3] | 33.7 [19.0, 43.5] | 71.2 | 47 |
| 14 | 17.9 [15.3, 20.5] | 7.2 [5.5, 8.9] | 7.1 [3.1, 7.3] | −0.1 [−3.3, 3.0] | 7.6 [3.8, 10.8] | 39.6 [24.3, 50.6] | 87.0 | 46 |
| 16 | 21.3 [17.8, 24.8] | 8.4 [6.5, 10.4] | 5.6 [2.6, 5.6] | −0.2 [−3.0, 2.6] | 8.2 [4.1, 11.6] | 43.4 [28.1, 55.0] | 99.4 | 44 |
| 18 | 23.9 [19.7, 28.0] | 9.6 [7.5, 11.8] | 4.7 [2.2, 4.8] | −0.2 [−2.6, 2.2] | 7.9 [3.9, 11.3] | 45.9 [30.7, 58.1] | 110.3 | 42 |
| 20 | 25.3 [20.5, 30.0] | 10.8 [8.5, 13.2] | 4.1 [2.0, 4.4] | 1.3 [−2.0, 4.6] | 8.3 [4.2, 11.8] | 49.8 [33.2, 64.1] | 123.1 | 40 |
| 22 | 27.0 [21.6, 32.4] | 11.4 [8.4, 14.4] | 3.6 [1.8, 3.9] | 2.8 [−1.4, 7.1] | 9.0 [4.5, 12.6] | 53.8 [34.9, 70.4] | 128.8 | 42 |
| 24 | 28.0 [22.2, 33.8] | 12.0 [8.3, 15.5] | 3.2 [1.6, 3.5] | 2.9 [−1.3, 7.2] | 9.7 [4.9, 13.5] | 55.7 [35.6, 73.5] | 136.9 | 41 |
| 26 | 28.5 [22.4, 34.7] | 12.5 [8.2, 16.7] | 2.9 [1.5, 3.1] | 1.5 [−1.7, 4.6] | 11.3 [5.6, 15.4] | 56.7 [35.9, 74.5] | 145.0 | 39 |
| 28 | 29.0 [22.6, 35.6] | 12.8 [8.1, 17.3] | 2.6 [1.3, 2.8] | 0.1 [−2.2, 2.3] | 11.9 [5.9, 16.1] | 56.4 [35.8, 74.0] | 145.6 | 39 |
| 30 | 29.5 [22.8, 36.3] | 12.8 [8.1, 17.3] | 2.3 [1.2, 2.5] | 0.1 [−2.2, 2.3] | 12.2 [6.1, 16.4] | 56.8 [36.0, 74.7] | 159.8 | 36 |

**Model and observed curves.** The model comparison in Section 5.4 preserves the model's municipal-to-taxable yield ratio and applies it to the observed corporate yield. Subtracting the observed household benchmark then puts the model gap on the same basis as the empirical adjustments. With bars denoting sample averages and superscripts $m$ for model and $d$ for data,

$$\widetilde{y}_T^M = \frac{\bar{y}_T^{M,m}}{\bar{y}_T^{C,m}} \bar{y}_T^{C,d}, \qquad \widetilde{\mathcal{G}}_T = \widetilde{y}_T^M - B_T^d.$$

Here $B_T^d = \bar{\xi}_H(T)\bar{y}_T^{C,d}$ is the observed household benchmark of (113). All yields use nominal semiannual units, and dated after-tax yields are formed before the averaging of Appendix C.1.1. This preserves the observed benchmark without rescaling the model's own averaged benchmark. Figure 8 and Table 4 report the normalized comparison. For reference, the gaps on the model's own taxable curve are −0.2, 36.2 and 107.4 basis points at two, ten and thirty years. Section 5.4 explains the distinction.

The normalized root-mean-square unexplained residual is 7.53 bp (12.94 bp unnormalized), and the maximum absolute residual is 10.72 bp at sixteen years. Fourteen of fifteen absolute residuals are below 10 bp. The signed residuals are −7.83 bp at twelve years, −6.72 at twenty and +5.62 at thirty.

## C.2 Paired issues: Build America Bonds and tax-exempt bonds

Build America Bonds (BABs) provide taxable debt from the same issuers as tax-exempt municipal bonds. Comparing the two checks whether the widening gap survives when issuer identity is held fixed and the taxable reference is also municipal debt. Multiplying BAB offering yields by 0.65, one minus the period's top federal household rate, puts them on the household after-tax scale used in Section 3.8.

**Sample and construction.** The issue-level Mergent sample, one observation per CUSIP, includes bonds issued in 2009Q2–2010Q4 and not subject to the AMT. Table 32 uses AA-and-above issues (agency score $\leq 3$) and A-or-above issues ($\leq 7$). Figure 3 uses the AA-and-above issues and averages yields by bond type and integer maturity $n \in \{1, \ldots, 30\}$.

**Regression specification.** As in equation (4), $y_b$ is the tax-exempt offering yield or 0.65 times the BAB offering yield, $EXMP_b$ indicates a tax-exempt bond, and $n_b$ is maturity in years. I estimate

$$y_b \;=\; \alpha_1\, EXMP_b \;+\; \beta\,(EXMP_b \times n_b) \;+\; \gamma\,(0.65\, y^{C}_{n_b,t_b}) \;+\; \lambda_{j(b)\times n_b\times q_b} \;+\; \varepsilon_b. \tag{142}$$

The corporate control $0.65\, y^{C}_{n_b,t_b}$ varies across issue months within matched quarters, identifying $\gamma$. Equation (4) uses the unscaled control, so only its coefficient, $\gamma_C = 0.65\gamma$, and its standard error change.

**Results.** Table 32 shows a positive maturity gradient in every column. All maturity interactions are significant at one percent, and Panel B bin coefficients rise with maturity. Call matching lowers the longest-bin coefficient by about fifteen percent in the AA-and-above sample but leaves the $[5, 10)$ coefficient unchanged.

**Call features.** Optional calls appear on 49 percent of BABs and 39 percent of tax-exempt bonds. The call-matched columns therefore compare only bonds with the same call status.

**Original-issue premium and discount.** Exact after-tax internal rates of return would treat OID, acquisition premium, interest income and basis separately for BABs and exempt bonds. I therefore check the 0.65 approximation on bonds issued near par. Restricting the call-matched AA-and-above sample to offering prices of 99–101 leaves a maturity gradient of 2.1 bp per year (s.e. 0.6), against 3.1 in the full call-matched sample. The 98–102 restriction gives 2.7 bp (s.e. 0.4). Both remain significant at one percent.

**Interpretation.** The paired comparison of Section 3.8 holds issuer identity fixed, with the BAB yield replacing the corporate benchmark. Under the 0.65 approximation, and with comparable security design and issuance premia, the widening gap supports the participation mechanism. The scaling uses the 35% top federal rate, which is below the 41.1% average household wedge of the AA-and-above curve (Appendix Figure 28). A lower after-tax share would widen the gap. The BAB gaps are closest to that curve's gaps at the shortest and longest maturities and farthest at $[15, 20)$ years.

## C.3 The municipal puzzle on unlimited-tax general-obligation bonds

This check asks whether the declining implied tax persists when the sample is restricted to high-grade bonds backed by the issuer's general taxing power.

**Sample.** Unlimited-tax GO bonds pledge the issuer's full faith and credit and unlimited taxing power, while revenue bonds pledge project revenues. Figure 21 restricts the comparison to AA-and-above unlimited-tax GO bonds[30] passing the main call-pricing and calendar screens (Appendix C.1.2). To test whether the maturity pattern persists among high-grade GO bonds, this check applies only the call and issuance-premium adjustments. Appendix C.1.6 separately assesses the aggregate default adjustment, which uses state-government CDS as a proxy.

**Result.** After call and issuance-premium adjustments, GO implied tax falls from 43.1% at two years to 26.0% at thirty. This 17.1-point decline matches the A-or-above full sample and exceeds the 14.6-point decline of the AA-and-above sample. The screens exclude 4,029 GO bonds (0.86%). Holding the issuance premium fixed, the screens change the decline by 0.49 points relative to the unscreened call-adjusted GO series, almost entirely at the short end. The result agrees with earlier evidence that the long-end residual persists after default and liquidity adjustments and in high-grade or GO samples (Longstaff, 2011; Fleckenstein and Longstaff, 2025; Chalmers, 1998).

## C.4 The single-tax-rate benchmark of Green (1993)

Sections 2 and 3.7 refer to the single-tax-rate explanation of Green (1993).

[30] Mergent security code "K".

Table 32: Paired issues: the tax-exempt-versus-BAB gap widens with maturity.

| | AA and above (score $\leq 3$) | | A or above (score $\leq 7$) | |
|---|---|---|---|---|
| | (1) base | (2) call-matched | (3) base | (4) call-matched |
| *Panel A. Continuous maturity* | | | | |
| Tax-exempt | $-0.102^{***}$ | $-0.091^{**}$ | $-0.062^{**}$ | $-0.043$ |
| | (0.034) | (0.036) | (0.027) | (0.030) |
| Tax-exempt $\times\, n_b$ | $0.034^{***}$ | $0.031^{***}$ | $0.033^{***}$ | $0.030^{***}$ |
| | (0.003) | (0.003) | (0.002) | (0.003) |
| After-tax corporate yield | $0.379^{***}$ | $0.376^{***}$ | $0.264^{*}$ | $0.261^{*}$ |
| | (0.138) | (0.135) | (0.134) | (0.130) |
| *Panel B. Five-year maturity bins* | | | | |
| Tax-exempt | $-0.048$ | $-0.050$ | $-0.007$ | $-0.008$ |
| | (0.033) | (0.033) | (0.024) | (0.024) |
| Tax-exempt $\times\, [5, 10)$ y | $0.192^{***}$ | $0.191^{***}$ | $0.192^{***}$ | $0.190^{***}$ |
| | (0.022) | (0.022) | (0.017) | (0.017) |
| Tax-exempt $\times\, [10, 15)$ y | $0.396^{***}$ | $0.364^{***}$ | $0.368^{***}$ | $0.342^{***}$ |
| | (0.033) | (0.033) | (0.026) | (0.030) |
| Tax-exempt $\times\, [15, 20)$ y | $0.450^{***}$ | $0.390^{***}$ | $0.419^{***}$ | $0.403^{***}$ |
| | (0.041) | (0.039) | (0.029) | (0.039) |
| Tax-exempt $\times\, [20, 25)$ y | $0.576^{***}$ | $0.540^{***}$ | $0.593^{***}$ | $0.510^{***}$ |
| | (0.061) | (0.066) | (0.071) | (0.066) |
| Tax-exempt $\times\, [25, 30]$ y | $0.755^{***}$ | $0.639^{***}$ | $0.996^{***}$ | $0.804^{***}$ |
| | (0.079) | (0.039) | (0.074) | (0.060) |
| After-tax corporate yield | $0.380^{***}$ | $0.378^{***}$ | $0.264^{*}$ | $0.262^{**}$ |
| | (0.138) | (0.135) | (0.134) | (0.130) |
| Issuer $\times$ maturity $\times$ quarter FE | Yes | Yes | Yes | Yes |
| $\times$ call status | No | Yes | No | Yes |
| State cluster | Yes | Yes | Yes | Yes |
| Observations | 29,839 | 28,890 | 61,816 | 59,886 |

*Notes.* FE denotes fixed effects. Offering yields in percent, 2009Q2–2010Q4, excluding AMT bonds. BAB yields are multiplied by 0.65. The corporate-yield row reports $\gamma$, the coefficient on the after-tax control in equation (142). The omitted bin is $[1, 5)$ years. Standard errors are clustered by issuing jurisdiction (51 in columns (1)–(2), 52 in columns (3)–(4)). Call-matched columns add the optional-call flag to the absorbed issuer-by-maturity-by-quarter cell. $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

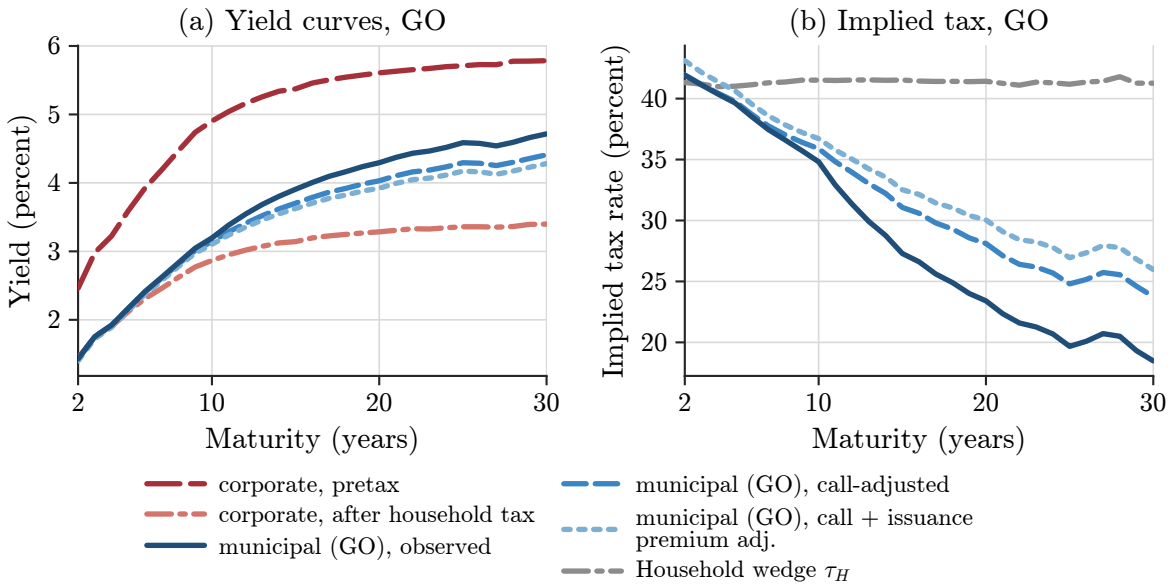



Figure 21: AA-and-above unlimited-tax GO bonds: 465,258 bonds pass the main call-pricing and calendar screens, and curves use the sample's own weights and household benchmark. Calls use the calibrated model. Because GO offering-week matches are scarce, the issuance-premium adjustment uses the all-bond institutional estimate (Appendix C.1.3). Adjusted implied tax falls 17.1 points from two to thirty years, matching the A-or-above full sample.

**Green's mechanism.** Green explains rising municipal-to-taxable yield ratios by the timing of tax on taxable bonds. Buying two five-percent market-discount bonds and shorting one ten-percent par bond of equal maturity defers tax. After purchase, the package's only net cash flow is one dollar of principal at maturity. At rate $\tau$, the package costs $P_T$ and pays $1-\tau(1-P_T)$ after tax at maturity. With $d_T$ the price of a dollar paid at $T$, the package is worth $P_T = (1-\tau)d_T/(1-\tau d_T)$. Asymmetric taxation of long and short positions limits the strategy, because borrowing and short-sale deductions cannot exceed the investment income they support. Green's high-tax investor prices both municipal bonds and the package. Dealers face no such limit and tie taxable par-bond prices to the package, lowering long taxable yields below fully taxed returns. One investor can thus price municipal bonds at every maturity while the yield ratio rises. Fama (1977), by contrast, requires different price-setting investors at different maturities.

**Relation to participation.** Deep market-discount bonds may be scarce, annual OID accrual removes the deferral and shorts consume deduction capacity. Households forgoing the package can accept the household benchmark $\xi_H y_T^C$, below Green's relation. With sufficient capacity, these higher-benefit households set municipal prices, and investors using the strategy do not. Raising municipal yields to Green's relation would require shorting municipal bonds, which faces the frictions in Appendix D.5. His relation therefore need not bind. Nor does it account for observed long yields. Green's own 1987–1990 estimates leave thirty-year municipal yields about twenty taxable-equivalent basis points above his relation, and the gap increases with maturity.

**Green's critique of clientele explanations.** Green rejects earlier clientele explanations because term premia persisted after mid-1980s reforms disrupted bank holdings. Here participation depends on exemption benefits and the duration benefit (Section 4.3), so the clientele that prices each maturity can change. Households alone hold short maturities and set their yields. Institutions with smaller exemption benefits also hold long maturities, so long yields reflect these institutions' tax treatment too. Implied taxes respond to supply and risk capacity. Section 3.4 documents ownership across benefits. I keep $\xi_H y_T^C$ and measure tax timing separately in the capital-gains adjustment (Appendix C.1.4). To distinguish participation

from tax timing alone, P1 tests household state-tax pass-through and P2 tests bank and insurer benefits (Sections 6.1 and 6.2).

## C.5 State-tax pass-through: full regression tables

These tables check whether the declining state-tax pass-through in P1 depends on the yield construction or state credit risk. They retain the specification of Section 6.1 and first compare offering and zero-coupon yields, then add CDS controls on a common sample.

**Sample.** The main P1 regressions use federally exempt AA-and-above issuance (agency score $\leq 3$) from 2000–2017. They exclude zero-maturity bonds and group maturities into $[0, 5)$, $[5, 10)$, $[10, 15)$, $[15, 20)$ and $[20, 30]$ years. The zero-yield columns have 681,838 complete-control observations. Six missing offering yields leave 681,832 in the offering-yield columns.

**Call-adjusted outcome.** The main outcome uses call-adjusted zero-coupon yields. The P1 call adjustment prices each bond on the fifteenth of its issue quarter's final month, while fixed effects use actual issue months. The option-adjusted spread is fitted on the calibrated model's supply-weighted curve for the issue date, with a regular semiannual cash-flow grid. As in the main adjustment, each bond starts at the observed rate state of its issue quarter. The rate loadings of the short-rate tree reproduce the curve's local rate sensitivity at every maturity. Observed call dates map to the nearest grid date within 0.27 years. The call is then exercisable on every grid date from the first mapped date, whereas the main adjustment follows the contractual call calendar (Appendix C.1.2). Callable bonds are excluded when the first call date lies more than 0.27 years before the first semiannual grid date or the call schedule lacks a usable date and price.

**Controls and fixed effects.** All P1 tables retain equation (25)'s six lagged state controls and the rounded average rating, each interacted with all five maturity bins, plus log offering par and security-type indicators. Stock and issuance controls use all observed federally exempt bonds, regardless of rating. Outstanding par equals the offering amount until the date of the reported outstanding amount, and that amount afterward. Fixed effects are cells defined by issue year, issue month, rounded maturity and callability. Observations with missing controls are excluded without imputation, and singleton cells are then removed.

**Estimation and inference.** Unweighted OLS is the main estimator. For the IV comparison, the 1990 rate interacted with each maturity bin instruments the corresponding current-rate interaction. The first stage gives a Kleibergen–Paap (KP) rk Wald $F$ of 10.4 and Sanderson–Windmeijer (SW) conditional statistics of 76.1, 115.9, 104.4, 155.7 and 230.4. Standard errors are clustered by state, with $t(G-1)$ critical values for $G$ states ($t(46)$ in the main sample). Each table's decline $p$-value tests the endpoint equality $\beta_1 = \beta_5$.

### C.5.1 Comparison across yield constructions

Table 33 compares Mergent offering yields, which may be yields to worst for premium callable bonds, with zero-coupon yields from offering prices, with and without call adjustment. Both OLS and IV reject endpoint equality at one percent for every construction.

### C.5.2 State credit-risk controls

Table 34 keeps pre-2018 observations that have a five-year state-government CDS quote in the preceding month and adds that quote interacted with the maturity bins. On this common sample of 345,877 bonds from 25 states, the CDS controls raise the longest-bin OLS estimate from 0.335 to 0.416 and the IV estimate from 0.322 to 0.368. Both estimators retain a declining profile and reject endpoint equality at five percent in every row.

Table 33: P1 across yield constructions, 2000–2017.

| | Offering yield | | Call-adjusted zero | | Unadjusted zero | |
|---|---|---|---|---|---|---|
| | OLS | IV | OLS | IV | OLS | IV |
| $\beta_1$: $[0, 5)$ years | $0.985^{***}$ | $1.084^{***}$ | $1.069^{***}$ | $1.088^{***}$ | $1.080^{***}$ | $1.102^{***}$ |
| | (0.238) | (0.249) | (0.195) | (0.225) | (0.196) | (0.227) |
| $\beta_2$: $[5, 10)$ years | $0.635^{***}$ | $0.648^{***}$ | $0.677^{***}$ | $0.694^{***}$ | $0.683^{***}$ | $0.693^{***}$ |
| | (0.075) | (0.093) | (0.069) | (0.086) | (0.069) | (0.087) |
| $\beta_3$: $[10, 15)$ years | $0.414^{***}$ | $0.410^{***}$ | $0.484^{***}$ | $0.447^{***}$ | $0.515^{***}$ | $0.479^{***}$ |
| | (0.053) | (0.075) | (0.096) | (0.118) | (0.100) | (0.117) |
| $\beta_4$: $[15, 20)$ years | $0.268^{***}$ | $0.257^{***}$ | $0.390^{***}$ | $0.326^{***}$ | $0.428^{***}$ | $0.360^{***}$ |
| | (0.060) | (0.067) | (0.102) | (0.116) | (0.108) | (0.121) |
| $\beta_5$: $[20, 30]$ years | $0.193^{***}$ | $0.152^{**}$ | $0.305^{***}$ | $0.229^{***}$ | $0.335^{***}$ | $0.255^{***}$ |
| | (0.052) | (0.057) | (0.075) | (0.081) | (0.078) | (0.085) |
| Decline $\beta_1 - \beta_5$ | $0.791^{***}$ | $0.933^{***}$ | $0.764^{***}$ | $0.859^{***}$ | $0.745^{***}$ | $0.847^{***}$ |
| $p$-value of decline | 0.004 | $< 0.001$ | 0.002 | $< 0.001$ | 0.003 | 0.001 |
| Issue month × maturity × call FE | Yes | Yes | Yes | Yes | Yes | Yes |
| State cluster | Yes | Yes | Yes | Yes | Yes | Yes |
| Observations | 681,832 | 681,832 | 681,838 | 681,838 | 681,838 | 681,838 |
| First-stage $F$ (Kleibergen–Paap) | — | 10.4 | — | 10.4 | — | 10.4 |

*Notes.* FE denotes fixed effects. Outcomes are implied state taxes (24) from Mergent offering yields or from zero-coupon yields with and without the call adjustment. The call-adjusted zero is the main outcome. Controls, fixed effects and state-clustered inference (47 states) follow Appendix C.5, 2000–2017. IV instruments each bin's current rate with its 1990 rate. The decline row reports $\beta_1 - \beta_5$. $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

Table 34: P1 sample and credit controls, 2000–2017.

| Sample | Obs. | States | $\beta_1$, OLS | $\beta_5$, OLS | Decline $p$ | $\beta_1$, IV | $\beta_5$, IV | Decline $p$ |
|---|---|---|---|---|---|---|---|---|
| Main sample | 681,838 | 47 | $1.069^{***}$ | $0.305^{***}$ | 0.002 | $1.088^{***}$ | $0.229^{***}$ | $< 0.001$ |
| | | | | (0.075) | | | (0.081) | |
| Observed CDS sample | 345,877 | 25 | $1.210^{***}$ | $0.335^{***}$ | 0.009 | $1.349^{***}$ | $0.322^{***}$ | 0.009 |
| | | | | (0.081) | | | (0.098) | |
| CDS sample + CDS × bin | 345,877 | 25 | $1.400^{***}$ | $0.416^{***}$ | 0.011 | $1.391^{***}$ | $0.368^{***}$ | 0.025 |
| | | | | (0.068) | | | (0.087) | |

*Notes.* Controls, fixed effects and state-clustered inference follow Appendix C.5, 2000–2017. The observed-CDS sample keeps bonds with a prior-month five-year state-government CDS quote, without imputation, and the last row adds that quote interacted with the maturity bins. IV uses bin-interacted 1990 rates. Parentheses give standard errors of $\beta_5$, and decline $p$-values test $\beta_1 = \beta_5$. $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

## C.6 State-tax pass-through: instrument vintages, weighting and further checks

These checks ask whether P1 depends on the historical tax-rate instrument, the weight assigned to each state or any single state. They retain the controls, fixed effects and state-clustered inference of Appendix C.5, then examine state fixed effects and the post-2018 tax regime separately.

**Instruments.** The baseline IV instrument is the state's 1990 top income-tax rate. The exclusion restriction is that, conditional on the controls and fixed effects, the historical rate affects current implied taxes only through the current rate. Tables 35 and 36 compare instruments from different years (vintages) in 1945–2000 on the same pre-2018 sample.

**Continuous-maturity specification.** Table 35 replaces the five current-rate bin interactions in (25) with the current rate and its interaction with maturity in years, $n_b$. The interaction coefficient is the

pass-through gradient per year, and IV uses the historical rate and its maturity interaction. OLS and all instrument vintages give negative gradients, unweighted and with equal total weight per state.

Table 35: Continuous-maturity P1 checks, 2000–2017.

| | | Unweighted | | | | Equal per state | |
|---|---|---|---|---|---|---|---|
| Vintage | Corr. | Gradient | s.e. | KP $F$ | SW $F$ | Gradient | KP $F$ |
| OLS | — | $-0.0392^{***}$ | (0.0131) | — | — | $-0.0326^{**}$ | — |
| 1945 | 0.40 | $-0.0535^{**}$ | (0.0251) | 8.0 | 19.2 | $-0.0440$ | 2.8 |
| 1950 | 0.31 | $-0.0607^{**}$ | (0.0273) | 7.8 | 18.4 | $-0.0878$ | 1.8 |
| 1960 | 0.38 | $-0.0620^{**}$ | (0.0254) | 9.0 | 28.8 | $-0.0997^{*}$ | 2.5 |
| 1970 | 0.51 | $-0.0574^{***}$ | (0.0201) | 11.5 | 45.2 | $-0.0795^{***}$ | 7.3 |
| 1980 | 0.76 | $-0.0424^{**}$ | (0.0164) | 10.6 | 103.4 | $-0.0385^{**}$ | 26.9 |
| 1990 | 0.82 | $-0.0440^{***}$ | (0.0127) | 30.1 | 404.1 | $-0.0307^{**}$ | 25.9 |
| 2000 | 0.88 | $-0.0399^{***}$ | (0.0124) | 65.1 | 998.5 | $-0.0258^{*}$ | 36.5 |

*Notes.* Controls, fixed effects and state-clustered inference follow Appendix C.5, 2000–2017. The gradient is the coefficient on the current rate times maturity $n_b$, the change in pass-through per year. IV instruments the current rate and its maturity interaction with the vintage's rate and its maturity interaction. KP $F$ is the Kleibergen–Paap statistic and SW $F$ the Sanderson–Windmeijer conditional statistic for the interaction. Corr. is the cross-state correlation between each state's mean current top rate over the estimation sample and its vintage-year top rate. $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

**Binned profiles.** Table 36 compares the five-bin IV profiles with OLS, and Figure 22 plots them. Every vintage rejects endpoint equality at five percent.

Table 36: P1 instrument-vintage checks, 2000–2017.

| Estimator/vintage | $\beta_1$ | $\beta_2$ | $\beta_3$ | $\beta_4$ | $\beta_5$ | Decline $p$ | KP $F$ |
|---|---|---|---|---|---|---|---|
| OLS (main) | $1.069^{***}$ | $0.677^{***}$ | $0.484^{***}$ | $0.390^{***}$ | $0.305^{***}$ | 0.002 | — |
| | (0.195) | (0.069) | (0.096) | (0.102) | (0.075) | | |
| 1945 | $1.257^{***}$ | $0.868^{***}$ | $0.551^{***}$ | $0.363^{**}$ | $0.245^{**}$ | 0.027 | 3.1 |
| | (0.421) | (0.158) | (0.158) | (0.143) | (0.103) | | |
| 1950 | $1.453^{***}$ | $0.942^{***}$ | $0.622^{***}$ | $0.428^{***}$ | $0.313^{***}$ | 0.024 | 3.1 |
| | (0.473) | (0.181) | (0.153) | (0.143) | (0.116) | | |
| 1960 | $1.483^{***}$ | $0.942^{***}$ | $0.627^{***}$ | $0.432^{***}$ | $0.332^{***}$ | 0.014 | 3.6 |
| | (0.439) | (0.167) | (0.147) | (0.140) | (0.113) | | |
| 1970 | $1.394^{***}$ | $0.858^{***}$ | $0.566^{***}$ | $0.400^{***}$ | $0.325^{***}$ | 0.006 | 4.4 |
| | (0.350) | (0.125) | (0.143) | (0.141) | (0.114) | | |
| 1980 | $1.084^{***}$ | $0.756^{***}$ | $0.528^{***}$ | $0.374^{***}$ | $0.280^{***}$ | 0.010 | 3.8 |
| | (0.273) | (0.083) | (0.121) | (0.120) | (0.084) | | |
| 1990 | $1.088^{***}$ | $0.694^{***}$ | $0.447^{***}$ | $0.326^{***}$ | $0.229^{***}$ | $< 0.001$ | 10.4 |
| | (0.225) | (0.086) | (0.118) | (0.116) | (0.081) | | |
| 2000 | $1.072^{***}$ | $0.678^{***}$ | $0.480^{***}$ | $0.378^{***}$ | $0.280^{***}$ | 0.001 | 23.5 |
| | (0.215) | (0.079) | (0.095) | (0.096) | (0.069) | | |

*Notes.* Controls, fixed effects and state-clustered inference follow Appendix C.5, 2000–2017. IV instruments each bin's current rate with the vintage's top rate in that bin. KP $F$ is the Kleibergen–Paap statistic, and decline $p$-values test $\beta_1 = \beta_5$. $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

**Weighting.** Every weighting in Table 37 retains a declining OLS profile, with endpoint differences significant at one percent.

**State fixed effects.** Adding state fixed effects to the pre-2018 OLS specification, on 681,838 bonds, gives a short-bin coefficient of 0.665 (s.e. 0.201) and a long-bin coefficient of $-0.112$ (0.212). The decline is

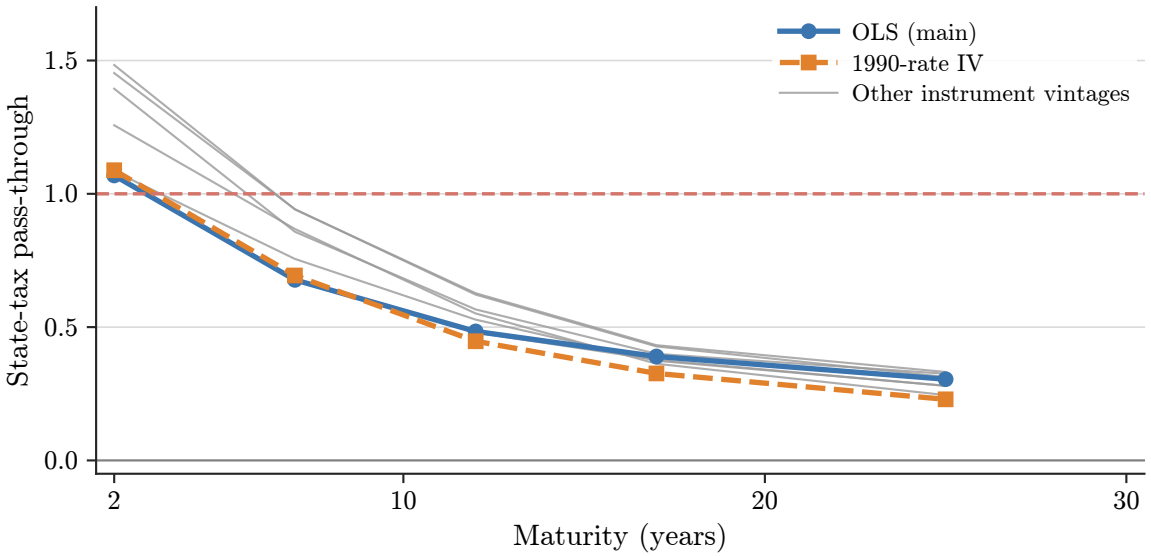



Figure 22: State-tax pass-through using historical tax-rate instruments, 2000–2017. Blue shows the main OLS profile. Orange shows IV using the 1990 rate. Gray lines use the 1945, 1950, 1960, 1970, 1980 and 2000 rates. The dashed red line is the representative-tax benchmark. Table 36 reports state-clustered uncertainty.

Table 37: P1 OLS estimates under four weighting schemes, 2000–2017.

| Weighting | $\beta_1$ | $\beta_2$ | $\beta_3$ | $\beta_4$ | $\beta_5$ | Decline $p$ |
|---|---|---|---|---|---|---|
| Unweighted (main) | $1.069^{***}$ | $0.677^{***}$ | $0.484^{***}$ | $0.390^{***}$ | $0.305^{***}$ | 0.002 |
| | (0.195) | (0.069) | (0.096) | (0.102) | (0.075) | |
| Equal per state | $0.863^{***}$ | $0.543^{***}$ | $0.374^{***}$ | $0.308^{***}$ | $0.234^{**}$ | 0.006 |
| | (0.191) | (0.082) | (0.101) | (0.111) | (0.088) | |
| Equal per state-year | $0.911^{***}$ | $0.509^{***}$ | $0.317^{***}$ | $0.261^{***}$ | $0.211^{***}$ | 0.002 |
| | (0.190) | (0.069) | (0.080) | (0.090) | (0.074) | |
| $1/\sqrt{n_{\text{state}}}$ | $0.990^{***}$ | $0.617^{***}$ | $0.420^{***}$ | $0.344^{***}$ | $0.282^{***}$ | 0.006 |
| | (0.202) | (0.080) | (0.109) | (0.121) | (0.091) | |

*Notes.* OLS with the controls, fixed effects and state-clustered inference of Appendix C.5, 2000–2017. The equal-weight rows give each state or state-year the same total weight, and $n_{\text{state}}$ is the bond's state observation count, all measured after the complete-case and singleton restrictions. Decline $p$-values test $\beta_1 = \beta_5$. $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

0.777 (0.220, $p = 0.001$). State effects absorb fixed tax levels. The five excluded instruments, the 1990 rate interacted with each maturity bin, sum to the state's 1990 rate. Because that rate is collinear with the state effects, IV is not identified.

**Leave-one-state-out.** All 47 single-state omissions retain a positive short-to-long decline under both OLS and IV (Figure 23). OLS long-bin estimates span 0.232–0.482, with declines 0.452–0.889. IV long-bin estimates span 0.116–0.315, with declines 0.669–1.004. The largest endpoint-test $p$-values are 0.062 for OLS and 0.036 for IV.

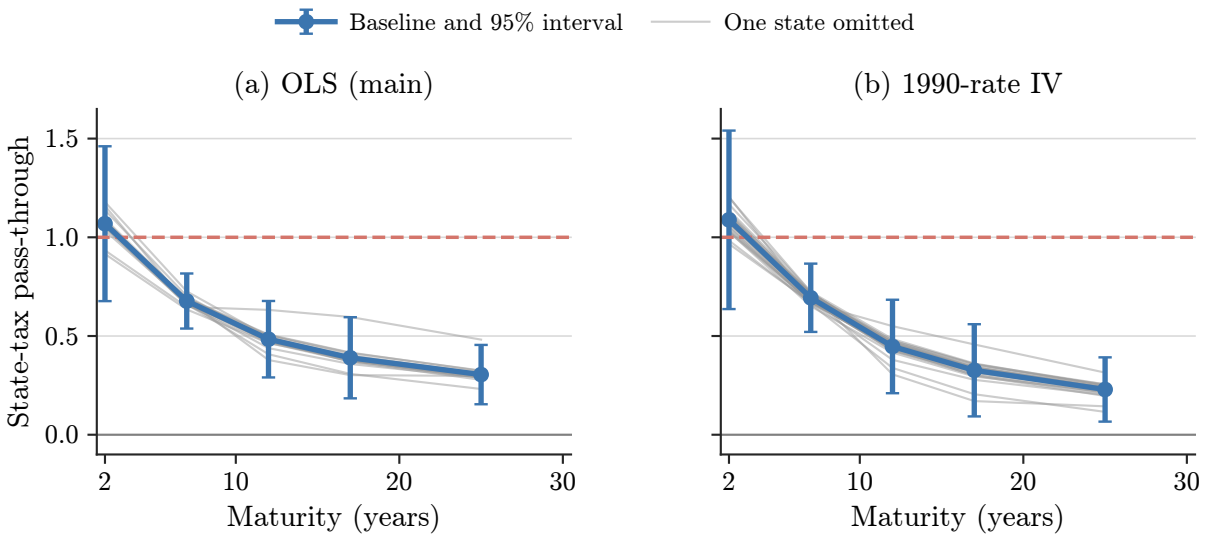



Figure 23: Leave-one-state-out P1 profiles for 2000–2017. Gray lines show each of the 47 state omissions, and blue shows the baseline estimates with state-clustered 95% intervals. Panel (a) uses OLS and panel (b) uses the 1990-rate instruments. The dashed line is the representative-tax benchmark.

**The two tax regimes.** The main P1 sample ends in 2017, so its product of federal and state after-tax shares matches the six pre-2018 model economies in Figure 9. From 2018, the TCJA capped SALT deductions at $10,000 per return. When the cap binds, an additional dollar of state tax is no longer federally deductible. The combined after-tax share on taxable interest is then $\xi_t^{\text{fed}} - \tau_s^{\text{st}}$ (notation of Section 6.1). The implied state tax for this regime is

$$\mathit{ImpliedStateTax}_{s,t}^{\text{cap}}(T) \;\equiv\; \xi_t^{\text{fed}} - \frac{y_{s,t}^{M}(T)}{y_t^{C}(T)}. \tag{143}$$

For a household whose cap binds, this measure equals the statutory state rate under representative-tax indifference.[31] Table 38 estimates each regime separately. Both estimators reject endpoint equality at one percent in each period.

## C.7 The bank-qualified discontinuity: construction, robustness and placebos

Bank qualification changes banks' tax benefit without changing households' exemption. It therefore tests whether an institution-specific benefit has a larger yield effect at long maturities, as P2 predicts. This subsection derives the capitalization benchmarks, explains the issuance-cap design and funding-cost measure, and checks alternative windows, periods and placebo cutoffs. Table 39 reports the full five-bin estimates of (27), extending Section 6.2.

### C.7.1 Benchmarks and the designation premium

**The two benchmarks.** Both benchmarks of (26) compare annual after-tax carry per dollar of basis on matched bonds, in the notation of Section 6.2. I suppress dates, and $y_k^C > 0$ is the taxable yield at maturity $k$.

[31] See the IRS, 2017 Schedule A instructions, and its summary of the 2018 changes.

Table 38: P1 under the two federal/state tax regimes.

| | 2000–2017 (main) | | 2018–2020 | |
|---|---|---|---|---|
| | OLS | IV | OLS | IV |
| $\beta_1$: $[0, 5)$ years | $1.069^{***}$ | $1.088^{***}$ | $0.948^{***}$ | $0.909^{***}$ |
| | (0.195) | (0.225) | (0.182) | (0.246) |
| $\beta_2$: $[5, 10)$ years | $0.677^{***}$ | $0.694^{***}$ | $0.834^{***}$ | $0.879^{***}$ |
| | (0.069) | (0.086) | (0.121) | (0.143) |
| $\beta_3$: $[10, 15)$ years | $0.484^{***}$ | $0.447^{***}$ | $0.446^{***}$ | $0.466^{***}$ |
| | (0.096) | (0.118) | (0.061) | (0.074) |
| $\beta_4$: $[15, 20)$ years | $0.390^{***}$ | $0.326^{***}$ | $0.312^{***}$ | $0.323^{***}$ |
| | (0.102) | (0.116) | (0.078) | (0.105) |
| $\beta_5$: $[20, 30]$ years | $0.305^{***}$ | $0.229^{***}$ | 0.088 | 0.082 |
| | (0.075) | (0.081) | (0.091) | (0.109) |
| Decline $\beta_1 - \beta_5$ | $0.764^{***}$ | $0.859^{***}$ | $0.860^{***}$ | $0.827^{***}$ |
| $p$-value of decline | 0.002 | $< 0.001$ | $< 0.001$ | 0.008 |
| Issue month × maturity × call FE | Yes | Yes | Yes | Yes |
| State cluster | Yes | Yes | Yes | Yes |
| Observations | 681,838 | 681,838 | 146,174 | 146,174 |
| First-stage $F$ (Kleibergen–Paap) | — | 10.4 | — | 7.9 |

*Notes.* FE denotes fixed effects. Controls, fixed effects and state-clustered inference (47 states) follow Appendix C.5. IV uses bin-interacted 1990 rates. The outcome is (24) before 2018 and, for a household whose SALT deduction cap binds, (143) from 2018. $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

The representative-tax benchmark has $y_k^N = (1 - \tau_H) y_k^C$, where $\tau_H$ is the top household rate. At this yield, a designated bond earns the bank $y_k^N - 0.2\,\tau_B c$ after the tax cost of its disallowed interest deduction. A taxable bond earns $(1 - \tau_B) y_k^C$, so the taxable bond has an advantage of

$$\underbrace{(1 - \tau_B) y_k^C}_{\text{taxable alternative}} - \underbrace{\left[(1 - \tau_H) y_k^C - 0.2\,\tau_B c\right]}_{\text{qualified bond at the household yield}} = (\tau_H - \tau_B) y_k^C + 0.2\,\tau_B c > 0. \tag{144}$$

It is positive whenever $\tau_H > \tau_B - 0.2\,\tau_B c / y_k^C$, which holds in particular when household rates are at least banks' and funding costs are positive. Banks still prefer taxable bonds and designation leaves household taxes unchanged, so $y_k^Q = y_k^N$ and $\varpi_k = 0$ at every maturity.

In the bank-only benchmark, each municipal bond gives banks the same after-tax carry as taxable bonds:

$$\begin{aligned} y_k^N - \underbrace{\tau_B c}_{\text{full disallowance}} &= (1 - \tau_B) y_k^C, \\ y_k^Q - \underbrace{0.2\,\tau_B c}_{\text{qualified disallowance}} &= (1 - \tau_B) y_k^C. \end{aligned} \tag{145}$$

Subtracting the two conditions cancels the taxable yield and gives $\varpi_k = 0.8\,\tau_B c = X$. With constant funding costs, the premium in basis points is flat in maturity. The expected annual benefit $\bar{X}_b$ (Appendix C.7.3) evaluates this benefit at expected annual funding costs.

**Derivation of the designation premium.** The derivation compares matched bonds of maturity $k$ in supply $q_k > 0$, with outside opportunities and non-tax characteristics fixed. The designation benefit $u$ is zero for a non-designated bond and $X$ for a designated one. At yield $y$, household demand is $d_k^H(y)$ and bank demand $d_k^B(y + u)$, both reflecting baseline funding costs, taxes and risk compensation. For each

Table 39: Bank-qualified designation and yield premia by maturity.

| | (1) | (2) | (3) |
|---|---|---|---|
| Dependent variable | BQ designation | Offering yield (bp) | Yield / expected benefit |
| Estimator | OLS | 2SLS | 2SLS |
| Regressor, interacted with each maturity bin | Below-cap indicator | BQ designation | BQ designation |
| Reported effect | Designation increase | Yield reduction | Scaled yield reduction |
| 1–2 years | 0.42*** | 1.8 | 0.17 |
| | (0.015) | (2.2) | (0.11) |
| 2–5 years | 0.42*** | 8.7*** | 0.31*** |
| | (0.015) | (2.2) | (0.08) |
| 5–10 years | 0.44*** | 13.8*** | 0.41*** |
| | (0.015) | (2.2) | (0.06) |
| 10–20 years | 0.45*** | 18.4*** | 0.52*** |
| | (0.017) | (2.3) | (0.07) |
| Over 20 years | 0.43*** | 18.2*** | 0.47** |
| | (0.027) | (6.3) | (0.19) |
| 10–20 minus 1–2 years | | 16.5*** | 0.34*** |
| | | (2.7) | (0.11) |
| Fixed effects | Yes | Yes | Yes |
| Issuer cluster | Yes | Yes | Yes |
| Observations | 355,372 | 355,372 | 355,372 |

*Notes.* Each column is a separate pooled regression containing all five maturity interactions. Column (1) reports the designation increase associated with below-cap status, estimated by OLS. Columns (2)–(3) instrument designation by below-cap status within each bin, as in (27), and report $-\beta_g$, so positive numbers mean lower yields on designated bonds. Column (3) divides each bond's yield by its expected annual benefit $\bar{X}_b = 0.8\,\tau_{B,t}\hat{c}_b$ before estimation (Appendix C.7.3). Contrasts use unrounded estimates. Parentheses give standard errors clustered by issuer (15,889 issuers). $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

*Fixed effects in every column.* Maturity bucket × year, state × year, state × offering year–month, rating category × year, security type × year, use of proceeds, coupon bin × maturity bucket and callability × maturity bucket × year. Other controls are log tranche par, maturity-bin indicators and separate linear issuance slopes on each side of the cap within each bin.

*Sample.* Tax-exempt new issues in 2000–2021 within \$0.5–\$4 million of the statutory cap (\$30 million in 2009–10, \$10 million otherwise). Maturity is measured at issuance. Appendix C.7.2 defines the fixed-effect buckets.

$u \in [0, X]$, assume that a clearing yield exists, that demands are continuous and nondecreasing with finitely many continuously differentiable regimes and that total demand has a strictly positive slope at clearing. Then

$$d_k^H(y_k(u)) + d_k^B(y_k(u) + u) = q_k, \qquad y_k^N = y_k(0), \quad y_k^Q = y_k(X). \tag{146}$$

Differentiating (146) on a smooth segment gives

$$-\frac{dy_k(u)}{du} = \frac{\partial d_k^B/\partial y}{\partial d_k^H/\partial y + \partial d_k^B/\partial y} \equiv w_{Bk}(u) \in [0, 1], \tag{147}$$

with demand slopes at the clearing net returns, so $w_{Bk}(u)$ is the banks' share of the total demand response. A larger benefit raises bank demand, and the yield falls to offset it. Integrating continuously across participation boundaries gives

$$\varpi_k = y_k(0) - y_k(X) = \int_0^X w_{Bk}(u)\,du = \theta_k X, \qquad \theta_k = \frac{1}{X}\int_0^X w_{Bk}(u)\,du. \tag{148}$$

If banks never participate, $\partial d_k^B/\partial y = 0$ and $\varpi_k = 0$, the representative-tax benchmark. If only banks hold the bond, $\partial d_k^H/\partial y = 0$ and $\varpi_k = X$, the bank-only benchmark. When both respond, capitalization lies between them. The capitalization share $\theta_k$, the yield reduction as a fraction of the benefit, is distinct from banks' share of holdings.

The prediction of a flat premium $X$ concerns annual yield-equivalent carry and makes no claim about a fixed price premium. The benefit applies per dollar of tax basis, so mapping it into offering yields uses comparable cash flows and the ratio of basis to price.

### C.7.2 Design

**Running variable.** The running variable $R_{jt}$ of Section 6.2 measures issuance as all tax-exempt par within each six-digit CUSIP issuer prefix and offering calendar year, before sample restrictions. Statutory issuance aggregates subordinate entities and includes bank loans and leases, mostly absent from Mergent. It excludes refundings of designated bonds, which Section 265(b)(3)(D) deems designated, and private-activity bonds other than 501(c)(3) bonds. This measurement noise motivates the fuzzy design.

Among bonds \$0.5–4 million above the cap, 21% are designated. In that window, 4% of single-sale issuer-years designate, and issuer-years with multiple sales designate 44% of their issues. The statute tests reasonably anticipated issuance at each sale, so a March designation stands even when an unanticipated October sale takes the year over the cap. Refundings are 46% of the bonds in that window, and measuring issuance net of refundings reduces designation in the corresponding window to 9%. Designation falls from 24% at \$0.5–1 million above the cap to 15–16% at \$4–8 million and 5% beyond \$10 million, consistent with measurement noise. Designation, refunding status, security type, bond insurance and credit enhancement come from the Mergent bond master files.

**Controls.** In (27), $W_b$ contains the fixed effects listed under Table 39, the log par amount of the maturity tranche (floored at \$1,000) and a linear function of $R_{jt}$ on each side of the cap for each maturity bin. The seven maturity buckets appear only in these fixed effects and differ from the five maturity bins $g$ of the coefficients. Their upper limits are 2, 5, 10, 15, 20 and 30 years, with the seventh bucket beyond thirty.

**The window.** The baseline window keeps issuer-years with $0.5 \leq |R_{jt}| \leq 4$. Its half-million donut excludes managed bunching (2,081 issuer-years within \$0.25 million below the cap versus 572 above). It also excludes 1,034 issuer-years with 12,076 bonds exactly at the \$10 million cap (\$30 million in 2009–10). At-cap issuers are eligible (Section 6.2) and designate 79.6% of their bonds, the below-cap rate, against 27% at \$1–50 thousand above the cap. The below-cap indicator is always $E_{jt} = \mathbf{1}\{R_{jt} \leq 0\}$. A strict inequality would misclassify at-cap issuers, lowering the no-donut first stage from 0.39 to 0.30 but barely changing 2SLS, because dividing by the first stage offsets the attenuation.

**Balance at the cap.** Local linear covariate regressions in the baseline window, with maturity-bucket-by-year and state-by-year effects, find callability and the number of maturities per issue balanced in every bin. Ratings, observed for 62% of bonds, differ only at 1–2 years (0.17 notch, $t = 2.1$). Coupons, set at pricing and covered by coupon-bin-by-maturity-bucket fixed effects, are 0.09–0.14 percentage points lower below the cap in the four bins beyond two years. Log par, always controlled, is lower below the cap at 1–2 years ($-0.26$) and higher at 10–20 years ($+0.12$).

### C.7.3 Funding costs and the benefit measures

**Funding costs.** Section 265(b) allocates a bank's interest expense across its assets in proportion to their average adjusted bases, without matching each bond to borrowing of the same maturity. With bank $n$'s annual interest expense $I_{nt}$ and average asset basis $\bar{A}_{nt}$, a dollar of municipal basis held throughout the year is allocated $c_{nt} = I_{nt}/\bar{A}_{nt}$ of interest expense. Designation reduces the nondeductible fraction from one to 0.2, saving $X_{nt} = 0.8\,\tau_{B,t}c_{nt}$ at the bank's federal tax rate $\tau_{B,t}$ (35% before 2018 and 21% thereafter). Here $I_{nt}$ is the year-to-date interest expense in banks' December reports, already an annual total, and $\bar{A}_{nt}$ is average assets, both from the BankFind database of the Federal Deposit Insurance Corporation (FDIC). This expense-to-assets ratio is neither a deposit rate nor a maturity-matched funding cost.

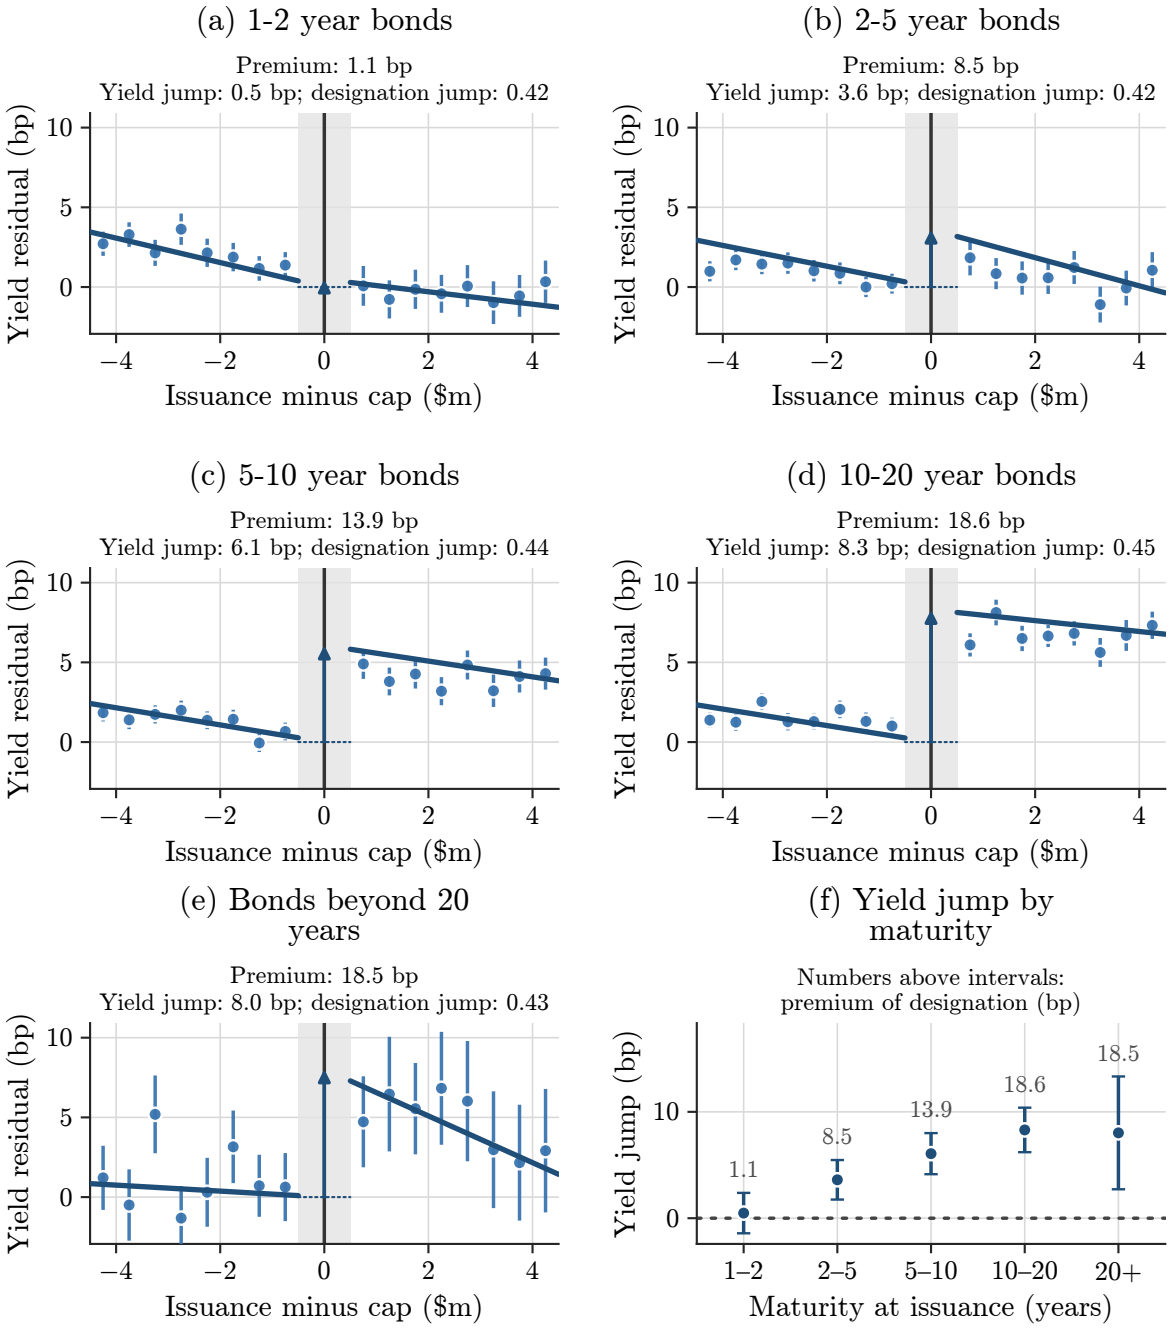



Figure 24: Offering-yield discontinuities by maturity. Panels (a)–(e) show residualized $0.5 million issuance-bin means relative to the below-cap fitted value at the cap, with local linear fits, donut shading and 95 percent bars. Arrows show above-cap-minus-below-cap jumps, and headers divide each yield jump by the designation jump to obtain the premium. Panel (f) shows jumps and premia, the binwise counterparts to Table 39. The common scale covers the central 96 percent of binned means.

Banks are weighted by their municipal-security holdings: state and local securities held to maturity at amortized cost plus those available for sale at fair value. These cover all municipal securities, bank-qualified or not, exclude loans and have no bank-size cutoff.[32] With $M_{nt}$ these holdings and $\mathcal{B}_t$ the retained banks with positive holdings, the holdings-weighted funding cost and the current-year benefit are

$$c_t = \frac{\sum_{n\in\mathcal{B}_t} M_{nt}(I_{nt}/\bar{A}_{nt})}{\sum_{n\in\mathcal{B}_t} M_{nt}}, \qquad X_t = 0.8\,\tau_{B,t}c_t. \tag{149}$$

The retained banks have positive average assets, nonmissing interest expense and an expense-to-assets ratio from zero to below 15%. Mean $c_t$ is 2.34% in 2000–04, 2.31% in 2005–08, 1.09% in 2009–10, 0.51% in 2011–13, 0.31% in 2014–17 and 0.66% in 2018–19. The corresponding current-year benefits $X_t$ are approximately 66, 65, 30, 14, 9 and 11 bp, with the 2009–10 value under the ordinary rule.

**Scope of the measure.** The Call Report measure approximates the tax allocation with book assets and expense. The statute uses adjusted tax bases and otherwise-deductible interest, which are not observed bank by bank, and the IRS permits specified book-asset averaging conventions. $X_t$ measures the ordinary-rule benefit of a C-corporation bank from observable data and does not replicate any bank's tax return.[33] The weights include Subchapter S banks' municipal holdings. In 2009–10, the Section 265(b)(7) exception for holdings up to two percent of assets could reduce the disallowance on non-designated bonds and with it designation's incremental benefit. Modeling 2009–10 by the $30 million cap alone omits this provision, so I report those years separately.

Bank holdings of municipal securities rose from $97 billion in 2000 to $370 billion in 2017 and fell after the 2018 rate cut.

**The expected annual benefit.** A bond's benefit accrues over its life. For a bond $b$ issued in month $t$, the expected annual benefit is $\bar{X}_b = 0.8\,\tau_{B,t}\,\hat{c}_b$, where $\hat{c}_b$ is the expected annual funding cost over the bond's life. Regressing annual holdings-weighted funding costs on one-year expected rates over 2000–2021 ($R^2 = 0.75$) gives $\hat{c}_b = 0.38 + 0.55\,r_b$. Here $r_b$ is the issue-month Kim–Wright (Kim and Wright, 2005) expected average short rate at the bond's rounded whole-year maturity, capped at ten years. Both $\hat{c}_b$ and $r_b$ are in percent per year. The expected rate is the fitted yield less its term premium, which keeps a term premium rising with maturity out of the denominator of the scaled outcome. Bin means of $\bar{X}_b$ are 33.9, 41.5, 50.5, 56.8 and 53.3 bp.

### C.7.4 Scaled and interacted specifications

**Premia and the current-year benefit.** The current-year benefit $X_t$ describes today's carrying cost, whereas $\bar{X}_b$ averages expected benefits over the bond's life. The current-year benefit falls from 104 bp in 2000 to an average of 9 bp in 2013–17, but premia do not track it. Ratios of the 5–20 year premium to $X_t$ exceed two in some low-cost periods and are at most about 0.3 before 2009 (Appendix C.7.6). Such ratios depend on the years pooled. The next specification instead scales each bond's outcome by its expected benefit over its life.

**Outcome scaled by the expected benefit.** For column (3) of Table 39, I divide each offering yield by its own $\bar{X}_b$ before any pooling and then estimate (27) by 2SLS. The premia in expected-benefit units are 0.17, 0.31, 0.41, 0.52 and 0.47, and the 10–20 minus 1–2 year contrast is 0.34 (s.e. 0.11). The Adrian–Crump–Moench (Adrian et al., 2013) decomposition removes larger term premia and gives 0.17, 0.35, 0.46, 0.55 and 0.48 (contrast 0.39).

[32] Interest expense is BankFind series `EINTEXP` (Schedule RI, item 2.e, `RIAD4073`), average assets `ASSET5` and municipal holdings `SCMUNI` (Schedule RC-B, item 3, columns A/D, 8496/8499 with applicable prefixes), as defined in the FDIC's Financial API definitions and the FFIEC December 2017 Call Report, Schedules RI and RC-B. BankFind harmonizes the historical regulatory-report series.

[33] See the IRS, Bank Qualified Bonds: Section 265, pp. 13-11–13-12, for the allocation and 20% disallowance. The annual, per-dollar-of-basis benefit is distinct from the present value of deductions over a bond's remaining life.

Long-minus-short contrasts remain positive within benefit regimes: 0.57 in low-benefit years, 0.37 in middle years and 0.27 in high-benefit years. In the high-benefit years 2000–02 and 2005–08, expected benefits were nearly flat in maturity (62.8–71.2 bp) and yield curves often inverted, so the benefit's maturity structure does not drive the gradient. Excluding the strong bank-demand years 2020–21 gives 0.18, 0.25, 0.35, 0.39 and 0.36 (contrast 0.21).

**Yield-level benefit interactions.** A second specification keeps offering yields in basis points and interacts designation with the expected annual benefit:

$$\begin{aligned} y_b &= \sum_{g=1}^{5} \beta_g B_b \bar{X}_b \mathbf{1}\{n_b \in g\} + \widetilde{W}_b'\delta + u_b, \\ B_b\bar{X}_b\mathbf{1}\{n_b \in g\} &\quad \text{instrumented by } E_{jt}\bar{X}_b\mathbf{1}\{n_b \in g\}, \qquad g = 1, \ldots, 5. \end{aligned} \tag{150}$$

Here $\widetilde{W}_b$ adds $\bar{X}_b$ by bin to the fixed effects and running-variable controls of (27), and $u_b$ is the regression error. Table 40 reports the yield reductions $-\beta_g$ and the joint first-stage strength. Table 39 divides the outcome by $\bar{X}_b$. This specification multiplies designation and its instruments by $\bar{X}_b$ and adds benefit controls, so the estimated levels need not coincide. Both give rising profiles. The statement of Section 6.2 that about half of the benefit is priced at 10–20 years refers to the outcome-scaled specification. Excluding 2009–10 leaves the interacted 10–20 minus 1–2 year contrast at 0.113 (s.e. 0.032).

Table 40: Bank designation interacted with the expected annual benefit.

| Maturity (years) | 1–2 | 2–5 | 5–10 | 10–20 | > 20 |
|---|---|---|---|---|---|
| Yield reduction per benefit unit | 0.058** | 0.096*** | 0.154*** | 0.175*** | 0.233*** |
| | (0.024) | (0.025) | (0.027) | (0.031) | (0.089) |
| 10–20 minus 1–2 | 0.117*** | | | | |
| | (0.032) | | | | |

*Notes.* Entries are $-\beta_g$ in equation (150), the offering-yield reduction in basis points per basis point of expected annual benefit $\bar{X}_b$. Each designation–benefit interaction is instrumented by the below-cap indicator times the benefit in that bin. Controls add the benefit by bin to the fixed effects, log tranche par and running-variable slopes of Table 39, on its 2000–2021 sample within \$0.5–\$4 million of each year's cap. Observations: 355,372, issuer clusters: 15,889, joint Kleibergen–Paap $F$: 122.0. Issuer-clustered standard errors are in parentheses. $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

### C.7.5 Robustness

**Donut, bandwidth and window.** Donut half-widths from zero to \$1.5 million and bandwidths from \$2 million to \$8 million (Table 41) give 10–20 minus 1–2 year contrasts of 13.5–18.4 bp. The 1–2 year premia are never significant. In the \$5–\$15 million window of Dagostino (2025), without a donut, the contrast is 15.4 (1.8) bp, one of the most precise estimates.

The last row of Table 41 drops the region that her bunching estimation treats as contaminated, \$8.5–\$13.5 million at the \$10 million cap. Within $|R| \leq 8$, it compares issuers at least \$1.5 million below the cap with issuers at least \$3.5 million above it. First stages rise to 0.47–0.54. The longer extrapolation gives premia of 7.6 bp at 5–10 years and 10.0 bp at 10–20 years, with a contrast of 9.8 (4.2).

**Construction and the relation to Dagostino (2025).** Dagostino (2025) measures issuer-year issuance with Ipreo MuniIC data and documents the same bunching and first stage at the \$30 million cap in 2009–10. Her regressions on \$1 million issuance bins control for size, maturity, insurance, rating and use of proceeds, with year and issuer effects and standard errors clustered by year. In that study, tax-adjusted spreads over treasuries jump about 10 bp at the cap for general-obligation (GO) bonds issued before 2008. In her Call Report data for 2000–2010, designated bonds are about 4% of the average bank's assets, rising to 5% by 2010, about nine times the non-designated share. Table 42 compares her construction choices with

Table 41: Bank-qualified discontinuity: sensitivity to the donut, the bandwidth and the window.

| | first stage | 1–2y | 2–5y | 5–10y | 10–20y | 20y+ | 10–20 minus 1–2 |
|---|---|---|---|---|---|---|---|
| *Donut half-width, bandwidth $4 million* | | | | | | | |
| none | 0.39*** | 1.9 (1.8) | 7.5*** (1.8) | 13.9*** (1.8) | 17.7*** (1.8) | 17.7*** (4.5) | 15.7*** (2.0) |
| $0.25 million | 0.40*** | 1.6 (2.0) | 7.6*** (2.0) | 13.4*** (2.0) | 16.7*** (2.2) | 16.3*** (5.5) | 15.1*** (2.4) |
| $0.5 million (baseline) | 0.42*** | 1.8 (2.2) | 8.7*** (2.2) | 13.8*** (2.2) | 18.4*** (2.3) | 18.2*** (6.3) | 16.5*** (2.7) |
| $1 million | 0.44*** | 1.2 (2.9) | 8.2*** (2.7) | 13.3*** (2.8) | 19.7*** (3.1) | 20.6*** (8.0) | 18.4*** (3.5) |
| $1.5 million | 0.44*** | 2.0 (4.4) | 7.3* (3.8) | 11.8*** (4.0) | 15.5*** (4.2) | 19.5* (10.2) | 13.5*** (5.2) |
| *Bandwidth, donut $0.5 million* | | | | | | | |
| $2 million | 0.38*** | 0.8 (4.7) | 7.9 (4.9) | 13.3*** (4.8) | 16.9*** (5.0) | 21.0 (14.1) | 16.1*** (4.8) |
| $3 million | 0.40*** | 2.7 (3.0) | 8.9*** (3.0) | 15.2*** (2.9) | 19.2*** (3.1) | 19.9*** (7.6) | 16.5*** (3.3) |
| $5 million | 0.42*** | −0.9 (1.9) | 5.7*** (1.8) | 11.4*** (1.8) | 14.9*** (1.9) | 15.9*** (4.9) | 15.8*** (2.3) |
| $6 million | 0.43*** | −0.5 (1.7) | 5.4*** (1.5) | 10.8*** (1.6) | 13.2*** (1.6) | 11.4*** (4.1) | 13.6*** (2.1) |
| $8 million | 0.45*** | 0.0 (1.5) | 6.3*** (1.2) | 10.1*** (1.3) | 13.5*** (1.3) | 12.3*** (3.5) | 13.5*** (1.8) |
| *Other windows* | | | | | | | |
| $5–15 million, no donut | 0.40*** | 0.1 (1.6) | 5.6*** (1.5) | 12.1*** (1.6) | 15.5*** (1.5) | 16.7*** (3.8) | 15.4*** (1.8) |
| $8.5–13.5 million dropped, $\lvert R\rvert \le 8$ | 0.47*** | 0.2 (3.5) | 6.0** (3.0) | 7.6** (3.1) | 10.0*** (3.4) | −1.8 (9.5) | 9.8** (4.2) |

*Notes.* Cells report the 2SLS premium of designation in basis points by maturity bin, pooled over 2000–2021 at each year's cap, with the baseline fixed effects and issuer-clustered standard errors in parentheses. The bandwidth and the donut half-width are the largest and smallest $|R_{jt}|$ retained, and the first-stage column refers to the 1–2 year bin. The $5–15 million window is the one used in the spread comparison of Dagostino (2025). The last row drops the region her bunching estimation excludes. $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

mine and nets refunding par from the running variable. Every row retains a larger long-maturity premium: issuer effects steepen it, while netting refundings shifts issuer-years and widens confidence intervals.

Table 42: Bank-qualified discontinuity: construction checks.

| | 1–2y | 2–5y | 5–10y | 10–20y | 20y+ | 10–20 minus 1–2 |
|---|---|---|---|---|---|---|
| Baseline (Table 39) | 1.8 (2.2) | 8.7*** (2.2) | 13.8*** (2.2) | 18.4*** (2.3) | 18.2*** (6.3) | 16.5*** (2.7) |
| Bond insurance and enhancement × year | 0.9 (2.2) | 8.0*** (2.2) | 13.0*** (2.2) | 17.7*** (2.3) | 17.3*** (6.2) | 16.8*** (2.6) |
| General-obligation bonds only | 1.7 (2.3) | 7.6*** (2.2) | 13.3*** (2.3) | 15.2*** (2.5) | 15.8** (7.3) | 13.5*** (2.9) |
| Issuer fixed effects | −5.0* (2.6) | 1.8 (2.5) | 7.5*** (2.6) | 17.3*** (2.8) | 16.8*** (4.1) | 22.3*** (2.4) |
| Tax-adjusted spread over treasuries | 5.3 (3.3) | 12.6*** (3.2) | 21.4*** (3.2) | 26.9*** (3.5) | 22.8*** (8.6) | 21.6*** (3.8) |
| Insurance controls, year clusters (22) | 0.9 (2.9) | 8.0*** (2.6) | 13.0*** (3.6) | 17.7*** (3.4) | 17.3** (7.6) | 16.8*** (3.7) |
| Running variable net of refundings | 0.8 (3.1) | 4.9 (3.1) | 8.1*** (3.1) | 15.7*** (3.6) | 11.0 (8.2) | 14.8*** (3.7) |
| and refunding bonds excluded | 0.9 (3.1) | 6.3* (3.2) | 8.9*** (3.2) | 16.0*** (3.6) | 10.9 (7.5) | 15.1*** (3.6) |

*Notes.* Cells report the 2SLS premium of designation in basis points, pooled over 2000–2021 at the year's cap in the window $0.5 to $4 million, with the baseline fixed effects and issuer-clustered standard errors unless stated. First stages are 0.40 to 0.51 in every bin of every row. Insurance and enhancement: adds bond insurance × year and credit enhancement × year (47% of window bonds are insured). The general-obligation, issuer-fixed-effect, tax-adjusted-spread and year-clustered rows also include these controls. General obligation: unlimited-tax, limited and double-barreled general-obligation bonds, 70% of the window. Issuer fixed effects: identified from issuers that cross the cap across years. Tax-adjusted spread: $100\,[y/(1-\tau_{B,t}) - y^{\mathrm{Tsy}}]$, the outcome of Dagostino (2025). Net of refundings: the running variable excludes refunding par, and the last row also drops refunding bonds from the estimation sample. $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

Applying her $1 million issuance-bin estimator to her window, with maturity-bucket, state and rating effects by year and use-of-proceeds controls, gives flat below-cap bins and a step above the cap, matching her Figure 3. Pooled across maturities, the step is 4.2 bp in the offering yield and 5.6 bp in the tax-adjusted spread, 1.1 and 1.4 bp at 1–5 years and 7.8 and 10.2 bp at 10–20 years. In her pooled-spread setting for GO bonds issued before 2008, with her controls but with state fixed effects and clustering by issuer, the reduced-form jump at the cap is 10.4 bp (s.e. 2.5), against her approximately 10. With her issuer fixed effects and clustering by year, which leaves eight clusters, the jump is 6.3 (3.8), because within-issuer variation across the cap is thin before 2008.

### C.7.6 Placebos and periods

**Placebo cutoffs and the ARRA cap.** At a placebo cutoff of $5 million, first stages are 0.00–0.05 and all yield reductions are within two bp of zero, with a 10–20 minus 1–2 year contrast of 0.4 (1.2). At $15 million, first stages are −0.01 to 0.02. Yield reductions are −0.1, 0.6, 1.7, 2.4 and 8.6 bp and are not jointly significant ($p = 0.17$). The 10–20 minus 1–2 year contrast is 2.5 (1.6), against 7.8 (1.2) in the reduced form at the cap. Around the $30 million cap of the American Recovery and Reinvestment Act (ARRA), a $0.5–$8 million window covers 16 thousand bonds issued in 2009–10 by issuers with $22–$38 million of issuance. With the baseline fixed effects, the first stages are 0.19, 0.20, 0.19, 0.16 and 0.26, jointly significant at five percent. At that cutoff, first stages are 0.01–0.04 in 2004–08 and −0.05 to 0.01 in 2011–15. At $10 million in 2009–10, they are 0.01–0.06 and not jointly significant. The 2009–10 reduced-form yield reductions at the $30 million cap are 0.5, 3.2, 1.6, 10.9 and 21.5 bp. The average reduction beyond ten years exceeds the 1–2 year reduction by 15.7 (6.7), significant at five percent.

**By period.** Period estimates for 2000–04, 2005–08, 2009–10, 2011–13, 2014–17 and 2018–19 use the baseline fixed effects and issuer clustering. They use the $30 million cap and a $0.5–8 million window in the ARRA period and the $10 million cap and $0.5–4 million window otherwise. Before 2009, funding costs $c_t$ are 2.3 to 2.4% and $X_t$ is 65–66 bp. Premia are at most about 0.3 of $X_t$ at every maturity, so the deduction is weakly capitalized when the carrying penalty on non-designated bonds approaches a percentage point.

The 2009–10 $30 million cap has smaller designation jumps, 0.16–0.26, so for this period the reduced form is the informative estimate. Its contrast between bonds beyond ten years and 1–2 year bonds, reported above, is significant at five percent.

In 2011–17, funding costs are 0.31–0.51% and $X_t$ is 9–14 bp. Premia are 2.4–2.5 bp at 1–2 years and 14–28 bp at 5–20 years, at or above $X_t$.

After the 2018 cut in the bank tax rate, premia are 6.9 (4.3), 6.4 (5.0), 7.7 (5.7), 14.9 (7.7) and 3.7 (24.2) bp. A difference-in-discontinuities estimate pools 2014–2019 and interacts designation with a post-2018 indicator. It gives post-2018 changes of 3.1 (6.5), −9.1 (6.7), −19.2 (7.6), 3.6 (10.0) and −1.3 (28.7) bp, with unchanged first stages. Subchapter S banks hold 6–14% of banks' municipal holdings over 2000–2021. Their income is taxed to their shareholders, and the TCJA affected them less.

Annual reduced-form yield reductions at the $10 million cap, pooled across maturities with the baseline fixed effects, are largest in 2016 and 2019–21, at 11–18 bp. They are 6–8 bp in 2011–13 and within about 3 bp of zero in 2000–01, 2004–07 and 2018. First stages are 0.30–0.62 in every year except 2009–10, when the cap was $30 million. Premia do not simply track the current-year benefit, and the pattern is consistent with changing bank participation.

## C.8 Investor holdings by tax status

Portfolio shares vary with exemption benefits (Section 3.4 and Table 2). In 2000–2017, exempt municipal bonds account for 34–48% of the average P&C insurer's bond portfolio, compared with 6–9% for life insurers (Figure 25). In the public-pension sample, mean portfolio shares never exceed 0.13% for exempt and 0.78% for taxable municipals (Figure 26). After the corporate-rate cut, the P&C exempt share fell from 38.0% in 2017 to 28.6% in 2019, while the share of other taxable bonds rose from 50.2% to 57.5%. This reallocation is the quantity counterpart of the repricing in the insurer test of P2 (Appendix C.9).

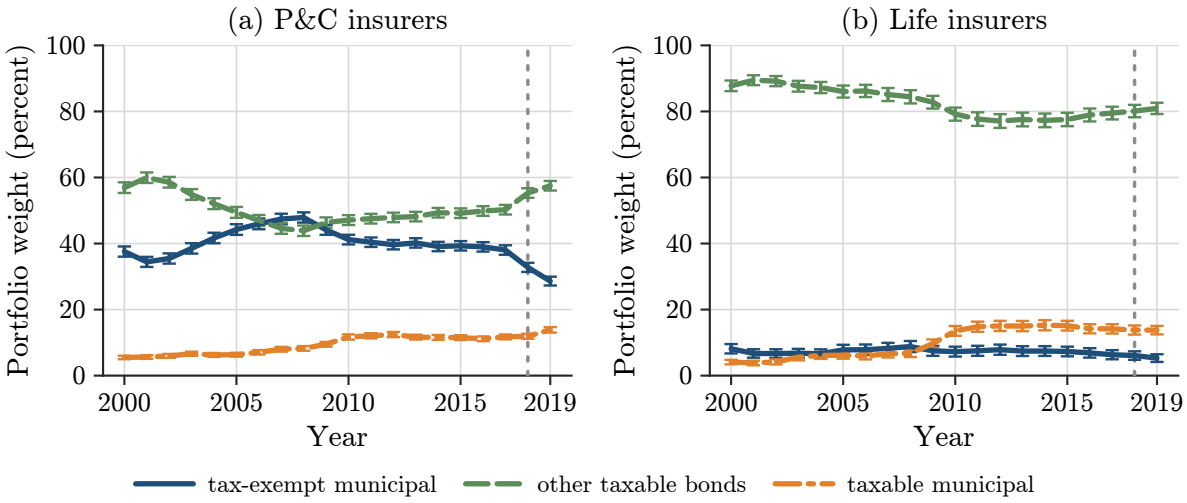



Figure 25: Tax-exempt municipals, taxable municipals and other taxable bonds as shares of insurers' total bond holdings: P&C in panel (a) and life in panel (b), 2000–2019, NAIC Schedule D. The dotted vertical line marks the TCJA corporate-rate cut. Bars are 95 percent intervals for yearly mean insurer portfolio shares, with standard errors clustered by insurer.

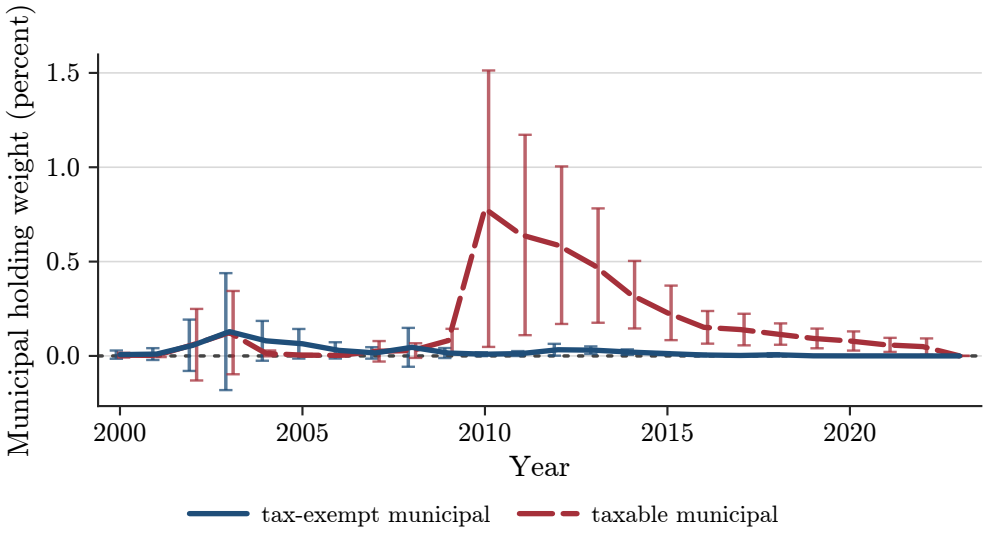



Figure 26: Municipal holdings by federal tax status as shares of total bond portfolios in the public-pension sample obtained by Freedom of Information Act (FOIA) requests (Section 3.1), 2000–2023. All thirty plans are covered from 2015. Bars are 95 percent intervals for yearly mean plan portfolio shares, with standard errors clustered by plan. Other taxable bonds are omitted. Taxable municipals include Build America Bonds.

## C.9 An additional institutional tax test: the 2017 corporate-rate cut

This additional test of P2 uses the Tax Cuts and Jobs Act (TCJA), which reduced P&C insurers' exemption benefit. It asks whether yield-implied taxes consequently fell more at long maturities in states with greater pre-reform insurer ownership. The difference-in-differences (DiD) design compares these states' maturity profiles before and after the reform. The event study examines timing, and Appendices C.9.1 and C.9.2 address inference and the concurrent household tax change.

**Prediction.** Appendix Proposition A.9 gives zero pass-through of the lower-tax investor's exemption benefit at the riskless short end, where the household's after-tax cash return anchors the yield. Pass-through is positive at the long end. Through rebalancing, a change in that benefit also moves yields at risky maturities shorter than those the lower-tax investor holds. P&C insurers are lower-tax institutions in this sense. Since the TCJA lowered their exemption benefit, P2 predicts that after the reform the implied tax fell more at longer maturities in states whose municipal debt P&C insurers held more heavily. The same reform also changed household taxes, which the model predicts move short-maturity implied taxes. Quarter effects and the event-time interactions with maturity and with the state tax rate absorb these changes where they are common across states or vary with the state tax rate. The test therefore concerns the exposure gradient $\beta_1$, defined below.

**The 2017 tax change.** The TCJA cut the corporate rate from 35% to 21% effective 2018Q1, and the Section 832 proration rule amplified the resulting decline in P&C insurers' exemption benefit. The disallowed fraction of their tax-exempt interest rose from 15 to 25 percent, so the effective tax on municipal interest stayed at 5.25% ($35\% \times 15\%$ before, $21\% \times 25\%$ after). The after-tax share of taxable interest rose from 0.65 to 0.79. The P&C relative after-tax share $\xi_P$ rose from 0.686 to 0.834, and the exemption benefit $1 - \xi_P$ narrowed from 0.31 to 0.17.

The top household ordinary rate also fell, from 39.6% to 37%, and the cap on SALT deductions changed the value of in-state exemptions. Life insurers' relative after-tax share rose from $0.65/0.825 \simeq 0.788$ to $0.79/0.937 \simeq 0.843$ (Appendix D.1), and banks' taxable after-tax share and the tax cost of their financing changed. The design compares states with different P&C exposure, and quarter effects absorb yield changes common to all states.

**Treatment and sample.** The treatment is a state's exposure to P&C insurers, $PCShare_s$: the market value of their holdings of state $s$'s municipal bonds relative to the state's par outstanding. It is measured at the end of 2015, before the 2016Q3 start of the estimation window and two years before the reform. Across the 50 states and the District of Columbia in the sample, it averages 0.103 with a standard deviation of 0.030. Each bond $b$ takes the exposure of its issuing state, $PCShare_{s(b)}$. The sample is the 1,733,881 customer trades reported to the MSRB over 2016Q3–2019Q1 in non-callable municipal bonds rated AA and above and not subject to the AMT. Restricting to non-callable bonds rules out direct repricing of embedded options. The outcome is the implied tax of (2) for each trade, $ImpliedTax_{b,t} = 1 - y^M_{b,t}/y^C_t(n_b)$. It uses the trade yield and the corporate zero yield for the last month of the trade's calendar quarter, at a maturity equal to the bond's maturity year minus the trade year.

**Specification.** I estimate

$$\begin{aligned} ImpliedTax_{b,t} = \beta_0\big(PostTCJA_t \times PCShare_{s(b)}\big) + \beta_1\big(PostTCJA_t \times PCShare_{s(b)} \times n_b\big) \\ + \alpha_{\lfloor n_b \rfloor \times Purchase} + \alpha_{issuer \times Purchase} + \alpha_t + \Gamma' Z_{b,t} + \varepsilon_{b,t}, \end{aligned} \tag{151}$$

where $PostTCJA_t = \mathbf{1}\{t \geq 2017Q4\}$ starts the post period at the enactment quarter and $n_b$ is the bond's remaining maturity in years. The fixed effects compare trades of the same issuer and of the same whole-year maturity, purchases with purchases and sales with sales, and $\alpha_t$ absorbs each quarter's level. $Purchase_{b,t}$ is the MSRB dealer-side indicator for a customer purchase. The controls $Z_{b,t}$ include $PCShare_{s(b)} \times n_b$ and event time interacted with maturity and with the state tax rate. They also include the issuing state's resident tax rate, the pre-period life-insurer holding share and the bond rating, each alone and interacted with maturity.

The event-time-by-maturity interactions let the common maturity gradient change each quarter, so $\beta_1$ is identified from cross-state differences in *PCShare* within a quarter.

The post-reform change in the effect of exposure at maturity $n_b$ is $\beta_0 + \beta_1 n_b$. The model predicts $\beta_1 < 0$, and the representative-tax benchmark predicts $\beta_0 = \beta_1 = 0$.

**Instruments.** Pre-reform P&C ownership may be correlated with state conditions that also move the maturity profile of implied taxes. The IV specifications therefore instrument $PCShare_s$ with the holdings of state-domiciled P&C insurers across all issuers, relative to the state's municipal debt:

$$IV1_s = \log\Big(\frac{\sum_{i \in D_s} \text{Municipal Assets}_i}{\text{Municipal Outstanding}_s}\Big), \qquad IV2_s = \log\Big(\frac{\sum_{i \in D_s^{\text{act}}} \text{Bond Assets}_i}{\text{Municipal Outstanding}_s}\Big), \tag{152}$$

where $D_s$ is the set of P&C insurers domiciled in state $s$ and $D_s^{\text{act}}$ the subset holding any municipal bonds. From insurer $i$'s Schedule D filings with the NAIC, $\text{Municipal Assets}_i$ is its par holding of municipal bonds and $\text{Bond Assets}_i$ its par holding of all bonds. $\text{Municipal Outstanding}_s$ is the state's municipal par outstanding. $IV2_s$ uses total bond holdings, so it does not depend on which municipal bonds the domiciled insurers chose to hold. The interactions of each instrument with $n_b$, with $PostTCJA_t$ and with their product instrument the three corresponding exposure terms. Relevance comes from insurer home bias: locally domiciled insurers disproportionately hold the state's bonds. The Kleibergen–Paap first-stage $F$ statistics are 25.7 for $IV1_s$ and 17.5 for $IV2_s$.

The exclusion restriction is that, conditional on the fixed effects and controls, the domiciled insurers' total holdings affect the post-TCJA change in the maturity gradient of implied taxes only through P&C insurers' exposure to the state's bonds. Insurer domicile largely predates the reform. Quarter effects and event-time-by-maturity controls absorb the common component of the federal tax changes. Event-time-by-state-tax interactions absorb level changes associated with state tax rates, including SALT-induced ones. State-tax-by-maturity controls capture a persistent maturity gradient. Neither absorbs a post-reform change in the gradient that varies with state tax rates. Appendix C.9.2 states the covariance conditions under which omitting such a SALT gradient biases the OLS estimate toward zero.

**Results.** Table 43 reports the static DiD estimates. Under OLS, the maturity gradient $\beta_1$ is $-0.017$ with an issuer-clustered standard error of 0.005. Under the two IV specifications it is $-0.038$ (0.010) and $-0.040$ (0.015). Consider two states whose pre-reform P&C shares differ by one standard deviation, 3.0 percentage points. After the reform, the implied tax in the more exposed state falls by an additional 0.05 percentage point per year of maturity under OLS and 0.11 under the first IV specification. Between a five- and a twenty-five-year bond, the first IV estimate cumulates to about 2.3 percentage points of implied tax. At a corporate yield of four percent, that is about 9 basis points of municipal yield relative to the taxable yield.

Table 43: TCJA difference-in-differences: implied-tax response to insurer exposure. The sample is MSRB trades in non-callable, AA-and-above, non-AMT municipal bonds in 2016Q3–2019Q1. The outcome is the implied tax of equation (2), and *PCShare* is the end-2015 P&C ownership share. $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

| | (1) OLS | (2) IV1 | (3) IV2 |
|---|---|---|---|
| $PostTCJA \times PCShare$ | $-0.122^{***}$ | $-0.129$ | $-0.017$ |
| | (0.029) | (0.084) | (0.100) |
| $PostTCJA \times PCShare \times n_b$ | $-0.017^{***}$ | $-0.038^{***}$ | $-0.040^{***}$ |
| | (0.005) | (0.010) | (0.015) |
| $\lfloor n_b \rfloor \times Purchase$ FE | Yes | Yes | Yes |
| $Issuer \times Purchase$ FE | Yes | Yes | Yes |
| Quarter FE | Yes | Yes | Yes |
| Issuer cluster | Yes | Yes | Yes |
| Observations | 1,733,881 | 1,733,881 | 1,733,881 |
| First-stage $F$ (Kleibergen–Paap) | | 25.7 | 17.5 |

*Notes.* FE denotes fixed effects. Columns (2) and (3) instrument *PCShare* with $IV1_s$ and $IV2_s$ of equation (152). All columns include the controls of equation (151), among them the rating and its interaction with $n_b$. Standard errors are clustered by issuer (six-digit CUSIP, 9,903 issuers) and reported in parentheses.

The maturity gradient persists on the wider A-or-above sample (Appendix C.11.2 and Table 49).

**Event study.** The event study replaces the post-reform indicator in (151) with quarter indicators, omitting 2017Q3, so $\hat{\beta}_{1,q}$ measures the maturity gradient in quarter $q$ relative to 2017Q3. The IV versions use the corresponding quarter interactions of the instruments, with the fixed effects and clustering of Table 43. These estimates locate the change at the reform. Across the three estimators, eleven of the twelve pre-TCJA coefficients are individually insignificant at five percent. In every column, the coefficients turn significantly negative at 2018Q1 and stay negative. Five of the six post-reform quarters exclude zero under IV1 and four under IV2 (Figure 27; Table 44 reports each quarter).

### C.9.1 State-clustered inference

As a check for shocks shared within states, I also cluster the insurer regressions above by issuing state (51 clusters: the 50 states and the District of Columbia). Point estimates are unchanged. State-clustered standard errors are 0.0129, 0.0245 and 0.0342 for the OLS, IV1 and IV2 maturity coefficients, and their 95% intervals include zero.

### C.9.2 SALT-induced household demand and the insurer-channel regression

A binding SALT cap can raise high-tax households' exemption benefit and lower municipal yields where these households are active. The conditions below determine when this household channel attenuates the OLS insurer coefficient.

Write $Y_{b,t} = 1 - y^M_{b,t}/y^C_t(n_b)$ for the implied tax of equation (2), $X_{b,t} = PostTCJA_t PCShare_{s(b)} n_b$ for the insurer interaction and $W_{b,t} = PostTCJA_t Tax_{s(b)} n_b$ for a household-SALT interaction. The state tax rate $Tax_s$ proxies SALT exposure. Here $X_{b,t}$ is unrelated to the bank benefit $X$ of Appendix C.7. Consider

$$Y_{b,t} = \beta_1 X_{b,t} + \beta_2 PostTCJA_t Tax_{s(b)} + \beta_3 W_{b,t} + \text{controls} + \text{FE} + \varepsilon_{b,t}. \tag{153}$$

The insurer mechanism implies $\beta_1 < 0$. Stronger household demand at the short end implies $\beta_2 > 0$ and $\beta_3 < 0$.

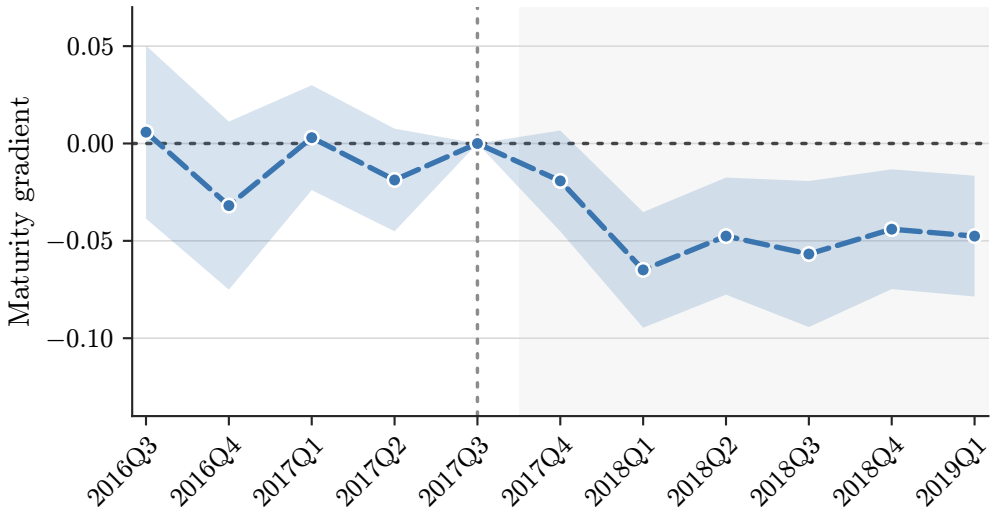



Figure 27: IV1 event-study estimates of the quarter-specific maturity gradient, the coefficient on $\mathbf{1}\{t = q\} \times PCShare \times n_b$, relative to 2017Q3 (vertical dotted line). Shading marks the post-reform quarters from 2017Q4. Bands are 95% intervals clustered by six-digit CUSIP issuer. Appendix C.9.1 reports further inference checks.

Suppose the insurer regression controls for the SALT level term and omits $W$, and the remaining disturbance is orthogonal to residualized $X$. Tildes denote residuals after projecting on all included controls and fixed effects. The omitted-variable formula is

$$\operatorname{plim} \widetilde{\beta}_1 = \beta_1 + \beta_3 \frac{\operatorname{Cov}(\widetilde{X}, \widetilde{W})}{\operatorname{Var}(\widetilde{X})}. \tag{154}$$

Consequently,

$$\beta_3 < 0, \quad \operatorname{Cov}(\widetilde{X}, \widetilde{W}) < 0 \quad \Longrightarrow \quad \operatorname{plim} \widetilde{\beta}_1 > \beta_1. \tag{155}$$

Equation (154) uses residualized variables, so it accounts for controls, fixed effects and weights. An upward bias attenuates a negative coefficient as long as the estimate keeps its sign.

IV identifies the insurer coefficient when the instruments are relevant and orthogonal to the remaining household-demand term, under the exclusion restriction stated in Appendix C.9.

Table 44: TCJA dynamic event study: full triple-interaction coefficients $\hat{\beta}_{1,q}$ on $\mathbf{1}\{t = q\} \times PCShare \times n_b$, by quarter. Standard errors in parentheses are clustered by issuer (six-digit CUSIP) as in Table 43, and 2017Q3 is the omitted base quarter. FE denotes fixed effects. $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

| Quarter $q$ | (1) OLS | (2) IV1 | (3) IV2 |
|---|---|---|---|
| 2016Q3 | $0.019^{**}$ | 0.006 | $-0.020$ |
| | (0.010) | (0.023) | (0.036) |
| 2016Q4 | $-0.014$ | $-0.032$ | $-0.029$ |
| | (0.009) | (0.022) | (0.036) |
| 2017Q1 | 0.001 | 0.003 | 0.009 |
| | (0.007) | (0.014) | (0.020) |
| 2017Q2 | $-0.005$ | $-0.019$ | $-0.021$ |
| | (0.006) | (0.013) | (0.020) |
| 2017Q4 | $-0.001$ | $-0.019$ | $-0.021$ |
| | (0.006) | (0.013) | (0.020) |
| 2018Q1 | $-0.028^{***}$ | $-0.065^{***}$ | $-0.063^{***}$ |
| | (0.007) | (0.015) | (0.023) |
| 2018Q2 | $-0.027^{***}$ | $-0.048^{***}$ | $-0.039^{*}$ |
| | (0.007) | (0.015) | (0.023) |
| 2018Q3 | $-0.031^{***}$ | $-0.057^{***}$ | $-0.063^{**}$ |
| | (0.009) | (0.019) | (0.027) |
| 2018Q4 | $-0.011$ | $-0.044^{***}$ | $-0.051^{**}$ |
| | (0.008) | (0.016) | (0.023) |
| 2019Q1 | $-0.014^{*}$ | $-0.048^{***}$ | $-0.065^{***}$ |
| | (0.007) | (0.016) | (0.024) |
| $\lfloor n_b \rfloor \times$ *Purchase* FE | Yes | Yes | Yes |
| *Issuer* $\times$ *Purchase* FE | Yes | Yes | Yes |
| Quarter FE | Yes | Yes | Yes |
| Issuer cluster | Yes | Yes | Yes |
| Observations | 1,733,881 | 1,733,881 | 1,733,881 |

## C.10 P3: municipal-yield response to treasury yields

P3 tests whether municipal yields respond more to a given movement in maturity-matched treasury yields as maturity rises. The comparison concerns this relative response, so longer bonds' greater duration alone does not explain the prediction. This subsection constructs the trade-pair changes used in Section 6.3, then checks aging, trade windows, clustering and the reference market.

**Transaction sample.** Each MSRB customer trade records the CUSIP, trade date, par, dollar price, implied yield to maturity and dealer side. Estimation trades, from January 2005 to August 2020, are in non-callable, federally exempt bonds rated AA and above (agency score $\leq 3$, Section 6). Noncallability removes call-option repricing. I exclude trades less than 92 days after issue, when offering yields still anchor prices, and trades less than 365 days before maturity, where the fitted treasury curve is erratic.

**Matched treasury changes.** At each trade, I interpolate the daily treasury par-yield curve of Gürkaynak, Sack and Wright (GSW) linearly between integer maturities at the bond's remaining maturity $n_b(t) = d_{b,t}/365$, where $d_{b,t}$ is its days to maturity, holding the curve flat beyond thirty years. Within each CUSIP

Table 45: P3: observed coupon-yield response to treasury par yield by remaining maturity.

| | (1) Continuous $\Delta y^M_{b,t}$ | (2) Bin $\Delta y^M_{b,t}$ |
|---|---|---|
| $\Delta y^{\text{Tsy}}_{n_b,t}$ | $0.346^{***}$ | $0.351^{***}$ |
| | (0.016) [0.012] | (0.013) [0.006] |
| $\Delta y^{\text{Tsy}}_{n_b,t} \times n_b$ | $0.007^{***}$ | |
| | (0.002) [0.002] | |
| $\Delta y^{\text{Tsy}}_{n_b,t} \times Long_b$ | | $0.075^{***}$ |
| | | (0.012) [0.009] |
| $Long_b$ | | $0.016^{***}$ |
| | | (0.002) [0.002] |
| $n_b$ | $0.002^{***}$ | |
| | (0.000) [0.000] | |
| Trade-date cluster | Yes | Yes |
| Observations | 952,874 | 952,874 |
| Weights | Inverse pair count | Inverse pair count |

*Notes.* The sample is pairs of consecutive same-side MSRB customer trades within 90 days in federally tax-exempt, non-callable bonds rated AA and above (agency score $\leq 3$), January 2005 to August 2020, over which the maturity-matched treasury yield moved at least 10 basis points in either direction. The treasury yield is interpolated on the daily GSW curve at the bond's exact remaining maturity on each transaction date. Observations are weighted by the inverse count of their CUSIP's base-sample trade pairs, computed before observations with missing or trimmed outcomes are dropped. The dependent variable is trimmed at its first and ninety-ninth percentiles. $Long_b = \mathbf{1}\{n_b \geq 5\}$. First differencing absorbs time-invariant bond characteristics. Parentheses give trade-date-clustered standard errors (3,864 clusters). Brackets give state-clustered standard errors (54 clusters: the 50 states, the District of Columbia, Puerto Rico, Guam and the U.S. Virgin Islands). Stars are identical under either clustering. $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

and dealer side, consecutive trades form the municipal yield change $\Delta y^M_{b,t}$ and the matched treasury change $\Delta y^{\text{Tsy}}_{n_b,t} = y^{\text{Tsy}}_{n_b(t),t} - y^{\text{Tsy}}_{n_b(t-1),t-1}$, where $t-1$ denotes the trade preceding $t$. Comparing purchases with purchases and sales with sales limits bid–ask and intraday noise. I keep pairs with gaps of 1–90 days and treasury moves of at least 10 bp, with the weights and trimming of Table 45. Over $\Delta t$ elapsed days, aging contributes $-(\Delta t/365)\partial_m y^{\text{Tsy}}_{m,t} + O((\Delta t/365)^2)$ to the matched change, an effect measured below.

**Specifications.** Column (1) of Table 45 estimates the continuous interaction in (28), including maturity but no fixed effects. Column (2) replaces maturity by $Long_b = \mathbf{1}\{n_b \geq 5\}$, so its short and long responses are $\beta_0$ and $\beta_0 + \beta_1$.

**Holding the benchmark maturity fixed.** Part of the matched treasury change comes from the bond becoming shorter between trades. This aging component is $y^{\text{Tsy}}_{n_b(t),t} - y^{\text{Tsy}}_{n_b(t-1),t}$, holding the later date's curve fixed. Fixing the benchmark at the first trade's maturity, $\Delta y^{\text{Tsy,fixed}}_{b,t} = y^{\text{Tsy}}_{n_b(t-1),t} - y^{\text{Tsy}}_{n_b(t-1),t-1}$, removes it. On upward-sloping curves, aging lowers the reference yield by an average of 2.6 bp below five years and 1.6 bp beyond, against absolute changes near 26 bp. Removing this component lowers the short response more than the long response, increasing the gradient (Table 46). Municipal aging does not offset treasury aging because the two curves have different slopes. The third row adds elapsed time and its interaction with maturity to control for linear municipal aging. The gradient rises further. The headline specification keeps the benchmark matched at each trade, the more conservative choice.

Table 46: P3 observed coupon-yield responses and treasury-curve aging.

| | Bin specification | | Interaction | |
|---|---|---|---|---|
| Treasury benchmark | $< 5$y | $\geq 5$y | Bin | Continuous |
| Benchmark matched at each trade's maturity (Table 45) | $0.351^{***}$ | $0.426^{***}$ | $0.0749^{***}$ | $0.0070^{***}$ |
| Benchmark held at the first trade's maturity | $0.340^{***}$ | $0.423^{***}$ | $0.0825^{***}$ | $0.0084^{***}$ |
| plus elapsed time and elapsed time $\times$ maturity | $0.334^{***}$ | $0.419^{***}$ | $0.0850^{***}$ | $0.0086^{***}$ |

*Notes.* Each row estimates (28) on the 952,874 pairs of Table 45, with its weights and trade-date clustering. The first two columns are the bin specification's responses below and from five years, $\beta_0$ and $\beta_0 + \beta_1$, and the last two are the interactions with $Long_b$ and $n_b$. $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

**Sensitivity to the filter thresholds.** Across the treasury-move cutoffs and trade windows of Table 47, all observed maturity-interaction estimates are positive. Every binned cell is significant at one percent, and every continuous cell is significant at five percent, with twenty-seven of the thirty significant at one percent.

Table 47: P3 sensitivity: observed maturity interaction across the treasury-move cutoff and trade-gap window (weighted by inverse CUSIP base-pair counts).

| *Panel A. Continuous interaction* $\Delta y^{\text{Tsy}}_{n_b,t} \times n_b$ | | | | | |
|---|---|---|---|---|---|
| Cutoff \ window | 15d | 30d | 60d | 90d | 120d |
| $\geq 0.05$ | $0.0053^{***}$ | $0.0057^{***}$ | $0.0051^{**}$ | $0.0071^{***}$ | $0.0090^{***}$ |
| $\geq 0.075$ | $0.0061^{***}$ | $0.0057^{***}$ | $0.0047^{**}$ | $0.0068^{***}$ | $0.0087^{***}$ |
| $\geq 0.10$ | $0.0065^{***}$ | $0.0060^{***}$ | $0.0049^{**}$ | $0.0070^{***}$ | $0.0087^{***}$ |
| $\geq 0.15$ | $0.0080^{***}$ | $0.0071^{***}$ | $0.0059^{***}$ | $0.0077^{***}$ | $0.0091^{***}$ |
| $\geq 0.20$ | $0.0088^{***}$ | $0.0082^{***}$ | $0.0067^{***}$ | $0.0081^{***}$ | $0.0094^{***}$ |
| $\geq 0.25$ | $0.0104^{***}$ | $0.0083^{***}$ | $0.0074^{***}$ | $0.0087^{***}$ | $0.0097^{***}$ |
| *Panel B. Binned interaction* $\Delta y^{\text{Tsy}}_{n_b,t} \times Long_b$ | | | | | |
| Cutoff \ window | 15d | 30d | 60d | 90d | 120d |
| $\geq 0.05$ | $0.049^{***}$ | $0.061^{***}$ | $0.070^{***}$ | $0.074^{***}$ | $0.088^{***}$ |
| $\geq 0.075$ | $0.056^{***}$ | $0.063^{***}$ | $0.068^{***}$ | $0.074^{***}$ | $0.086^{***}$ |
| $\geq 0.10$ | $0.062^{***}$ | $0.066^{***}$ | $0.069^{***}$ | $0.075^{***}$ | $0.086^{***}$ |
| $\geq 0.15$ | $0.079^{***}$ | $0.077^{***}$ | $0.079^{***}$ | $0.083^{***}$ | $0.092^{***}$ |
| $\geq 0.20$ | $0.094^{***}$ | $0.090^{***}$ | $0.085^{***}$ | $0.088^{***}$ | $0.097^{***}$ |
| $\geq 0.25$ | $0.111^{***}$ | $0.100^{***}$ | $0.096^{***}$ | $0.096^{***}$ | $0.102^{***}$ |

*Notes.* Each cell re-estimates column (1) (Panel A) or (2) (Panel B) of Table 45 on pairs whose maturity-matched treasury move is at least the cutoff, in percentage points, and whose trades are at most the window apart. The headline sample uses 0.10 and 90 days. Standard errors are clustered by trade date as in Table 45. $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

**State clustering.** Brackets in Table 45 give standard errors clustered by issuing state, which allow correlation among a state's bonds across dates. The headline trade-date clustering allows common shocks across states on the same date. Both headline interactions remain significant at one percent, and binned-grid significance levels are unchanged. In the continuous grid of Table 47, the 0.05, 0.075 and 0.10 cutoffs give $p < 0.05$ at thirty days and $p < 0.10$ at sixty days, and the 0.15 cutoff gives $p < 0.05$ at sixty days. All other significance levels are unchanged, so twenty-seven of the thirty continuous cells and all binned cells remain significant at five percent.

**Taxable-market comparison.** I also compare municipal changes with monthly corporate-curve changes. The corporate and treasury curves use monthly averages at the bond's remaining maturity in each

trade month. Keeping cross-month pairs whose absolute monthly treasury change is at least 10 bp gives 517,864 pairs, with the original weights. The maturity interaction is 0.0132 (s.e. 0.0026) against corporate changes and 0.0072 (s.e. 0.0025) against treasury changes on these pairs, both with trade-date clustering.

## C.11 Robustness to the rating cut

This section checks whether the yield pattern and empirical tests depend on the rating cutoff. Both samples use the unrounded mean agency score on the twenty-two-notch scale, on which 1 = AAA, 3 = AA and 7 = A−. The curves in Sections 3 and 5 use scores of seven or better, the A-or-above sample, matching the corporate benchmark's credit band. The Build America Bond (BAB) comparison in Section 3 and the P1, TCJA and P3 tests in Section 6 use scores of three or better, the AA-and-above sample.

The curve pattern survives the narrower cut, and the BAB, P1, TCJA and P3 tests retain their gradient signs on the wider cut. P1's pre-2018 state-tax pass-through declines across all five bins under OLS and IV (Table 48), with endpoint differences significant at five percent under OLS and one percent under IV. In the TCJA insurer test, the maturity gradients are negative and significant under all three estimators with issuer clustering, and Appendix C.9.1 reports state-clustered inference for the AA-and-above sample. The wider sample gives larger P3 interactions and a larger OLS TCJA gradient, with similar BAB, P1 and IV TCJA gradients.

### C.11.1 Curve construction at AA and above

Figure 28 recomputes the AA-and-above curves with the three-stage averaging of Appendix C.1.1: equal weights within cells of maturity, issue quarter and rating, then across ratings, then across quarters. At each maturity, the implied tax is one minus the ratio of these mean municipal to taxable yields. Net of calls and issuance premia, the implied tax falls from 40.1% at two years to 25.5% at thirty, a decline of 14.6 points against 17.1 on the A-or-above sample.

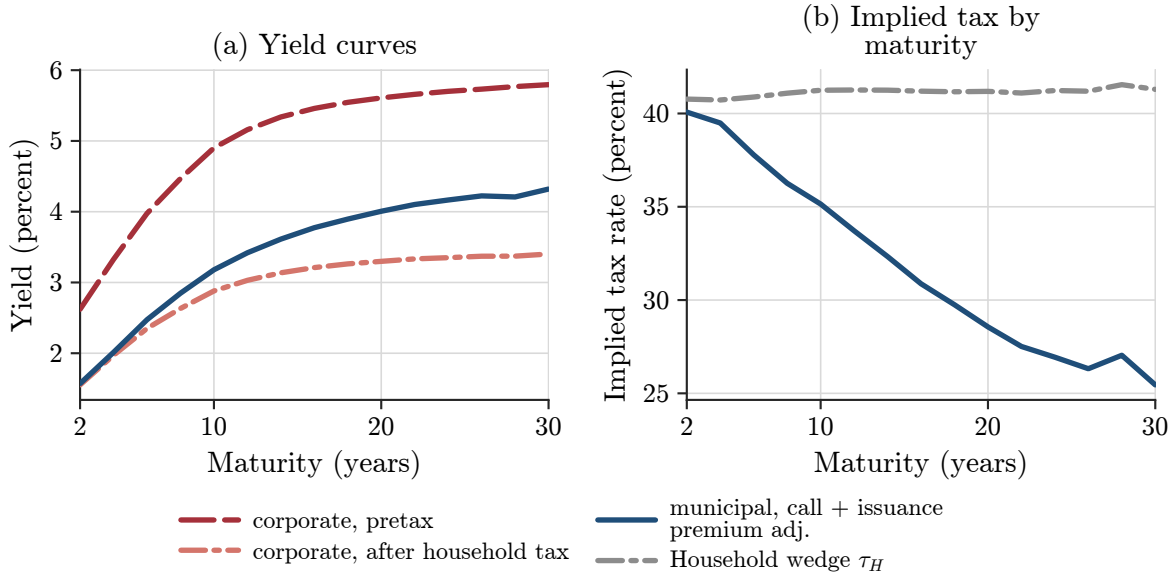



Figure 28: The AA-and-above municipal puzzle, 2000–2020, with this sample's own weights and household benchmark. Panel (a) shows pretax and household after-tax corporate zero-coupon yields and the municipal yield after removing callability and the issuance premium, and panel (b) shows the implied tax $1 - y_T^M / y_T^C$ and the household wedge $1 - \bar{\xi}_H(T)$, labeled $\tau_H$ in the legend. Calls use dated calibrated prices, and the sample's all-size issuance premium is scaled by the full sample's institutional-to-all-size ratio (see text).

The model comparison repeats the accounting of Section 5.4, applying this sample's model ratio of municipal to taxable mean yields to observed corporate yields (the corporate normalization of (22)). Call,

liquidity and default estimates are sample-specific, and capital gains use the common method. The issuance adjustment scales this sample's premium from all trade sizes by the full-sample institutional-to-all-size ratio, giving 11.5 bp at thirty years. The normalized model gap exceeds the residual before the model at every maturity beyond two years, 96.2 against 89.9 basis points at thirty years. The unexplained residuals beyond two years are therefore negative. Fourteen of fifteen absolute residuals are below ten basis points, the thirty-year residual is −6.3 and the root-mean-square residual is 7.1 basis points.

### C.11.2 Empirical tests on the A-or-above sample

Table 32 (Appendix C.2) compares the two BAB rating cuts in columns (1) and (3), with call matching in columns (2) and (4). Table 48 repeats the P1 specification of Appendix C.5, so its AA-and-above columns reproduce the headline P1 estimates. Table 49 repeats the TCJA difference-in-differences of Table 43 with issuer clustering. It omits the headline rating and rating-by-maturity controls in both samples, so its AA-and-above columns differ slightly. Tables 50 and 51 repeat the P3 specification of Table 45 and the binned grid of Table 47 with state-clustered standard errors. The headline P3 estimates cluster by trade date. Their AA-and-above coefficients equal the headline ones.

Table 48: P1 by rating cut, 2000–2017.

| | OLS | | IV | |
|---|---|---|---|---|
| | Score $\leq 3$ | Score $\leq 7$ | Score $\leq 3$ | Score $\leq 7$ |
| $\beta_1$: $[0, 5)$ years | 1.069*** | 0.986*** | 1.088*** | 1.080*** |
| | (0.195) | (0.255) | (0.225) | (0.233) |
| $\beta_2$: $[5, 10)$ years | 0.677*** | 0.607*** | 0.694*** | 0.665*** |
| | (0.069) | (0.100) | (0.086) | (0.117) |
| $\beta_3$: $[10, 15)$ years | 0.484*** | 0.380*** | 0.447*** | 0.384*** |
| | (0.096) | (0.083) | (0.118) | (0.098) |
| $\beta_4$: $[15, 20)$ years | 0.390*** | 0.285*** | 0.326*** | 0.264*** |
| | (0.102) | (0.086) | (0.116) | (0.095) |
| $\beta_5$: $[20, 30]$ years | 0.305*** | 0.240*** | 0.229*** | 0.214*** |
| | (0.075) | (0.077) | (0.081) | (0.079) |
| Decline $\beta_1 - \beta_5$ | 0.764*** | 0.746** | 0.859*** | 0.866*** |
| $p$-value of decline | 0.002 | 0.011 | $< 0.001$ | $< 0.001$ |
| Issue month × maturity × call FE | Yes | Yes | Yes | Yes |
| State cluster | Yes | Yes | Yes | Yes |
| Observations | 681,838 | 1,397,768 | 681,838 | 1,397,768 |
| First-stage $F$ (Kleibergen–Paap) | — | — | 10.4 | 12.6 |

*Notes.* FE denotes fixed effects. Controls, fixed effects and state-clustered inference (47 states) follow Appendix C.5, 2000–2017, with call-adjusted zero yields. Score $\leq 3$ is the AA-and-above sample and score $\leq 7$ the A-or-above sample. IV uses bin-interacted 1990 rates. $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

Table 49: P2: TCJA insurer difference-in-differences, by rating cut

| | OLS | | IV1 | | IV2 | |
|---|---|---|---|---|---|---|
| | $\leq 3$ | $\leq 7$ | $\leq 3$ | $\leq 7$ | $\leq 3$ | $\leq 7$ |
| *Post* × *PCShare* | $-0.118^{***}$ | $-0.161^{***}$ | $-0.124$ | $-0.255^{**}$ | $-0.015$ | $-0.175$ |
| | (0.029) | (0.035) | (0.085) | (0.110) | (0.102) | (0.131) |
| *Post* × *PCShare* × $n_b$ | $-0.018^{***}$ | $-0.036^{***}$ | $-0.037^{***}$ | $-0.039^{***}$ | $-0.038^{**}$ | $-0.037^{***}$ |
| | (0.005) | (0.005) | (0.011) | (0.009) | (0.015) | (0.012) |
| $\lfloor n_b \rfloor$ × *Purchase* FE | Yes | Yes | Yes | Yes | Yes | Yes |
| *Issuer* × *Purchase* FE | Yes | Yes | Yes | Yes | Yes | Yes |
| Quarter FE | Yes | Yes | Yes | Yes | Yes | Yes |
| Issuer cluster | Yes | Yes | Yes | Yes | Yes | Yes |
| Observations | 1,733,881 | 3,043,035 | 1,733,881 | 3,043,035 | 1,733,881 | 3,043,035 |
| First-stage $F$ | | | 25.6 | 45.5 | 17.5 | 34.6 |

*Notes.* FE denotes fixed effects. *Post* is the post-reform indicator *PostTCJA* of equation (151). Standard errors are clustered by issuer (six-digit CUSIP), with 9,903 issuers at score $\leq 3$ and 17,634 at score $\leq 7$. First-stage $F$ is Kleibergen–Paap. Both samples omit Table 43's rating and rating-by-maturity controls. $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

Table 50: P3: observed coupon-yield response to treasury par yields, by rating cut

| | (1) continuous | | (2) bin <5y vs ≥5y | |
|---|---|---|---|---|
| | score $\leq 3$ | score $\leq 7$ | score $\leq 3$ | score $\leq 7$ |
| $\Delta y^{\text{Tsy}}_{n_b,t}$ | $0.346^{***}$ | $0.281^{***}$ | | |
| | (0.012) | (0.011) | | |
| $\Delta y^{\text{Tsy}}_{n_b,t} \times n_b$ | $0.007^{***}$ | $0.013^{***}$ | | |
| | (0.002) | (0.002) | | |
| $\Delta y^{\text{Tsy}}_{n_b,t} \times \mathit{Long}_b$ | | | $0.075^{***}$ | $0.113^{***}$ |
| | | | (0.009) | (0.006) |
| State cluster | Yes | Yes | Yes | Yes |
| Observations | 952,874 | 2,023,637 | 952,874 | 2,023,637 |
| Weights | Inverse pair count | | | |

*Notes.* The estimates use Table 45's first-differenced specification for coupon yields and par yields, weighted by inverse CUSIP pair counts and clustered by state, with $\mathit{Long}_b = \mathbf{1}\{n_b \geq 5\}$. $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

Table 51: P3 observed-response sensitivity grid, binned specification, by rating cut

| | AA and above (score $\leq 3$) | | | | | A− and above (score $\leq 7$) | | | | |
|---|---|---|---|---|---|---|---|---|---|---|
| cutoff | 15d | 30d | 60d | 90d | 120d | 15d | 30d | 60d | 90d | 120d |
| $\geq 0.05$ | $0.049^{***}$ | $0.061^{***}$ | $0.070^{***}$ | $0.074^{***}$ | $0.088^{***}$ | $0.078^{***}$ | $0.099^{***}$ | $0.098^{***}$ | $0.114^{***}$ | $0.121^{***}$ |
| $\geq 0.075$ | $0.056^{***}$ | $0.063^{***}$ | $0.068^{***}$ | $0.074^{***}$ | $0.086^{***}$ | $0.080^{***}$ | $0.098^{***}$ | $0.096^{***}$ | $0.112^{***}$ | $0.118^{***}$ |
| $\geq 0.1$ | $0.062^{***}$ | $0.066^{***}$ | $0.069^{***}$ | $0.075^{***}$ | $0.086^{***}$ | $0.088^{***}$ | $0.101^{***}$ | $0.097^{***}$ | $0.113^{***}$ | $0.119^{***}$ |
| $\geq 0.15$ | $0.079^{***}$ | $0.077^{***}$ | $0.079^{***}$ | $0.083^{***}$ | $0.092^{***}$ | $0.095^{***}$ | $0.110^{***}$ | $0.105^{***}$ | $0.117^{***}$ | $0.123^{***}$ |
| $\geq 0.2$ | $0.094^{***}$ | $0.090^{***}$ | $0.085^{***}$ | $0.088^{***}$ | $0.097^{***}$ | $0.108^{***}$ | $0.118^{***}$ | $0.111^{***}$ | $0.122^{***}$ | $0.128^{***}$ |
| $\geq 0.25$ | $0.111^{***}$ | $0.100^{***}$ | $0.096^{***}$ | $0.096^{***}$ | $0.102^{***}$ | $0.120^{***}$ | $0.126^{***}$ | $0.119^{***}$ | $0.128^{***}$ | $0.134^{***}$ |

*Notes.* Cells report the observed interaction $\Delta y^{\text{Tsy}}_{n_b,t} \times \mathit{Long}_b$ in the regression of coupon-yield changes on par-yield changes, with state-clustered standard errors. Rows give minimum absolute treasury moves in percentage points. Columns give maximum days between consecutive trades, as in Table 47. Every wider-sample coefficient is larger, with the largest proportional gains at the lowest cutoffs. $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

# D Institutional detail and data

This appendix connects the institutional setting of Section 3 to the model and its policy calculations. It first explains federal and state tax treatment, then follows taxation from interest accrual to resale. The fiscal section shows how these rules enter receipts and financing costs under the two reforms. The final sections examine financing and hedging restrictions and list the data sources. Throughout, statutory rules are distinguished from the assumptions used in the quantitative model.

## D.1 Investor-by-investor federal tax treatment since 2000

The federal exemption is worth more to some investors than others. Households face different marginal rates, insurers lose part of the benefit through proration, and banks can lose deductions on financing costs. This subsection describes those rules before adding state taxes and resale taxation below.

Here the shares $\xi_i^C = 1-\tau_i^C$ and $\xi_i^M = 1-\tau_i^M$ are absolute, before the normalization of Section 4.2: they are the fractions of a dollar of taxable and municipal interest that investor $i$ keeps. Pricing uses their ratio $\xi_i = \xi_i^C/\xi_i^M$. The exemption benefit $1-\xi_i = (\tau_i^C - \tau_i^M)/(1-\tau_i^M)$ therefore equals $\tau_i^C$ only when municipal interest is fully exempt. Financing deductions add a separate cost, described for banks below and for other investors in Appendix D.5. Table 2 summarizes the pre-2018 treatment, and Table 52 gives the 2018–2020 comparison.

Table 52: Investor heterogeneity in the municipal market after the 2017 Act

| Investor type | Federal treatment | State treatment | Net exemption benefit |
|---|---|---|---|
| High-tax household | Interest exempt | Exempt on in-state bonds only | 0.35–0.41 |
| Middle-tax household | Interest exempt | Exempt on in-state bonds only | 0.22–0.32 |
| Low-tax household | Interest exempt | Exempt on in-state bonds only | 0.10–0.12 |
| P&C insurer | Exempt, less proration | Premium taxes, no distinction | 0.17 |
| Life insurer | Exempt, less proration | Premium taxes, no distinction | 0.16 |
| Bank | Exempt, funding cost disallowed | Varies by state | 0.16–0.18 |
| Other (pensions, etc.) | Varies | Varies | $< 0.05$ |

*Notes.* Household values are federal marginal tax rates in three groups (35% and above, 22% to 32%, below 22%), and the top value includes the 3.8% net investment income tax. The $10,000 cap on the state and local tax deduction makes the home-state exemption more valuable to households whose cap binds. Institutions face the 21% corporate rate. P&C proration rises to 25% of exempt interest, giving $\xi = 0.79/0.9475$, and the life insurer's policyholders' share is set at 30%, giving $\xi = 0.79/0.937$. The bank value is the break-even implied tax of Table 2 at the 21% rate, with the 2018–2020 financing-tax cost of 0.117% a year and 2018–2020 taxable yields. The Act narrows the difference between the P&C and life-insurer exemption benefits.

**Households.** Municipal interest is generally federally exempt, and taxable bond interest is taxed as ordinary income. A household's federal after-tax shares are therefore approximately $(\xi_i^M, \xi_i^C) = (1, 1-\tau_{oi})$, where $\tau_{oi}$ is its ordinary-income tax rate. Households in both high- and lower-income brackets report exempt interest (Figure 1), so the exemption benefit varies across municipal-holding households.

Table 53 gives top federal rates. The 2001 cuts, accelerated in 2003, lowered the top ordinary rate to 35% in 2003. The rate rose to 39.6% in 2013 and fell to 37% in 2018. The top long-term capital-gains rate fell to 15% in 2003 and rose in 2013. Since 2013, the 3.8% NIIT has applied to taxable interest and many gains, but not to exempt interest. The calibration computes federal and state rates for each Survey of Consumer Finances (SCF) tax unit under each year's law (Appendix B.1).

Table 53: Top federal tax rates relevant to municipal pricing, 2000–2020

| Period | Top ordinary rate | + NIIT (on taxable interest) | Top corporate rate |
|---|---|---|---|
| 2000 | 39.6% | 39.6% | 35% |
| 2001 | 39.1% | 39.1% | 35% |
| 2002 | 38.6% | 38.6% | 35% |
| 2003–2012 | 35.0% | 35.0% | 35% |
| 2013–2017 | 39.6% | 43.4% | 35% |
| 2018–2020 | 37.0% | 40.8% | 21% |

**Banks.** Banks' municipal coupons are exempt, but financing those holdings can be costly after tax. Let $c$ be the annual interest cost allocated to a dollar of holdings and $\tau_B$ the bank's tax rate, as in Section 6.2. Interest expense that finances non-bank-qualified exempt bonds is generally nondeductible. It costs $c$ after tax, against $c(1-\tau_B)$ if deductible, a penalty of $\tau_B c$. Section 265(b)(3) excepts bank-qualified bonds, but Section 291 cuts their deduction by 20%, leaving 80% and an after-tax cost of $c(1-0.8\tau_B)$. Interest expense is allocated to exempt bonds in proportion to adjusted tax basis.[34]

The 15% disallowance under the Tax Equity and Fiscal Responsibility Act became 20% in 1984. The 1986 Act generally denied financial institutions' interest deductions allocable to exempt obligations acquired after August 7, 1986. The regular bank-qualified exception covers designated small-issuer obligations under the $10 million annual limit and other conditions. Temporary 2009–2010 rules expanded eligibility. In the model, each dollar of bank municipal holdings bears a financing-tax cost of $\tau_B$, times the measured ratio of interest expense to assets, times $(1-0.8q_{BQ})$. Here $q_{BQ}$ is the bank-qualified share of those holdings. The aggregate calibration sets $q_{BQ}=0$. Appendix C.7.1 derives how much of the deduction on designated bonds is capitalized in yields.

**Property and casualty insurers.** P&C insurers pay the corporate rate $\tau_c$ on taxable bond interest. Proration cuts their losses-incurred deduction by a percentage of exempt interest, so that percentage of municipal interest is in effect taxed at $\tau_c$. Before the 2017 Tax Cuts and Jobs Act, the percentage was 15%, giving an effective municipal rate of $0.15\times0.35=0.0525$. The Act replaced it with 5.25% divided by the highest corporate rate, or 25% at $\tau_c=21\%$. The municipal after-tax share thus stays at 0.9475 while the taxable share rises from 0.65 to 0.79. The relative after-tax share rises from 0.686 to 0.834 and the exemption benefit falls from 31.4% to 16.6%. Cutting the corporate rate while keeping 15% proration would not reproduce the 2018 law.

The calibration applies the top corporate rate to tax-paying insurers. Before 2018, marginal corporate rates were 15–34% below $100,000 and 34–39% from $100,000 to $18,333,333, including bubble rates that recapture the lower brackets' benefit, and a flat 35% above. Since 2018 the rate is 21% at every taxable-income level.

**Life insurers and other institutions.** Section 812 divides a life insurer's net investment income into company and policyholder shares, and exempt interest reduces its reserve deductions. The pre-2018 calibration uses a one-half policyholders' share, which at a 35% corporate rate gives an effective municipal rate of 17.5%. Absolute after-tax shares are 0.825 on municipal and 0.65 on taxable interest, and the relative after-tax share is 0.788. The post-2017 calibration's 30% policyholders' share and 21% corporate rate give an effective municipal rate of 6.3%. After-tax shares are 0.937 on municipal and 0.79 on taxable interest, and the relative after-tax share is 0.843. Tax-exempt institutions and tax-deferred accounts gain little from the exemption.

[34] See the IRS, Bank Qualified Bonds: Section 265, pp. 13-9 and 13-11. Special rules apply to certain 2009–2010 issues and grandfathered obligations.

## D.2 State exemptions and their implementation in the model

State exemptions add a benefit that depends on both the issuer's state and the investor's residence. This is the source of variation in P1. The distinction between residents' own-state and other-state interest also determines which tax rule changes in the home-state repeal.

**Resident households.** The quantitative model applies a common home-state exemption in every income-tax market. The P1 sample instead excludes the four states with exceptional rules: Illinois, Iowa, Wisconsin and Oklahoma. Thus the empirical comparison uses the jurisdictions where the exemption applies, while the model maintains a common rule across its income-tax markets.

For a household resident in state $s$, let $\tau_s^{\text{st}}$ be its state tax rate and $\xi^{\text{fed}}$ its federal relative after-tax share. The baseline normalizes the home-state municipal share to one and sets the out-of-state municipal share to $1-\tau_s^{\text{st}}$. It gives no federal deduction credit for state tax on that interest, a convention varied below. Through 2017, the share on taxable interest is $\xi^{\text{fed}}(1-\tau_s^{\text{st}})$, formed tax unit by tax unit. After 2017, the model treats the state-and-local-tax deduction cap as binding. The share is then $\xi^{\text{fed}}-\tau_s^{\text{st}}$, as in (143). These are the pricing shares, before the conversion to tax dollars in Appendix D.4.

**Funds and institutions.** Municipal funds generally pass the federal exemption through to shareholders as exempt-interest dividends. State treatment depends on the shareholder's residence and the share of interest from that state's bonds. A single-state fund preserves the home-state exemption for resident shareholders. A national fund passes it through only on the eligible share, which weakens the household's state-tax incentive. Among institutions, tax-deferred accounts and many pensions owe no current state tax on interest. Insurers commonly pay premium taxes, which do not distinguish municipal from taxable interest. Banks' and other corporations' state taxes depend on the tax base and business apportionment. P1 exploits the household exemption that depends jointly on issuer and holder residence (Section 6.1). The model assigns this home-state wedge to households and uses federal proration and financing deductions for insurers and banks.

### D.2.1 Sensitivity to a federal deduction for state tax

Before 2018, itemizers could deduct state income taxes on federally exempt interest against ordinary federal income tax. Not every household gets the deduction's full benefit. Taking the standard deduction, paying the alternative minimum tax or choosing a larger sales-tax deduction can eliminate the added saving.[35]

To test sensitivity, I grant every tax unit the full deduction, ignoring eligibility, valued at the ordinary-income component of its baseline federal marginal rate, which excludes NIIT. I form products tax unit by tax unit before aggregation. I re-solve the six pre-2018 economies, holding all other inputs and calibrated coefficients fixed. The post-2017 binding-cap convention is unchanged.

For the 2015–2017 California high-tax group, the effective after-tax shares with the deduction are 0.498 for taxable interest and 0.919 for out-of-state municipal interest. At a riskless taxable yield of 5%, a municipal yield of 2.71% equates the two after-tax returns, compared with 2.87% for the high-tax group in states without an income tax. The taxable alternative already incorporates the baseline deduction treatment, and the added municipal deduction lowers the yield needed to match it. Residents of high-tax states can therefore accept lower out-of-state yields than residents of states without an income tax. This compresses cross-state yield differences and reduces state-tax pass-through even at the riskless maturity.

With the main sample's 2000–2020 weights and corporate normalization, the thirty-year model gap falls only 0.23 basis points, from 97.39 to 97.16. The model's P1 pass-through, averaged over the six economies with its original taxable-yield normalization, falls from 1.001 to 0.904 at the two-year riskless maturity. At thirty years it falls from 0.460 to 0.416. The deduction lowers the level of pass-through but preserves its declining profile and the large long-maturity gap.

[35] IRS Publication 550 (2017), "State income taxes"; 2017 Instructions for Form 6251, line 3; and 2017 Instructions for Schedule A, line 5. 2016 Instructions for Form 8960, line 9b, restrict NIIT deductions to taxes attributable to net investment income, which excludes exempt interest.

## D.3 Taxation over a bond's life: accrual, basis and resale

The capital-gains adjustment in Appendix C.1.4 requires a distinction between interest and gains on resale. Accrued interest changes the bond's tax basis, so the same accrual is not taxed again as a price gain. This subsection follows that accounting from issuance to maturity and then to a secondary-market sale.

The model instead applies the investor's interest-income after-tax share to the whole return $\mathcal{R} = \mu + u$ in (32). That approximation does not separately track coupons, accrual and resale gains.[36]

**Buying at issue and holding to maturity.** For individuals, federal law separates interest accrual from gains realized at disposition (Sections 171, 1272–1278 and 1288). Consider a bond with face value 100, annual coupon payment $c$ in dollars, original price $\text{AIP}_0$ and annual effective issue yield $y_0$. Under constant-yield accrual, its adjusted issue price is the original buyer's tax basis after each payment and satisfies

$$\text{AIP}_t = (1 + y_0)\text{AIP}_{t-1} - c, \qquad t = 1, \ldots, n, \qquad \text{AIP}_n = 100. \tag{156}$$

A discount, $\text{AIP}_0 < 100$, is OID, and a premium, $\text{AIP}_0 > 100$, is original-issue premium (OIP). Annual interest $y_0\text{AIP}_{t-1} = c + (\text{AIP}_t - \text{AIP}_{t-1})$, including the noncash accrual, is taxed at the ordinary-income rate $\tau_{oi}$:

$$\text{tax}_t = \begin{cases} \tau_{oi} y_0 \text{AIP}_{t-1}, & \text{taxable bond,} \\ 0, & \text{tax-exempt bond.} \end{cases} \tag{157}$$

The basis change $\text{AIP}_t - \text{AIP}_{t-1}$ prevents taxing accrued interest again at redemption. OID raises basis, including for zero-coupon bonds, and OIP lowers it. Amortizing premium on a taxable bond requires a Section 171 election, which the capital-gains adjustment assumes. Amortization on an exempt bond gives no deduction. At maturity, basis equals face, so there is no capital gain. For taxable bonds, de minimis OID below a quarter point per full year can be treated as zero, with the discount recognized as capital gain at maturity. This OID exception generally does not apply to tax-exempt obligations.[37]

**Selling before maturity.** A sale at price $P_h$ at date $h$ realizes the gain $G_h = P_h - \text{basis}_h$. This is the price relative to adjusted basis, not relative to the original purchase price. A holding period of more than one year makes a capital gain long-term, taxed at rate $\tau_{cg}$.

**The next buyer's basis.** A secondary buyer starts with basis equal to the purchase price $P$. If $P$ exceeds face, the premium amortizes at the buyer's yield, mandatorily for exempt bonds and by election for taxable bonds. On an OID bond bought at date $s$ for $\text{AIP}_s < P \leq 100$, the acquisition premium $P - \text{AIP}_s$ reduces the remaining OID by the fraction $(P - \text{AIP}_s)/(100 - \text{AIP}_s)$, leaving basis 100 at maturity.

**Market discount and the tax at sale.** Market discount is the shortfall of the purchase price below $\text{AIP}_s$ for OID bonds or below face for par and premium bonds. It accrues ratably, or at constant yield by election. Unless included in income annually, accrued market discount $D_h$ is taxed as ordinary income at sale or redemption, up to any positive gain. This rule also covers exempt bonds bought after April 30, 1993. Below the de minimis threshold (126), the discount is treated as zero and the gain is capital income. The sale tax therefore separates the accrued market discount from the remaining gain:

$$\text{tax at sale} = \tau_{oi} \min\{\max(G_h, 0), D_h\} + \tau_{cg}[G_h - \min\{\max(G_h, 0), D_h\}], \qquad G_h = P_h - \text{basis}_h. \tag{158}$$

[36] The after-tax share multiplies both the log of the expected gross return, $\log \mathbb{E}[1+R]$, and the centered log-return innovation. Taxing a simple return $R$ gives the after-tax gross return $1+\xi R$, with $\log \log(1+\xi R)$. For a deterministic two-year simple return of 5.4 percent and $\xi = 0.62$, this log differs from $\xi \log(1+R)$ by 3.2 basis points per two-year period (1.6 a year) and zero for a fully exempt leg. Appendix A.1.1 reports the effect of taxing simple returns exactly in the solved model. The return convention matches the empirical implied tax and the after-tax yield mappings of Appendix C.1.4.

[37] IRS Publication 550 (2017), "Original Issue Discount (OID)"; Publication 1212, "De minimis rule." The market-discount threshold in (126) applies separately.

The first term taxes the positive gain up to accrued market discount at $\tau_{oi}$. The second applies $\tau_{cg}$ to the residual gain or loss.

**Losses, institutions and the adjustment.** Individual capital losses offset gains and up to $3,000 of ordinary income a year, with carryforwards (Sections 1211–1212). The capital-gains adjustment assumes full offsets at the capital-gains rate, as in (127). Corporations pay their ordinary rate on gains but offset capital losses only against gains. Banks treat gains and losses on debt as ordinary income (Section 582(c)). The model's institutional after-tax shares apply one rate to the full return. Appendix C.1.4 uses the household accrual and basis distinctions above, with its stated federal rates and loss-offset convention, to calculate the yield adjustment.

## D.4 Fiscal accounting and counterfactual evidence

### D.4.1 Pricing shares and the annual fiscal ledger

Portfolio choice depends on relative after-tax shares, but the tax receipts of Section 7 require absolute tax rates. For example, the pre-2018 P&C relative share of 0.6860 reflects a 5.25% tax on municipal interest and a 35% tax on taxable interest. Treating municipal interest as untaxed and taxing taxable interest at $1-0.6860$, or 31.4%, would preserve the relative share but misstate both receipts. Reforms change absolute tax rates, while preferences and liability hedges stay measured in each investor's baseline payoff units (Section 4.2).

For households, I keep the dated federal marginal rates, state-tax approximations and Social Security taxation proxy of Appendix B.1. State receipts go to the investor's state of residence. Before 2018, the federal tax per dollar of taxable interest equals total federal and state tax, one minus the combined after-tax share, less the state rate. This allocation is consistent with the multiplicative after-tax share and needs no new tax-unit calculation.

**Interest income and tax receipts.** The ledger covers interest on bonds and taxable cash. In the portfolio units of Appendix B.1, type $i$ holds dollar cash $m_i(W_i - \mathbf{1}'x_i)$, its investment budget less its bond positions. Cash interest faces the same absolute tax rates as taxable bond interest. Aggregate funding, bond supplies and the short rate are fixed, so aggregate cash interest is unchanged. Its tax receipts can nevertheless change when cash moves between investors with different tax rates. That change is offset by the change in investors' after-tax cash income, preserving identity (29).

Changes in the bank financing-tax cost are recorded separately as federal receipts and deducted from banks' income. They do not count interest expense or issuer payments a second time. Other liability flows remain outside the ledger.

**Annual financing costs and averaging.** Issuer costs measure annual replacement financing at the solved yields. I sum yield times debt outstanding, $y_T D_T$, across maturities at $z = 0$. This does not reset the coupons, original-issue discount schedules or redemption payments of existing bonds. For both reforms, I weight calendar years equally and divide mean dollar costs by mean dollar receipts. Tables 55 and 56 report the individual periods. Their annual fiscal calculations are distinct from yield changes averaged over each economy's monthly rate states.

### D.4.2 Portfolio reallocation under the two reforms

Table 54 reports how portfolios change after each reform. Its shares describe the solved allocations, while the following tables combine those holdings with tax rates and yields to calculate fiscal effects.

**Home-state repeal.** After home-state repeal, both home- and out-of-state municipal interest are state-taxable, so switching issuers does not avoid tax. High-tax residents increase their riskless share by 0.6–2.7 percentage points across the seven states. Lower-tax residents increase their taxable-bond share by 5.1–14.6 points and reduce their riskless share. For both groups together, the riskless share rises in New Jersey and Pennsylvania and falls in the other five states (Table 54).

Table 54: Portfolio allocation under the two tax reforms

| Panel A. Home-state repeal: average resident allocations, 2000–2017 | | | |
|---|---|---|---|
| | Home-state municipal | All risky | Riskless |
| California | 20.33 → 4.35 | 30.02 → 31.56 | 69.98 → 68.44 |
| New York | 19.65 → 3.92 | 29.97 → 31.18 | 70.03 → 68.82 |
| New Jersey | 18.63 → 1.32 | 31.34 → 31.12 | 68.66 → 68.88 |
| Massachusetts | 18.96 → 1.55 | 29.95 → 30.99 | 70.05 → 69.01 |
| Ohio | 19.10 → 1.09 | 29.75 → 31.00 | 70.25 → 69.00 |
| Pennsylvania | 17.42 → 2.15 | 30.76 → 30.58 | 69.24 → 69.42 |
| Minnesota | 20.19 → 0.64 | 29.60 → 31.38 | 70.40 → 68.62 |
| Panel B. Equal treatment: average investor allocations, 2000–2017 | | | |
| | Risky maturity (years) | Share of long bonds (%) | Riskless |
| High-tax households | 6.71 → 4.91 | 2.6 → 0.0 | 75.31 → 76.96 |
| Lower-tax households | 12.85 → 12.14 | 29.7 → 27.2 | 66.07 → 64.95 |
| Banks | 11.76 → 8.03 | 8.6 → 0.0 | 0.00 → 0.00 |
| P&C insurers | 9.93 → 6.83 | 5.3 → 1.7 | 21.70 → 0.65 |
| Life insurers | 11.81 → 12.07 | 14.6 → 14.0 | 0.38 → 0.68 |
| Tax-neutral investors | 7.73 → 8.53 | 39.2 → 57.1 | 27.90 → 29.56 |

*Notes.* Entries are baseline → reform values. Home-state, all-risky and riskless shares are percentages of each population's modeled financial assets, including cash and excluding outside assets and liabilities. The share-of-long-bonds column instead reports each group's percentage of the total long-bond market. Panel A aggregates resident households within each state and averages their portfolio shares across the six 2000–2017 economies. Each row is a separate repeal. Panel B averages each investor group's allocations across the same economies under 99.9% equal treatment. Both panels weight types by their mass $m_i$ and use rate state $z = 0$. Tax-neutral investors are the residual group of Table 3. Riskless holdings comprise cash and one-period municipal bonds (two years on the numerical grid). All longer bonds are risky. In Panel A, home-state municipal holdings overlap the other columns, and the last two columns sum to 100%. In Panel B, risky maturity is the value-weighted mean maturity in years of the group's risky municipal and taxable bonds. The long-bond market comprises all municipal and taxable bonds maturing in more than ten years, and its shares sum to 100% across groups. Risky municipal and taxable bonds of the same maturity become close substitutes, so Panel B reports their combined maturity profile (Appendix B.6.5).

The high-tax group's mean bond maturity rises by 0.4–0.9 years as it replaces riskless municipal bonds with cash, which is excluded from the bond-maturity average. Its risky bonds become slightly shorter. In California, mean bond maturity rises from 5.99 to 6.44 years, while the mean maturity of risky bonds falls from 7.05 to 6.52. The lower-tax group's mean risky-bond maturity there falls from 14.47 to 11.84 years. The added receipts thus accompany substitution across markets and tax groups, without a shift toward longer risky bonds.

**Equal treatment.** Panel B shows the different reallocation under the equal-treatment reform. The tax-neutral group's share of long municipal and taxable bonds rises from 39.2% to 57.1%, while the high-tax household and bank shares fall to almost zero. The lower-tax household and life-insurer shares decline slightly. These are shares of the long-bond market, not shares of each group's own portfolio. The other columns describe each group's own risky maturity and riskless allocation.

### D.4.3 Home-state repeal: financing costs and receipts

Table 55 separates the seven state experiments by period. Panel A compares own-state and all-issuer costs per dollar of added state receipts. Panel B explains the six-period average by separating tax and yield changes from portfolio reallocation.

Table 55: Home-state repeal: period financing ratios and the decomposition of costs and receipts

| Panel A. Own-state / all-issuer financing costs per added state receipt dollar | | | | | | | |
|---|---|---|---|---|---|---|---|
| State | 2000–02 | 2003–05 | 2006–08 | 2009–11 | 2012–14 | 2015–17 | Average |
| CA | 0.80 / 1.19 | 0.61 / 0.84 | 0.63 / 0.67 | 0.78 / 1.13 | 0.94 / 1.41 | 0.96 / 1.52 | 0.81 / 1.17 |
| NY | 0.74 / 1.15 | 0.64 / 0.89 | 0.59 / 0.49 | 0.78 / 1.11 | 0.94 / 1.38 | 0.92 / 1.41 | 0.78 / 1.08 |
| NJ | 0.82 / 1.35 | 0.79 / 1.23 | 0.80 / 1.11 | 0.83 / 1.18 | 1.00 / 1.49 | 0.98 / 1.46 | 0.88 / 1.30 |
| MA | 0.66 / 1.03 | 0.51 / 0.75 | 0.56 / 0.49 | 0.66 / 0.95 | 0.80 / 1.04 | 0.90 / 1.20 | 0.68 / 0.89 |
| OH | 0.73 / 1.17 | 0.57 / 0.83 | 0.60 / 0.57 | 0.70 / 1.04 | 0.88 / 1.19 | 0.99 / 1.39 | 0.73 / 0.99 |
| PA | 0.66 / 1.11 | 0.56 / 0.88 | 0.66 / 1.02 | 0.63 / 0.87 | 0.90 / 1.12 | 1.02 / 1.32 | 0.73 / 1.03 |
| MN | 0.69 / 1.07 | 0.52 / 0.74 | 0.53 / 0.50 | 0.68 / 0.95 | 0.86 / 1.22 | 0.83 / 1.16 | 0.68 / 0.93 |

| Panel B. Average annual changes ($million) | | | | | | | |
|---|---|---|---|---|---|---|---|
| State | Cost at fixed $y^0$ | Change in $y^0$ | Total cost | Tax at $y^0$ | Yield changes | Reallo-cation | Total receipts |
| CA | 798.7 | 17.2 | 815.9 | 902.5 | 1.0 | 100.8 | 1004.3 |
| NY | 455.4 | 6.6 | 462.0 | 526.0 | 0.1 | 65.2 | 591.4 |
| NJ | 138.4 | 0.2 | 138.6 | 141.8 | -0.2 | 15.2 | 156.9 |
| MA | 76.1 | 0.1 | 76.2 | 98.9 | -0.1 | 14.0 | 112.8 |
| OH | 75.1 | 0.1 | 75.1 | 90.1 | -0.1 | 12.7 | 102.6 |
| PA | 59.8 | -0.3 | 59.5 | 73.0 | -0.1 | 9.0 | 81.9 |
| MN | 59.5 | 0.1 | 59.6 | 75.8 | -0.1 | 11.3 | 87.0 |

*Notes.* Each state-period cell is a separate repeal from its dated baseline. All-issuer costs add other municipal and pooled taxable issuers to the state's own issuers, and receipts are the repealing state's added receipts. The average column divides six-period mean dollar costs by mean receipts. Where the all-issuer ratio is below the own-state ratio, as in several 2006–08 cells, residents' reallocation lowers other issuers' yields and hence their financing costs. Panel B splits the average own-state costs and receipts around each economy's pooled no-income-tax-state municipal yield $y_T^0$. Cost at fixed $y^0$ is $\sum_T D_T y_T^0(-\Delta\tau_T^{\text{impl}})$, with $D_T$ the state's debt and $\tau_T^{\text{impl}} = 1 - y_T^M/y_T^0$, and the remaining cost comes from the change in $y^0$. Tax at $y^0$ applies each resident group's state rate to its baseline home-state holdings at $y_T^0$. Yield changes are the rest of the tax at new yields on those holdings. Reallocation is the net effect of residents' portfolio changes, including cash. Components sum to totals, subject to rounding. All dollar calculations use continuous yields and holdings at $z = 0$ and include residents' taxable cash interest.

**Costs and receipts around the comparison yield.** Panel B of Table 55 splits each state's own-state cost and receipts into the fixed-yield comparison of Section 7.2 and adjustments for yield changes and reallocation. With $D_T$ the state's debt, $y_T^0$ the pooled no-income-tax-state municipal yield, $\tau_T^{\text{impl}} = 1 - y_T^M/y_T^0$ and primes marking post-repeal values, the cost increase is exactly

$$\Delta C = \sum_T D_T\, y_T^0 \left(\tau_T^{\text{impl}} - \tau_T^{\text{impl}\prime}\right) + \sum_T D_T \left(1 - \tau_T^{\text{impl}\prime}\right)\left(y_T^{0\prime} - y_T^0\right).$$

The first term holds $y^0$ fixed, and the second comes from the change in $y^0$. The tax at $y^0$ is $\sum_i \tau_i^{\text{st}} \sum_T H_{iT} y_T^0$, where $H_{iT}$ is resident group $i$'s baseline home-state holding and $\tau_i^{\text{st}}$ its state rate. Beyond the tax at fixed yields, almost all additional receipts come from residents' reallocation. At the comparison yield, own-state cost per dollar of tax ranges from 0.770 in Massachusetts to 0.976 in New Jersey. Reallocation raises receipts by another 10.8% to 14.9% of that tax, so costs stay below receipts in every state. California's additional annual receipts of $1.004 billion comprise $0.227 billion from home-state bonds and net increases of $0.300 billion from other states' bonds, $0.368 billion from taxable bonds and $0.110 billion from cash. Interest on these last three assets was already taxable. The increases are relative to baseline receipts and arise from changes in residents' holdings and the interest they earn, not from a new tax on those assets.

For California, Panel B's $100.8 million reallocation term counts purchases and sales at post-repeal yields. If every bond earned the comparison yield at its maturity and cash kept its taxable yield, the net portfolio effect would be $4.9 million. Taxable yields above this reference contribute $96.8 million and municipal yield differences −$0.9 million. Most of the increase therefore comes from substitution into higher-yielding taxable bonds. These terms decompose the solved allocation, rather than separate policy experiments.

**Differences across states.** At longer maturities, lower pass-through reduces the cost per resident-held dollar, but debt held by nonresidents and institutions adds cost without adding tax. New Jersey's residents hold only 49% of its 12–20 year debt. On that debt, cost is 1.38 times the tax at the comparison yield, which gives New Jersey the highest overall ratio. Minnesota's residents hold 80% of its 22–30 year debt, against 67% in New Jersey, and its lower-tax residents face a higher rate. Cost on those bonds is therefore 0.59 per dollar of tax, against 0.98 in New Jersey. Household composition, outside exposures and resident funding relative to local debt are common across states within each period, so these differences come from tax schedules and market sizes.

### D.4.4 Equal treatment and an issuer subsidy

Table 56 reports the equal-treatment reform and the issuer subsidy discussed in Section 7. The first two columns are yield changes, averaged over monthly rate states. The fiscal columns instead compare annual costs and receipts at $z = 0$. The average row divides mean dollar costs by mean dollar receipts, rather than averaging the period ratios.

The subsidy returns all added federal and state receipts to municipal issuers at the solved equal-treatment prices and portfolios. Per receipt dollar, municipal costs before subsidy are the all-issuer cost less the taxable-issuer cost. Subtracting one dollar then gives the last column. The average municipal cost remaining after subsidy is 0.063 per receipt dollar, and taxable issuers retain their 0.133 increase. The transfer changes issuers' net costs, while gross yields and portfolios remain those of the equal-treatment equilibrium.

Table 56: Equal treatment and a revenue-financed issuer subsidy across the six pre-2018 economies

| | Thirty-year yield increase (bp) | | Costs per receipt dollar | | |
|---|---|---|---|---|---|
| Period | Municipal | Taxable | All issuers | Taxable issuers | Municipal after subsidy |
| 2000–2002 | 199.8 | 38.2 | 1.434 | 0.282 | 0.152 |
| 2003–2005 | 136.9 | 9.3 | 1.125 | 0.089 | 0.036 |
| 2006–2008 | 146.5 | 11.4 | 1.077 | 0.073 | 0.004 |
| 2009–2011 | 152.7 | 11.2 | 1.113 | 0.043 | 0.070 |
| 2012–2014 | 178.9 | 33.9 | 1.245 | 0.170 | 0.075 |
| 2015–2017 | 208.6 | 34.7 | 1.242 | 0.174 | 0.068 |
| Average 2000–2017 | 170.6 | 23.1 | 1.196 | 0.133 | 0.063 |

*Notes.* Each row removes 99.9% of the municipal–taxable tax distinction from its own calibrated baseline. Yield changes average each economy's 36 monthly rate states. Fiscal ratios use continuous yields and holdings at $z = 0$, with bond and taxable cash interest in receipts. The last column subtracts a subsidy equal to all added federal and state receipts from municipal issuer costs. The other columns describe the reform before subsidy. The subsidy leaves gross yields and portfolios unchanged because supplies are fixed. The average row divides calendar-year-weighted mean costs by mean receipts.

## D.5 Financing, short sales and rate hedges

The model rules out short sales and borrowing (Sections 3.5 and 4.1). The institutional question is why investors cannot costlessly eliminate the difference in their valuations of rate risk. Financing a municipal purchase, hedging a long bond and shorting a municipal bond are distinct transactions with different tax and market costs. The nine cases below are grouped accordingly. I then relate the remaining rate exposure to participation and test household swap access in equilibrium.

Household tax rules matter because high-income households hold much of the market. In custodial data, municipal bonds are wealthy households' second-largest liquid risky asset, and their portfolio share rises with wealth (Gabaix et al., 2026). In tax years 2008, 2016 and 2019, returns with adjusted gross income (AGI) of $1 million or more, about 0.3 percent of all returns, reported 26–33 percent of tax-exempt interest.[38]

[38] IRS Statistics of Income, Individual Income Tax Returns, Table 1.4, all returns by size of adjusted gross income, tax years 2008, 2016 and 2019. Tax-exempt interest includes exempt-interest dividends from funds.

Define the after-tax coupon gap $g = c_M - (1-\tau)c_C$, where $c_M$ and $c_C$ are municipal and corporate par coupons of the same maturity and $\tau$ is the investor's marginal tax rate. This coupon measure differs from the 118-basis-point thirty-year puzzle in zero-coupon yields (Section 5.4) and from the model's one-period return gap $g_T$. Carry is computed for equal principal in each leg, before fees, margin and collateral costs. The legs need not match in duration or price risk.

### D.5.1 Financing a municipal position

**(i) Financing municipals with debt or a corporate short.** Let $c_F$ be the interest on a dollar borrowed to buy a dollar of municipal principal. If that interest were deductible, carry would be $c_M-(1-\tau)c_F$. Section 265(a)(2) denies deductions for interest on debt incurred or continued to purchase or carry tax-exempt obligations, so carry becomes

$$c_M - c_F = \{c_M - (1-\tau)c_F\} - \tau c_F. \tag{159}$$

At after-tax parity, $c_M = (1-\tau)c_F$, the denial turns zero carry into a loss of $\tau c_F$. Financing the position with a corporate short does not restore the tax shield. Section 265(a)(5) treats the short seller's compensatory coupon payments $c_C$ as interest under Section 265(a)(2), which leaves carry $c_M - c_C$. This carry is negative at every maturity because the sample-average municipal zero curve, and hence the implied par yield, lies below its corporate counterpart. At thirty years with annual compounding, the municipal and corporate par yields are 4.498% and 5.443%.[39] Municipal fund disclosures tell investors that interest on debt allocable to exempt-interest dividends is nondeductible.[40] While the deduction was available, depository institutions financed municipal holdings at scale, holding 51.8% of municipal bonds in 1972 and $235.1 billion (27.3%) in 1985. After the 1986 denial (Appendix D.1), their holdings fell to $120.4 billion in 1990 and $95.4 billion (7.5%) in 1995. Municipal debt outstanding meanwhile grew from $862 billion in 1985 to $1.27 trillion in 1995.[41]

**(ii) Pairing corporate positions.** If the proceeds of a corporate short buy a corporate bond, the investor's own capital finances the municipal bond, and the remaining trade is the hedge in (iii). If they finance the municipal bond, the denial in (i) applies. Revenue Procedure 72-18 tests the purpose of borrowing through tracing of the proceeds, exempt collateral or the facts and circumstances. Its presumption that business borrowing is not incurred to carry exempt bonds lapses when the borrowing exceeds business needs. If the positions form an offsetting straddle, Section 263(g) also capitalizes interest and carrying charges, net of the positions' taxable income, deferring the deductions.[42]

### D.5.2 Corporate and derivative hedges

Hedging a bond's rate risk is distinct from financing its purchase. An investor can use its own funds to buy the municipal bond and take an offsetting rate position elsewhere. The relevant costs then include the hedge's tax treatment, access requirements and residual risks.

**(iii) Hedging duration with a corporate short.** An unlevered hedge of a municipal bond's duration shorts a long-maturity corporate bond and buys a short-maturity corporate bond with the proceeds. Let $c_\ell$ and $c_s$ be the long and short corporate coupons, so that $g = c_M - (1-\tau)c_\ell$. Let $\delta \in [0,1]$ be the

[39] 26 U.S.C. §265(a)(2), (5) and (b), 2018 U.S. Code. Section 265(a)(5)(B) excepts specified short-sale payments when the taxpayer posts cash collateral that earns no material return. That exception does not remove the opportunity cost of the collateral. See also Hines and Horton (2024) on banks' financing-tax restrictions.

[40] BlackRock Municipal Bond Fund, Inc. (2025), Part II, "Dividends and Taxes," p. II-131; iShares Trust (2026), "Taxes," p. 210; Nuveen Municipal Trust (2026), "Tax Matters," p. S-80.

[41] Board of Governors of the Federal Reserve System, Financial Accounts of the United States (Z.1), release of September 8, 2023, Table L.212, year-end levels. Depository institutions comprise U.S.-chartered depository institutions, foreign banking offices in the U.S. and banks in U.S.-affiliated areas.

[42] Rev. Proc. 72-18, 1972-1 C.B. 740, §§3.01–3.04 and 6.01, as summarized in Revenue Ruling 2004-47 and in IRS Field Service Advice 200129009, which is nonprecedential. On straddles, see 26 U.S.C. §263(g)(1)–(2) and §1092(c)(1)–(2). Positions are offsetting when holding one substantially diminishes the risk of loss from holding the other.

currently deductible fraction of the short payment. With equal principal in each leg, annual carry relative to holding the short corporate bond alone is

$$\Delta = \{c_M + (1-\tau)c_s - (1-\delta\tau)c_\ell\} - (1-\tau)c_s = g - (1-\delta)\tau c_\ell. \tag{160}$$

Denial under Section 265 gives $\delta = 0$ and $\Delta = c_M - c_\ell$, negative at every maturity and a loss of 95 basis points a year at thirty years. Section 263(g) capitalization and Section 163(d) carryforwards defer deductions. Current carry then approaches $c_M - c_\ell$, and later deductions bear a timing and liquidity cost.

Even where neither Section 265 nor Section 263(g) applies, Section 163(d) lets individuals deduct investment interest only if they itemize and up to net investment income, carrying the excess forward. Coupons on the short-maturity corporate bond add to this deduction capacity, while wages and exempt coupons do not. With principal $q$ on both corporate legs, coupon difference $k = c_\ell - c_s > 0$, unused deduction capacity $C \geq 0$ and no prior carryforwards, the hedge's current cash flow is

$$\Pi(q) = -(1-\tau)qk - \tau[qk - C]_+, \tag{161}$$

where $[x]_+ = \max\{x, 0\}$, so $\delta < 1$ once $qk$ exceeds $C$.[43]

Even with $\delta = 1$, the hedge removes only common rate risk. Bond-specific risk (Appendix A.1.1) and credit, liquidity and call differences (Appendix C) remain. Shorting requires available bonds and incurs fees, recall risk and margin demands against an illiquid callable long position.[44] Fees for borrowing corporate bonds are 10–20 bp in the major-lender sample of Asquith et al. (2013).

**(iv) Hedging with swaps.** A pay-fixed, receive-floating swap hedges duration without borrowing or shorting a bond, but household access is restricted. A person that is not an eligible contract participant may trade a swap only on, or subject to the rules of, a designated contract market. An individual qualifies only with more than $10 million invested on a discretionary basis, or more than $5 million when the swap manages the risk of an asset owned or a liability incurred.[45] Ordinary swap payments are not automatically interest on indebtedness, so Section 265(a)(2) alone does not deny their deduction. A household's Section 212 investment expenses attributable to exempt interest fall under Section 265(a)(1).[46] Whatever the classification, let $X$ be a net taxable amount and $d \in [0, 1]$ the denied fraction of the tax saving on its losses, so that its after-tax value is

$$A_d(X) = (1-\tau)X - d\tau X^-, \qquad X^- := \max\{-X, 0\}. \tag{162}$$

Symmetric offsets give $d = 0$. Taxing gains and denying losses gives $d = 1$, which lowers expected proceeds by $\tau\mathbb{E}[X^-]$.

Netting gives the swap a tax advantage when the corporate short payment is disallowed. Suppose both trades have matched periodic receipts $F \geq 0$ and payments $K \geq 0$. When the corporate short payment is wholly nondeductible, the bond trade's after-tax cash is $V_C = (1-\tau)F - K$. Under the rule that taxes net gains and denies losses, the swap's is $V_S = F - K - \tau[F - K]_+$. Hence

$$V_S - V_C = \tau \min\{F, K\} \geq 0. \tag{163}$$

If the corporate payment is fully deductible, the two coincide for $F \geq K$. For $F < K$, the corporate trade benefits from a loss offset that the swap calculation denies. Additional deduction capacity therefore favors the corporate short.

[43] 26 U.S.C. §163(d)(1)–(3), which includes otherwise deductible short-sale expenses in investment interest; IRS Publication 550 (2016), Chapter 3, "Interest Expenses" and "Interest and carrying charges on straddles"; IRS Form 4952 and instructions (2016). Net capital gain and qualified dividends count as investment income only through an election that relinquishes their preferential rates to that extent.

[44] FINRA Rule 4210, paragraph (e)(2), sets margin requirements for bond shorts.

[45] Commodity Exchange Act §2(e) and §1a(18)(A)(xi), 7 U.S.C. §§2(e), 1a(18)(A)(xi); see also the CFTC's 2012 final-rule summary.

[46] Treasury Regulations §1.446-3 and §1.265-1(a)(2), (c). IRS Field Service Advice 200037034 treats the ordinary swap payments before it as non-interest in the absence of a loan. The advice concerns a corporate taxpayer and is nonprecedential.

A pay-fixed swap has the rate exposure of a short fixed-rate bond and a long floating-rate bond. It avoids the corporate short's credit and borrowing exposures but carries its own collateral and counterparty exposure. In the mean–variance objective, the swap weakly dominates any corporate hedge that delivers the same exposure with only an additional nonnegative cost and uncompensated risk orthogonal to the existing portfolio and the swap.[47] An outside dealer can also absorb swap risk without trading domestic bonds, whereas physical shorts change the gross long positions required to clear those markets.

Because the wealthiest households can meet the eligibility thresholds, the equilibrium test of Appendix D.5.5 grants every household a thirty-year pay-fixed swap from an outside dealer with no fee, margin or notional limit. The test taxes swap gains and denies losses ($d = 1$). Physical bond shorts remain excluded, institutions' hedging is unchanged and both bond markets are re-solved. The calculation and results are presented after the residual-risk condition below.

**(v) Treasury futures.** Regulated futures avoid the swap eligibility rules but require cash margin and annual marking to market, with gains and losses treated as 60% long-term and 40% short-term. Individuals offset these losses fully against capital gains, deduct at most $3,000 against ordinary income and carry the excess forward. Net losses on these Section 1256 contracts may be carried back three years against net Section 1256 gains. Section 1092 defers losses to the extent of unrecognized gains on offsetting bonds. Gains can be taxed before losses yield tax benefits. In present value, this timing and liquidity cost equals a partial loss denial, $0 \le d < 1$ in (162).[48]

### D.5.3 Municipal shorts, financing structures and other entities

**(vi) Shorting municipals.** The substitute payment a short seller makes for the exempt coupon loses its exemption. A lender with tax rate $\tau_L$ who requires an after-tax payment equal to the coupon $c_M$ must therefore receive $L$ satisfying

$$(1-\tau_L)L = c_M, \qquad L - c_M = \frac{\tau_L}{1-\tau_L} c_M. \tag{164}$$

A short also needs a willing lender of the specific issue and must meet the contract's collateral terms.[49] The reverse trade shorts a municipal bond and buys an equal-principal corporate bond. Let $\tau_S$ be the short seller's tax rate and $\delta$ the currently deductible fraction of the substitute payment, as in (iii). Under this lender compensation, the trade's coupon carry is

$$\pi_{M\to C} = (1-\tau_S)c_C - (1-\delta\tau_S)\frac{c_M}{1-\tau_L}. \tag{165}$$

Even with full deduction, $\delta = 1$, equal borrower and lender rates $\tau$ give $\pi_{M\to C} = (1-\tau)c_C - c_M$. This carry is nonpositive when $c_M \ge (1-\tau)c_C$, and the carry $c_M - c_C$ of (i) is nonpositive when $c_M \le c_C$, so both coupon trades have nonpositive carry throughout the band

$$(1-\tau)c_C \le c_M \le c_C, \tag{166}$$

and long-end municipal coupons lie in it: above the household after-tax benchmark and, on the sample-average curves, below the corporate coupon. A municipal curve trade, which shorts long-maturity and buys short-maturity municipal bonds, pays the gross-up in (164) on its short leg. To the extent its proceeds buy the short-maturity municipal bond, that payment is nondeductible as in (i).

[47] In common payoff units, let the corporate hedge be $B = \alpha S - \varsigma + \varepsilon$, where $S$ is the swap payoff, $\alpha > 0$, $\varsigma \ge 0$, $\mathbb{E}\varepsilon = 0$, and $\mathrm{Cov}(\varepsilon, W) = \mathrm{Cov}(\varepsilon, S) = 0$ for the investor's existing risk index $W$. With $\mathrm{CE}(Y) = \mathbb{E}Y - \gamma\,\mathrm{Var}(Y)/2$, one has $\mathrm{CE}(W + qB) - \mathrm{CE}(W + q\alpha S) = -q\varsigma - \gamma q^2\,\mathrm{Var}(\varepsilon)/2 \le 0$ for $q \ge 0$, provided the replacement is feasible at no greater cost of financing the position.

[48] 26 U.S.C. §1256(a)(1), (a)(3), (b) and (g)(1); §1211(b); §1212(b)–(c); §1092(a)(1)(A). Section 1256(a)(4) removes Sections 1092 and 263(g) only for straddles consisting entirely of Section 1256 contracts.

[49] IRS Publication 550 (2016), "Reporting Substitute Payments"; Treasury Regulation §1.6045-2(a)(4).

**(vii) Tax-exempt floaters.** A tender-option-bond partnership divides a long municipal bond into variable-rate equity and an inverse interest that receives the residual. Exempt income passes through to both classes, so the structure needs no interest deduction. The inverse holder receives the municipal curve's slope and bears its duration, so capturing $g$ still requires one of the hedges in (iii)–(v). Demand for the floaters limits the scale of the strategy.[50]

**(viii) The de minimis presumption.** For insubstantial holdings, Revenue Procedure 72-18 ordinarily presumes that borrowing is not incurred to carry exempt bonds, absent tracing or exempt collateral. A corporation qualifies only up to 2% of average business assets. Direct evidence overrides the presumption, and larger holdings face the facts-and-circumstances test. Sections 163(d), 263(g), 1092 and 1211(b) still constrain deductions and losses.[51]

**(ix) Other entities.** Financial institutions outside the bank-qualified exception face the pro rata disallowance of Appendix D.1, without tracing. Revenue Procedure 72-18 applies a similar ratio to dealers' business borrowing that cannot be traced. Insurer proration reduces the exemption regardless of financing, through P&C insurers' loss deductions (Section 832(b)(5)(B)) and life insurers' reserve deductions (Sections 807 and 812). Partnerships pass the exempt character of municipal interest through to partners under Section 702(b), but their interest expense remains subject to Section 265(a)(2) under Section 703(a). Over 2000–2020, dealers held on average 0.85% of municipal bonds, with a maximum of 1.94%. Hedge funds, reported from 2012, held at most 0.36%.[52]

### D.5.4 The model constraint and risk after hedging

**The model constraint.** With a concave objective and linear shorting costs, a zero short position is optimal when the risk-adjusted benefit of the first dollar shorted does not exceed the tax, borrowing and financing cost. The quantitative model imposes this corner: bond positions are nonnegative and limited by measured investment budgets, and banks bear the measured financing-tax cost. The household swap extension below tests the mechanism when households can hedge rate risk while these bond-market constraints remain.

**Residual risk and participation after hedging.** The mechanism depends on investors' costs of bearing additional rate risk after they use available hedges. It therefore allows derivative trading. To isolate this tradeoff, suppose the cash constraint is slack and bond-specific risk is zero. Let a hedge remove $q_i$ units of rate exposure at an expected charge $f_i q_i$ and a differentiable convex additional cost $\psi_i(q_i)$, all in the investor's normalized payoff units. With residual exposure $\widehat{D}_i^{\text{hedge}} = \widehat{D}_i - q_i$, where $\widehat{D}_i$ is the net rate exposure (10) before hedging, the hedge-dependent part of the objective is $-f_i q_i - \psi_i(q_i) - \gamma_i \sigma_r^2 (\widehat{D}_i^{\text{hedge}})^2/2$. An interior hedge choice therefore satisfies

$$\lambda_i := \gamma_i \sigma_r^2 \widehat{D}_i^{\text{hedge}} = f_i + \psi_i'(q_i). \tag{167}$$

Here $\lambda_i$ is the unit cost of rate risk at the post-hedge exposure. A binding hedge limit turns the equality into a one-sided condition. When only the high-tax household holds municipal bonds, with hedges and other holdings optimized, the lower-tax investor's municipal slack is its required return on the $T$-period municipal bond minus the offered return. Superscript hh denotes this household-only municipal allocation, giving

$$\nu_{I,T}^{M,\text{hh}} = (\xi_I - \xi_H) r - A_{T-1}^M (\lambda_H^{\text{hh}} - \lambda_I^{\text{hh}}). \tag{168}$$

[50] Revenue Procedure 2003-84, Section 2; 26 U.S.C. §702(b).

[51] Rev. Proc. 72-18, §3.05, as summarized in Revenue Ruling 2004-47 and Field Service Advice 200129009.

[52] Rev. Proc. 72-18, §§5.03 and 7, as summarized in Revenue Ruling 2004-47; 26 U.S.C. §265(b)(1)–(2), §832(b)(5)(B)–(C), §807(a)–(b) and §812. Holdings: Financial Accounts of the United States (Z.1), release of September 8, 2023, Tables L.212 and B.101.f, quarterly levels.

This is the participation comparison (59) with risk costs measured after hedging. Even when both investors hedge, differences in liabilities sustain unequal valuations of duration through unequal hedge access, binding limits or hedging costs that rise with the hedge.

If unlimited hedging in either direction spans rate risk, $\psi_i = 0$ and both investors face a common normalized charge $f$, then $\lambda_H = \lambda_I = f$. At a positive short rate, the slack reduces to the positive tax difference $(\xi_I - \xi_H)r$, which excludes lower-tax municipal holdings. This conclusion requires a common after-tax hedging opportunity. Zero dealer fees alone do not suffice. Observed insurer holdings (Figure 25) therefore require a departure from this benchmark, which any of the documented frictions can sustain.

Insurers already hedge. At year-end 2015, interest-rate swaps accounted for 64% of the notional value of U.S. insurers' interest-rate hedges. Insurers that used derivatives held 87% of life-insurer assets but 19% of P&C assets.[53]

### D.5.5 Household swaps: incentives and equilibrium

The baseline model gives households no separately traded rate hedge. This exercise asks how much household access to a pay-fixed swap changes the mechanism. It has two stages. First, I evaluate a small swap at fixed baseline bond portfolios in the 2015–2017 calibrated economy at $z = 0$, for all 630 household types (14 in each of 45 markets). Second, I let households choose swaps and bonds jointly and re-solve both bond markets. The swap is the pay-fixed, receive-floating contract of case (iv), which gains when rates rise and can offset losses on long bonds.

**Quotes, pricing and taxes.** Swap quotes come from the 95 dates with a two-year rate within 10 bp of the model's mean rate, 3.0206% a year, all in 2001–2008, outside the 2015–2017 block. Fixed rates at ten, twenty and thirty years are 4.8217%, 5.0976% and 5.3736%, and the twenty-year rate is interpolated.[54]

Floating payments are quarterly and fixed payments semiannual. Liquidation after two years uses an approximate calendar. A single curve discounts payments and projects floating rates. Under this pricing rule, the rate state is Gaussian with quarter-by-quarter drifts fitted to the quotes, and the same rule values the swap at liquidation. Quarterly physical rate dynamics preserve the calibrated two-year persistence and variance. Annual net gains and the second-year liquidation payoff are taxed as ordinary income. Losses are denied in the baseline and offset symmetrically in a sensitivity check. Municipal funds likewise distribute swap gains as taxable ordinary income or capital gains and allocate net payments pro rata against taxable and exempt income.[55] Separately held swap cash earns taxable interest when positive and pays nondeductible interest when negative. Existing cash does not change in this fixed-portfolio calculation.

**Stage 1: the incentive at fixed bond portfolios.** Let $\widetilde{W}_i$ be the calibrated log-return risk index, including outside exposure $-h_i z'$, and $M_i^0$ its fixed baseline mean. Here $h_i = -\ell_i$ is the outside rate exposure of Appendix B.1 (positive values add rate risk). The absolute home-municipal after-tax share $\xi_{i,\mathrm{abs}}^M$ is $\xi_i^M$ of Appendix D.1 on home-state bonds. Multiplying it by the relative after-tax share $\xi_i$ gives the swap's absolute after-tax share. Home-municipal after-tax shares are 0.999905 for high-tax and 0.990429 for lower-tax households, and swap tax rates span 21.49%–50.25%. For a pay-fixed notional $q \geq 0$, after-tax cash $A_i$

[53] NAIC, "Update on the Insurance Industry's Use of Derivatives and Exposure Trends," year-end 2015, Tables 1 and 8 and the accompanying discussion. Notional shares measure instrument use and are silent on how much liability risk is eliminated.

[54] Federal Reserve H.15 DSWP10 and DSWP30 are mid-market fixed rates against three-month London Interbank Offered Rate (LIBOR) and differ from household quotes.

[55] BlackRock Municipal Bond Fund, Inc. (2025), Part II, pp. II-20, II-126 and II-132; iShares Trust (2026), "Taxes," p. 213.

per unit notional (unrelated to the rate loading $A_T^a$) changes normalized wealth by $qA_i/\xi^M_{i,\text{abs}}$. The objective and the score $b_i$, the objective's annualized slope at zero notional in basis points, are

$$\mathcal{W}_i(q) = \xi^M_{i,\text{abs}} \left[ M_i^0 + \frac{q\mathbb{E}A_i}{\xi^M_{i,\text{abs}}} - \frac{\gamma_i}{2} \operatorname{Var}\left( \widetilde{W}_i + \frac{qA_i}{\xi^M_{i,\text{abs}}} \right) \right], \tag{169}$$

$$b_i = \frac{10^4}{2} \mathcal{W}_i'(0+) = \frac{10^4}{2} \left[ \mathbb{E}A_i - \gamma_i \operatorname{Cov}(\widetilde{W}_i, A_i) \right]. \tag{170}$$

The score compares the swap's expected after-tax cash payoff with its contribution to portfolio risk. A positive score means that a small pay-fixed position improves the objective at the fixed bond portfolio. It is a marginal benefit per unit notional, not a return on the household's whole portfolio. The score uses the model's log-return approximation. Independent bond-specific shocks remain in baseline risk but have zero covariance with a rate-only swap.

Table 57: Fixed-portfolio swap benefits with gains taxed and losses denied.

| Household group | Years | Mean payoff | Risk benefit | Total | Median | Positive |
|---|---|---|---|---|---|---|
| High tax | 10 | −408.6 | 852.8 | 444.3 | −14.4 | 135/315 |
| High tax | 20 | −473.5 | 996.0 | 522.6 | −13.1 | 135/315 |
| High tax | 30 | −473.9 | 988.2 | 514.2 | −17.2 | 135/315 |
| Lower tax | 10 | −387.9 | 675.6 | 287.7 | −28.8 | 135/315 |
| Lower tax | 20 | −450.0 | 787.6 | 337.7 | −31.5 | 135/315 |
| Lower tax | 30 | −449.5 | 783.3 | 333.9 | −33.3 | 135/315 |

*Notes.* Annualized scores $b_i$ of equation (170), in basis points per unit notional. Years is the swap tenor. Mean payoff and risk benefit are the score's two terms and total is their sum, each averaged within the group with portfolio-asset weights. The median and the count of positive scores weight types equally. Entries are rounded separately.

**Results.** At every tenor (Table 57), 360 types have negative and 270 positive scores. The positive-score types, the top three outside-duration bins of both groups, hold 42.9% of household portfolio assets and 9.8% of household municipal holdings. Individual scores reach 3,100 bp, but a finite trade also bears the quadratic term $-\gamma_i q^2 \operatorname{Var}(A_i)/(2\xi^M_{i,\text{abs}})$. Most types have negative scores, giving a negative median. Because the positive-score types gain more, the mean weighted by portfolio assets is positive.

Relative to symmetric offsets, loss denial lowers expected swap cash for every type. Through the covariance term, it lowers the risk component of the score for 450 types and raises it for 180. Overall, loss denial lowers 585 of the 630 scores. Doubling pricing resolution and draws preserves signs (maximum change 2.57 bp), as does using terminal dollar payoffs and their covariance. Widening the quote window to 25 bp gives 285 positive scores at ten years, 302 at twenty and 275 at thirty, some near zero.

**From a marginal incentive to a finite hedge.** At fixed bond prices and holdings, let $c_i \geq 0$ be any additional two-year certainty-equivalent cost per unit notional, in $A_i$'s absolute units. Maximizing (169) net of this cost gives

$$q_i^*(x; c_i) = \frac{\xi^M_{i,\text{abs}}}{\gamma_i \operatorname{Var}(A_i)} \left[ \mathbb{E}A_i - \gamma_i \operatorname{Cov}\left( \widetilde{W}_i(x), A_i \right) - c_i \right]_+. \tag{171}$$

Here $[u]_+ = \max\{u, 0\}$, as in (161), and $\widetilde{W}_i(x)$ is the risk index of bond portfolio $x$. Costs reduce hedge demand at given bond holdings, and margin and position limits restrict feasible hedges.

**Stage 2: equilibrium with household hedge access.** With the 2015–2017 calibration fixed, households jointly choose bonds and nonnegative thirty-year pay-fixed swap notionals. Outside dealers quote swaps without trading the modeled bonds. A new swap has zero upfront value at the current rate state's par rate. Every contract is priced with the drifts fitted to the historical quotes, counted from its start date, and liquidation uses the same rule. Bond prices in both markets and all portfolios adjust across rate states, with nonnegative bonds and cash, taxed gains, denied losses and separate interim financing.

Swap access reduces the thirty-year gap by 9.48 bp, or 9.1% (Table 58). These are the 2015–2017 model curves, before the full-sample corporate normalization. Thirty-year municipal yields fall 15.95 bp and taxable yields 12.04 bp. At $z = 0$, 270 types hedge, but over 99.99% of their municipal holdings remain riskless two-year bonds, because these types mainly hedge outside exposure. Portfolio responses at other rate states affect continuation prices.

Table 58: Household swap access and the equilibrium municipal gap.

| Maturity | Baseline | Swap access | Change |
|---|---|---|---|
| 2 years | 0.00 | 0.00 | 0.00 |
| 10 years | 23.68 | 22.24 | −1.44 |
| 20 years | 72.63 | 65.42 | −7.20 |
| 30 years | 104.27 | 94.79 | −9.48 |

*Notes.* Entries are the supply-weighted municipal yield minus the household benchmark, in basis points, averaged over the 36 monthly rate states of 2015–2017 with continuous yields. After-tax shares, parameters and dollar supplies are fixed; both yields adjust. Changes are computed before rounding. Refining the simulation fourfold moves either yield curve or the gap by less than 0.003 bp, and price and optimality residuals are below $10^{-10}$.

**Scope of the equilibrium test.** The test includes asymmetric taxes and nondeductible interest on negative swap cash but excludes dealer fees, bid–ask spreads, margin, collateral opportunity costs and notional limits. The largest mean-state notional is 6.88 times the type's budget, so the test substantially expands household access.

These frictions oppose risk transfer and are expected to dampen yield responses. The equilibrium gap depends on both yields, and the 9.1% reduction excludes these costs.

## D.6 Data sources and analysis coverage

Table 59 records what each source measures and its available calendar coverage. Bond characteristics identify the securities and their contractual terms, holdings and balance sheets describe investors, and transactions and curves supply pricing observations. The table's coverage is not a common estimation window. Appendix B explains the calibration inputs, and Appendix C gives the matched samples for the adjustments and empirical tests. Each analysis states its own estimation window.

Table 59: Data sources, variables and coverage.

| Source | Variables | Coverage |
|---|---|---|
| Municipal issues: Mergent Municipal Bond Database (FTSE Russell) | CUSIP, issue & maturity dates, coupon, callability, tax status (federal/state), issuer state, purpose | 1986–2023 |
| Corporate issues: Mergent FISD (FTSE Russell) | CUSIP, issue & maturity dates, coupon, callability, rating, issuer industry, seniority | 1986–2023 |
| Insurers: NAIC Schedule D (annual, Q1–Q3 supplements) | CUSIP-level par, book, and market values for U.S. P&C and life insurers | 1991–2023 |
| Banks: FDIC BankFind (Call Reports RI, RC-B) | Annual interest expense, average assets and municipal securities holdings by bank | 2000–2021 |
| Public pensions: FOIA requests and annual financial reports | CUSIP-level municipal/taxable par and market values, 30 state/local plans (coverage broadens from 2015) | 2000–2023 |
| Sector holdings: Financial Accounts (Z.1, L.212, B.101.f) | Quarterly municipal holdings: households, funds, banks, P&C, life, state/local pensions, other | 1972–2022 |
| Households: SCF, Federal Reserve | Tax-unit income and filing status, financial assets and debt | 2001–2019 waves |
| Household taxes: IRS SOI | Filer counts, taxable and tax-exempt interest, effective federal rates by AGI bracket | 2000–2022 |
| Municipal trades: MSRB RTRS | Trade-level price, par volume, customer-side flag, trade-condition codes | 2005–2023 |
| Corporate trades: FINRA TRACE Enhanced (Academic) | Trade-level price, par volume, customer-side flag, condition codes | 2002–2023 |
| Treasury curve: GSW, Federal Reserve | Daily fitted par yields, 1–30y, used in P3 | 2000–2023 |
| Short rates: TB3MS, FRED | Monthly three-month treasury bill rates for the rate process | 1952–2024 |
| Treasury supply: Monthly Statement of the Public Debt | Security-level amounts and maturities for the treasury maturity profile | 2001–2020 |
| Corporate zeros: HQM (FRED), Mergent FISD robustness | Monthly A-or-above zero curve, 1–30y | 2000–2023 |
| Credit risk: Markit/IHS | State-government and AAA, AA and A corporate CDS at 1–30y tenors, coverage of the matched adjustment sample | 2004–2020 |
| Expected rates: Kim–Wright, Federal Reserve | Monthly fitted yields and term premia, expected-average-short-rate component capped at 10y | 2000–2021 |
| Swaps: Federal Reserve H.15 (FRED) | Par interest-rate swap rates for household swap pricing | 2001–2008 |
| State taxes: Tax Foundation | Top marginal rate, bracket thresholds, treatment of out-of-state municipal interest | 2000–2022 |
| State income: BEA | Quarterly state-level personal income | 2000–2022 |
| Historical state rates: Institute on Taxation and Economic Policy (ITEP) | Top statutory rate by state, the P1 instruments | 1945–2000 vintages |